

When dendritic structure helps local credit assignment

Houman Safaai^{1,*}, Maceo Richards¹, and Bernardo L. Sabatini^{1,2,*}

¹Kempner Institute for the Study of Natural and Artificial Intelligence at Harvard University, Cambridge, MA, USA

²Department of Neurobiology, Howard Hughes Medical Institute, Harvard Medical School, Boston, MA, USA

*Correspondence: houman_safaai@harvard.edu; bernardo_sabatini@hms.harvard.edu

Abstract

Learning requires a neuron to assign a circuit-level error to individual synapses, but the contribution of its dendritic tree to this computation is unknown. We derive exact gradients for conductance-based trees and show that each synaptic update factorizes into local eligibility and a transported compartment error. This factorization yields a credit-operator theory that predicts when a restricted route improves learning: retained task signal must outweigh representation error, gain mismatch and admitted noise. Controlled simulations show that neuron identity supplies most of the useful feedback, whereas subtree addresses and serial dendritic depth help only at intermediate bandwidth or when they match the task hierarchy; depth alone is costly. Reconstructed arbors provide sparse, predominantly coarse addresses, while focal shunting regulates their gain only in permissive electrotonic regimes; measured visual responses provide no evidence that these anatomical routes are aligned with endogenous task credit. Dendritic structure therefore does not generically solve credit assignment. It supplies a conditional routing resource whose value is jointly determined by coordinate precision, address resolution, conductance state and task-morphology alignment.

Introduction

A biological neuron must distribute credit for its output errors across thousands of synapses arrayed on a branching dendritic tree, a problem that the point units of standard neural networks do not face [1–3]. Each synapse samples presynaptic drive, local voltage, reversal potential and input resistance. By increasing local conductance, inhibition can shunt excitation and alter the gain of every dendritic path through that compartment [4, 5].

Backpropagation applies the chain rule from an objective to every parameter [6]. It is therefore the exact optimization reference, not itself a biological mechanism: a literal implementation requires parameter-specific derivatives, coordinated backward Jacobians and suitable forward-feedback relationships. Synapses instead access presynaptic activity, local postsynaptic state and modulatory signals of uncertain spatial precision [7, 8]. Credit assignment in a neuron consequently has two spatial levels: identifying the neuron that should change and distributing that coordinate through its arbor. We therefore ask when a restricted signal can approximate the exact dendritic error field.

Biologically motivated approaches relax different requirements of backpropagation. Three-factor rules combine local eligibility with a modulator [9]; feedback alignment replaces transposed weights with alternative pathways, requiring positive functional alignment rather than exact symmetry [10, 11]; predictive, equilibrium and prospective schemes infer desired neural states [12–14]; and eligibility-propagation methods separate online synaptic traces from neuron-specific learning signals [15]. Dendritic-error and burst models separate forward and teaching signals through compartments, interneurons or temporal codes [16–20]. Most deliver a neuronal coordinate or use a few named compartments, leaving open what a full arbor contributes after credit reaches one neuron.

We first derive the exact gradient required by a conductance-based dendritic network. Each gradient factorizes into *local eligibility* $\times$ *transported error*: eligibility depends on presynaptic activity, driving force and input resistance, whereas a neuronal teaching coordinate is transported through path-specific gains. The factorization turns dendritic credit assignment into a concrete approximation problem: which spatial error fields can a restricted biological signal represent, and when do those fields improve learning?

We formalize a coordinate–address–gain hierarchy inside a common credit-operator theory. Teaching signals provide signed neuronal coordinates; dendritic ancestry provides nested candidate addresses; and conductance determines route gain. The theory predicts a phase boundary: a restricted route helps stochastic learning only when the useful gradient signal it retains compensates for its representation bias, coefficient error and admitted noise. Here, “phase” means an operational regime in alignment–bandwidth–noise space—not a temporal training phase or a thermodynamic phase transition. The theory also distinguishes three possible roles of depth—additional address scales, nonlinear forward composition and the cost of transporting a signal along more variable paths.

The experiments walk this hierarchy in order. Controlled simulations separate coordinate bandwidth, neuron-to-tree assignment and task-dependent routing between sibling subtrees; a hierarchy-depth experiment tests the predicted signal–noise crossover; and an experiment that adds physical dendritic stages at fixed synaptic, parameter and state budgets asks when serial nonlinear stages improve forward learning. MICrONS reconstructions then test route capacity, focal shunting and task alignment; an external six-animal dataset tests for signed neuron-specific coordinates. The evidence is conditional: neuron identity supplies most useful information, dendrites offer sparse addresses, shunting regulates gain only in suitable electrotonic regimes, and anatomy helps only when its routes align with task credit (Fig. 1).

Results

We proceed from the error available to a point neuron to the structured field required by its synapses. A global scalar conveys only error sign or magnitude; a neuron-indexed coordinate identifies the receiving cell; a within-tree address selects a subtree; and route gain sets the strength delivered along that path. Theory defines these levels, controlled networks isolate them, and anatomical and functional data test whether real trees supply and use the predicted routes. Study chronology, numerical gates and exclusions are reported in Methods.

Exact transported compartment error is used only as an information ceiling, never as a proposed biological signal. Every routed advantage reported here is a resource claim rather than a claim of dendrite-exclusive computation: a point model explicitly supplied with the same routing fields and gains matches it by construction, so the question is what anatomy makes available, and at what wiring cost.

Core notation. Indices u , n and i denote a neuron, a compartment and a synapse, respectively. K is the number of communicated feedback channels; H is task-hierarchy depth; D_r is the number of address levels used by a feedback route; and D_p is the number of physical dendritic stages in a trained conductance model. D1, D2 and D3 therefore mean one, two and three serial dendritic stages, not three cell types. Bracketed labels give branching factors from soma to tips: D1 [8], D2 [2, 3] and D3 [2, 1, 2] all contain eight nonsomatic branch units. The symbols δ_u and $\delta_{u,k}$ denote a neuron-level teaching coordinate and its route- k refinement; $\Phi_{u,K}$ is the corresponding K -route dictionary and P_Φ its projector. Bold $\mathbf{g} = \nabla \mathcal{L}$ denotes a gradient, whereas scalar g_i denotes a conductance. The Supplementary Information contains the complete symbol table.

Exact dendritic gradients factor into eligibility and transported error

For a point neuron u with drive $s_u = \sum_i w_i x_i$, activity $V_u = f(s_u)$ and error $\delta_u = \partial \mathcal{L} / \partial V_u$, the chain rule gives

$$\frac{\partial \mathcal{L}}{\partial w_i} = \underbrace{x_i f'(s_u)}_{\text{synapse-local eligibility}} \underbrace{\delta_u}_{\text{one coordinate shared by neuron } u}. \quad (1)$$

Backpropagation computes δ_u from the downstream Jacobian, whereas biological models approximate it. A point unit has no subcellular state: all its synapses share this coordinate and differ only through eligibility. We ask how a dendritic tree extends that credit.

69 The local factorization is not restricted to one feedforward tree. Let the steady state of a compartmental neuron
 70 satisfy

$$\mathbf{F}(\mathbf{V}, \mathbf{g}, \mathbf{x}) = \mathbf{0}, \quad J_V = \frac{\partial \mathbf{F}}{\partial \mathbf{V}}, \quad J_V^\top \mathbf{q} = \nabla_V \mathcal{L}. \quad (2)$$

71 This fixed-point adjoint is the steady-state member of the Almeida–Pineda recurrent-backpropagation lineage
 72 [21, 22]. Intuitively, q_n measures how much the loss would change if a small current were injected at compartment
 73 n . When J_V is nonsingular, so that the steady state varies differentiably with its parameters, implicit differentiation
 74 gives the exact parameter gradient

$$\frac{\partial \mathcal{L}}{\partial g_i} = -\mathbf{q}^\top \frac{\partial \mathbf{F}}{\partial g_i}. \quad (3)$$

75 For a conductance synapse at compartment n , using the residual convention $\mathbf{F} = \mathbf{G}\mathbf{V} - \mathbf{b}$, this becomes

$$\frac{\partial \mathcal{L}}{\partial g_i} = \underbrace{x_i(E_i - V_n)}_{\text{synapse-local factor}} \underbrace{q_n}_{\text{compartment adjoint}}. \quad (4)$$

76 The directed path product below is the triangular-tree corollary (an acyclic directed tree whose Jacobian can be ordered
 77 triangularly) of Eq. 2. In a reciprocal passive cable, $J_V = G$ is symmetric positive definite and $G^{-\top} = G^{-1}$, giving
 78 the Green’s-function adjoint used for reconstructed cells. For a general nonsymmetric or recurrent compartmental
 79 system the inverse transpose is essential. For the triangular tree system used below, $q_n = R_n^{\text{tot}} \partial \mathcal{L} / \partial V_n$, so Eq. 4 and
 80 the compartment factorization of Eq. 6 are the same identity with the input-resistance factor moved between the
 81 local and transported terms. Thus dendritic credit is always local eligibility multiplied by a compartment-specific
 82 adjoint; only the construction of that adjoint changes.

83 We replace the point unit by a directed, feedforward rooted tree [23] whose compartment n receives synaptic
 84 conductances g_i with presynaptic activities x_i and reversal potentials E_i . Let $\mathcal{S}_n$ and $\mathcal{C}_n$ denote its synapses and child
 85 compartments, respectively, and let g_n^L and E_L denote leak conductance and reversal potential. At the conductance
 86 stage, its steady-state voltage is

$$V_n = R_n^{\text{tot}} \left(\sum_{i \in \mathcal{S}_n} g_i x_i E_i + \sum_{c \in \mathcal{C}_n} g_{c \rightarrow n}^{\text{den}} a_c + g_n^L E_L \right), \quad R_n^{\text{tot}} = (g_n^{\text{tot}})^{-1}, \quad (5)$$

87 where $a_c = f_c(V_c)$ is the activity transmitted by child c . Here $g_n^{\text{tot}} = g_n^L + \sum_{i \in \mathcal{S}_n} g_i x_i + \sum_{c \in \mathcal{C}_n} g_{c \rightarrow n}^{\text{den}}$ includes leak,
 88 active synaptic conductance, and child coupling. Differentiating with respect to a synaptic conductance gives

$$\frac{\partial \mathcal{L}}{\partial g_i} = \underbrace{x_i R_n^{\text{tot}} (E_i - V_n)}_{\text{synapse-local eligibility}} \underbrace{\frac{\partial \mathcal{L}}{\partial V_n}}_{\text{compartment error}}. \quad (6)$$

89 Presynaptic activity, driving force, and input resistance are local to the synapse and its compartment. All remaining
 90 task dependence is contained in $\partial \mathcal{L} / \partial V_n$.

91 Because a tree has one route from each compartment to the soma, its exact compartment error is

$$\frac{\partial \mathcal{L}}{\partial V_n} = \delta_{0,u} \tilde{\alpha}_n, \quad \tilde{\alpha}_n = \prod_{(i \rightarrow k) \in \text{path}(n \rightarrow 0)} f'_i(V_i) R_k^{\text{tot}} g_{i \rightarrow k}^{\text{den}}, \quad (7)$$

92 where $\delta_{0,u}$ is the somatic error for neuron u . Exact reconstruction from Eqs. 6–7 matched automatic differentiation
 93 across 100 diagnostic runs: median relative gradient error was 7.43×10^{-8} and the maximum was 1.88×10^{-7}
 94 (Fig. 2a).

95 Equation 7 applies to the directed artificial network used in the regular-tree experiments. The reconstructed-cell
 96 analyses below instead use a reciprocal passive cable matrix grounded in classical dendritic cable theory [24]. In that
 97 model the compartment error is an exact Green’s-function adjoint rather than this simple product, but the synaptic
 98 gradient retains the same separation between local sensitivity and a soma-derived compartment error.

99 This identity defines what a local learning rule must approximate. Our three-factor rule retains the exact eligibility
 100 in Eq. 6 and replaces the compartment error by an available teaching field. Additional four- and five-factor variants
 101 use slowly estimated reliability preconditioners, but do not alter the fast local factorization. Exact path transport is
 102 used only as an information oracle.

103 **From exact gradients to restricted credit operators.** The factorization identifies the object that differs among
 104 local rules. Let $\hat{\mathbf{g}} = \mathbf{g} + \boldsymbol{\xi}$ be an unbiased stochastic gradient with covariance Σ , and let a fixed credit operator M
 105 encode the available coordinate, route assignment and gain. The update is $\mathbf{w}^+ = \mathbf{w} - \eta M \hat{\mathbf{g}}$. If the population loss
 106 has L -Lipschitz gradient, the descent lemma gives

$$\mathbb{E}\mathcal{L}(\mathbf{w}^+) \leq \mathcal{L}(\mathbf{w}) - \eta \mathbf{g}^\top M \mathbf{g} + \frac{L\eta^2}{2} \left(\|M\mathbf{g}\|_2^2 + \text{tr}(M\Sigma M^\top) \right). \quad (8)$$

107 When $\mathbf{g}^\top M \mathbf{g} > 0$, optimizing this bound over η yields the phase utility

$$U(M) = \frac{[\mathbf{g}^\top M \mathbf{g}]^2}{2L\{\|M\mathbf{g}\|_2^2 + \text{tr}(M\Sigma M^\top)\}}. \quad (9)$$

108 Equation 8 is an upper guarantee for a general L -smooth loss and is exact for the corresponding quadratic when
 109 its curvature equals L on the update directions; comparisons of $U(M)$ otherwise order guarantees, not necessarily
 110 realized losses. Here and throughout, a *phase* means a regime separated by a quantitative crossover—for example,
 111 whether rejected noise is larger or smaller than discarded signal. It does not denote a thermodynamic phase transition.
 112 Thus more address dimensions are predicted to help only when their retained task signal outweighs finite-step gain
 113 and admitted gradient noise. Exact full-batch backpropagation has $M = I$ and $\Sigma = 0$; at a matched infinitesimal update
 114 norm no alternative direction can provide greater first-order descent. A projected stochastic gradient can nevertheless
 115 outperform an unprojected stochastic gradient because it rejects noise. For an orthogonal projector P on an identity-
 116 curvature quadratic at the common step $\eta = 1/L$, the exact crossover is $\text{tr}[(I - P)\Sigma] > \|(I - P)\mathbf{g}\|_2^2$: rejected
 117 noise must exceed discarded signal (at smaller common steps the rejected-noise threshold rises; Supplementary
 118 Information). Backpropagation followed by the same projection is therefore the relevant upper control for any
 119 projected local rule.

Scope of the credit-operator prediction. One-step geometry predicts within-family progress and large contrasts caused by retained signal, rejected noise or gross misalignment. It does not predict every advantage accumulated along a training trajectory; in particular, the trained $K = 4$ ancestry advantage is outside the initialization bound’s scope.

Capture vocabulary. *Field capture* is the fraction of exact mean-field energy retained by a dictionary projection; *spectral capture* averages this fraction over a credit covariance; *checkpoint capture* applies it to a trained gradient; and *capture per wiring* divides field capture by feedback-connection density. We use an explicit qualifier whenever the estimand is not field capture.

121 For a route dictionary Φ , topology controls the span $\text{col}(\Phi)$, whereas learning must also estimate the route
 122 coefficients. The total weighted field error separates exactly into the irreducible projection residual and the
 123 within-span coefficient-estimation error (Supplementary Information, “Stochastic credit-operator phase theory”).
 124 Multiplying fixed route columns by nonzero scalar shunting gains leaves their span and spectral capture unchanged,
 125 although it can strongly alter Gram conditioning. By contrast, state- or branch-dependent gain can act as a reliability
 126 preconditioner. For disjoint branches with signal energy S_b , noise energy N_b and fixed step $\eta = c/L$, Eq. 8 is
 127 optimized over attenuations $0 \leq a_b \leq 1$ by $a_b^* = \min\{1, S_b/[c(S_b + N_b)]\}$; the reliability factor $S_b/(S_b + N_b)$ is the
 128 special case $c = 1$, paralleling the forward reliability weighting derived for conductance-based dendrites [25]. A point
 129 model supplied with the identical branch gates must match; the biological question is whether dendritic conductance
 130 can make such localized gain available.




Figure 1: **From neuronal coordinates to dendritic credit addresses.** **A**, Local synaptic states are distributed across a branching tree, whereas the task supplies an error at the neuronal output. **B**, A population layer of dendritic excitatory and inhibitory units supplies a task readout; output error first arrives at each neuronal soma. **C**, Once a neuron-level coordinate δ_u arrives, dendritic subtrees provide distinct addresses and path conductances scale route gain. The following experiments test these levels in that order before asking when additional route resolution is useful. **D**, The central factorization as a teaching cartoon: local synaptic eligibility is multiplied by a neuron-level error transported through a route-specific gain. **E**, Evidence path from factorization and signal–noise theory through trained route tests, anatomical capacity and the functional alignment boundary.

Per-neuron bandwidth dominates restricted feedback

A global scalar is low bandwidth but discards which neuron produced an error. The matched-width/scalar-fallback rule used in our initial experiments retains the somatic vector only at a dendritic stage of identical width and otherwise averages it to one scalar per example. In contrast, neuron-indexed feedback repeats the appropriate somatic coordinate across all compartments belonging to that neuron’s tree. It communicates one value per neuron, not one value per compartment.

This scalar-versus-indexed distinction is a feedback-bandwidth effect, not a new scaling principle: global perturbation estimates are known to become noisy as network size grows [10, 26]. Our question is what is added by neuron-to-tree ownership and within-tree address after neuronal identity has been supplied.

We compared these fields in regular $[3, 3]$ trees, which contain 13 modeled compartments per neuron. Across 15 paired MNIST seeds, neuron-indexed feedback raised three-factor accuracy from $90.54 \pm 0.54\%$ to $97.16 \pm 0.11\%$ in shunting networks and from $91.76 \pm 0.78\%$ to $97.14 \pm 0.11\%$ in width- and connectivity-matched additive networks. All 15 pairs improved in each architecture. At the same neuron-indexed-trained checkpoints, the cosine between local and exact dendritic gradients rose from 0.012 ± 0.059 to 0.708 ± 0.027 in shunting trees and from -0.001 ± 0.047

145 to 0.625 ± 0.035 in additive trees (Fig. 2b,c). The increase was positive at all 15 matched checkpoints in each
146 architecture.

147 Supplying exact transported compartment errors tests the derived local eligibility without feedback approximation.
148 With five paired shunting seeds, three-factor learning with a local decoder reached $97.18 \pm 0.12\%$, compared with
149 $97.03 \pm 0.08\%$ for a same-architecture backpropagation reference. Adding the empirical five-factor preconditioner or
150 a backpropagated decoder provided no material gain under exact transport (Fig. 2d). Thus the poor scalar-feedback
151 result is not a failure of the derived local eligibility. The large scalar-to-neuron-indexed gap results from loss of
152 neuronal identity; the smaller remaining gap to exact transport reflects omitted path information.

153 Shunting changes the conductance component of path gain. In a separate five-seed diagnostic cohort with
154 five inhibitory synapses per branch, the unweighted path-gain coefficient of variation across compartments within
155 examples was 0.229 for shunting and 1.161 for additive trees, with the same ordering in all five paired seeds. This
156 statistic describes the electrical operating point; it does not establish better gradient direction. Indeed, a 15-seed
157 comparison crossing architecture with analytical or occupancy-calibrated branch-output-nonlinearity initialization
158 showed no general shunting advantage: both architectures reached a mean of 90.92% under analytical initialization
159 (± 0.42 and ± 0.51 s.d. for shunting and additive; paired difference -0.00 , $P = 0.982$), whereas occupancy calibration
160 modestly favored the additive model. We therefore interpret conductance as a route-gain mechanism, not as a universal
161 optimizer (Supplementary Fig. S4).

162 Broader archived controls define the boundary of this conclusion (Supplementary Figs. S1–S4). Under restricted
163 feedback, shunting trees were more stable than matched additive trees in depth and teaching-noise stress tests, but
164 the advantage was task dependent and did not isolate a backward shunting mechanism. Increasing feedback structure
165 consistently improved learning, including on a flattened CIFAR-10 mechanism control, whereas unmatched local
166 errors and forward nonlinearities did not replace the required teaching field. These experiments support a conditional
167 robustness claim rather than a general accuracy advantage.

168 We prospectively trained 640 models crossing two tasks, two neuron models, one to four dendritic stages, three
169 feedback fields and matched backpropagation, with ten paired seeds per condition; all local conditions used the
170 same factorization-derived eligibility and training protocol. An input-validity audit, defined without reference to
171 outcomes, excluded all 160 synthetic-noise shunting runs, whose signed transfer output is not a valid drive for a
172 positive-conductance shunting denominator (Methods), leaving 480 retained runs comprising both models on MNIST
173 and the additive model on the noise task.

174 Neuron-indexed feedback increased held-out accuracy over a scalar broadcast by 4.92–6.99 percentage points
175 on MNIST and 9.76–14.16 in the additive noise model, with all 120 retained paired seeds improving (Fig. 2d and
176 Supplementary Fig. S19a,b). Exact transport placed mean three-factor accuracy within 0.14 points of matched
177 backpropagation in all 12 retained task–model–depth conditions (Supplementary Fig. S19c).

178 A fresh Fashion-MNIST ladder tested whether the dominant identity effect was specific to MNIST (Fig. 2g,h).
179 Without retuning the architecture or learning protocol, neuron-indexed feedback improved over the scalar fallback by
180 4.46 points in shunting trees (95% paired-seed interval 4.09–4.88) and 3.78 points in additive trees (3.43–4.14), with
181 all ten paired seeds positive in both architectures (exact sign-flip $P = 0.002$). Exact path transport added no reliable
182 benefit beyond neuron identity: -0.11 points in shunting trees (-0.21 to -0.01 ; 2/10 positive; $P = 0.096$) and 0.13
183 points in additive trees (-0.08 to 0.35; 8/10 positive; $P = 0.314$). This second dataset strengthens the bandwidth
184 conclusion while leaving the distinct within-tree-address question to the credit-reversal experiments below.

185 A bandwidth-matched control fixed coordinate count and values but deranged the coordinate-to-tree map. Across
186 the six retained conditions, correct assignment improved accuracy by 0.15–0.70 points and was favored in 58/60
187 pairs (Fig. 2e); five of six contrasts survived correction. Bandwidth therefore explains most of the scalar-to-neuron-
188 indexed gain, while correct coordinate ownership adds a smaller effect. This comparison does not assign different
189 task-dependent coordinates to different subtrees of one neuron and is not a trained test of within-tree addressing.

190 We therefore froze a separate two-stream credit-reversal task before its outcomes were inspected. Both sibling
191 streams were active on every example; context selected the stream controlling the somatic output, while the other
192 carried an anticorrelated distractor. Context and label were exactly balanced, so context alone could not predict the
193 target. The same scalar somatic error therefore required different branch updates in paired contexts. The one-neuron
194 field captured 0.509 of the eligibility-weighted exact-gradient energy at initialization, below the prespecified 0.90




Figure 2: Neuron identity is the dominant feedback bottleneck, with a smaller cost for incorrect route ownership. **A**, MNIST accuracy under one scalar feedback signal or a neuron-indexed coordinate. **B**, Exact- gradient cosine at matched checkpoints. **C**, Accuracy under exact path transport, local rules and backpropagation (BP); exact transport is an information oracle. **D**, Neuron-indexed-minus-scalar accuracy across physical depth in the prospective MNIST cohort. **E**, Gain from correct rather than deranged neuron-to-tree ownership at matched bandwidth. **F**, Held-out accuracy in the two-stream credit-reversal task, where the same neuronal error requires different updates in sibling subtrees. **G,H**, Fashion-MNIST replication of the scalar, neuron-indexed and exact- path ladder and their paired contrasts. Faint lines or small points are paired seeds; large symbols and bars are means and 95% seed-bootstrap intervals. Exact transport, gated-point and correct-ancestry values coincide in the shallow representation-matched condition in **F**.

rejection threshold, whereas two correct subtree coordinates captured 1.000.

Across ten paired seeds, correct two-subtree routing reached 80.63% held-out accuracy, compared with 76.24% for one shared neuronal coordinate (difference 4.39 percentage points, 95% paired bootstrap interval 3.33–5.48; 10/10 seeds) and 19.63% after a fixed within-neuron derangement (a permutation assigning every route to the wrong sibling) (difference 61.00 points, 59.93–62.09; 10/10 seeds; Fig. 2f). A dense random rank-two map reached 58.47% and was below correct routing in all ten seeds. Gradient capture tracked norm-matched one-step progress (Supplementary

Fig. S19h). After a single-context adaptation block, the shared rule lost 55.66 points on the previous context, whereas the correct route left the inactive sibling unchanged (difference in forgetting, 55.66 points, 53.84–57.67; Supplementary Fig. S19i).

Correct subtree routing, exact transport and analytic backpropagation were identical on this shallow linear branch model. Crucially, a point-neuron implementation with the same two context-gated parameter groups also matched them exactly. The experiment therefore demonstrates the value of the address, not a unique dendritic computational advantage: the same solution can be paid for by explicit feature gating in a point model. It executes the decisive $K = 2$ support contrast on one forward architecture. We next completed the broader bandwidth and representation factorial.

The same audit excluded the 400-run inhibitory-dose family from inference because its shunting-minus-additive endpoint depends on the invalid signed-shunting cells. An independent clean-source 320-run exact-transport/backpropagation audit instead tested implementation agreement across ten paired seeds, four depths, two tasks and two cores; the synthetic shunting cells used an explicit ReLU transfer before positive conductances were formed (Supplementary Fig. S19e). All 320 runs passed the configuration, log, metric and finite stage-complete checkpoint audit. After averaging depths within seed, exact-minus-backpropagation accuracy was 0.005 percentage points on additive MNIST (95% paired bootstrap interval -0.028 – 0.041), -0.013 on shunting MNIST (-0.047 – 0.022), -0.112 on additive noise (-0.465 – 0.247) and -0.098 on shunting noise (-0.653 – 0.390). All four intervals included zero. Across all 160 raw method pairs, the mean difference was -0.055 points and the largest absolute pair difference was 5.90 points. These are empirical agreement checks, not formal equivalence tests or pointwise identity claims.

A retained 240-run spatial-map control improved accuracy by 0.36–2.57 points, including under backpropagation and in the additive randomly projected noise model (Supplementary Fig. S7). At 336 contacts, spatial branches covered 336 unique inputs versus 276.2 for random branches, identifying a forward-connectivity rather than task-specific credit-routing effect.

Of 320 fixed-budget runs, the 160 additive cells provide the valid synthetic comparison at 16 terminal branches and 960–968 contacts per soma. Depth four underperformed depth one in all 40 paired seed contrasts: -7.92 points under scalar feedback, -1.06 under neuron-indexed feedback, and -1.37 under both exact transport and backpropagation (Supplementary Fig. S19f; Supplementary Fig. S8). Deeper trees retain more trainable parameters, but scalar feedback clearly amplified depth stress.

At 120 retained backpropagation checkpoints, mean gradient cosine was 0.027 for scalar and 0.586 for neuron-indexed feedback. Norm-matched updates descended at 75/120 and 120/120 checkpoints, respectively, and cosine predicted retained one-step progress ($\rho = 0.893$, 95% checkpoint-bootstrap interval 0.860–0.923; Supplementary Fig. S9). This connects field geometry to loss change without replacing the trained-learning factorial.

Structured subtree addresses help at an intermediate feedback budget

We froze an eight-context hierarchical credit-reversal task and crossed five parameter-matched representations, feedback budgets $K \in \{1, 2, 4, 8\}$, six route families and three references. Twenty paired independent seeds produced all 2,700 planned fits. Context and label were exactly balanced; every route field had unit norm per example; initial weights, data and optimization were paired. The dendritic, explicitly gated point, flat-compartment and grouped-point implementations had 64 forward parameters, 64 active input contacts and one nonlinearity. They received identical routed coefficient fields. A degree- and depth-matched rewired tree preserved resource counts while breaking task-topology alignment.

Correct ancestry routing displayed a non-monotonic advantage. Mean held-out accuracy increased from 0.187 at $K = 1$ to 0.449, 0.801 and 0.810 at $K = 2, 4, 8$, respectively (Fig. 3b). It strongly exceeded fixed within-neuron route derangement at $K = 2$ (difference 0.264, 95% paired bootstrap interval 0.243–0.285) and $K = 4$ (0.611, 0.601–0.620). Against the best matched non-anatomical control selected separately within each seed, however, correct ancestry lost at $K = 1$ and $K = 2$, won at $K = 4$ by 0.0127 (0.0059–0.0197; 15/20 seeds; two-sided Wilcoxon $P = 0.0048$, Benjamini–Hochberg-adjusted $P = 0.0064$ across the four budgets), and tied at full rank $K = 8$ (Fig. 3c). An unrestricted learned rank- K oracle and random rank- K routes therefore exploited low-dimensional task structure

more efficiently at the smallest budgets; nested anatomy became advantageous only when four coordinates were available.

The correct matched tree exceeded the rewired tree by 0.231 (0.210–0.254) at $K = 2$ and 0.050 (0.042–0.060) at $K = 4$, with all 20 seeds positive in both contrasts (Fig. 3d). At $K = 1$ the route is shared, and at $K = 8$ every leaf is addressed, so rewiring had no effect. Exact transport and analytic backpropagation were identical. The four implementation-equivalent dendritic, gated-point, flat and grouped-point controls also matched exactly in every deterministic and paired-stochastic condition (maximum accuracy difference 0; Fig. 3e). Initial exact-gradient capture alone did not determine final accuracy across route families: fields with similar capture but different route signs and task alignment learned differently (Fig. 3f). We made that boundary explicit by deterministically reconstructing all 2,700 fits and evaluating the shared-coefficient partition residual at initialization and at the trained checkpoint (Supplementary Fig. S23). For the matched tree, trained ancestry routes captured 0.162, 0.315, 0.547 and 1.000 of the weighted exact coefficient field at $K = 1, 2, 4, 8$, whereas ownership derangement captured zero for $K > 1$. Across the broad route ladder, capture and held-out accuracy were positively associated within every seed (mean within-seed Spearman $\rho = 0.824$, 95% seed-bootstrap interval 0.810–0.838; 20/20 seeds). The association remained positive after excluding the exact $K = 8$ boundary ($\rho = 0.639$, 0.610–0.665; 20/20). However, correct ancestry, depth bins and random sparse groups had similar capture at a given K despite different trained outcomes. The residual therefore explains the large bandwidth and ownership failures, but not the smaller support-specific advantage at $K = 4$; that distinction requires the signal–noise and trajectory analyses below. Thus the result establishes a structured-address resource advantage at an intermediate bandwidth, not an intrinsic superiority of dendritic material: point models reproduce it when explicitly supplied with the same parameter groups and routing fields.

A credit-operator phase theory predicts when restricted routes help

We next froze a controlled experiment directly from Eqs. 8–9. Sixteen coordinates were represented in an orthonormal balanced-tree Haar basis. Fifty independent confirmatory seeds generated 10,800 seed–condition rows crossing route rank, task–tree alignment, task hierarchy, routed model depth, signal and noise retention, and branch-reliability heterogeneity. The experiment holds the forward coordinate budget fixed and isolates stochastic credit geometry; it is not a positive-conductance forward-neuron simulation.

First, the benefit of ancestry routes depended continuously on spectral task–tree alignment $\rho \in [0, 1]$, where $\rho = 0$ is the seed-specific unaligned covariance and $\rho = 1$ is tree-Haar-aligned. At full alignment and $K = 4$, ancestry captured 0.552 more task-gradient energy than a random rank-matched route (95% paired bootstrap interval 0.535–0.569; 50/50 seeds; two-sided Wilcoxon $P = 1.78 \times 10^{-15}$; Fig. 4b). The advantage was near zero under random alignment, increased with alignment, never exceeded the Ky–Fan dense principal-subspace upper bound—the maximum covariance capture available to any rank- K orthogonal projector—and vanished at full rank. At full alignment and $K = 4$, ancestry attained the rank-matched upper bound to numerical precision. This attainment is the generic case for hierarchical credit: the K -route ancestry span equals the span of the K coarsest tree-Haar modes, so anatomy attains the rank- K Ky–Fan bound whenever the task covariance is diagonalized by the tree’s wavelet basis with a coarse-dominant spectrum—and, in general, exactly when the leading rank- K eigenspace lies inside the laminar span—while its shortfall for a general covariance is bounded by the leading-eigenspace energy outside that span (Supplementary Section S4). A deterministic endpoint reanalysis used the affine covariance mixture to obtain an exact ancestry-versus-random alignment threshold within each seed. Nineteen of 50 random-route draws already favored ancestry at zero alignment; all 50 favored it at full alignment. Counting the former as requiring zero alignment, the median $K = 4$ threshold was 0.0173 and the mean was 0.0523 (95% seed-bootstrap interval 0.0336–0.0730). This is a threshold relative to sampled random routes, not a universal biological constant.

Second, routed depth was optimal when it matched task hierarchy. This match was imposed by the controlled generator—signal occupied levels through H , whereas finer levels carried additional noise—so the test is a quantitative validation of predicted crossover locations, not a discovery that $D_r = H$ in arbitrary tasks. For task depths $H = 1, 2, 3, 4$, the lowest mean final population loss occurred at model depth $D_r = H$ (Fig. 4c,d). Across the 200 seed–task pairs, matched depth reduced loss relative to the best mismatched depth by 0.0957 (95% interval 0.0782–0.1134 after averaging the four task-depth contrasts within each independent seed; 48/50 seeds; two-sided




Figure 3: The benefit of a structured subtree address depends on feedback bandwidth and task-topology alignment. **A**, Eight leaves are addressed by one shared coordinate or $K = 2, 4, 8$ nested route fields. **B**, Held-out accuracy across route families in the dendritic-tree implementation. Curves and error bars show means and 95% seed-bootstrap intervals over 20 paired seeds; overlapping points are fanned slightly, with lines meeting at the shared values. The learned rank- K condition is an unrestricted task-derived upper bound. **C**, Correct ancestry minus the best of four matched non-anatomical controls within each seed; positive values favor ancestry. **D**, Correct routing on the task-matched tree minus the same rule on a degree- and depth-matched rewired tree. Small points are paired seed differences. **E**, Correct-route learning is numerically identical for dendritic, explicitly gated point, flat-compartment and grouped-point implementations when each receives the same routed fields. **F**, Initial exact-gradient capture versus final accuracy for each route family and budget; marker size increases with K , and overlapping points are fanned for visibility. Capture alone does not determine learning when route signs and task structure differ. Every condition matched forward parameter count, active contacts, nonlinearity count, initialization, data and optimizer.

Wilcoxon $P = 3.00 \times 10^{-13}$). Insufficient depth left target components outside the route span; excess depth admitted fine-scale stochastic gradient energy. At $D_r = 4$ the projector is full rank and the update is identical to the common stochastic-backpropagation reference. Depth therefore helps here by matching the address resolution of the task, not by being large in isolation.

Third, independently varying signal and noise retention recovered the predicted boundary. Backpropagation followed by the same route projection had lower mean one-step population loss than unprojected stochastic backpropagation in exactly those of the 25 condition means for which rejected noise exceeded discarded signal (Fig. 4e).

304 The routed local estimator carried an additional specified coefficient-noise cost and never exceeded this projected-
 305 backpropagation control. Branch-specific reliability gains then improved progress only when they were aligned with
 306 branch signal-to-noise ratio. At maximum heterogeneity, reliability-aligned gain exceeded the best single global gain
 307 by 1.241 in one-step loss reduction (95% interval 1.116–1.373; 50/50 seeds; $P = 1.78 \times 10^{-15}$), whereas shuffling or
 308 reversing the gains removed or reversed the benefit (Fig. 4f). An explicit point gate receiving the same gains matched
 309 exactly (maximum difference 0).

310 We next tested whether a physical conductance could realize this oracle attenuation in a trained branch model.
 311 Eight branches received strictly nonnegative input rates and excitatory conductances. A nonnegative shunt to reversal
 312 zero reduced branch input resistance, while an oracle compensating current $\kappa_b V_b$ held the forward voltage identical;
 313 this isolates the gain multiplying local eligibility. Fifty fresh confirmatory seeds crossed five reliability-heterogeneity
 314 levels and nine matched controls (Supplementary Fig. S13). At maximum heterogeneity, reliability-aligned shunting
 315 increased one-step test-loss reduction relative to the best global shunt by 0.001817 (95% paired bootstrap interval
 316 0.001618–0.002020; 49/50 seeds; two-sided Wilcoxon $P = 3.55 \times 10^{-15}$). Its mean one-step advantage over no
 317 shunt was 0.001339 (0.000594–0.002096; 33/50 pairs; $P = 0.0030$).

318 After 40 updates, aligned shunting reduced test loss by 0.02040 relative to the best global shunt (-0.02222
 319 to -0.01871 for aligned minus global; 50/50 seeds favored aligned, $P = 1.78 \times 10^{-15}$), and also outperformed
 320 shuffled and anti-aligned shunts. Its final loss differed from the unshunted noisy rule by 0.000425 (-0.002456 –
 321 0.003071 ; $P = 0.154$), providing no reliable final-loss advantage in either direction. Exact clean backpropagation
 322 remained the lowest-loss reference. A state-matched additive control was identical to no shunt, and the explicit
 323 point gate receiving the same sample-dependent factors was computed through an independent update path and
 324 agreed with the conductance trajectory to numerical precision. Thus positive conductance approximates the predicted
 325 reliability ordering under a state clamp, but the immediate benefit does not establish autonomous shunt learning,
 326 dendrite-exclusive computation or uniform training superiority.

327 Finally, static nonzero gain rescaling changed numerical conditioning without changing the route span: nested
 328 indicators, their tree-Haar basis and a gain-scaled version had identical capture to 1.33×10^{-15} , while the condition
 329 number ranged from 1 for the Haar basis to 583 for the illustrated scaled parameterization (Fig. 4g,h). We then
 330 prospectively trained noisy route coefficients in a separate 50-seed, exactly same-span experiment (Supplementary
 331 Fig. S14). The three rank-eight coordinates had condition numbers 1, 15 and 388.5 and zero target address residual.
 332 The preregistered unconditional Haar advantage was falsified at effective sample size four: raw nested and scaled
 333 nested coordinates reduced final loss by 0.277 and 0.191 relative to Haar because their slow modes filtered update
 334 noise. At effective sample size 256 the ordering reversed, with Haar reducing loss by 0.0256 and 0.1709, respectively,
 335 in 50/50 paired seeds. An exact finite-time bias–variance calculation predicted median crossovers at effective sample
 336 sizes 38.5 and 8.25, and an oracle Gram-preconditioned control made every same-span field trajectory identical to
 337 1.16×10^{-15} . Thus conditioning is neither a pure defect nor new capacity: it trades finite-time bias against estimator
 338 variance.

339 The phase utility was also diagnostic in the already completed 2,700-fit factorial. Reconstructing 64 initialization
 340 minibatches for each of 20 seeds gave 540 dendritic-route conditions; $U(M)$ predicted norm-matched one-step
 341 progress with Spearman $\rho = 0.937$ (seed-block 95% interval 0.9336–0.9412) and final held-out accuracy with
 342 $\rho = 0.916$ (0.905–0.929). Spectral capture alone predicted one-step progress less well ($\rho = 0.877$), with a positive
 343 seed-block difference of 0.060 (0.050–0.070; Fig. 4i). This secondary reanalysis explains immediate geometry at the
 344 fixed initialization; it is not an independent trained experiment.

345 The bound also locates optima, with an instructive boundary (Supplementary Fig. S16). Because every tested
 346 field is a projector on one nested ladder, utility rises with feedback bandwidth only while each added route scale
 347 contributes proportionally more signal than admitted noise, so an interior optimum is predicted whenever the signal
 348 spectrum decays faster than the noise (Supplementary Section S4). In the depth experiment, the per-seed argmax of
 349 the initialization utility coincided with the trained loss minimum in 135 of 200 seed–task pairs, and the group-mean
 350 argmax matched exactly at $H = 1, 2, 3$; at the full-rank boundary $H = 4$ the one-step bound was conservative,
 351 favouring $D_r = 2$. In the factorial, the same utility reproduced the exact ties at $K = 8$, where all spans coincide, but
 352 did not predict the trained $K = 4$ ancestry advantage over the best control (6 of 20 seeds): that advantage accrues
 353 along the training trajectory and is not visible in one-step initialization geometry, consistent with capture alone not

354 determining learning (Fig. 3f).




Figure 4: A signal-noise phase theory predicts when dendritic credit structure is useful. **A**, A fixed credit operator M maps an unbiased stochastic gradient into a routed local update; its guaranteed utility balances retained signal against finite-step gain and admitted noise. **B**, Mean ancestry-minus-random spectral capture across alignment and rank ($n = 50$ seeds); the advantage grows with alignment and vanishes at full rank. **C,D**, Mean final population loss for crossed task hierarchy H and routed depth D_r . Black outlines mark $D_r = H$; shading in **D** is the 95% seed-bootstrap interval, and coincident markers in **D** are slightly offset. $D_r = 4$ is full rank and equals stochastic backpropagation. **E**, Expected one-step loss of backpropagation plus the route projection minus unprojected stochastic backpropagation. Negative values favor projection; the dashed line is the analytical $n = r$ crossover. **F**, Simulated one-step loss decrease for SNR-aligned, best global, shuffled and anti-aligned gains (means and 95% intervals over 50 seeds). The explicit point gate is identical to the aligned condition, with coincident markers offset at zero heterogeneity. **G,H**, Static nonzero column gains preserve spectral capture but can alter the positive-spectrum Gram condition number. Markers in **G** are horizontally offset to reveal exact overlap. **I**, Prespecified signal-noise-curvature utility versus observed norm-matched one-step progress for 540 dendritic conditions from the 2,700-fit factorial; the interval resamples 20 seed blocks. Panels **B–H** are a fixed-state stochastic quadratic validation, not a positive-conductance forward model or evidence of uniform superiority to exact full-batch backpropagation.

Task-aligned physical depth improves learning at fixed budgets

The fixed-state phase experiment tests address resolution, not whether serial dendritic dynamics contribute to forward computation. We therefore constructed a separate positive-rate task—all inputs are nonnegative rates driving positive conductances—in DendriNet, the conductance-based population-network implementation used throughout (Methods). A distal class signal was modulated by nested fine, coarse and global gains, while separate nonnegative inhibitory streams measured those factors. This is a mechanism-matched existence test built around divisive normalization, not a generic task benchmark [27, 28]. D1 [8], D2 [2, 3] and D3 [2, 1, 2] contained the same eight nonsomatic branch units per soma, 14,336 active contacts, 21,760 candidate slots, 66,178 trainable parameters and 2,944 persistent state scalars. Only the number and placement of serial divisive stages changed (Fig. 5a,b).

With backpropagation and aligned gain sensors, mean test accuracy increased from 0.6096 at D1 to 0.6752 at D2 and 0.9182 at D3. The paired D3-minus-D1 effect was 30.86 percentage points (95% seed-bootstrap interval 30.45–31.30; 10/10 positive pairs; exact sign-flip $P = 0.002$; Fig. 5c,d). The operating point was selected transparently by exploratory signal, coupling and gain ladders and then fixed before these ten inference seeds (Methods and Supplementary Fig. S15); the 31-point contrast must therefore be read as a replicated effect at a calibrated operating point. Physical depth was flat when the sensors were independent (0.02 points) or trial shuffled (–0.01 points), and decreased under exact-resource reversal of fine and global tree placement (–0.43 points, –0.67 to –0.15). The corresponding aligned-minus-control interactions were 30.84, 30.87 and 31.28 points, respectively, all positive in 10/10 paired seeds. Thus the effect requires correspondence between task factors, branch-local sensors and their physical ordering; stage count alone is insufficient.

The nonlinearity also mattered at this operating point. Aligned raw-additive BP decreased from 0.5733 at D1 to 0.5383 at D2 and 0.5190 at D3, whereas the matched shunting series increased (Fig. 5f). This control supports a divisive hierarchical-composition mechanism in this task, not a general impossibility for additive or expanded point networks.

A same-task architecture ladder sharpened that boundary (Fig. 5g,h and Supplementary Fig. S18a). An all-active grouped-star model retained every branch module, mask, parameter, contact and persistent state but removed serial child-to-parent composition. It matched D1 (–0.005 points), remained near 0.610 at every depth, and therefore fell behind the serial tree by 6.54 points at D2 and 30.81 points at D3 (30.41–31.26; 10/10 positive). Reversing fine and global placement reversed the D3 contrast to –0.43 points, giving a 31.24-point alignment interaction (30.65–31.77; 10/10 positive). Thus the large effect requires serial composition in the model, rather than merely possessing the same collection of branch units. It is not an unconstrained capacity advantage: active- and total-parameter-matched point MLPs reached 0.9861 and 0.9902 and exceeded the serial D3 tree by 6.79 and 7.20 points, respectively, in every paired seed (Supplementary Fig. S18b). The tree supplies a task-aligned structured resource, not a computation that a flexible point network cannot learn.

Local learning retained the aligned depth benefit. A single decoder-derived soma coordinate rose from 0.6083 at D1 to 0.7778 at D3 (16.95 points, 16.37–17.44; 10/10 positive), while exact path transport reached 0.8873 at D3 (27.90 points, 27.24–28.53; 10/10 positive; Fig. 5e). Exact-resource reversal removed both effects: the aligned-minus-reversed depth interactions were 17.72 points (17.27–18.18; 10/10 positive) for the shared coordinate and 28.68 points (28.07–29.27; 10/10 positive) for path transport. A separate replay on the 30 aligned BP checkpoints reproduced autograd with exact path transport (minimum gradient cosine 0.9999999999999994), whereas mean shared-soma cosine fell from 0.946 at D1 to 0.767 at D3. The local-rule performance and checkpoint diagnostic together show that physical hierarchy can remain accessible to local eligibility when a sufficiently specific error is transported; they do not establish that exact path transport is itself biologically available.

An autograd teaching-coordinate ladder localized the remaining separation (Fig. 5i,j). Replacing each compartment’s loss derivative by its owning soma derivative while retaining autograd eligibility and the BP optimizer cost only 0.64 points at D3 (0.15–1.17; 8/10 positive; $P = 0.033$). Because this arm and LocalCA initially used different optimizer groups, a transparently post-outcome diagnostic repeated soma broadcast with the frozen LocalCA optimizer. It then matched shared-soma LocalCA (–0.06 points, –0.23 to 0.09; $P = 0.492$), whereas changing only the broadcast arm’s optimizer accounted for 13.47 points (12.89–14.05). Within LocalCA, path transport recovered 10.96 points over the shared coordinate (10.34–11.58); full BP retained a 3.09-point advantage over path LocalCA

(2.48–3.79). In this calibrated task, the dominant BP–local gaps are therefore optimizer regime and within-tree coordinate specificity, not the tested local eligibility expression itself.

An alignment dose response tested whether the depth effect changed before the fully matched endpoint (Supplementary Fig. S18c–e). Intermediate sensor alignments $\alpha \in \{0.25, 0.50, 0.75\}$, with $\alpha = 0$ denoting independent sensors and $\alpha = 1$ matched sensors, were registered after the $\alpha = 0$ and 1 outcomes were known, so this is a prospective interpolation rather than an independent confirmation. The D3-minus-D1 effect increased from 0.02 points at $\alpha = 0$ to 0.21, 0.82 and 3.13 points at the intermediate levels, before rising to 30.86 points at full alignment. Every seed had a positive within-seed linear slope (mean 25.84 points per unit α , 95% interval 25.43–26.26; exact sign-flip $P = 0.002$), but the curve was strongly nonlinear: the $\alpha = 0.75$ minus 0.25 increase was only 2.92 points (2.65–3.17), whereas the final 1.00 minus 0.75 increment was 27.73 points (27.35–28.11). The result therefore localizes a steep alignment boundary near the fully matched task, rather than supporting a generic linear dose law.

We then tightened the point control and changed task hierarchy (Supplementary Fig. S18f–k). The earlier grouped-star control retained serial-shaped child aggregators as independent arms. A literal grouped-point emulation instead replaced each aggregator by a parameter-matched direct projection to the soma: every nonsomatic level kept its local branch module, mask and shunting computation, but received the external streams independently and never inherited another branch state. The replacement copied the corresponding initialized raw conductances, making D1 an exact initialization control. On the existing $H = 3$ task, grouped-point accuracy remained near 0.610 across D1–D3 (aligned D3-minus-D1, -0.01 points, -0.06 to 0.04), while serial-minus-grouped-point reached 30.87 points at aligned D3 (30.48–31.31; 10/10 positive; exact sign-flip $P = 0.002$). The same contrast was -0.44 points after placement reversal (-0.69 to -0.16), giving a 31.31-point alignment interaction (30.73–31.84; 10/10 positive). The earlier grouped star and the literal grouped point differed by only 0.06 points at aligned D3 (0.01–0.13; 7/10 positive; $P = 0.084$), so the conclusion does not depend detectably on which nonserial parameter-matched readout is used.

An independent $H = 2$ hierarchy then changed the sensor partition to two 32-feature tiers, the fixed branch inventory to 4/4 and the serial morphologies to D1 [8] and D2 [4, 1], while retaining the calibrated task family and all other optimizer and stopping settings. Ten fresh paired seeds 10300–10309 were used. Serial BP increased from 0.6588 to 0.9672, a 30.84-point D2-minus-D1 effect (30.46–31.23; 10/10 positive; $P = 0.002$). Reversing the two sensor tiers changed that effect to -1.56 points, yielding a 32.40-point depth-by-placement interaction (32.10–32.70; 10/10 positive). The literal grouped point stayed flat in both placements (aligned, 0.02 points, -0.02 to 0.06 ; reversed, -0.02 points, -0.06 to 0.02), and serial BP exceeded it at aligned D2 by 30.82 points (30.41–31.23), with a 32.36-point architecture-by-placement interaction. Shared-soma LocalCA retained a 21.20-point aligned depth effect (19.34–23.00) and exact-path LocalCA retained 31.38 points (30.95–31.83); placement reversal made both effects negative, producing 22.93- and 33.36-point interactions, respectively. Every positive interaction was ordered in 10/10 paired seeds. This is an independent hierarchy-depth replication within the same synthetic mechanism-matched task family, not a second dataset or evidence that natural tasks generally favor serial trees.

Together, the new same-task controls separate four explanations that the original depth comparison could not distinguish. The parameter-matched point networks reject a dendrite-exclusive expressivity claim. The serial-minus-star and serial-minus-literal-point interactions identify ordered divisive composition, rather than branch count, parameter count or a particular nonserial readout, as the source of the aligned forward advantage; the $H = 2$ cohort shows that this conclusion is not unique to the original three-stage inventory. The optimizer-matched soma-broadcast comparison finds no detectable penalty from the tested local eligibility expression, whereas path-specific transport recovers 10.96 percentage points over a shared neuronal coordinate; the principal local-credit limitation in this task is therefore the teaching coordinate and its optimization regime. Finally, the alignment interpolation shows that the forward depth advantage appears only near a strongly matched task–sensor hierarchy. The experiments consequently support a conditional role for dendritic structure: it provides a useful structured computation and address system when the task matches the tree, but neither depth alone nor dendritic material guarantees better learning.

Reconstructed arbors provide sparse, mostly coarse route capacity

The evidence so far comes from synthetic trees with hand-specified topology. We next asked whether reconstructed cortical arbors provide such addresses. We constructed ancestry dictionaries from eight layer 2–5 intratelencephalic

neurons in MICrONS materialization 1822, mapping 28,669 classified inputs to collapsed cable segments from one mouse. Electrical parameters were normalized passive proxies, and reciprocal axial coupling was solved globally. The analysis complements prior connectomic proposals for anatomical credit assignment and the larger MICrONS morphology map [29, 30].

For neuron u , we define a K -channel feedback dictionary $\Phi_{u,K} \in \mathbb{R}^{N_u \times K}$, with rows for its synaptic or compartment locations and columns for available feedback fields. An *address* is one such column: its support and gain specify which locations can share one task-dependent coefficient. A point neuron has only the shared column **1**; depth bins supply coarse columns; an ancestry dictionary supplies nested subtree columns; and an unrestricted dense rank- K subspace is a descriptive upper bound. This definition separates neuronal ownership from within-tree address and makes clear that dendrites supply structured, not arbitrary, credit bandwidth.

This distinction can be made quantitative. For a fixed family of K -column dictionaries $\Phi_{u,K}$ and a distribution of exact task-dependent coefficient fields $\mathbf{c}_u$, let $P_{\Phi_{u,K}}$ be the W -orthogonal projector onto $\text{col}(\Phi_{u,K})$. For tolerated unexplained field-energy fraction $0 < \varepsilon < 1$, define the ε -address dimension

$$d_{\varepsilon,u} = \min \left\{ K : \frac{\mathbb{E} \|P_{\Phi_{u,K}} \mathbf{c}_u\|_W^2}{\mathbb{E} \|\mathbf{c}_u\|_W^2} \geq 1 - \varepsilon \right\}. \quad (10)$$

A point neuron with one neuronal teaching coefficient has rank-one task-dependent address space, irrespective of its number of synapses. A K -route dendritic dictionary has address rank at most $\text{rank}(\Phi_{u,K}) \leq K$ and can express K structured coefficient fields without assigning those routes to K different neurons. This is a resource statement, not a claim that dendrites create arbitrary feedback bandwidth: a point-neuron system can emulate the same field if it is supplied with K separately routed teaching coefficients or an equivalent gated feature expansion (Supplementary Information).

For an excitatory site i and candidate inhibitory route k , the direct topological route matrix was

$$\Gamma_{ik} = A_{ik} \beta_k, \quad (11)$$

where A_{ik} marks route ancestry and $\beta_k = g_k^I / g_k^{\text{tot}}$ is normalized inhibitory gain. This is a route-sensitivity proxy, not the exact passive-cable derivative used below. Its nested domains addressed a median 2.67% of mapped excitatory input and had mean participation rank 20.81, versus rank one for a scalar broadcast (Fig. 6a–c).

We measured the fraction of a modeled field retained by K feedback channels. For weighted field $\bar{\mathbf{g}} = W^{1/2} \mathbf{g}$ and dictionary $\bar{\Phi}_{T,K} = W^{1/2} \Phi_{T,K}$, let $\bar{P}_\Phi$ be the Euclidean projector onto $\text{col}(\bar{\Phi}_{T,K})$. Field capture is

$$\mathcal{C}_K(\mathbf{g}; \Phi, W) = \frac{\|\bar{P}_\Phi \bar{\mathbf{g}}\|_2^2}{\|\bar{\mathbf{g}}\|_2^2} = 1 - \frac{\|\bar{\mathbf{g}} - \bar{P}_\Phi \bar{\mathbf{g}}\|_2^2}{\|\bar{\mathbf{g}}\|_2^2}. \quad (12)$$

Here W contains summed mapped excitatory area, so Eq. 12 is first an anatomical field-energy measure. Only if $\bar{\mathbf{g}}$ is an exact gradient in explicitly transformed coordinates $\bar{\mathbf{w}} = W^{-1/2} \mathbf{w}$ and the transformed loss has L_W -Lipschitz gradients does $\mathcal{C}_K$ also equal the ratio of projected and full one-step descent guarantees at step $1/L_W$ (Supplementary Information).

We led with an independent generator: exact state-matched responses to focal shunts in the reciprocal cable model. At eight channels, ancestry routes captured 0.458, versus 0.363 for random paths and 0.249 after row shuffling, with both advantages positive in all eight cells. Differences from depth bins (0.014) and degree- and depth-matched surrogate trees (0.035) were inconclusive: the cell-bootstrap interval for the latter excluded zero, whereas the primary paired Wilcoxon test did not ($P = 0.078$). We therefore infer sparse support beyond random assignment but not a demonstrated advantage of fine ancestry beyond coarse depth and degree (Fig. 6e,f).

As a model-matched ceiling, we then generated fields directly from K .

At eight channels, morphology-selected routes captured 0.487 of the energy in the model-derived route-perturbation fields using 6.92% of the connections in a dense feedback matrix. This was 85.0% of the field energy captured by an unconstrained dense eight-channel oracle. Random paths captured 0.285, depth bins 0.222, and ancestry-shuffled routes 0.198. The morphology advantage was positive in every cell for each control; hierarchical cell-and-stream intervals excluded zero (Fig. 6c,d). Normalizing by wiring makes the resource statement

493 explicit: per unit of feedback wiring, morphology-selected routes captured 14.2 times more field energy than the
494 dense oracle (95% cell-bootstrap interval 10.7–18.2), whereas the ancestry-shuffled control, matched at the identical
495 6.9% wiring density, reached 5.3-fold; the anatomy-specific advantage at matched density is therefore 2.7-fold, with
496 the remainder contributed by sparsity itself (Fig. 6g,h). The ordering also survived an analysis restricted to directly
497 typed presynaptic E/I inputs: morphology captured 0.696, compared with 0.347 for random paths, 0.447 for depth
498 bins, and 0.308 after ancestry shuffling (Supplementary Fig. S20h).

499 In 47 additional cells, morphology captured 0.756 with 8.08% of dense wiring, versus 0.302–0.441 for structural
500 controls; every within-cell contrast was positive (Supplementary Fig. S10). The cohort is nucleus-disjoint but from
501 the same mouse and a convenience-derived historical reconstruction.

502 All comparisons matched channel count. They establish anatomical compression of modeled fields, not a task
503 gradient or route use by the animal.




Figure 5: Matched nonlinear hierarchy makes exact-resource physical dendritic depth useful. **A**, D1–D3 have $D_p = 1, 2, 3$ stages. Bracket entries are soma-to-distal branching factors: [8], [2, 3] and [2, 1, 2] give 8, 2 + 6 and 2 + 2 + 4 nonsomatic branch units, with identical contact, candidate-slot, parameter and state budgets. **B**, A distal positive-rate signal is multiplied by fine, coarse and global gains; independent sensors, trial-shuffled sensors and reversed physical placement are matched controls. **C**, Test accuracy under shunting backpropagation. Lines in **C, E, F** are ten-seed means and error bars 95% seed-bootstrap intervals; coincident flat-control markers are offset. **D**, Prespecified paired D3-minus-D1 effect and aligned-minus-control interactions; bars are 95% paired-seed bootstrap intervals, all excluding zero. **E**, Aligned and reversed-placement local credit assignment (LocalCA) with one shared-soma coordinate or exact path transport; the coincident reversed curves are drawn with offset markers. **F**, Aligned shunting and raw-additive BP. **G**, Architecture controls that separate a flexible point multilayer perceptron, a nonserial grouped star and the serial tree. **H**, Serial-tree and grouped-star accuracy across physical depth. **I**, Credit-coordinate ladder from full BP through soma-broadcast autograd to shared- and path-transport LocalCA. **J**, Paired D3 decomposition of the BP–local gap into coordinate, optimizer, eligibility and path-specificity contrasts. All inferential runs use paired seeds 10200–10209; lines are seed means and error bars are 95% seed-bootstrap intervals.




Figure 6: Reconstructed arbors provide a sparse, mainly coarse routing basis for modeled perturbations. **A**, Example reconstructed tree and mapped excitatory and inhibitory contacts. **B**, Two nested ancestry-defined addresses on the same arbor. **C,D**, Capture of the model-derived perturbation field across channel budgets and at eight channels for a dense oracle, ancestry routes and matched structural controls. **E,F**, The same comparison for an independently generated reciprocal- cable response field. **G**, Eight-channel capture versus wiring density; dashed curves mark equal capture per unit wiring. **H**, Capture per unit wiring across budgets. At eight channels the model-matched efficiency is 14.2-fold relative to the dense oracle, whereas the density-matched anatomy-specific factor is 2.7-fold. Lines summarize eight reconstructed cells and bands or intervals are 95% cell-bootstrap intervals. These are modeled fields on measured anatomy, not observed task gradients; wiring counts feedback connections, not cable length or metabolic cost.

The route analysis above used inferred inhibitory distributions. We next asked whether synapse-resolved inhibitory connectivity is organized in a manner compatible with branch-level control. This question builds on physiological and modeling work on tunable dendritic inhibition and inhibitory plasticity [31, 32]. We joined the disjoint reconstruction cohort to the published MICrONS inhibitory census [33]. Twenty of 47 target cells could be joined by stable nucleus identity across releases; five had different root IDs after re-segmentation. The source tables contained 116,401 input locations and 2,789 known contacts from 116 inhibitory cells. Of these, 114,775 inputs and 2,637 contacts from 114 inhibitory cells mapped within $5\mu\text{m}$. This is a structural analysis of one mouse; it does not identify a learning signal, and the packaged census does not provide a per-target proofreading-status field.

For each inhibitory cell–target connection, we retained one representative contact per independent axonal clump. The 402 connections with at least two clumps provided a within-axon test. Repeated contacts from the same inhibitory axon were $12.0\mu\text{m}$ closer along the target tree than compartment- and path-distance-matched input locations (95% bootstrap interval -21.5 to $-4.2\mu\text{m}$; 15 of 20 target means were negative; one-sided Wilcoxon $p = 0.0032$). They also shared 0.0357 more normalized path (0.0091 – 0.0766 ; 15 of 20; $p = 0.0028$) and addressed descendant input domains with 0.0322 greater weighted cosine overlap (0.0060 – 0.0743 ; 14 of 20; $p = 0.0053$; Supplementary Fig. S12d–f). These target-level comparisons do not by themselves separate dendritic organization from repeated sampling of the same axons or local tissue geometry. Averaging first within each of 92 presynaptic axons retained the tree-distance effect ($p = 0.0008$), whereas the shared-path and domain-overlap effects were not reliable ($p = 0.104$ and $p = 0.393$). With candidate locations jointly matched for path distance and three-dimensional proximity, none of the three contrasts was reliable ($p = 0.108, 0.826$, and 0.54 , respectively). Thus repeated contacts are spatially co-targeted, but these data do not establish geometry-independent organization over descendant credit domains.

Inhibitory classes exposed different spatial scales. After retaining one contact per axonal clump, distal dendrite-targeting contacts addressed a smaller fraction of target-cell input than perisomatic-targeting contacts in all 20 cells. The mean within-target difference of the class medians was -0.0162 (-0.0231 to -0.0105 ; one-sided $p = 9.5 \times 10^{-7}$; Supplementary Fig. S12g). This class effect should not be confused with a global placement advantage. Across all inhibitory contacts, descendant-domain size did not differ reliably from matched generic input locations (mean difference 0.0016 , $p = 0.546$; Supplementary Fig. S12h).

Finally, we measured the internal compressibility of the actual inhibitory route dictionary. At eight sites, its highest-leverage columns captured 0.743 of that same weighted dictionary, close to a dense singular-vector oracle (0.774), compared with 0.379 for random actual sites, 0.397 for depth bins, and 0.369 after depth-preserving ancestry shuffling. Each selected-route contrast was positive in all 20 targets (one-sided $p = 9.5 \times 10^{-7}$ for each; Supplementary Fig. S12i,j). Selection by leverage is an anatomical compression oracle: it measures redundancy within a dictionary using columns of that dictionary, not independent credit routing or a claim that the animal chose those sites. Together, the census establishes class-dependent spatial scales and available nested domains, while the axon-clustered and 3D-matched controls prevent a stronger functional-routing interpretation.

Shunting regulates route gain only in permissive electrotonic regimes

Topology-based reconstruction does not show that conductance can selectively control the routes. At the textbook passive calibration ($R_m = 15,000\Omega\text{cm}^2$), the primary contrast was effectively null (-0.0019 , two of eight cells positive); localization emerged only in permissive normalized or higher-conductance regimes. We therefore specified a focal intervention before examining its outcome. This adjoint-credit question extends prior analyses of spatial inhibitory domains in dendrites [4]. The passive conductance matrix of each reconstructed tree was denoted by G , with voltage $\mathbf{V} = G^{-1}\mathbf{b}$, where $\mathbf{b}$ is the current-source vector. Let $\mathbf{e}_s$ and $\mathbf{e}_k$ be standard basis vectors at the soma and focal compartment k , respectively. For a somatic loss, the adjoint error is $\mathbf{q} = G^{-1}\mathbf{e}_s\partial\mathcal{L}/\partial V_s$. A shunt of conductance κ_k at compartment k changes the transport operator according to

$$\mathbf{q}' = \mathbf{q} - \frac{\kappa_k q_k}{1 + \kappa_k (G^{-1})_{kk}} G^{-1} \mathbf{e}_k. \quad (13)$$

Writing $\mathbf{c}_k = G^{-1}\mathbf{e}_k$ and $r_k = (G^{-1})_{kk}$ gives $\Delta\mathbf{q} = -a_k\mathbf{c}_k$, where $|a_k| = \kappa_k|q_k|/(1 + \kappa_k r_k)$. Thus dose controls the response amplitude, which increases monotonically and saturates at $|q_k|/r_k$, whereas the cable column $\mathbf{c}_k$ controls its

spatial spread. For descendant set $D(k)$ and a matched comparison set $C(k)$, the transport-only selectivity criterion is

$$S_k = \frac{\text{median}_{i \in D(k)} |c_{k,i}|}{\text{median}_{j \in C(k)} |c_{k,j}|} > 1. \quad (14)$$

Equation 14 separates electrotonic selectivity from perturbation strength. It is necessary for descendant enrichment of absolute adjoint change under this median contrast, but is not sufficient for the synaptic-gradient statistic, which also depends on baseline adjoints and local driving forces. A first-order-current-matched additive perturbation changes the voltage-driving term but not G^{-1} . This yields a direct test: after a state-matching compensatory current injection restores somatic voltage and the scalar output error, conductance shunting should alter exact synaptic gradients more strongly among descendants than among depth-matched unrelated sites. Reciprocal coupling still changes credit outside the descendant set.

We tested 101 depth-stratified focal sites across the eight cells; the gradient-change profile across topological relations is shown in Fig. 7b. At the unit dose, the descendant-minus-depth-matched-unrelated localization index—descendant minus depth-matched unrelated median log gradient attenuation—was 0.104 for shunting and 0.035 for the matched additive perturbation. The cell-level paired difference was 0.069 log units (95% bootstrap interval 0.053–0.086), positive in all eight cells. The contrast rose monotonically with dose, from 0.018 at 0.25 local-conductance units to 0.128 at two units. True focal relations also exceeded relations sampled from other focal sites, stratified by topological-depth quartile when another site was available, by 0.080 (0.059–0.098), again in all eight cells (Fig. 7b–d).

Because a synaptic gradient is the product of adjoint error and local driving force, we next froze either factor algebraically. Keeping the full post-shunt voltage and driving force but replacing the post-shunt adjoint with its baseline value gave localization 0.026. Restoring the post-shunt adjoint raised it to 0.104, an exact transport-specific difference of 0.079 that was positive in all eight cells (Fig. 7e). This comparison holds the post-shunt driving-force field fixed and therefore isolates the contribution of changed error transport without relying on first-order current matching. The complementary adjoint-only state gave localization 0.130. A two-factor Shapley sensitivity analysis assigned 0.087 of the planned shunt-minus-additive contrast to the adjoint change and -0.017 to the driving-force change, reconciling exactly to the 0.069 net contrast at every site (Supplementary Section S8). These factor substitutions are algebraic controls, not separate biological interventions.

The state-matching residual at the soma was numerically zero, conductance matrices remained positive definite, and adjoint gradients agreed with finite differences to a maximum relative error of 1.03×10^{-6} . The principal contrast also survived changes in normalized conductance scale and inhibitory reversal potential (Supplementary Fig. S21f). Restricting both synaptic weighting and focal sites to directly typed inputs gave a shunt-minus-additive contrast of 0.101 (0.075–0.129), positive in all eight cells (Supplementary Fig. S21g). In the disjoint sensitivity cohort, the same contrast was 0.138 (0.116–0.162) across 235 sites in 45 eligible cells and was positive in every cell (Supplementary Fig. S10f). These results identify a conductance-specific gradient-localization mechanism at fixed somatic state in the normalized proxy model.

We next calibrated axial and leak conductances from reconstructed radii and lengths using explicit specific resistances. With $R_a = 150 \Omega \text{cm}$ and $R_m = 15,000 \Omega \text{cm}^2$, the median axial-to-leak ratio was 68.0. The shunt-minus-additive contrast was then -0.0019 in the original cohort (two of eight cells positive) and 0.0024 in the v661 cohort (45 of 45 positive). Lower effective membrane resistance, representing a progressively higher-conductance state, restored a larger effect: at $R_m = 1,000 \Omega \text{cm}^2$ the original eight-cell contrast was 0.0164, and at $R_m = 300 \Omega \text{cm}^2$ it was 0.0674 (eight of eight positive in both conditions; Fig. 7f). Low input resistance and conductance-driven gain modulation are established features of in-vivo-like high-conductance states [34, 35]. These values use an explicit relative calibration of synapse size because MICrONS synaptic area is not a conductance measurement. The focal signature is therefore an electrotonic-regime-dependent model prediction, not a universal property of the reconstructed arbors or evidence that inhibitory synapses carried endogenous teaching signals in the animal.

We then froze a passive sensitivity matrix that removes the relative-dose confound. Across the same 101 sites and eight cells, it crossed three membrane resistances, three levels of distributed leak-like background conductance, and relative-local, fixed-absolute and input-conductance-normalized doses (16,362 site–condition–intervention rows). At a fixed 0.5-nS shunt, the shunt-minus-additive localization contrast was 0.0104 at $R_m = 300 \Omega \text{cm}^2$, 0.0073 at

597 $R_m = 1,000 \Omega \text{cm}^2$, and -0.0039 at $R_m = 15,000 \Omega \text{cm}^2$; the first two were positive in all eight cells, whereas only
 598 two of eight were positive in the last regime. At the high-resistance setting, adding one and four times the baseline
 599 leak as distributed background conductance raised the input-conductance-normalized contrast from -0.0004 to
 600 0.1210 and 0.2587 , respectively. Every tested descendant gradient was attenuated in magnitude; none was enhanced
 601 or changed sign. Thus passive shunting supplies graded attenuation, not signed reversal, and its conductance-specific
 602 spatial advantage requires both sufficient dose and a selective electrotonic column (Supplementary Fig. S11).

603 Finally, we asked whether the same localization survives weak active conductances rather than a purely passive
 604 operator. The ensemble was evaluated at the permissive electrotonic operating point identified above ($R_m =$
 605 $1,000 \Omega \text{cm}^2$ with one baseline-leak unit of distributed background conductance; Methods); it extends the localization
 606 result within that regime rather than restoring it at the high-resistance calibration. For each of the eight cells,
 607 we solved 64 channel-density draws containing voltage- dependent Na, K, Ca, HCN and NMDA currents, then
 608 linearized the exact nonlinear steady-state residual to obtain its local Jacobian. All 512 equilibria were accepted;
 609 every Jacobian remained positive definite and the maximum equilibrium residual was 9.6×10^{-11} . Draws were
 610 nested sensitivities, not independent cells. At the unit input-conductance-normalized dose, focal shunting produced
 611 descendant localization 0.529 (95% cell-bootstrap interval $0.513\text{--}0.541$), attenuated every modeled descendant
 612 gradient, produced no sign flips and reduced descendant gradient energy to 0.330 of baseline (Fig. 7g–j). Its
 613 localization exceeded the state-matched additive control by 0.382 ($0.356\text{--}0.409$), positive in all eight cells (two-sided
 614 Wilcoxon $P = 0.0078$). The effect rose monotonically across the three doses. This active steady-state ensemble
 615 strengthens the local-Jacobian prediction but is not a kinetic-channel or spiking simulation.

616 **Measured responses reveal a boundary for morphology-specific alignment**

617 Structural capacity used ancestry-kernel perturbations and the focal result used modeled adjoints. We next tested
 618 whether the routes align with measured functional inputs and response-derived task gradients. A MICrONS/DANDI
 619 join identified seven reconstructed targets and 136 repeated visual conditions per scan. The original one-scan-per-
 620 target cohort contained 69 connected, imaged excitatory partners. Functional similarity did not increase consistently
 621 with shared ancestry after controlling for Euclidean separation and depth: the mean target-level partial rank effect was
 622 -0.069 (95% target-bootstrap interval -0.255 to 0.108), positive in three of seven cells (Fig. 8a,b). This estimand
 623 asks whether within-tree position adds organization beyond geometry and complements the broader like-to-like
 624 wiring principle reported in the same MICrONS volume [36]. Using the deterministic rule “all automatic-conservative
 625 scans with at least five connected imaged partners” expanded the analysis to 13 scan observations. After averaging
 626 scans within target, the partial shared-path effect was -0.048 (-0.224 to 0.118), positive in three of seven cells. The
 627 scan-complete analysis therefore removed the selection ambiguity without changing the null conclusion.

628 We trained a branch surrogate for each postsynaptic response and projected its exact error into four-channel
 629 dictionaries. Each partner mapped through its dominant contact to one nonlinear state on the first major branch
 630 below the soma. Sites on that branch therefore shared one task-dependent error. With only 3–8 occupied major
 631 branches per target, this cannot resolve within-branch credit or strongly test fine topology. Stimulus identities were
 632 held out. Mean held-out field capture was 0.999 for a training-only dense oracle, 0.475 for morphology paths, 0.564
 633 for nonempty random paths, 0.621 for depth bins, 0.446 for ancestry-shuffled paths, and 0.222 for a scalar broadcast
 634 (Fig. 8c). Ancestry routes minus ancestry-shuffled capture was 0.028 (95% interval -0.146 to 0.206), with three of
 635 seven target cells positive (Fig. 8c).

636 Learning outcomes gave the same boundary. Mean normalized held-out error was 0.778 for exact backpropa-
 637 gation, 0.781 for the dense oracle, 0.791 for depth bins, 0.801 for scalar broadcast, 0.834 for random paths, 0.837
 638 for morphology paths, and 0.842 after ancestry shuffling. Thus morphology was close to its ancestry-shuffled
 639 control but did not outperform the simpler depth-bin or scalar dictionaries on this task. The paired morphology
 640 improvement was 0.004 normalized-error units (95% interval -0.021 to 0.027), positive in four of seven targets
 641 (Fig. 8d). A pooled capture–performance association vanished after centering within target ($\rho = -0.150$; label-
 642 permutation $P = 0.419$; Supplementary Fig. S22g). Higher within-target capture predicted lower error at one
 643 channel ($\rho = -0.621$, $P = 0.0006$) and eight channels ($\rho = -0.480$, $P = 0.0104$), but not two or four (Source Data).
 644 Channel-wise morphology-minus-shuffle contrasts (Supplementary Fig. S22h), reliability thresholds and a matched




Figure 7: Focal shunting produces a descendant-enriched modeled-credit signature in defined electrotonic regimes. **A**, Focal intervention with state-matching somatic current, shown as a shunt-off/shunt-on tree pair; the faded descendant subtree marks where the shunt attenuates transported credit. **B**, Absolute gradient change across tree relations. **C**, Within-cell localization for the matched additive perturbation, route reassignment and focal shunting. **D**, Dose response. **E**, Exact factor-freeze comparison isolating changed adjoint transport. **F**, Shunt-minus-additive localization across the cable-calibrated axial-to-leak ratio; the standard high-resistance calibration is effectively null. **G**, Active-channel steady-state design. **H,I**, Active- model dose response and paired shunt-minus-additive localization. **J**, Descendant outcomes at unit dose: attenuation, enhancement, sign reversal and residual gradient energy. Points are reconstructed-cell means after averaging sites and channel draws; lines, bands and intervals are 95% cell-bootstrap summaries. The reconstructed cell is the inferential unit. Active validation uses the permissive electrotonic regime and does not restore the effect at the standard passive calibration.

subset did not establish morphology-specific benefit.

We then replaced the coarse branch surrogate with an exact passive solve on each complete compressed reconstructed tree. Every imaged partner supplied a trainable non-negative conductance at its mapped segment; the exact local rule used the trial-specific soma-to-segment adjoint. Restricted rules used only scalar output error, local eligibility and a fixed four-channel feedback vector calibrated from topology before learning. All 13 eligible scans and ten held-out-stimulus splits per scan completed without exclusion, giving 520 method fits. Split replicates were averaged within scan, then scans within each of seven target cells.

Exact compartment error achieved mean normalized held-out MSE 0.788. The topology-matched, random-route and site-shuffled rules reached 0.832, 0.833 and 0.839, respectively (Fig. 8g). Topology improved MSE over site shuffling by 0.0070 units (95% target-bootstrap interval -0.0207 to 0.0321) and over random routes by 0.0013 (-0.0090 to 0.0113), neither a reliable advantage. At the common exact checkpoint, topology captured 0.232 of the eligibility-weighted update field, below site-shuffled routes (0.439) and random routes (0.358; Fig. 8h). Restricted dictionaries had identical effective rank within every scan-split run, and exact conductance gradients passed finite differences with maximum relative error 3.1×10^{-7} .

Thus route availability did not imply route use in either the coarse or segment-resolved task. Four dense dimensions nearly matched exact learning in the coarse model, but anatomical routes did not align reliably with measured- response credit on the full tree. The dense-oracle capture curve indicates why the coarse-model null is expected under the credit-operator account: a single channel already captured 0.965 of this response family's credit energy (effective rank 1.07), so a four-channel budget exceeds the task's credit rank and every equal-rank dictionary spans it—the regime in which the theory predicts the observed tie (Fig. 8k). Oracle calibration makes each restricted method a capacity ceiling. The null does not exclude other response windows, higher input coverage or electrophysiologically fitted conductances.

Measured responses showed no reliable morphology-specific alignment; the conditional claim, however, predicts that the same dictionary should become useful exactly when task credit rotates into its span. To test the conditional statement directly, we generated alignment-controlled quadratic objectives on each reconstructed tree. Exact-gradient fields were rotated from the span of the morphology dictionary into its orthogonal complement while channel count, dictionary, gradient energy, and loss curvature were held fixed. The same fields were evaluated with morphology, random-path, depth-bin and ancestry-shuffled dictionaries before norm-matched or iterative projected descent (Fig. 8e,f and Supplementary Fig. S5).

The morphology dictionary captured exactly the imposed aligned energy: from 0 at complete orthogonality to 1 at complete alignment. At 40% alignment it captured 0.400 and supported 0.615 of the loss reduction from a norm-matched full-gradient step. Its capture advantage over the three controls ranged from 0.203 to 0.228 and was positive in all eight cells. At complete orthogonality, morphology lost to every control in all eight cells; at complete alignment, it exceeded the controls by 0.710–0.799 and won in all eight. Thus the same fixed anatomy changes from an unhelpful to a useful feedback basis as only its alignment with task credit changes. Across the three control dictionaries and alignment levels, initial capture tracked 20-step progress within each cell (median Spearman $\rho = 0.990$, range 0.969–0.996). The one-step relation is shown in Fig. 8f and the 20-step relation in Supplementary Fig. S5.

This constructive sufficiency test imposes alignment, fits oracle projection coefficients and uses a quadratic objective. It links field capture to optimization under controlled conditions but does not show that MICrONS cells used these gradients; independently measured responses showed no reliable morphology-specific benefit.

The theory assumes that a neuron-specific teaching coordinate is available. We tested that assumption using published source data from a six-animal brain-computer-interface experiment in which the causal signs of two selected neurons, P+ and P-, were fixed by the experimental mapping [37]. We defined one contrast before inspecting its sign: standardized dendritic population residual during error reduction minus the residual during error increase. If the dendritic signal follows causal credit rather than one common scalar error, the contrast should be positive around P+ and negative around P-.

The P+ contrast was positive in five of six animals (mean 0.0753; 95% bootstrap interval 0.0295–0.1214; exact one-sided random-sign permutation $P = 0.0313$). The P- contrast was negative in all six (mean -0.1498 ; -0.2054 to -0.0866 ; $P = 0.0156$). Their signed separation was positive in all six animals (mean 0.2251; 0.1416–0.3108; $P = 0.0156$; Fig. 8i). Decomposing the six animal contrast vectors showed that 83.7% of pooled squared projection energy lay in the signed P+/P- mode, compared with 16.3% in a common scalar mode (animal-bootstrap 95% interval for signed-mode energy, 60.6–95.7%; leave-one-animal-out range, 76.5–90.2%) (Fig. 8j); descriptive neuron-level residual distributions from the published source table are shown in Supplementary Fig. S17d. With six animals, these exact permutation tests sit at or one step above their resolution floor. The analysis independently supports the first level of the theory—signed neuron identity—but does not determine how the signal is transported within a dendritic tree or uniquely implicate shunting.




Figure 8: Measured responses define the biological boundary, while imposed alignment validates the conditional prediction. **A,B**, Seven-cell measured-response cohort and the geometry-controlled structure–function effect. **C,D**, Held-out task-field capture and learning for exact, dense, ancestry and matched control dictionaries; projection coefficients are oracle ceilings. **E,F**, Rotating task credit into the fixed ancestry span increases its capture, and capture predicts one-step progress. **G,H**, Complete-tree learning and ancestry-minus-control contrasts across all eligible scans; neither shows a morphology-specific benefit. **I**, Signed neuron-coordinate contrast in six animals: P+ and P- are defined by the causal experimental mapping. **J**, Most contrast energy lies in the signed P+/P- mode rather than a common scalar mode. **K**, Alignment-by-bandwidth synthesis across trained and controlled experiments. Filled markers denote an ancestry win or tie; open markers denote a loss. Shaded regions summarize the credit-operator regimes and are not fitted to the displayed outcomes. Target cells or animals are the inferential units; nested splits and simulations are averaged within them.

Together, the experiments occupy a common alignment-by-bandwidth plane (Fig. 8k). Low bandwidth loses neuronal identity, whereas high bandwidth can make equal-rank spans coincide. Between these limits, ancestry routes help only when task credit falls sufficiently inside their span. The trained $K = 4$ interior optimum, the spectral and controlled- alignment sweeps, the credit-reversal task and both measured-response nulls all fall on the corresponding side of this boundary. The shaded regions are theory-derived regimes, not a fitted classifier of the experiments.

707 Discussion

708 Each exact synaptic gradient separates into locally available eligibility—presynaptic activity, driving force and
709 input resistance—and a transported compartment error. Biological learning in a dendrite is therefore a well-posed
710 approximation problem: recover a defined credit field from restricted signals. Point-neuron theories ask how an
711 error reaches a neuron or gates eligibility; a dendritic theory must additionally ask which subtree receives that
712 coordinate and how strongly it is transported. Learning can fail by losing neuron identity, addressing the wrong tree
713 or descendants, or omitting route gain.

714 The prospective experiments separate these failures quantitatively. One coordinate per neuron closes most
715 of the scalar-feedback gap, while deranging the neuron-to-tree map at matched bandwidth produces a smaller
716 but reproducible cost. A 2,700-fit credit-reversal factorial then shows that the value of within-neuron structure
717 is bandwidth dependent. Correct ancestry is inferior to unrestricted low-rank routes at one and two coordinates,
718 acquires a small but reliable advantage over the best matched non-anatomical control at four coordinates, and ties at
719 full rank. It also outperforms a degree- and depth-matched rewired tree at the intermediate budgets. Dendritic, flat,
720 grouped-point and explicitly gated point implementations are identical when given the same routed fields, locating
721 the benefit in structured address resources rather than dendritic material per se. Exact transported feedback makes
722 the derived local eligibility sufficient to approach matched backpropagation, and the detached clean audit found no
723 material exact-transport-backpropagation difference, bounding the claim to empirical agreement rather than formal
724 equivalence. Gradient alignment predicts immediate descent.

725 The credit-operator bound now unifies these observations. Route quality is not capture alone: it is retained
726 task signal divided by finite-step gain and admitted stochastic noise. This utility predicted the existing factorial
727 better than spectral capture alone. In the independent phase experiment, ancestry became useful as task covariance
728 rotated into its span; routed depth was best at the task's hierarchy depth; and a projection improved a noisy gradient
729 only below the rejected-noise/discarded-signal boundary. These are not routes by which a local rule can beat exact
730 full-batch backpropagation on the same objective. They explain how a restricted stochastic rule can improve relative
731 to another noisy or bandwidth-limited rule, with backpropagation plus the same projection remaining the relevant
732 upper control. The interior-optimum reanalysis sharpens this division of labour: the bound's optimum located the
733 trained depth optimum wherever one existed, whereas the trained $K = 4$ anatomical advantage lies beyond one-step
734 geometry (Supplementary Fig. S16). The results therefore order into four comparisons. A restricted rule never beats
735 exact full-batch backpropagation at matched step norm; it beats stochastic backpropagation exactly when rejected
736 noise exceeds discarded signal; it beats a bandwidth-matched non-anatomical rule only when task covariance aligns
737 with the anatomical span; and any within-tree address beats scalar or neuron-only feedback whenever exact credit
738 varies across the tree in an addressable way. Reliability-weighted route gains improve each comparison when branch
739 signal-to-noise is heterogeneous. Figure 8k places every bandwidth-varied experiment of this study in the resulting
740 alignment-by-bandwidth plane, including the measured-response null, which falls where the theory predicts a null.
741 In operator terms, topology sets the span of addressable fields, conductance sets the metric that weights their use,
742 and the task determines whether either matters.

743 Depth can affect learning in at least three distinct ways. First, nested branch points reuse route coordinates
744 across exponentially many leaves and provide progressively finer addresses; the matched-depth phase result isolates
745 this role. Second, serial dendritic stages can compose nonlinear forward operations, a role deliberately absent
746 from the quadratic phase experiment. The exact-resource positive-rate experiment now isolates one such case:
747 aligned divisive composition improved accuracy by roughly 31 percentage points in both $H = 2$ and $H = 3$ variants,
748 whereas literal grouped-point models stayed flat and reversed placement made serial depth harmful. Independent and
749 shuffled-sensor controls removed the H3 gain, and the raw-additive series worsened with depth. This is evidence for
750 matched hierarchical shunting within the calibrated task family, not a general law that deeper trees are better. Third,
751 a deeper physical route transports a shared coordinate through more differentiated branch states. Consistent with that
752 cost, shared-soma gradient cosine fell from 0.946 at D1 to 0.767 at D3, whereas exact path transport reproduced
753 autograd at every checkpoint. The earlier additive fixed-contact control likewise showed no intrinsic depth benefit
754 when terminal branches and contacts were fixed. Thus the dominant requirement remains a signed coordinate for
755 the correct neuron, followed conditionally by enough task-aligned address and computational depth—not depth

in isolation. The new controls make the point–dendrite boundary explicit: the serial tree beat a resource-identical nonserial star only under the aligned hierarchy; it also beat the stricter literal grouped-point model at both H2 and H3, yet ordinary parameter-matched point MLPs beat the tree by about seven points. Dendritic depth is therefore a structured inductive and wiring resource in this task, not an absolute expressive separation. The credit ladder draws the complementary BP–local boundary: soma-broadcast autograd nearly matched BP under a common optimizer, optimizer choice accounted for the largest raw separation from shared LocalCA, and path specificity recovered most of the remaining local-rule deficit. Stated as an accounting at matched resources, the tree’s contribution separates into a forward-composition term isolated by the backpropagation depth contrast, an address-resolution term isolated by the bandwidth factorial, and a transport cost isolated by the falling shared-coordinate cosine, with the local-rule retention measuring their interaction.

Real dendrites then constrain how that coordinate could be distributed. Reconstructed arbors compress model-matched and independently generated cable fields into sparse, ancestry-defined supports using a small fraction of dense feedback wiring—in wiring-normalized terms, an order-of-magnitude advantage in captured energy per feedback connection over a dense oracle, of which a 2.7-fold factor is anatomy-specific at matched density (Fig. 6g,h). A dendritic tree, on this account, multiplexes K addressable fields within one neuron’s arbor. An equivalent point implementation requires the same K separately addressable teaching coordinates; its wiring cost depends on how those coordinates are routed. Yet fine morphology does not reliably outperform controls matched for depth and parent degree. The synapse-resolved inhibitory census similarly supports coarse co-targeting and class-specific spatial scales, while shared-path effects do not survive axon-level aggregation and joint geometric matching. These one-mouse measurements establish available addresses, not their use for endogenous learning.

Shunting has a precise and bounded role within this hierarchy. Unlike models in which inhibitory circuits compute an error [38–40], focal conductance here changes the sensitivity of routes crossing a compartment. Exact factor freezing assigns the positive localization effect in the normalized proxy to altered adjoint transport, partly opposed by the change in synaptic driving force. Physical calibration shows that the effect can become very small in passive regimes with high axial-to-leak coupling and re-emerge in higher-conductance states. This spatial boundary is determined by the focal Green’s-function column and cable length constant: strong axial coupling spreads the perturbation and reduces the descendant-versus-matched off-path contrast. The passive inverse-operator corollary is used only as a global magnitude bound, $\|\Delta G^{-1}\|_2 \leq \kappa \|G^{-1}\|_2^2 / [1 + \kappa / \lambda_{\max}(G)]$; it does not predict spatial selectivity. Shunting is therefore a state-dependent route-gain mechanism, not a universal optimizer and not evidence that inhibition itself carries a teaching signal. The active-channel ensemble extends this conclusion beyond a passive operator: exact linearization at 512 accepted Na/K/Ca/HCN/NMDA steady states preserved descendant-enriched attenuation without gradient sign reversal. It remains a local steady-state result, not a claim about channel kinetics or spiking. The reliability phase identifies a possible computational use for that attenuation: fixed-step branch gains $\min\{1, S_b / [c(S_b + N_b)]\}$ suppress unreliable credit and outperform one global gain when reliability differs across branches. The state-matched positive-conductance experiment joins that prediction to a physical input-resistance mechanism in a trained branch model: aligned shunts improved the immediate credit step and final loss over a global shunt, but not final loss over unshunted noisy learning. The independently computed point-gate control agrees to numerical precision and proves that the operation is not mathematically exclusive to dendrites. This differs from the sibling gain–load study, which asks when branch-local inhibition improves forward signal robustness; here the forward state is matched and conductance weights backward credit [41]. Moreover, branch reliabilities and shunt placement were supplied by an oracle, and the compensating current removed forward state changes. The experiment therefore validates a conductance-based, state-dependent approximation to a reliability factor, not autonomous inference of that factor or a completed biological causal chain.

Anatomical capacity is useful only when it matches the task. Neither all 13 eligible visual-response scans nor 520 learning fits on the complete compressed segment trees reveal a reliable morphology-specific advantage for the derived objective. By contrast, rotating exact task credit into or out of the same anatomical span reverses its learning benefit. The six-animal reanalysis supports the availability of signed neuron-specific coordinates, but does not identify their transport within a tree. Route overlap and off-route leakage then predict one-step interference (Supplementary Fig. S6). Together, these results replace the claim that more dendritic structure is inherently better with a conditional statement: route capacity matters when the task requires credit fields that the arbor can express.

806 The alignment requirement has a population-level analogue: motor learning is faster for perturbations within an
807 existing neural manifold than outside it [42]. Our claim is the subcellular counterpart, not an identification of the
808 population mechanism: a pre-existing ancestry span helps only when the task-gradient field lies sufficiently within it.

809 These results relate to, and delimit, three strands of prior work. Studies of dendritic nonlinearity establish
810 that branch-local processing expands a single neuron’s forward capacity [43–46]; here forward resources were
811 deliberately matched or budgeted, and the tree’s contribution was measured instead in feedback bandwidth, so the two
812 accounts are complementary rather than competing. Dendritic-error and segregated-compartment learning models
813 deliver a teaching signal to one or a few named compartments [17–19, 47]; the present framework asks how far such
814 a signal can usefully be distributed within a full arbor, and answers with an explicit bandwidth–alignment condition
815 rather than an architectural preference. Observations of branch-specific plasticity and structured synaptic placement
816 [48–50] demonstrate that learning is spatially organized in vivo; our analysis specifies what that organization must
817 deliver to constitute credit routing rather than forward gain control—nested addresses whose span aligns with
818 task-gradient covariance—and provides the perturbation signature that distinguishes the two.

819 Bio-inspired machine-learning systems likewise use explicit context gates, active dendritic segments or localized
820 dendritic updates to allocate learning capacity [51–53]. Those algorithms establish useful engineered routing
821 schemes; the present theory instead predicts when a restricted biological route should help or fail at a fixed
822 bandwidth.

823 The same accounting suggests why credit might be routed through trees at all. With a fixed budget of routed
824 teaching coefficients, a circuit can spend them as many neurons carrying one coordinate each or as fewer neurons
825 with finer within-tree addresses; the implementations are mathematically interchangeable, but their wiring is not.
826 Per-neuron coordinates require long-range feedback axons, whereas per-branch routes reuse local inhibitory synapses
827 on existing arbors, so the wiring-efficient allocation depends on whether task credit varies mostly across neurons or
828 hierarchically within them. On this view the tree is the inexpensive way to buy feedback bandwidth—a question of
829 circuit economics rather than computational exclusivity.

830 The framework makes a branch-resolved experimental prediction. After matching somatic voltage, output and
831 global error, focal dendrite-targeting inhibition during learning should alter subsequent plasticity more strongly for
832 spines below the perturbed branch than for equidistant spines on sister subtrees. The contrast should scale with local
833 input resistance and perturbation conductance, and a current-matched additive perturbation should not reproduce it.
834 Simultaneous dendritic recording, focal inhibition and longitudinal spine measurements could distinguish credit
835 routing from forward gain control [37, 48].

836 The theory assumes steady-state, non-spiking trees with unique directed paths; the reconstructed-cell analyses
837 are sensitivity models rather than cell-specific physiological fits. The primary, sensitivity and inhibitory cohorts
838 come from one mouse, the functional cohort contains seven targets, projection and transport coefficients are oracles,
839 and the external animal analysis is retrospective. The trained within-neuron factorial uses one synthetic eight-context
840 hierarchy, analytically controlled route fields and implementation-equivalent linear parameterizations. The active
841 extension linearizes steady states and does not simulate channel kinetics. Temporal eligibility, spikes, recurrent paths
842 and long-term network adaptation remain outside the current model. In particular, branch plateau potentials and
843 behavioral-timescale synaptic plasticity provide documented dendritic teaching signals beyond this steady-state
844 theory [54, 55]. The hierarchy-depth phase uses a fixed-state quadratic family and prescribed task covariance.
845 The nonlinear physical-depth bridge now tests $H = 2$ and $H = 3$, two fixed sensor inventories and one calibrated
846 operating-point family; it does not cross the full $D_r \times H \times K \times \rho$ design, infer sensor alignment or establish prevalence
847 beyond this synthetic task. The positive-conductance reliability experiment trains excitatory branch weights but fixes
848 oracle shunts and uses a compensating current clamp; it does not demonstrate autonomous estimation or learning of
849 shunt placement. We establish no memory, latency or benchmark advantage over backpropagation. Within these
850 limits, the result is an information hierarchy for cellular learning: coordinates identify neurons, ancestry partitions
851 synapses, conductance regulates routes and task alignment determines utility.

852 Methods

853 Directed conductance-tree network

854 Each compartment combines nonnegative excitatory and inhibitory synaptic conductances, leak, and child-to-parent
855 dendritic coupling. Excitatory and inhibitory reversal potentials were fixed within an experiment. Raw trainable
856 parameters were transformed to nonnegative conductances. A compartment could apply a monotone branch output
857 nonlinearity $a_n = f_n(V_n)$ before transmitting activity to its parent; this mechanism is named “reactivation” in the
858 implementation, and f_n was the identity when it was disabled. The full forward equations, raw-parameter derivatives,
859 and temporal order of the local update are provided in Supplementary Section S1.

860 Writing $p(n)$ for the parent of compartment n , the compartment-error recursion follows from the chain rule on
861 the tree:

$$\frac{\partial \mathcal{L}}{\partial V_n} = \frac{\partial \mathcal{L}}{\partial V_{p(n)}} f'_n(V_n) R_{p(n)}^{\text{tot}} g_{n \rightarrow p(n)}^{\text{den}}. \quad (15)$$

862 Multiplying this recursion along the unique path yields Eq. 7. The local derivative of Eq. 5 with respect to g_i yields
863 Eq. 6. Analytical gradients were checked against PyTorch automatic differentiation on fixed batches and checkpoints.
864 Parameter tensors, rather than individual elements, were the units of the reconstruction diagnostic.

865 Local learning and feedback fields

866 The three-factor update was the batch average of

$$\widehat{\nabla}_{g_i} \mathcal{L} = x_i R_n^{\text{tot}} (E_i - V_n) \widehat{\delta}_n. \quad (16)$$

867 The tested $\widehat{\delta}_n$ fields included a global scalar, the matched-width/scalar-fallback implementation, neuron-indexed
868 sharing, low-rank random feedback, selected paths, and exact path transport. The neuron-indexed field repeats one
869 somatic coordinate over all descendants of that soma. It does not reveal an independent error at each compartment.
870 When $\widehat{\delta}_n = \partial \mathcal{L} / \partial V_n$, the three-factor rule equals the exact conductance gradient; the exact-transport experiment
871 therefore isolates feedback approximation from local eligibility. Four- and five-factor variants multiply the three-
872 factor update by bounded, slowly estimated stage-level scalar reliability terms. They were treated as preconditioners
873 rather than theoretically required teaching signals.

874 Regular-tree feedback comparisons used MNIST classification with matched [3,3] morphologies, optimizers,
875 decoders, schedules, and paired initializations within architecture. The feedback-definition experiment used 15
876 independently initialized seeds per architecture and condition. Exact transport used five matched shunting seeds per
877 rule and decoder-update condition. Test data remained hidden during checkpoint selection. Full hyperparameters and
878 provenance are reported in Supplementary Section S5 and the accompanying machine-readable configurations. This
879 feedback-definition cohort is a re-execution under the released implementation with the width- and connectivity-
880 matched additive architecture; the original conference cohort gave $90.75 \pm 1.01 \rightarrow 97.15 \pm 0.12$ (shunting) and
881 $91.58 \pm 0.61 \rightarrow 97.14 \pm 0.10$ (additive), consistent with the values reported here [23].

882 The Fashion-MNIST replication changed only the dataset and used the same regular [3,3] architecture, 128
883 somata, contact counts, three-factor rule, local decoder, architecture-specific initialization, optimizer and validation-
884 based checkpoint selection. Images were flattened to 784 inputs without normalization or augmentation. Within
885 each fresh seed 10200–10209, the 60,000-image training set was split into 48,000 training and 12,000 validation
886 examples; the official 10,000-image test set was used only after selection. Sixty fits crossed two architectures with
887 scalar fallback, neuron-indexed feedback and exact path transport. The primary replication criterion required the
888 paired neuron-indexed-minus-scalar interval to exclude zero and at least 8/10 positive seed signs in each architecture.
889 Intervals used 50,000 paired-seed bootstrap draws and exact two-sided sign-flip tests enumerated all 2^{10} assignments.
890 The artifact gate required complete finite results and checkpoints, the registered seeds and no fallback warning; all
891 gates passed. A frozen-source guard divided each architecture across two source manifests after a concurrent edit.
892 The exact diff only added an optional rule override that no resolved configuration enabled, and the audit retains both
893 manifests.

894 Regular-tree regime and one-step analyses

895 The task, inhibition-dose, depth, broadcast-noise, rule-family, error-source, mechanism-control, feedback-ladder,
896 and CIFAR-10 analyses use the frozen source tables and configurations from the original regular-tree study. Unless
897 stated otherwise, each condition contains five independent training seeds. The three-factor, four-factor, five-factor
898 comparison and the local-mismatch control contain three seeds and are treated as descriptive. The context-gating and
899 noise-resilience tasks are synthetic classification tests described in Supplementary Section S5. CIFAR-10 images
900 were flattened; no convolutional front-end, augmentation pipeline, or large-scale benchmark claim was used.

901 The prospective depth-by-feedback cohort was frozen before outcome inspection. It crossed two tasks, two
902 neuron models, four dendritic depths, and three local feedback fields at ten seeds, together with one matched
903 backpropagation model per task–core–depth–seed combination (480 local-learning and 160 backpropagation runs).
904 The local conditions used the factorization-derived three-factor rule. Configurations fixed the decoder, optimizer,
905 training horizon, reactivation initialization policy, 40 excitatory and 20 inhibitory contacts per distal branch, and
906 one population of 128 modeled neurons. The single-population design prevented equal neuron indices in different
907 feedforward populations from sharing a teaching coordinate. Nominal depth changed compartment and parameter
908 count; it was therefore treated as a depth stress test rather than a capacity-matched comparison.

909 Each frozen run had to contain its resolved configuration, seed record, checkpoint, complete learning curves,
910 resource record, and finite held-out metrics. The artifact audit also required every expected dendritic stage and
911 excluded critical log errors, non-finite parameters, degenerate coupling stages, and transport-shape fallbacks. A
912 subsequent outcome-independent input- validity rule excluded the 160 synthetic-noise shunting runs because a
913 signed transfer output cannot parameterize positive conductances. This left 480 publication-facing runs. Condition
914 means used the ten paired seeds. Intervals for paired contrasts used 20,000 bootstrap resamples, and exact two-sided
915 Wilcoxon tests were computed from the ten seed differences.

916 The bandwidth-matched routing cohort was frozen separately and crossed the two tasks and neuron models with
917 depths two and four, correct neuron-indexed routing or a fixed derangement, and seeds 42–51. The derangement
918 maps every somatic coordinate to another neuron and has no fixed point. It preserves the number and per-example
919 distribution of teaching coordinates, the forward network, and all training settings. The same validity rule removed
920 40 synthetic-noise shunting runs, leaving 120 runs and six task–model–depth conditions. Correct-minus-deranged
921 accuracy was paired by seed. We report paired bootstrap intervals, exact two-sided Wilcoxon tests, and Benjamini–
922 Hochberg correction across the six retained contrasts.

923 Trained within-neuron subtree-address experiment

924 The two-stream support-task configuration, task parameters, ten confirmatory seeds and artifact gates were frozen
925 before the canary outcome was inspected. Each example contained two 12-dimensional sibling-stream inputs, a
926 binary context and a balanced binary label. For the context-selected stream, the input mean was the label sign times
927 a seed-specific unit teacher vector; the other stream had the opposite sign at 0.75 times the amplitude. Independent
928 Gaussian noise with standard deviation 1.15 was then added. Every train and test split contained identical counts
929 for all four context–label combinations. Thus context had zero label correlation, both streams were active, and the
930 inactive stream was a structured distractor rather than a zero mask.

931 For branch weights $\mathbf{w}_A, \mathbf{w}_B$, the somatic logit was

$$z = \mathbb{I}[c = 0] \mathbf{w}_A^T \mathbf{x}_A + \mathbb{I}[c = 1] \mathbf{w}_B^T \mathbf{x}_B. \quad (17)$$

932 All conditions used the same paired initialization, mini-batches and forward model. The one-neuron and depth-shared
933 fields applied the logistic output error to both branches. Correct $K = 2$ feedback applied it only to the selected
934 sibling, while fixed derangement applied it to the other sibling. The random rank-two control rotated the two context
935 coordinates by a seed-fixed orthogonal matrix. Exact transport and analytic backpropagation used the selected-branch
936 derivative. The gated point emulation used the same two parameter groups and context-gated input eligibility; it
937 therefore represents the explicit point-model resource that reproduces the dendritic address.

938 Each seed used 2,048 training and 2,048 independent test examples, mini-batch SGD with batch size 64 and
939 learning rate 0.3 for 60 epochs, and Gaussian weight initialization with standard deviation 0.03. Switch interference

was measured after a further 12 epochs on 1,024 balanced-label examples from context 1 only, without resetting optimizer or weights. Gradient cosine and scaled capture used the concatenated per-example parameter gradients. For one-step progress, each candidate batch gradient was norm matched to the exact gradient before applying a fixed step of 0.25. Canary seed 4099 was used only to enforce balance, nondegenerate capture, minimum exact-route accuracy and exact-BP equivalence. The unchanged code and configuration then ran confirmatory seeds 4100–4109. The independent seed was the paired inferential unit; intervals used 20,000 paired bootstrap resamples and exact two-sided Wilcoxon tests are reported in Source Data. This experiment tests $K = 2$ on one shallow matched forward architecture and served as the support check for the larger factorial.

The full subtree-address protocol used eight contexts assigned to the leaves of a balanced binary hierarchy and eight eight-dimensional input streams. Every context–label cell was exactly balanced. The selected stream carried unit label signal; non-selected streams carried an opposite-sign distractor whose amplitude increased with tree distance (0.15 for a sibling, 0.45 within the same half and 0.75 across halves), followed by independent Gaussian noise of standard deviation 1.1. Each seed supplied 2,048 train, 2,048 test and 1,024 switch examples. Full-batch gradient descent used learning rate 0.3 for 80 epochs and 20 switch epochs. Confirmatory seeds 4200–4219 were fixed before outcome inspection; canary seed 4199 was excluded from inference.

Five representations were crossed with $K = 1, 2, 4, 8$: the task-matched dendritic tree; an explicitly context-gated point model; a flat compartment model; grouped point subunits; and a degree- and depth-matched tree with a seed-fixed leaf permutation. The first four implementations deliberately received identical routed coefficient fields and are computationally equivalent controls. Every representation contained 64 trainable forward weights, 64 active input contacts, one output nonlinearity and one scalar decoder. Route fields were normalized to unit Euclidean norm for each context. The six families were correct contiguous ancestry subtrees, a fixed within-neuron route derangement, depth-interleaved bins, random sparse groups, an unrestricted random rank- K projector and a training-derived rank- K PCA upper bound. Shared neuronal feedback, exact compartment transport and analytic backpropagation supplied references. Random controls were keyed only by seed, family and bandwidth, ensuring identical stochastic realizations across implementation-equivalent representations.

The frozen canary required exact context–label balance, zero context–label correlation, neuron-shared capture below 0.20, exact capture above 0.999999999, exact-backprop endpoint difference below 10^{-12} , exact implementation equivalence below 10^{-12} and zero fallbacks. All gates passed before the 2,700 confirmatory fits ran with unchanged script and configuration hashes. Outcomes were paired by seed. Intervals used 20,000 paired bootstrap draws; two-sided Wilcoxon tests used the 20 seed differences. The per-seed best matched non-anatomical control was prespecified across route derangement, depth-interleaved, random-sparse and random-rank conditions; the learned PCA route remained an explicitly labeled upper bound.

Credit-operator phase experiment and frozen-factorial reanalysis

The phase experiment used 16 Euclidean coordinates and the orthonormal Haar basis of a balanced binary tree, ordered from the global mode (level 0) to single sibling contrasts (level 4). Its configuration, 50 confirmatory seeds 7000–7049, two artifact-only canary seeds and numerical tolerances were frozen before confirmatory outcomes. The canary required basis orthogonality, projector idempotence, static-gain span invariance and exact agreement of the explicit point gate. Confirmatory execution required unchanged SHA-256 hashes of the configuration, runner and numerical reference. Every expected row was present: 3,750 spectral, 1,800 depth, 3,750 projection and 1,500 reliability rows.

For the spectral screen, task covariance assigned total energy (0.15, 0.35, 0.30, 0.20, 0) to tree levels 0–4. An alignment coefficient $p \in \{0, 0.25, 0.5, 0.75, 1\}$ mixed this covariance with an independently rotated covariance having the same eigenvalues. Ancestry, random rank-matched and dense principal-eigenspace projectors were compared at $K \in \{1, 2, 4, 8, 16\}$. The dense value was also evaluated as the Ky–Fan sum of the leading K eigenvalues divided by total covariance energy. Because the two endpoint covariances had equal trace, ancestry-minus-random capture was affine in alignment; endpoint differences supplied the exact per-seed crossover, with already favorable zero-alignment seeds assigned threshold zero. Capture was $\text{tr}(P\Sigma_g)/\text{tr}(\Sigma_g)$. The same-span diagnostic compared raw nested indicators, their tree-Haar basis and indicators multiplied by fixed nonzero lognormal column gains; rank,

coherence, positive-spectrum Gram condition number and capture were computed directly.

For the depth screen, a unit-norm target was sampled from modes at levels 0 through H , with $H \in \{1, 2, 3, 4\}$. Gradient-noise standard deviation in level ℓ was $0.65[1 + 1.5 \max(\ell - H, 0)]$. Starting from zero, 80 common-noise updates of size 0.12 used the projector onto levels 0 through $D_r \in \{1, 2, 3, 4\}$; a fixed leaf permutation supplied the rewired control and the identity supplied stochastic backpropagation. The projection screen independently set retained signal r and retained noise n to $\{0.25, 0.5, 0.75, 0.9, 1\}$ and compared an unprojected stochastic gradient, backpropagation plus the same projection and a routed estimator with an added 0.05 total coefficient-noise fraction. At the unit step and identity-curvature quadratic, their expected losses were calculated analytically and stochastic losses were recorded from paired draws.

For the reliability screen, eight disjoint blocks had signal and noise energies proportional to $\exp(hx_b)$ and $\exp(-hx_b)$, respectively, with equally spaced $x_b \in [-1, 1]$ and heterogeneity $h \in \{0, 0.5, 1, 1.5, 2\}$. We compared no attenuation, the best common gain, $S_b/(S_b + N_b)$, seed-shuffled and reversed gains, and a point gate receiving the identical branch factors. Intervals resampled the 50 paired seeds with 20,000 draws; paired tests were two-sided Wilcoxon signed-rank tests.

The secondary analysis reconstructed the frozen training data, initialization and route matrices for all 2,700 previously completed subtree-factorial fits. For each of 20 seeds, 64 independently sampled minibatches of size 64 estimated gradient-noise energy at initialization. The exact full-data gradient supplied signal, the largest logistic Hessian eigenvalue supplied local smoothness, and Eq. 9 supplied utility. Correlations were restricted to the 540 dendritic-tree route conditions and uncertainty resampled whole seed blocks. This is an initialization diagnostic on an existing experiment, not a second confirmatory training cohort.

Positive-rate nonlinear physical-depth experiment

To connect the fixed-state hierarchy result to the original rate-based network, we used the registered `hierarchical_gain_load` generator and the production population-network DendriNet. Inputs were ordered as nonnegative excitatory and inhibitory rate streams. For binary label y , the distal excitatory signal had the schematic form

$$E_j = (e_0 + y\Delta + \varepsilon_j) \exp \left\{ \sum_{\ell=1}^3 \sigma z_{\ell, a_\ell(j)} - \frac{1}{2} \sum_{\ell=1}^3 \sigma^2 \right\}, \quad (18)$$

where $a_\ell(j)$ indexes the fine, coarse or global latent factor containing coordinate j . Three disjoint inhibitory blocks separately measured those factors. With sensor alignment α , their log gains were $\sigma[\alpha z_{\ell, a_\ell(j)} + \sqrt{1 - \alpha^2} z'_{\ell, a_\ell(j)}]$. Thus $\alpha = 1$ supplied matched branch-local nuisance sensors, whereas $\alpha = 0$ preserved the complete factor marginals with independent sensors. A trial-shuffled control cyclically permuted the sensor field after class balancing, preserving its marginal distribution while removing example-wise pairing.

At every physical stage the production shunting equation was

$$V_n = \frac{E_n + C_n}{1 + E_n + I_n + G_n + \varepsilon}, \quad (19)$$

where C_n and G_n are child current and child conductance. All rate and conductance terms were nonnegative. Reactivation was the identity, so the divisive branch equation was the only intermediate nonlinearity. D1 [8], D2 [2, 3] and D3 [2, 1, 2] had one, two and three nonsomatic stages, respectively, but each retained eight nonsomatic branch units per soma and the same 4/2/2 fine/coarse/proximal sensor inventory. Every model contained 64 somata, 16 excitatory and 12 inhibitory contacts per branch, 14,336 active synapses, 21,760 candidate slots, 66,178 trainable parameters and 2,944 persistent state scalars per sample. Reversing the fine and global feature ranges provided the same-resource tree-placement control.

The operating point used signal contrast $\Delta = 0.80$, matched per-factor train/test log-gain standard deviation 0.25, child conductance 16 and mechanism-neutral initialization. It was selected transparently after an original severe-shift canary failed at every depth and two-seed exploratory gain, coupling and signal ladders located an accessible non-ceiling boundary (Supplementary Fig. S15). No pilot seed entered inference. The fresh confirmatory

seeds 10200–10209 crossed aligned, zero-alignment, trial-shuffled and reversed-placement shunting BP; aligned raw additive BP; and aligned or reversed-placement shunting LocalCA with shared-soma or exact path transport, for 270 fits. Adam ran for at most 180 epochs with validation-based early stopping. Learning rate was 0.003 for BP and 0.0008 for LocalCA; all comparisons were paired by seed.

The prespecified primary effect was aligned D3 minus D1 accuracy. Topology specificity required the aligned depth effect minus the corresponding zero-alignment, shuffled-sensor or rewired-tree depth effect. A positive local claim separately required its depth effect and aligned-minus-rewired interaction. Intervals used 50,000 seed-paired bootstrap draws; exact two-sided sign-flip tests enumerated all 2^{10} paired sign assignments. The artifact gate required all 270 finite results, no LocalCA fallback message, exact resource equality, at least nine of ten aligned BP seeds above 0.55 at every depth and condition means below 0.98. A dedicated checkpoint replay measured branch voltage, total conductance, input resistance, activation derivative, path-gain dispersion and local-versus-autograd conductance-gradient cosine on a fixed held-out batch.

Point–dendrite and BP–local-credit controls

The same frozen generator, operating point, stopping rule and paired seeds were used for a 140-fit reviewer-requested extension. The all-active grouped-star model retained the hierarchical model’s complete state dictionary and resource counts but evaluated every nonsomatic level independently and pooled the resulting child aggregators only once at the soma. D1 is therefore an equality control; D2 and D3 remove serial composition without removing branch modules. The aligned and reversed-placement regimes were crossed with D1–D3. Two ordinary point MLPs were trained on the aligned task: one matched the active parameter count (15,055), and the nearest integer-width model matched the total count (66,200 versus 66,178 in the tree). Neither MLP was supplied the task hierarchy or branch grouping.

For the teaching-coordinate control, the forward pass and synaptic autograd derivatives were unchanged, but a backward hook replaced every nonsomatic output derivative by the derivative at its owning soma. This “soma-broadcast autograd” condition was first run with the standard BP optimizer, which makes its contrast with full BP coordinate-specific. A configuration audit performed after those outcomes revealed that it could not be compared mechanistically with LocalCA, which used split parameter groups and learning rates selected for its stability. Before viewing the amendment outcomes, we therefore froze 60 additional aligned and reversed D1–D3 fits that changed only soma-broadcast to the LocalCA optimizer groups (global rate 0.0008, branch-weight rate 0.0007 and decoder rate 0.0015). This amendment is identified as a post-outcome diagnostic, not a prospective rescue of the primary analysis.

The architecture, alignment-interaction and credit contrasts were computed within seed. Intervals used 50,000 paired-seed bootstrap draws, and exact two-sided sign-flip tests enumerated all 2^{10} assignments. The artifact gate required 200 finite new fits, no fallback pathway, seeds 10200–10209 and exact equality of trainable parameters, active contacts, candidate slots and persistent state between serial and star models. The implementation tests also required identical D1 forward logits and exact equality of forward logits between full and soma-broadcast autograd modes before training. All gates passed. Configurations, the dated primary contract and its amendment, seed-level outcomes and the machine-readable audit are included with the source data.

Physical-depth alignment dose response

The interpolation extension retained the same generator, shunting-BP model, operating point, early-stopping rule, resources and paired seeds as the physical-depth endpoint experiment. It crossed $\alpha \in \{0.25, 0.50, 0.75\}$ with D1–D3 for 90 new fits and combined them with the unchanged $\alpha = 0$ and 1 endpoint outcomes. The dated contract was frozen after the endpoint outcomes were available but before any intermediate run was inspected. Its primary summary was the ordinary least-squares slope, within seed, of D3-minus-D1 accuracy on alignment; the $\alpha = 0.75$ minus 0.25 change and adjacent-level changes diagnosed linearity. Intervals used 50,000 paired-seed bootstrap draws, and exact two-sided sign-flip tests enumerated all 2^{10} assignments.

A concurrent source-tree edit interrupted 26 of the first 90 scheduled jobs through the frozen-source guard.

Those jobs were rerun under the new source hash; the 64 completed pre-edit jobs were retained because the only code difference introduced an optional rule override that no resolved configuration enabled. A machine-readable two-manifest audit and the dated source-equivalence record accompany the source data. The final gate required exactly 90 unique, finite intermediate outcomes, no fallback pathway, the expected seeds and exactly equal resources across depth; all conditions passed.

Literal grouped-point and second-hierarchy experiments

After the original $H = 3$ endpoints, nonserial-star controls and alignment-dose outcomes were known, we froze a final 220-fit matrix before opening any new outcome. The literal grouped-point model retained each nonsomatic branch module, structured input mask, local shunting equation and trainable scalar, but replaced the serial child aggregator at every level by a direct parameter-matched projection to its owning soma. Each replacement contained the same number and initialized value of raw conductance parameters as the aggregator it replaced. Every nonsomatic level therefore operated on the raw external streams without receiving another branch state. Unit tests required finite trainable gradients, copied conductances and exact D1 logits against the serial model.

The H3 arm crossed aligned and reversed placement, D1–D3, BP and the existing paired seeds 10200–10209 for 60 new fits; the already observed serial cohort was retained unchanged as its paired reference. The H2 arm used two 32-feature sensor tiers, fixed inventory counts 4/4 and exact-resource morphologies D1 [8] and D2 [4, 1]. It crossed aligned and reversed placement with serial BP, literal grouped-point BP and serial LocalCA using shared-soma or exact-path transport over ten fresh paired seeds 10300–10309, for 160 fits. Task-family, signal, gain, coupling, learning-rate, early-stopping and all other selected operating-point settings were inherited without retuning from H3.

Primary estimands were the H3 aligned serial-minus-grouped-point contrast at D3 and its placement interaction; the H2 serial-BP D2-minus-D1 effect and its placement interaction; the H2 serial-minus-grouped-point contrast at D2 and its placement interaction; and the H2 shared- and path-LocalCA depth effects and placement interactions. Intervals used 50,000 paired-seed bootstrap draws, and exact two-sided sign-flip tests enumerated all 2^{10} assignments. A directional positive claim required an interval excluding zero and at least eight of ten positive differences. All 220 fits completed with finite metrics, no fallback or nonfinite alert, the expected seeds and exact resource equality. Frozen manifests, run-level outcomes, every paired contrast and the machine-readable audit are distributed with Source Data.

Fourth-hierarchy extension and clean-source replication

We next froze an $H = 4$ extension to distinguish a hierarchy-matched depth prediction from the simpler possibility that the deepest offered model always wins. The positive-rate task, optimizer, conductance operating point, train-validation protocol and stopping rule were inherited without retuning from the $H = 3$ cohort. The 64 signal features were divided into four 16-feature tiers with inventory counts [4, 2, 1, 1]. Aligned placement assigned the tiers in fine-to-global physical order; the reversed control assigned them in the opposite order. D1 [8], D2 [2, 3], D3 [2, 1, 2] and D4 [1, 1, 2, 2] each contained eight nonsomatic branch units per soma (8, 2 + 6, 2 + 2 + 4 and 1 + 1 + 2 + 4). Each resolved preflight contained 66,178 trainable parameters.

The 360-fit factorial used ten fresh paired seeds 10400–10409. It crossed D1–D4 for aligned and reversed serial shunting BP, aligned and reversed literal grouped-point shunting BP, aligned and reversed serial LocalCA with shared-soma or exact-path transport, and aligned serial raw-additive BP. Frozen primary estimands were aligned serial-BP D4-minus-D3 and D4-minus-D1 effects; their placement interactions; the serial-minus-grouped D4 contrast and its placement interaction; shared- and path-LocalCA D4-minus-D3 effects and placement interactions; and the shunting-minus-additive interaction in D4-minus-D3. Failure of D4 to exceed D3 was defined in advance as falsification of the simple $D_p = H$ trained crossover at $H = 4$. Every contrast is retained regardless of direction. Intervals use 50,000 paired-seed bootstrap draws; exact two-sided sign-flip tests enumerate all 2^{10} assignments, and the frozen primary family receives Benjamini–Hochberg correction. A positive directional statement additionally requires at least eight of ten paired differences in the predicted direction.

All physical-depth extension jobs were generated and submitted from immutable modeling commit a99c3a7;

each job exported that checkout explicitly and refused a changed commit or tracked diff. Before training, serial shunting, literal grouped-point, raw-additive and both LocalCA D4 configurations passed model construction and resource audits. For D1, serial and literal grouped-point models received an identical 32-example batch and produced bitwise-identical initial logits (maximum absolute difference zero). The source commit, config hashes and gate result are supplied in a machine-readable audit.

Because the earlier H2/H3 arrays had been guarded by a base commit plus a tracked-diff hash rather than a standalone immutable commit, we also froze a 430-fit source-concordance replication before inspecting any rerun outcome. It reuses the original configurations and paired seeds without changing any scientific design: all 220 literal-point/H2 fits and the 210 load-bearing H3 serial shunting, LocalCA and raw-additive fits were rerun from a99c3a7. This is a same-seed implementation replication, not an independent biological or task replication. Each clean-source outcome is paired to its historical seed and semantic condition; the full distribution and maximum absolute held-out-accuracy change are reported. Scientific contrasts are recomputed at the seed level. If a load-bearing contrast changes sign or loses its frozen directional gate, the clean-source endpoint replaces the historical value rather than being selectively discarded.

State-matched positive-conductance reliability experiment

The locked mechanism experiment used eight branches with four nonnegative rate inputs per branch, 512 training examples, 2,048 test examples and one learned positive excitatory conductance per input. Branch voltages were $V_b = g_b^E / (g_b^L + g_b^E)$. A nonnegative shunt κ_b with reversal zero was paired with the oracle current $\kappa_b V_b$, so the post-intervention voltage was exactly $(g_b^E + \kappa_b V_b) / (g_b^L + g_b^E + \kappa_b) = V_b$, while the local eligibility was attenuated by $(g_b^L + g_b^E) / (g_b^L + g_b^E + \kappa_b)$. This state clamp separated physical input-resistance attenuation from forward feature suppression.

Initial full-data gradients defined branch signal energy S_b . Independent coefficient noise was scaled to give $N_b = S_b \exp(-2hx_b)$ for eight equally spaced $x_b \in [-1, 1]$ and $h \in \{0, 0.5, 1, 1.5, 2\}$. With the actual step $\eta = c/L$ and $c = 0.5$, shunts were fixed from the initial oracle attenuation $\min\{1, S_b/[c(S_b + N_b)]\}$ or its step-consistent best global, shuffled or reversed control; they were not learned. Other controls were clean exact backpropagation, noisy credit without a shunt, a state-matched additive current, an explicit point gate receiving the identical factors, and aligned attenuation with clean credit. Forty common-noise full-data updates used one half of the numerically estimated initial smoothness limit. This correction was frozen after identifying that the preceding pilot had paired this half step with the $c = 1$ gain. Two new artifact-only canary seeds preceded 50 fresh paired confirmatory seeds 8100–8149. The frozen gates required nonnegative inputs and shunts, voltage matching and independently computed point/additive equivalence within 10^{-12} , and relative finite-difference errors below 10^{-5} for both the unshunted objective and shunted frozen-clamp surrogate. Observed maxima were 2.22×10^{-16} for state mismatch, 4.00×10^{-15} and 1.31×10^{-14} for the equivalence gates, and 3.60×10^{-8} for the physical finite-difference error. Intervals resampled paired seeds with 20,000 draws; paired tests were two-sided Wilcoxon signed-rank tests. Configuration, runner hashes, all 2,250 seed-condition rows and the prespecified contract are supplied with Source Data.

Same-span noisy coefficient-learning experiment

The rate-based coefficient experiment fixed a 16-coordinate balanced-tree space and the rank-eight span through depth three. Three parameterizations shared this projector: eight orthonormal Haar columns, 15 redundant raw nested indicators and the same indicators multiplied by fixed positive gains spanning $\exp(-3)$ to $\exp(3)$. Their positive Gram condition numbers were 1, 15 and 388.52. Each seed supplied a unit-norm target in the common span and 80 common noisy target estimates. Observation standard deviation was $2/\sqrt{n_{\text{eff}}}$ for $n_{\text{eff}} \in \{4, 16, 64, 256\}$. Vanilla coefficient descent used 0.2 of each dictionary’s stability limit. A separate oracle control multiplied the coefficient gradient by the pseudoinverse Gram matrix, producing the same projected field step in every parameterization.

The design and predictions were frozen before canary seeds 9098–9099 and 50 paired confirmatory seeds 9100–9149. All 4,800 checkpoint rows completed. Projectors agreed to 9.44×10^{-16} , target address residual was at

most 1.45×10^{-30} , and Gram-preconditioned trajectories agreed to 1.16×10^{-15} . The unconditional orthonormal-coordinate prediction failed in the low-data regime. We therefore report both the prespecified contrast and the exact post-result bias–variance decomposition rather than relabeling the outcome as confirmatory. Intervals resampled paired seeds with 20,000 draws and tests were two-sided Wilcoxon signed-rank tests.

The historical inhibitory-dose cohort used the noise-resilience task and crossed two neuron models, five inhibitory-contact counts per branch (0, 5, 10, 20 and 40), backpropagation or one of the three local feedback fields, and ten paired seeds, for 400 runs. Its prespecified endpoint was a shunting-minus-additive contrast. Because every shunting cell used the invalid signed-input convention, the complete family is retained in the run-validity ledger but excluded from figures and inference.

The replacement exact-transport/backpropagation audit crossed the two tasks, two cores, depths one to four and paired seeds 42–51, giving 320 runs and 160 method pairs. Training used detached tracked-clean source commit 74792ca and nested synthetic-data commit 4f3612a. The noise- task additive core retained signed inputs; the positive-conductance shunting core used an explicit ReLU transfer before conductances were formed. Acceptance required resolved configurations and seeds, complete learning summaries, finite held-out metrics, finite stage-complete checkpoints, current scheduler logs without a fatal error or transport-shape fallback, and clean source identity. Exact-minus-backpropagation effects were paired within task, core, depth and seed; the plotted task–core contrasts first averaged four depths within seed. Incomplete setup and resource attempts supplied no endpoint and were excluded independently of outcome.

The fixed-topology cohort contained 320 runs at depth four and crossed the two tasks, neuron models, backpropagation or three local feedback fields, random or spatially partitioned sparse input indices, and ten paired seeds. Both maps had 21 contacts per distal branch. The spatial map recursively divided the 28×28 feature plane into 16 disjoint rectangles per neuron; the random map sampled each branch independently from all 784 coordinates. A configuration-level audit quantified unique coordinate coverage and cross-branch overlap for the same ten topology seeds. The validity rule removed the 80 synthetic-noise shunting runs, leaving 240 comparisons. Because any spatial effect shared by backpropagation is a forward-connectivity effect, this cohort was not used to claim topology-specific credit transport.

The fixed-budget depth cohort used the synthetic task and contained 320 runs. The retained 160 additive runs crossed four depths, three local feedback fields and matched backpropagation at ten paired seeds while fixing 16 terminal branches and 960–968 active contacts per soma. This matches terminal branches and input contacts, not compartment or trainable-parameter count.

For the prospective fixed-checkpoint diagnostic, scalar, neuron-indexed, and exact teaching fields were evaluated at the 120 input-valid backpropagation checkpoints on one frozen test batch of 128 examples. Only non-somatic dendritic stages and their trainable branch parameters entered the field and gradient vectors. Scaled capture was one minus the squared residual after optimally rescaling a candidate vector to the exact vector. For the learning endpoint, each candidate gradient was normalized to the exact-gradient norm and applied without further training. Step size was $r \|\theta\| / \|\nabla \mathcal{L}\|$ for $r \in \{10^{-6}, 10^{-5}, 10^{-4}, 10^{-3}\}$; 10^{-5} was the primary plotted operating point and all four are supplied as source data. Relative progress was the candidate loss decrease divided by the loss decrease from the norm-matched exact update. The association interval resampled checkpoints while retaining the paired scalar and neuron-indexed fields.

MICrONS morphology processing

Cells were drawn from the functionally co-registered MICrONS cohort and resolved from stable nucleus identifiers to roots in materialization 1822 of minnie65_public. Skeletons were downloaded through the MICrONS skeleton service. Incoming synapses came from synapses_pni_2. The original eight cells were a convenience pilot selected under the earlier analysis contract, not a random or representative sample; every anatomical claim from this cohort is therefore explicitly single-animal and descriptive. Degree-two dendritic chains were collapsed into branch-to-branch segments. Synapses were assigned to the nearest dendritic cable node and retained within $5 \mu\text{m}$. Presynaptic coarse type supplied E/I identity where available; otherwise, a high-confidence target-structure prediction supplied a secondary proxy from the in-development synapse_target_predictions_ssa_v2 table. This table does not

currently have a stable public release identifier and is not required for the direct-type-only or disjoint-v661 analyses. The direct-type-only analysis discarded all proxy-classified contacts.

Raw axial coupling was proportional to radius squared divided by cable length; it was normalized by its nonzero cell median and clipped to 0.05–20. Raw leak was proportional to radius times the larger of cable length and radius; it was normalized by its nonzero median, and was clipped to 0.1–10. Excitatory and inhibitory synaptic areas were normalized by the median nonzero total synaptic area within cell. These quantities define sensitivity weights, not fitted physical conductances. The soma was the root of every compressed tree. Ancestry matrices were computed by explicit parent traversal and checked for acyclicity and a unique path to the soma.

For the disjoint sensitivity cohort, we froze a 55-cell eligibility universe before inspecting outcomes, mapped all candidates to the public MICrONS v661 static release, and excluded the original eight cells by stable nucleus identifier. All 47 remaining cells passed the structural criteria. Only direct presynaptic E/I annotations were used. Direct labels were available for 11,024 of 234,301 incoming synapses before spatial mapping (4.71%); 10,516 contacts mapped to retained cables (median 212 directly typed synapses per cell). The cohort was not a population-random sample and contained 34 L2 IT, 2 L3 IT, 10 L4 IT, and 1 L5 ET cell. The same structural-capacity procedure was applied without pooling the two cohorts.

MICrONS inhibitory-census analysis

The inhibitory analysis joined the 47-cell v661 cohort to the published v795 census by stable nucleus identifier, before evaluating any endpoint. All 20 overlapping targets were retained. We used the published cell-type table, inhibitory-synapse table, complete input-synapse table, and target SWCs [33]. Raw all-input coordinates were mapped into the published slanted micrometre frame with an affine transformation fit from 80,114 synapse locations represented in both coordinate systems. The maximum transform residual was $1.64 \times 10^{-12} \mu\text{m}$. Input and known inhibitory locations were mapped to the nearest target skeleton node and retained within $5 \mu\text{m}$. Inputs without a known presynaptic inhibitory cell are called generic rather than excitatory because the census does not label every incoming contact.

For a multisynaptic inhibitory connection, one largest contact was retained within each published axonal clump. Every retained contact was matched to 64 generic input segments in the same compartment with nearest log path distance to the soma. Five hundred independent matched assignments estimated the null within each connection. Pairwise outcomes were cable-tree distance, normalized shared path to the lowest common ancestor, and cosine overlap between the two input-size-weighted binary descendant domains. Outcomes were first averaged over inhibitory connections within each postsynaptic target. The target, not a contact or connection, was the inferential unit.

Two sensitivity analyses tested dependence and spatial matching. First, connection-level contrasts were averaged within each of 92 presynaptic axons before inference. Second, each candidate location was ranked jointly by its log path-distance and three-dimensional Euclidean proximity to the observed contact, using the sum of the two within-pool percentile ranks; the same 500 matched assignments were then repeated. These analyses ask whether a target-level contrast persists after repeated connections from one axon are clustered and after local tissue proximity is included in the null.

For route capacity, rows were segments with mapped input and columns were segments bearing at least one known inhibitory contact. The binary ancestry matrix was weighted by the square root of mapped input size by row and known inhibitory size by column. Actual sites were ranked by column norm. At each budget, we reconstructed the full weighted actual-route matrix from the selected columns. Controls were 200 random subsets of actual sites, depth bins, depth-preserving within-column ancestry shuffles, and an unconstrained SVD oracle. The leverage selection is a capacity diagnostic, not a biological learning rule. Because the selected columns and reconstruction target come from the same actual-route matrix, this endpoint measures internal dictionary compressibility rather than generalization to an independently generated credit field.

Credit dictionaries and capture

The weighted direct route matrix was $\Gamma_{ik} = A_{ik}\beta_k$. It was a normalized route-sensitivity proxy, not the exact derivative of the reciprocal passive cable solve. Perturbation fields were sampled by multiplying independent route amplitudes by K . At each channel budget, ancestry routes were ranked by β_k times the fraction of total mapped excitatory synaptic area among their descendants. Controls were an unconstrained principal-component basis, random nonempty real paths, bins of cable depth, and column-wise ancestry shuffles. Channel count and the number of evaluated examples were matched. Feedback wiring density was the fraction of nonzero entries in the site-by-channel dictionary.

For each dictionary, weighted least squares supplied the optimal projection after mapping both the field and dictionary into the synaptic-area-weighted coordinates of Eq. 12. Twenty independent Monte Carlo streams, each containing 384 perturbation fields and 64 random or shuffled dictionaries per cell and channel count, measured numerical sensitivity. Streams were not treated as biological replicates. Primary comparisons were averaged within cell before paired testing.

For the independent-generator control, the reconstruction target was not sampled from K . We added a small focal shunt at every inhibitory-bearing segment of the reciprocal passive-cable system, restored the baseline somatic voltage with an independently solved state-matching compensatory current, and finite-differenced the log magnitude of the exact excitatory-conductance gradient at every excitatory-bearing segment. The derivative dose was 10^{-3} times the focal segment's membrane-plus-synaptic conductance. This produced a weighted response operator H whose columns were exact focal responses. Expected field-energy capture for a dictionary Φ was evaluated without Monte Carlo sampling as

$$\frac{\|P_{W^{1/2}\Phi}W^{1/2}H\|_F^2}{\|W^{1/2}H\|_F^2}, \quad (20)$$

where W contained mapped excitatory synaptic area and P denoted the orthogonal projector onto the stated column space. In addition to the primary controls, 200 surrogate trees per cell reassigned child subtrees among parent slots at the preceding depth. This preserved every node and edge attribute, every node's topological depth, and every parent's out-degree while changing ancestral membership. A separate row shuffle preserved the selected route values and sparsity while breaking their assignment to excitatory segments.

Focal conductance perturbations

Before running the focal analysis, we fixed the site-selection rule, primary dose, comparison groups, and localization endpoint. Sites were selected by a depth-stratified rule and required at least three excitatory descendants and three comparison sites; at most 16 sites were sampled per cell across topological-depth quartiles. Primary excitatory and inhibitory synaptic scales were both 0.35, with normalized reversal potentials $E_E = 1$ and $E_I = -0.2$. Conductance doses were 0.25, 0.5, 1, and 2 times the local membrane-plus-synaptic conductance. The additive control injected the first-order current $\kappa_k(E_I - V_k)$ at the same node without changing G . An independently solved state-matching compensatory current injection restored baseline somatic voltage after either intervention. This is not the derivative of a continuously adjusted voltage-clamp controller.

The objective was squared somatic error, with the target voltage set one normalized unit below baseline. Excitatory-conductance gradients were obtained by an adjoint solve, treating the solved compensatory current as fixed during the post-intervention derivative, and were checked by finite differences. If $d_i = E_E - V_i$ and $d'_i = E_E - V'_i$, the post-shunt gradient $q'_i d'_i$ was decomposed algebraically into an adjoint-only substitution $q'_i d_i$ and a driving-force-only substitution $q_i d'_i$, alongside the full-shunt and matched-additive gradients. Two-factor Shapley values allocated the nonlinear full-shunt-minus-additive localization payoff between the adjoint and driving-force changes and reconciled exactly to the observed contrast at each site. The matched-additive reference is not a literal no-factor physical intervention, so an unperturbed-reference allocation was also computed as a sensitivity analysis. These substitutions are not independent biological interventions. Synapses were classified as descendants, ancestors, sister-subtree sites, or unrelated sites. Unrelated synapses within each template were matched to descendant cable depth. For baseline gradient g_i and perturbed gradient g'_i , we defined $m_i = |\log[(|g'_i| + \varepsilon)/(|g_i| + \varepsilon)]|$, where

1305 $\varepsilon = 10^{-15} \max(1, \max_i |g_i|)$. The frozen site-level localization index was the median m_i among descendants minus
 1306 the median among depth-matched unrelated synapses. Site effects were averaged within cell; the primary paired
 1307 contrast used the eight cell means. A topology control sampled 200 relation templates from other focal sites, stratified
 1308 by focal-site topological-depth quartile when another site was available; otherwise it used another site in the same
 1309 cell. Axial and leak normalization and clipping are specified in Supplementary Section S6.

1310 We separately calibrated the passive cable in physical units. Non-root segments were treated as cylinders
 1311 with axial conductance $\pi r^2/(R_a \ell)$ and leak conductance $2\pi r \ell/R_m$; the root used the spherical membrane area
 1312 $4\pi r^2$. Conductances were expressed in nS. Synaptic areas are not conductance measurements, so excitatory and
 1313 inhibitory synaptic conductances retained the prespecified relative scale and were referenced to the cell-median
 1314 physical leak conductance. We evaluated $R_a \in \{100, 150, 250\} \Omega \text{cm}$ at $R_m = 15,000 \Omega \text{cm}^2$ and, at $R_a = 150 \Omega \text{cm}$,
 1315 $R_m \in \{300, 1000, 3000, 5000, 15,000, 30,000\} \Omega \text{cm}^2$. The same sites, state-matching current injection, adjoint
 1316 gradient, unit relative dose, and additive comparison were used. Two zero-length non-root segments in the v661
 1317 source skeletons used their radius as the cable length; this fallback and all electrotonic ratios are recorded in source
 1318 data.

1319 The frozen passive focal-sensitivity matrix used the original eight cells at $R_a = 150 \Omega \text{cm}$. We crossed membrane
 1320 resistivities of 300, 1000 and $15,000 \Omega \text{cm}^2$ with distributed leak-like background conductance equal to 0, 1 or 4
 1321 times each segment's baseline leak. At every selected site we evaluated three dose schemes: 0.25, 1 and 2 times
 1322 local membrane-plus-synaptic conductance; fixed doses 0.05, 0.5 and 5 nS; and 0.25, 1 and 4 times local input
 1323 conductance $[(G^{-1})_{kk}]^{-1}$. Rank-one voltage and adjoint updates were computed exactly with Sherman–Morrison,
 1324 followed by the same state-matching somatic current. In addition to absolute localization, the frozen outputs included
 1325 signed log-magnitude change, attenuated and enhanced fractions, sign-flip fraction, descendant gradient-energy
 1326 ratio, local input resistance and S_k from Eq. 14. Site values were averaged within cell before eight-cell intervals and
 1327 shunt-minus-additive contrasts. A numerical canary required positive-definite matrices and somatic residual below
 1328 10^{-10} before the full matrix ran with unchanged code and configuration.

1329 For the active steady-state sensitivity, we added voltage-dependent Na, K, Ca, HCN and NMDA currents to the
 1330 physical passive system at $R_a = 150 \Omega \text{cm}$, $R_m = 1,000 \Omega \text{cm}^2$ and one baseline-leak unit of distributed background
 1331 conductance. Channel activation followed sigmoidal steady-state gates with channel-specific half-activation, slope,
 1332 power and reversal values fixed in the machine-readable protocol. Global and segment-level maximal-conductance
 1333 multipliers were independent log-normal draws. Newton iterations with residual-reducing line search solved the
 1334 full nonlinear equilibrium. Draws were accepted only when the voltage remained within the prespecified range,
 1335 the residual was below 10^{-10} and the symmetrized local Jacobian was positive definite above 10^{-8} . With purely
 1336 local voltage-dependent gates, the off-diagonal Jacobian entries are the reciprocal axial couplings, so the steady-
 1337 state Jacobian is symmetric and the symmetrized gate and the plain inverse used below are exact rather than
 1338 approximations.

1339 For every accepted equilibrium, the adjoint used the exact inverse of that local Jacobian. Focal rank-one shunts
 1340 and first-order-current-matched additive inputs were evaluated at 0.25, 1 and 4 times local input conductance,
 1341 followed by the same state-matching compensatory somatic current. We required 64 accepted draws in every cell, up
 1342 to 256 attempts; all 512 first attempts were accepted. The 64 channel draws and up to 16 focal sites were averaged
 1343 within each reconstructed cell before bootstrap intervals and paired tests. The active ensemble is a steady-state
 1344 local-linearization sensitivity, not a simulation of channel kinetics or spikes.

1345 The focal analysis was also repeated after removing all proxy-classified contacts. In the disjoint v661 sensitivity
 1346 cohort, 45 of 47 cells supplied 235 eligible focal sites for the primary shunt–additive comparison. Forty cells with at
 1347 least two eligible sites supplied the reassigned-relation control. Cells, rather than focal sites, remained the units of
 1348 inference.

1349 Measured-response analyses

1350 Functionally imaged presynaptic partners were linked to their anatomical contacts through the MICrONS/CAVE and
 1351 DANDI mappings. The original frozen one-scan cohort included 69 partners across seven targets. The scan-complete
 1352 analysis deterministically retained every automatic-conservative target–session–scan cohort with at least five unique

connected imaged presynaptic roots, giving 13 scan observations nested within the same seven targets. Each scan contained 464 trials, 280 stimulus identities and 136 repeated identities. Multiple contacts from one presynaptic partner were pooled; its analysis site was the segment with greatest summed contact area. Functional similarity was the Pearson correlation of condition-mean responses over repeated identities. For each target, the partial-rank analysis rank-transformed similarity and shared ancestry, separately residualized both against log Euclidean separation and path-depth difference, and correlated the residuals. Scan-level statistics were averaged within target before the target bootstrap and signed-rank test. The target neuron, not a scan or synaptic pair, was the inferential unit.

For task-derived credit, a small branch model predicted each postsynaptic response from its connected presynaptic responses. Complete visual-stimulus identities were assigned to held-out splits. Exact branch errors were computed on training trials, and dictionary projectors were fit without held-out identities. Matched models were then trained from the same initialization with exact or projected branch errors. Ten fixed split/initialization replicates were averaged within each of seven target cells. Because projection coefficients used exact errors, these analyses estimate the capacity of a feedback family rather than a biological encoder. Only input-weight errors were projected; branch-readout weights and bias received exact gradients. Normalized MSE was held-out MSE divided by the held-out error of the training-mean predictor. Complete model and optimizer specifications are in Supplementary Section S9.

For the segment-resolved analysis, each complete compressed reconstruction was instantiated as a reciprocal passive conductance system. A trainable non-negative excitatory conductance was placed at the dominant mapped contact of every imaged partner. Exact gradients of somatic squared prediction error used the trial-specific inverse conductance operator and factored into $x_i(E_E - V_i)$ times the segment adjoint. Restricted methods first projected the baseline soma-to-site transfer vector into a four-channel conductance-weighted ancestry dictionary, a column-value-preserving site shuffle or an equal-rank random set of nonempty anatomical routes. This vector was calibrated once without target responses and frozen during learning; each restricted update then used only scalar output error, the fixed site vector and local eligibility. Exact readout and bias gradients were used in all methods. Ten seed-fixed splits held out complete stimulus identities within every scan. Replicates were averaged within scan and all eligible scans within target before seven-cell inference. Finite differences audited one exact conductance coordinate per run, input hashes and runner hash were archived, and any absent route, rank mismatch, non-finite value or incomplete scan caused exclusion; none occurred.

The capture-learning association excluded the exact and dense oracles. Both variables were centered within target before computing Spearman's ρ . Its interval used 20,000 target-level bootstrap draws, and its two-sided null used 20,000 independent reassignments of learning outcomes among method labels within each target.

Alignment-controlled optimization

For each reconstructed cell, we rebuilt the conductance-weighted ancestry kernel and selected eight routes by inhibitory fraction times descendant excitatory synaptic-area fraction, before generating gradient fields. Segment coordinates were weighted by the square root of summed mapped excitatory synaptic area. If Q is an orthonormal basis for the selected routes, each unit-energy exact gradient was constructed as

$$\mathbf{g}(a) = \sqrt{a}\mathbf{u}_{\parallel} + \sqrt{1-a}\mathbf{u}_{\perp}, \quad (21)$$

where $\mathbf{u}_{\parallel} \in \text{span}(Q)$, $\mathbf{u}_{\perp} \perp \text{span}(Q)$, and $a \in \{0, 0.2, 0.4, 0.6, 0.8, 1\}$. Directions were paired across alignment levels. Eight random real paths, eight cable-depth bins, and row-permuted ancestry routes supplied controls. The permutation preserved column values and nonzero count while breaking ancestry.

Each cell contributed 40 nested Monte Carlo streams with 256 gradients per stream. The objective was a positive diagonal quadratic with an anatomy-independent curvature spectrum from 0.5 to 1.5, permuted within stream and fixed across methods and alignment levels. One-step updates had norm 0.1 times the exact-gradient norm. Iterative projected descent used learning rate 0.25 for 20 steps. Both outcomes were normalized by the corresponding full-gradient progress. Streams and gradient fields measured numerical sensitivity; cell-level means were the inferential units.

1397 External animal-data reanalysis

1398 Francioni et al. source data were obtained from the public Nature record associated with DOI 10.1038/s41586-026-
1399 10190-7. The archived workbook was verified by SHA-256 before analysis. Extended Data sheet 13E supplied one
1400 P+ and one P- contrast per animal, defined by the source study as the change in soma–dendrite residual between
1401 error-reduction and error-increase periods. We retained the source signs. Signed separation was P+ minus P- and
1402 was checked against the value reported in the workbook. The common basis was $(1, 1)/\sqrt{2}$ and the signed basis was
1403 $(1, -1)/\sqrt{2}$; squared projection norms were divided by total two-coordinate energy. Sheet 5E supplied neuron-level
1404 distributions for visualization only. Directional animal-level hypotheses followed the known causal mapping. We
1405 report exact random-sign and Wilcoxon tests together with 50,000-draw animal bootstrap intervals.

1406 Static task-interference calculation

1407 The task-interference identity follows from exact quadratic expansion around a Task-A optimum, where the linear
1408 term vanishes. For its equal-route illustration, each task occupied 1,000 coordinates, route overlap varied from
1409 zero to one, and the Task-B target opposed Task A on shared coordinates. Off-route leakage supplied an update of
1410 relative magnitude λ on Task-A-only coordinates. We evaluated a 101×101 grid over overlap and leakage at step
1411 size 0.2. Explicit vector updates were compared with the closed form at every grid point. This analysis contains no
1412 temporal or spiking neural dynamics and uses the published motor-learning study only for qualitative, cross-study
1413 interpretation.

1414 Interior-optimum and wiring-normalized reanalyses

1415 Two deterministic secondary analyses of frozen outputs support Supplementary Fig. S16 and Fig. 6g,h; neither
1416 alters or reruns any training outcome. For the interior-optimum analysis, the per-seed initialization utility of Eq. 9
1417 was taken from the archived operator metrics of the 2,700-fit factorial and from the frozen depth-phase tables. The
1418 predicted optimum is the argmax of that utility over routed depth, or of the ancestry-minus-best-control utility
1419 contrast over budget; the observed optimum is the argmin of final population loss, or the argmax of the corresponding
1420 held-out accuracy contrast. Before any new quantity was reported, the pipeline reproduced the published $K = 4$
1421 ancestry-minus-best-control accuracy contrast (0.0127; 15 of 20 seeds) from the frozen seed outcomes. For the
1422 wiring-normalized analysis, capture per unit wiring is weighted field-energy capture divided by feedback wiring
1423 density, computed per cell from the frozen routing-capacity curves after averaging Monte Carlo streams within cell;
1424 the published eight-channel capture and wiring anchors were reproduced before any ratio was formed. Intervals
1425 are 95% seed- or cell-bootstrap intervals with 20,000 resamples. The laminar-capture proposition and the marginal
1426 utility condition (Supplementary Section S4) are verified numerically in the released test suite.

1427 Statistics and reproducibility

1428 All tests were two sided unless a directional hypothesis had been frozen in advance. Effect sizes and 95% intervals
1429 accompany central comparisons. Reconstructed-cell intervals resampled cells first and, where appropriate, a Monte
1430 Carlo stream within cell. Textual focal-attribution intervals used a hierarchical bootstrap over cells and sites; Fig. 7h
1431 instead shows cell-bootstrap intervals after averaging sites within each cell. Measured-response intervals resampled
1432 target cells. Artificial-network means and standard deviations used independent initialization seeds. Exact Wilcoxon
1433 signed-rank tests were reported for small paired cell cohorts; no site, trial, epoch, or Monte Carlo stream was counted
1434 as an independent biological replicate. Null results are described as a lack of reliable evidence rather than proof of
1435 equality. Bootstrap intervals used 20,000 draws. The four-budget primary subtree scan used Benjamini–Hochberg
1436 correction; other parameter sweeps and sensitivity analyses are labeled secondary and are not treated as independent
1437 confirmation.

1438 Seed counts followed the evidential role of each cohort rather than one pooled power calculation: three to
1439 five seeds were reserved for numerical or mechanism diagnostics; ten for fresh paired confirmatory replications
1440 with interval and sign-agreement gates; 15–20 for inherited trained factorials; and 50 for phase boundaries and

1441 crossover distributions. Effect sizes were not known before the new simulation families, so no prospective parametric
1442 power calculation was used. Positive confirmatory claims required a paired 95% interval excluding zero and the
1443 prespecified fraction of concordant seed signs; diagnostic cohorts were not promoted to inferential evidence. A
1444 machine-readable table in Source Data maps every headline endpoint to its independent unit and identifies nested or
1445 descriptive observations.

1446 Analysis code records configuration files, seeds, software versions, source tables, and machine-readable sum-
1447 maries. Numerical source data are organized by figure and panel in the accompanying package. Tests cover tree
1448 construction, projection identities, and exact-gradient checks.

1449 **Reporting of computational assistance**

1450 OpenAI Codex was used to assist code development, statistical and consistency audits, figure preparation, literature
1451 checks, and language editing. It was not listed as an author and did not determine the scientific conclusions. The
1452 authors retain responsibility for the study design, interpretation, accuracy of the code and numerical results, citations,
1453 and final manuscript.

1454 **Ethics and data reuse**

1455 No new experiments involving animals or human participants were performed. All biological analyses reused
1456 public MICrONS, DANDI, and published animal source data; experimental collection and institutional approvals are
1457 described in the source studies [33, 37, 56].

1458 **Data availability**

1459 MICrONS anatomy and connectivity are available through the public CAVE datastack minnie65_public. The
1460 primary cohort used materialization 1822, skeleton-service reconstructions, and the synapses_pni_2 table. The
1461 disjoint sensitivity cohort used the public v661 static release. The secondary target-structure labels used only in
1462 the full-contact analysis came from the in-development synapse_target_predictions_ssa_v2 table and are
1463 therefore distributed only as derived labels subject to the source-data access terms; all principal structural findings
1464 were repeated without those labels. Volumetric imagery is archived in BossDB under DOI 10.60533/BOSS-2021-
1465 TOSY, and co-registered functional data are available as DANDI:000402; exact asset and table identifiers are recorded
1466 in the supplied manifest. The source publication describes the full data release [56]. Public benchmark datasets are
1467 available from their original repositories. The inhibitory-census tables and SWCs are available with the source record
1468 for Schneider-Mizell et al. [33]. The Francioni et al. workbook is available from the Source Data link associated
1469 with the Nature article [37]; its checksum and used sheet names are recorded in the manifest. Derived cell and
1470 exclusion manifests, numerical source data underlying every figure, and non-restricted summaries are supplied with
1471 this manuscript. Authentication tokens and externally hosted raw caches are not redistributed.

1472 **Code availability**

1473 The accompanying reviewer software archive contains the complete training implementation used here, exact-gradient
1474 and perturbation analyses, frozen configurations, tests, figure scripts and byte-identified snapshots of the critical
1475 training components under the MIT License. A separate Source Data archive contains the numerical data underlying
1476 every display item. Submission is gated on replacing this sentence with the public repository URL and versioned
1477 archival DOI; no public identifier is asserted before that record is created.

References

- [1] Christof Koch. *Biophysics of Computation: Information Processing in Single Neurons*. Oxford University Press, 1999. doi: 10.1093/oso/9780195104912.001.0001.
- [2] Nelson Spruston. Pyramidal neurons: dendritic structure and synaptic integration. *Nature Reviews Neuroscience*, 9:206–221, 2008. doi: 10.1038/nrn2286.
- [3] Blake A. Richards and Timothy P. Lillicrap. Dendritic solutions to the credit assignment problem. *Current Opinion in Neurobiology*, 54:28–36, 2019. doi: 10.1016/j.conb.2018.08.003.
- [4] Albert Gidon and Idan Segev. Principles governing the operation of synaptic inhibition in dendrites. *Neuron*, 75(2):330–341, 2012. doi: 10.1016/j.neuron.2012.05.015.
- [5] Matthew Larkum. A cellular mechanism for cortical associations: an organizing principle for the cerebral cortex. *Trends in Neurosciences*, 36(3):141–151, 2013. doi: 10.1016/j.tins.2012.11.006.
- [6] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning representations by back-propagating errors. *Nature*, 323:533–536, 1986. doi: 10.1038/323533a0.
- [7] Timothy P. Lillicrap, Adam Santoro, Luke Marris, Colin J. Akerman, and Geoffrey Hinton. Backpropagation and the brain. *Nature Reviews Neuroscience*, 21:335–346, 2020. doi: 10.1038/s41583-020-0277-3.
- [8] James C. R. Whittington and Rafal Bogacz. Theories of error back-propagation in the brain. *Trends in Cognitive Sciences*, 23(3):235–250, 2019. doi: 10.1016/j.tics.2018.12.005.
- [9] Nicolas Frémaux and Wulfram Gerstner. Neuromodulated spike-timing-dependent plasticity, and theory of three-factor learning rules. *Frontiers in Neural Circuits*, 9:85, 2016. doi: 10.3389/fncir.2015.00085.
- [10] Timothy P. Lillicrap, Daniel Cownden, Douglas B. Tweed, and Colin J. Akerman. Random synaptic feedback weights support error backpropagation for deep learning. *Nature Communications*, 7:13276, 2016. doi: 10.1038/ncomms13276.
- [11] Arild Nøkland. Direct feedback alignment provides learning in deep neural networks. In *Advances in Neural Information Processing Systems*, volume 29, pages 1037–1045, 2016.
- [12] Benjamin Scellier and Yoshua Bengio. Equilibrium propagation: bridging the gap between energy-based models and backpropagation. *Frontiers in Computational Neuroscience*, 11:24, 2017. doi: 10.3389/fncom.2017.00024.
- [13] James C. R. Whittington and Rafal Bogacz. An approximation of the error backpropagation algorithm in a predictive coding network with local Hebbian synaptic plasticity. *Neural Computation*, 29(5):1229–1262, 2017. doi: 10.1162/NECO_a_00949.
- [14] Yuhang Song, Beren Millidge, Tommaso Salvatori, Thomas Lukasiewicz, Zhenghua Xu, and Rafal Bogacz. Inferring neural activity before plasticity as a foundation for learning beyond backpropagation. *Nature Neuroscience*, 27:348–358, 2024. doi: 10.1038/s41593-023-01514-1.
- [15] Guillaume Bellec, Franz Scherr, Anand Subramoney, Elias Hajek, Darjan Salaj, Robert Legenstein, and Wolfgang Maass. A solution to the learning dilemma for recurrent networks of spiking neurons. *Nature Communications*, 11:3625, 2020. doi: 10.1038/s41467-020-17236-y.
- [16] Mathieu Schiess, Robert Urbanczik, and Walter Senn. Somato-dendritic synaptic plasticity and error-backpropagation in active dendrites. *PLOS Computational Biology*, 12(2):e1004638, 2016. doi: 10.1371/journal.pcbi.1004638.

- [17] Jordan Guerguiev, Timothy P. Lillicrap, and Blake A. Richards. Towards deep learning with segregated dendrites. *eLife*, 6:e22901, 2017. doi: 10.7554/eLife.22901.
- [18] João Sacramento, Rui Ponte Costa, Yoshua Bengio, and Walter Senn. Dendritic cortical microcircuits approximate the backpropagation algorithm. In *Advances in Neural Information Processing Systems*, volume 31, 2018.
- [19] Alexandre Payeur, Jordan Guerguiev, Friedemann Zenke, Blake A. Richards, and Richard Naud. Burst-dependent synaptic plasticity can coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24:1010–1019, 2021. doi: 10.1038/s41593-021-00857-x.
- [20] Benjamin Ellenberger, Paul Haider, Federico Benitez, Jakob Jordan, Kevin Max, Ismael Jaras, Laura Kriener, and Mihai A. Petrovici. Backpropagation through space, time and the brain. *Nature Communications*, 17:66, 2026. doi: 10.1038/s41467-025-66666-z.
- [21] Luis B. Almeida. A learning rule for asynchronous perceptrons with feedback in a combinatorial environment. In *Proceedings of the First IEEE International Conference on Neural Networks*, volume 2, pages 609–618, 1987.
- [22] Fernando J. Pineda. Generalization of back-propagation to recurrent neural networks. *Physical Review Letters*, 59(19):2229–2232, 1987. doi: 10.1103/PhysRevLett.59.2229.
- [23] Houman Safaai, Maceo Richards, and Bernardo L. Sabatini. Shunting inhibition and dendritic branching shape local credit assignment. *arXiv preprint arXiv:2607.03556*, 2026. URL <https://arxiv.org/abs/2607.03556>.
- [24] Wilfrid Rall. Theory of physiological properties of dendrites. *Annals of the New York Academy of Sciences*, 96(4):1071–1092, 1962. doi: 10.1111/j.1749-6632.1962.tb54120.x.
- [25] Jakob Jordan, João Sacramento, Willem A. M. Wybo, Mihai A. Petrovici, and Walter Senn. Conductance-based dendrites perform Bayes-optimal cue integration. *PLOS Computational Biology*, 20(6):e1012047, 2024. doi: 10.1371/journal.pcbi.1012047.
- [26] Justin Werfel, Xiaohui Xie, and H. Sebastian Seung. Learning curves for stochastic gradient descent in linear feedforward networks. *Neural Computation*, 17(12):2699–2718, 2005. doi: 10.1162/089976605774320539.
- [27] Matteo Carandini and David J. Heeger. Normalization as a canonical neural computation. *Nature Reviews Neuroscience*, 13:51–62, 2012. doi: 10.1038/nrn3136.
- [28] Gary R. Holt and Christof Koch. Shunting inhibition does not have a divisive effect on firing rates. *Neural Computation*, 9(5):1001–1013, 1997. doi: 10.1162/neco.1997.9.5.1001.
- [29] J. Kornfeld, Y. Wang, M. Januszewski, A. Rother, P. Schubert, M. Goldman, V. Jain, W. Denk, and M. S. Fee. An anatomical substrate of credit assignment in reinforcement learning. *bioRxiv*, 2020. doi: 10.1101/2020.02.18.954354. Version 2 posted 27 October 2025.
- [30] Marissa A. Weis, Stelios Papadopoulos, Laura Hansel, Timo Lüddecke, Brendan Celii, Paul G. Fahey, Eric Y. Wang, J. Alexander Bae, Agnes L. Bodor, Derrick Brittain, JoAnn Buchanan, Daniel J. Bumbarger, Manuel A. Castro, Forrest Collman, Nuno Maçarico da Costa, Sven Dorkenwald, Leila Elabbady, Akhilesh Halageri, Zhen Jia, Chris Jordan, Dan Kapner, Nico Kemnitz, Sam Kinn, Kisuk Lee, Kai Li, Ran Lu, Thomas Macrina, Gayathri Mahalingam, Eric Mitchell, Shanka Subhra Mondal, Shang Mu, Barak Nehoran, Sergiy Popovych, R. Clay Reid, Casey M. Schneider-Mizell, H. Sebastian Seung, William Silversmith, Marc Takeno, Russel Torres, Nicholas L. Turner, William Wong, Jingpeng Wu, Wenjing Yin, Szi-chieh Yu, Jacob Reimer, Philipp Berens, Andreas S. Tolias, and Alexander S. Ecker. An unsupervised map of excitatory neuron dendritic morphology in the mouse visual cortex. *Nature Communications*, 16(1):3361, 2025. doi: 10.1038/s41467-025-58763-w.

- [31] Matthew Lovett-Barron, Gergely F. Turi, Patrick Kaifosh, Peter H. Lee, Frédéric Bolze, Xiao-Hua Sun, Jean-François Nicoud, Boris V. Zeman, Scott M. Sternson, and Attila Losonczy. Regulation of neuronal input transformations by tunable dendritic inhibition. *Nature Neuroscience*, 15:423–430, 2012. doi: 10.1038/nn.3024.
- [32] Tim P. Vogels, Henning Sprekeler, Friedemann Zenke, Claudia Clopath, and Wulfram Gerstner. Inhibitory plasticity balances excitation and inhibition in sensory pathways and memory networks. *Science*, 334(6062):1569–1573, 2011. doi: 10.1126/science.1211095.
- [33] Casey M. Schneider-Mizell, Agnes L. Bodor, Derrick Brittain, JoAnn Buchanan, Daniel J. Bumbarger, Leila Elabbady, Clare Gamlin, Daniel Kapner, Sam Kinn, Gayathri Mahalingam, Sharmishta Seshamani, Shelby Suckow, Marc Takeno, Russel Torres, Wenjing Yin, Sven Dorkenwald, J. Alexander Bae, Manuel A. Castro, Akhilesh Halageri, Zhen Jia, Chris Jordan, Nico Kemnitz, Kisuk Lee, Kai Li, Ran Lu, Thomas Macrina, Eric Mitchell, Shanka Subhra Mondal, Shang Mu, Barak Nehoran, Sergiy Popovych, William Silversmith, Nicholas L. Turner, William Wong, Jingpeng Wu, Jacob Reimer, Andreas S. Tolias, H. Sebastian Seung, R. Clay Reid, Forrest Collman, and Nuno Maçarico da Costa. Inhibitory specificity from a connectomic census of mouse visual cortex. *Nature*, 640:448–458, 2025. doi: 10.1038/s41586-024-07780-8.
- [34] Frances S. Chance, L. F. Abbott, and Alex D. Reyes. Gain modulation from background synaptic input. *Neuron*, 35(4):773–782, 2002. doi: 10.1016/S0896-6273(02)00820-6.
- [35] Alain Destexhe, Michael Rudolph, and Denis Paré. The high-conductance state of neocortical neurons in vivo. *Nature Reviews Neuroscience*, 4:739–751, 2003. doi: 10.1038/nnr1198.
- [36] Zhuokun Ding, Paul G. Fahey, Stelios Papadopoulos, Eric Y. Wang, Brendan Celii, Christos Papadopoulos, Andersen Chang, Alexander B. Kunin, et al. Functional connectomics reveals general wiring rule in mouse visual cortex. *Nature*, 640:459–469, 2025. doi: 10.1038/s41586-025-08840-3.
- [37] Valerio Francioni, Vincent D. Tang, Enrique H. S. Toloza, Zilan Ding, Norma J. Brown, and Mark T. Harnett. Vectorized instructive signals in cortical dendrites. *Nature*, 652:1254–1263, 2026. doi: 10.1038/s41586-026-10190-7.
- [38] Julian Rossbroich and Friedemann Zenke. Dis-inhibitory neuronal circuits can control the sign of synaptic plasticity. In *Advances in Neural Information Processing Systems*, volume 36, pages 64059–64082. Curran Associates, Inc., 2023. doi: 10.52202/075280-2799.
- [39] Alessandro R. Galloni, Ajay Peddada, Yash Chennawar, and Aaron D. Milstein. Cellular and subcellular specialization enables biology-constrained deep learning. *Cell Reports*, 45(4):117159, 2026. doi: 10.1016/j.celrep.2026.117159.
- [40] Will Greedy, Heng Wei Zhu, Alexia Duriez, Joseph Pemberton, Patrick T. McCarthy, Kevin Kermani Nejad, and Rui Ponte Costa. Cell-type-specific cortical feedback coordinates hierarchical credit assignment. *bioRxiv*, 2026. doi: 10.1101/2026.06.16.732595.
- [41] Houman Safaai, Maceo Richards, Naeem Khoshnevis, and Bernardo L. Sabatini. When branch-local shunting helps: A gain-load-alignment principle for dendritic E/I networks. *arXiv preprint arXiv:2607.24990*, 2026. URL <https://arxiv.org/abs/2607.24990>.
- [42] Patrick T. Sadtler, Kristin M. Quick, Matthew D. Golub, Steven M. Chase, Stephen I. Ryu, Elizabeth C. Tyler-Kabara, Byron M. Yu, and Aaron P. Batista. Neural constraints on learning. *Nature*, 512:423–426, 2014. doi: 10.1038/nature13665.
- [43] Panayiota Poirazi, Terrence Brannon, and Bartlett W. Mel. Pyramidal neuron as two-layer neural network. *Neuron*, 37(6):989–999, 2003. doi: 10.1016/S0896-6273(03)00149-1.

- [44] Michael London and Michael Häusser. Dendritic computation. *Annual Review of Neuroscience*, 28:503–532, 2005. doi: 10.1146/annurev.neuro.28.061604.135703.
- [45] Tiago Branco and Michael Häusser. The single dendritic branch as a fundamental functional unit in the nervous system. *Current Opinion in Neurobiology*, 20(4):494–502, 2010. doi: 10.1016/j.conb.2010.07.009.
- [46] Spyridon Chavlis and Panayiota Poirazi. Dendrites endow artificial neural networks with accurate, robust and parameter-efficient learning. *Nature Communications*, 16:943, 2025. doi: 10.1038/s41467-025-56297-9.
- [47] Robert Urbanczik and Walter Senn. Learning by the dendritic prediction of somatic spiking. *Neuron*, 81(3): 521–528, 2014. doi: 10.1016/j.neuron.2013.11.030.
- [48] Joseph Cichon and Wen-Biao Gan. Branch-specific dendritic ca^{2+} spikes cause persistent synaptic plasticity. *Nature*, 520:180–185, 2015. doi: 10.1038/nature14251.
- [49] Erik B. Bloss, Mark S. Cembrowski, Bill Karsh, Jennifer Colonell, Richard D. Fetter, and Nelson Spruston. Structured dendritic inhibition supports branch-selective integration in CA1 pyramidal cells. *Neuron*, 89(5): 1016–1030, 2016. doi: 10.1016/j.neuron.2016.01.029.
- [50] Brendan A. Bicknell and Michael Häusser. A synaptic learning rule for exploiting nonlinear dendritic computation. *Neuron*, 109(24):4001–4017.e10, 2021. doi: 10.1016/j.neuron.2021.09.044.
- [51] Joel Veness, Tor Lattimore, David Budden, Avishkar Bhoopchand, Christopher Mattern, Agnieszka Grabska-Barwinska, Eren Sezener, Jianan Wang, Peter Toth, Simon Schmitt, and Marcus Hutter. Gated linear networks. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35, pages 10015–10023, 2021. doi: 10.1609/aaai.v35i11.17202.
- [52] Abhiram Iyer, Karan Grewal, Akash Velu, Lucas Oliveira Souza, Jeremy Forest, and Subutai Ahmad. Avoiding catastrophe: active dendrites enable multi-task learning in dynamic environments. *Frontiers in Neurobotics*, 16:846219, 2022. doi: 10.3389/fnbot.2022.846219.
- [53] Changze Lv, Jingwen Xu, Yiyang Lu, Xiaohua Wang, Zhenghua Wang, Zhibo Xu, Di Yu, Xin Du, Xiaoqing Zheng, and Xuanjing Huang. Dendritic localized learning: Toward biologically plausible algorithm. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 41682–41700. PMLR, 2025. URL <https://proceedings.mlr.press/v267/lv25c.html>.
- [54] Ken C. Bittner, Aaron D. Milstein, Christine Grienberger, Sandro Romani, and Jeffrey C. Magee. Behavioral time scale synaptic plasticity underlies CA1 place fields. *Science*, 357(6355):1033–1036, 2017. doi: 10.1126/science.aan3846.
- [55] Jeffrey C. Magee and Christine Grienberger. Synaptic plasticity forms and functions. *Annual Review of Neuroscience*, 43:95–117, 2020. doi: 10.1146/annurev-neuro-090919-022842.
- [56] MICrONS Consortium. Functional connectomics spanning multiple areas of mouse visual cortex. *Nature*, 640: 435–447, 2025. doi: 10.1038/s41586-025-08790-w.

Acknowledgements

We thank the MICrONS Consortium and the maintainers of the associated CAVE and DANDI resources. We also thank the authors of the published animal-learning studies for making their source data available. This work was supported in part by a gift from the Chan Zuckerberg Initiative Foundation to establish the Kempner Institute for the Study of Natural and Artificial Intelligence at Harvard University.

1638 **Author contributions**

1639 H.S. developed the computational framework, performed the analyses, and wrote the initial manuscript. M.R.
1640 contributed to the local-credit-assignment theory, implementation, and interpretation. B.L.S. supervised the work and
1641 contributed to the biological framing, study design, and manuscript. All authors reviewed and revised the manuscript.

1642 **Competing interests**

1643 The authors declare no competing interests.