

---

# Path-Gain Compressibility Mediates Local Credit Assignment in Conductance-Based Dendrites

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

Biologically plausible credit assignment requires replacing compartment-specific backpropagation errors with signals available to dendritic synapses. We ask when a low-dimensional somatic broadcast can be a faithful substitute in conductance-based dendritic trees. Exact gradients show that each synaptic update factorizes into a local eligibility term, presynaptic activity times driving force times input resistance, and a single non-local compartment error. Local learning is therefore an approximation to a path-specific error field. We show that shunting inhibition modulates input resistance along dendritic paths and makes the exact error field more compressible for rank-1 broadcast in the regimes studied here. Gradient diagnostics show that shunting improves local-backprop alignment and reduces gradient scale mismatch relative to matched additive controls. Post-training inhibition interventions show that learned inhibitory conductance carries sample-specific pathway structure, while transported-error oracle experiments show that supplying the exact path-gain-transported error largely closes the local-learning gap. Under nonnegative conductances and rank-1 broadcast, shunting LocalCA stays within about 5–6 percentage points of matched backpropagation ceilings on MNIST, Fashion-MNIST, and context gating. Harder-data and routed diagnostics reveal the boundary of this mechanism: when the exact error field is not rank-1 compressible, richer pathway-specific feedback is required. These results identify path-gain compressibility as a measurable condition under which simple local credit assignment is effective in dendritic networks.

## 1 Introduction

Credit assignment in deep networks usually depends on backpropagation: global error transport through symmetric pathways with no obvious biological substrate. Dendritic compartments suggest an alternative. Each synapse has access to rich local state such as driving force, membrane conductance, and branch voltage, while supervisory information could arrive as a low-bandwidth broadcast from the soma [14]. The question here is not whether such a broadcast can match unconstrained backpropagation in all settings, but when it is a faithful approximation to the exact dendritic credit signal.

Starting from conductance-based dendritic voltage equations [1], we derive exact gradients for dendritic trees (Theorem 1) and show that each synaptic gradient factorizes into purely local eligibility terms and a single non-local compartment error. Approximating that compartment error with a broadcast signal yields a theorem-facing local rule (3F), practical optimizer wrappers (4F/5F), and higher-rank feedback variants for tasks where different branches require different teaching signals.

The central finding is mechanistic: shunting can make low-bandwidth broadcast useful by improving the compressibility of the conductance-stage path-gain field. Dendritic branching creates multiple error paths from distal synapses to the soma. Excitatory and inhibitory synapse banks set the conductance state of each compartment. Shunting makes inhibitory conductance enter the denominator, so inhibition changes input resistance and therefore the gain of every upstream path that passes through that compartment. This does not mean inhibition always helps. It helps rank-1 broadcast when it attenuates high-gain paths more than low-gain paths, narrowing or low-rank compressing the exact error field that the broadcast must approximate.

**Related work.** This paper sits at the intersection of dendritic computation and credit assignment. The compartmental model follows classic accounts of dendritic integration, normalization, and branch-level nonlinear computation [1, 3, 4, 6, 30, 31]. The learning question connects to dendritic and cortical models in which segregated compartments, bursts, predictive signals, or prospective activity carry teaching information [5, 11–13, 16, 22–24, 44]. It also relates to feedback-alignment and direct-feedback approaches [9, 10], target-propagation methods [27], equilibrium and predictive objectives [15, 28, 29], forward-forward or perturbation-based learning [25, 26], and recent local-learning baselines [32–36]. Our contribution is narrower than these broad frameworks: we derive the exact conductance-tree credit object and test when a rank-1 dendritic broadcast approximates it.


We quantify this chain with three diagnostics. First, exact-gradient reconstruction verifies that the mathematical factorization matches autograd to numerical precision. Second, gradient-fidelity measurements compare local and backprop gradients at matched parameters. Third, transported-error oracle experiments replace the broadcast with the exact path-gain-transported compartment error, isolating broadcast fidelity from other architectural factors. In a checkpointed MNIST gradient-dynamics diagnostic, shunting improves final local-backprop alignment (cosine  $0.100 \pm 0.033$  vs.  $0.005 \pm 0.026$ ) and reduces scale mismatch ( $0.26 \pm 0.05$  vs.  $0.50 \pm 0.05$ ). On noise resilience under matched inhibition ( $N_I \geq 5$ ), transported error raises additive local learning from 81–84% under rank-1 broadcast to 92–94%, nearly matching shunting at 95–95.2%.

**Contributions.** We make three claims: (i) exact conductance-tree credit factorizes into synapse-local eligibility and a path-specific compartment error (Theorem 1, Corollary 1); (ii) path-gain compressibility is a measurable mediator of when rank-1 broadcast approximates that compartment error, and shunting provides a biophysical route to such compression by changing input resistance along dendritic paths; and (iii) gradient-fidelity, transported-oracle, competence, and morphology experiments validate this mechanism while identifying the failure modes that require richer feedback.


The empirical story follows this sequence. Section 2 derives the conductance-tree credit object, and Section 3 defines the local approximations. Figure 1 defines the conductance tree and broadcast taxonomy. Figure 2 verifies the exact factorization. Experiments in Section 4 then test the mechanism: Figure 3 links path-gain compressibility to broadcast fidelity and learning, and Figure 4 shows the practical supervised regimes, including morphology dependence and the rule-family ablation. Section 5 summarizes the boundary conditions. Harder-data experiments remain appendix diagnostics rather than headline benchmark claims.

## 2 Compartmental Voltage Model and Gradient Derivation

We use a steady-state conductance model derived from discretized passive cable dynamics [1, 2]. In normalized units, leak and inhibitory reversal are set to 0, the excitatory reversal is set to 1, and leak conductance is fixed at 1. Each compartment voltage is therefore a conductance-weighted average. Two properties of this form matter for what follows: local sensitivities depend on both driving force ( $E - V$ ) and input resistance  $R^{\text{tot}}$ , and shunting inhibition corresponds to adding conductance when  $E_{\text{inh}} \approx 0$ . Throughout the paper we write dendritic morphologies as rooted trees  $[b_1, b_2, \dots, b_D]$ , where each factor gives the fan-in from one dendritic stage to the next on the path to the soma. Thus  $[3, 3]$  means that each soma aggregates 3 proximal branches and each proximal branch aggregates 3 distal branches, for 9 distal leaves per soma. In code, each excitatory unit is a stack of branch layers, each branch layer carries explicit excitatory and optional inhibitory synapse banks plus a branch-to-parent aggregation map, and the headline feedforward regime keeps `somatic_synapses=false` so inputs terminate on dendritic branches rather than directly on the soma. The main experiments use `input_mode=1`: the transfer layer supplies nonnegative excitatory and inhibitory input streams,



**Figure 1: Model and credit assignment.** (A) Conductance-based dendritic tree. Nonnegative excitatory and inhibitory inputs drive distal compartments, which feed proximal branches and then the soma; shunting inhibition enters the conductance denominator and lowers input resistance. (B) Path gains through the tree. Exact backprop supplies path-specific compartment errors, while rank-1 LocalCA replaces that field with one shared broadcast. (C) Exact factorization and feedback choices. LocalCA keeps the exact synapse-local eligibility and replaces the path-specific compartment error  $\delta_n$  with a broadcast  $e_n$ ; rank-1, neuron-wise, rank- $K$ , and path-structured broadcasts are local approximations with increasing bandwidth, while transported error is an analysis oracle.



**Figure 2: Gradient-fidelity diagnostic.** (A) Exact-factorization sanity check across all theory-diagnostic conditions, reported as relative  $L_2$  reconstruction error and scale mismatch. Reconstructed gradients match autograd to numerical precision. (B) Checkpointed final-state weighted cosine similarity between local and backprop gradients. (C) Checkpointed final-state scale mismatch, defined as  $|\log_{10}(\|g^{\text{local}}\|/\|g^{\text{bp}}\|)|$ . Error bars denote  $\pm 1$  s.d. across 3 checkpointed seeds; the main learning and oracle sweeps use 5 seeds unless noted otherwise. (D) Full-gradient norms over learning. Additive local and backprop gradients remain finite, so near-zero additive cosine reflects directional mismatch rather than vanishing gradients.

and inhibition enters through learned branch-level I-to-E conductance banks. This is input-driven shunting inhibition, not a recurrent lateral inhibitory circuit. The additive and shunting cores are therefore architecture-matched at the tree level; they differ only in the branch-level voltage integration rule. These ingredients play distinct roles in the credit-assignment argument. Branching creates the path structure. E/I synapse banks determine the local conductance state. Shunting turns inhibitory conductance into a multiplicative change in input resistance. The local learning rule then asks whether a low-dimensional teaching signal can stand in for the exact path-specific error induced by that tree.

**Notation guardrails.** We use  $\alpha_n$  for the exact conductance-stage path gain from compartment  $n$  to the soma,  $r_n^{4F}$  for the empirical 4F branch-soma covariance proxy,  $\phi_n$  for the 5F confidence factor,  $\lambda_{\text{mix}}$  for the pathway-broadcast mixing weight, and  $\kappa_j$  for depth-dependent modulation in morphology extensions. All conductances and presynaptic activities in the headline regime are nonnegative; additive voltages may be signed comparators, but the conductance variables that define the shunting model are nonnegative.

## 2.1 Voltage Equation and Local Sensitivities

Consider compartment  $n$  with synaptic inputs  $j$  (activity  $x_j$ , reversal  $E_j$ , conductance  $g_j^{\text{syn}} \geq 0$ ) and dendritic inputs from children (voltage  $V_j$ , conductance  $g_j^{\text{den}} \geq 0$ ). The steady-state voltage is:

$$V_n = \frac{\sum_j E_j x_j g_j^{\text{syn}} + \sum_j V_j g_j^{\text{den}}}{\underbrace{\sum_j x_j g_j^{\text{syn}} + \sum_j g_j^{\text{den}} + 1}_{g_n^{\text{tot}}}}, \quad R_n^{\text{tot}} = 1/g_n^{\text{tot}}. \quad (1)$$

$V_n$  is a convex combination of reversal potentials, child voltages, and leak. Writing  $\mathcal{S}_n = \{0\} \cup \{E_j\}_j \cup \{V_j\}_j$ , we have  $\min \mathcal{S}_n \leq V_n \leq \max \mathcal{S}_n$  and  $0 < R_n^{\text{tot}} \leq 1$ . The local sensitivities follow directly:

**Proposition 1** (Local Sensitivities).

$$\frac{\partial V_n}{\partial g_i^{\text{syn}}} = x_i R_n^{\text{tot}} (E_i - V_n), \quad \frac{\partial V_n}{\partial V_i} = g_i^{\text{den}} R_n^{\text{tot}}, \quad \frac{\partial V_n}{\partial g_i^{\text{den}}} = R_n^{\text{tot}} (V_i - V_n). \quad (2)$$

Eq. (2) gives the eligibility factors used by the local rules below: presynaptic activity or voltage difference, driving force, and input resistance.

## 2.2 Shunting Inhibition as Divisive Gain Control

An inhibitory synapse with  $E_{\text{inh}} \approx 0$  contributes current  $(0 - V_n)x_j g_j^{\text{syn}}$  and increases  $g_n^{\text{tot}}$ . Its sensitivity is  $\partial V_n / \partial g_j^{\text{syn}} = -x_j R_n^{\text{tot}} V_n$ , which corresponds to multiplicative attenuation (divisive normalization). Shunting is divisive at the voltage level, but its effect on firing rates can be subtractive in some regimes [7]. Inhibitory plasticity can balance excitation dynamically [8]; our learned inhibitory conductances serve an analogous role. The same denominator also gives the direct connection to path gain. Writing  $G_n^E(x)$  and  $G_n^I(x)$  for the total excitatory and inhibitory synaptic conductance impinging on compartment  $n$ ,

$$R_n^{\text{tot}}(x) = (1 + G_n^E(x) + G_n^I(x) + G_n^{\text{den}})^{-1}.$$

Increasing  $G_n^I(x)$  lowers  $R_n^{\text{tot}}(x)$  and therefore multiplicatively suppresses every upstream path whose error must pass through compartment  $n$ . In this sense inhibition can gate credit flow by changing the path gain, even when the inhibitory drive is feedforward rather than lateral. Prop. 2 formalizes this statement after the tree path gain is defined.

## 2.3 Exact Gradients for Dendritic Trees

Let  $V_0$  be the somatic output and define the soma/core error as  $\delta_0 := \partial L / \partial V_0$ . For a linear decoder  $\hat{y} = W_{\text{dec}} V_0$ , this reduces to  $\delta_0 = W_{\text{dec}}^\top (\partial L / \partial \hat{y})$ ; for a nonlinear decoder,  $\delta_0$  is the corresponding decoder-input Jacobian product.

**Theorem 1** (Conductance-Stage Backpropagation on a Dendritic Tree). *For a rooted dendritic tree with soma at node 0 and unique parent  $p(n)$  for each non-somatic compartment, suppose the parent receives either the pre-reactivation voltage  $V_n$  or a local post-voltage activity  $a_n = f_n(V_n)$ . For identity upward transfer, the loss gradient satisfies*

$$\frac{\partial L}{\partial V_n} = \frac{\partial L}{\partial V_{p(n)}} R_{p(n)}^{\text{tot}} g_{n \rightarrow p(n)}^{\text{den}}, \quad (3)$$

with boundary condition  $\partial L / \partial V_0 = \delta_0$ . With post-voltage transfer  $a_n = f_n(V_n)$ , the local factor becomes  $f'_n(V_n) R_{p(n)}^{\text{tot}} g_{n \rightarrow p(n)}^{\text{den}}$ . Because the graph is a tree, each compartment has a unique path to the soma. Defining the conductance-stage path gain

$$\alpha_n = \prod_{(i \rightarrow k) \in \text{path}(n \rightsquigarrow 0)} R_k^{\text{tot}} g_{i \rightarrow k}^{\text{den}}, \quad \frac{\partial L}{\partial V_n} = \alpha_n \delta_0, \quad (4)$$

with  $\alpha_0 = 1$  under identity transfer. In the post-voltage case, the exact effective path gain is obtained by multiplying the same product by the local factors  $f'_i(V_i)$  on the transmitted child activities along the path.

135 *Proof.* Apply the chain rule on the tree-structured computation graph using Prop. 1.  $\square$

136 The one-step recursion in Eq. (3) and the path product in Eq. (4) separate local conductance transfer  
137 from the non-local somatic error.

138 **Proposition 2** (Inhibitory control of conductance-stage path gain). *Let  $\mathcal{A}(n)$  be the set of com-*  
139 *partments on the path from compartment  $n$  to the soma, excluding  $n$  and including the parent*  
140 *compartments through which the error must pass. Holding dendritic coupling conductances fixed, an*  
141 *added inhibitory conductance  $\Delta G_k^I(x) \geq 0$  at any  $k \in \mathcal{A}(n)$  changes the conductance-stage path*  
142 *gain by*

$$\frac{\alpha_n(x; \Delta G^I)}{\alpha_n(x; 0)} = \prod_{k \in \mathcal{A}(n)} \frac{g_k^{\text{tot}}(x)}{g_k^{\text{tot}}(x) + \Delta G_k^I(x)}, \quad \frac{\partial \log \alpha_n}{\partial G_k^I} = \begin{cases} -R_k^{\text{tot}}, & k \in \mathcal{A}(n), \\ 0, & k \notin \mathcal{A}(n). \end{cases} \quad (5)$$

143 *Proof.* Substitute  $R_k^{\text{tot}} = 1/g_k^{\text{tot}}$  into the path-gain product in Eq. (4). Adding inhibitory conductance  
144 changes only the denominator at compartments on the path from  $n$  to the soma, which gives the  
145 product ratio and the log derivative in Eq. (5).  $\square$

146 Prop. 2 makes two directed predictions. First, inhibition is a path gate: it attenuates all credit paths  
147 below the inhibited compartment. Second, inhibition is useful for rank-1 local learning only when this  
148 attenuation makes the distribution of  $\alpha_n$  more uniform or better aligned with the task. If inhibition  
149 is too strong, too heterogeneous, or generated by a poorly calibrated inhibitory population, it can  
150 suppress useful signals and hurt learning. In the additive control the same I input changes voltages  
151 and local eligibilities, but it does not enter the conductance-stage path recursion as a multiplicative  
152 resistance gate.

153 **When does shunting compress path gains?** The attenuation identity above is not by itself a  
154 concentration result. Let  $z_n = \log \alpha_n$  be the log path gain for compartment  $n$ . If added inhibitory  
155 conductance contributes a path-dependent attenuation

$$\beta_n(x) = \sum_{k \in \mathcal{A}(n)} \log \left( 1 + \frac{\Delta G_k^I(x)}{g_k^{\text{tot}}(x)} \right), \quad z'_n = z_n - \beta_n, \quad (6)$$

156 then, over compartments or examples,

$$\text{Var}(z') = \text{Var}(z) + \text{Var}(\beta) - 2 \text{Cov}(z, \beta). \quad (7)$$

157 Thus shunting narrows the log path-gain field exactly when

$$\text{Cov}(z, \beta) > \frac{1}{2} \text{Var}(\beta). \quad (8)$$

158 Eq. (6) follows by taking the logarithm of the product ratio in Prop. 2; Eq. (7) then applies the variance  
159 identity for  $z - \beta$ . Eq. (8) is an interpretive sufficient condition under the same averaging measure,  
160 not a headline empirical diagnostic. In finite trained networks, the covariance margin can disagree  
161 with CV and exact-error-rank summaries because these statistics average over different compartment  
162 and sample axes. We therefore use direct path-gain dispersion, exact-error compressibility, and  
163 inhibition-intervention measurements as the empirical tests, and report the covariance margin as a  
164 boundary check in the appendix.

165 **Corollary 1** (Local–Global Factorization). *The exact synaptic gradient at compartment  $n$  factorizes*  
166 *as:*

$$\frac{\partial L}{\partial g_i^{\text{syn}}} = \underbrace{x_i R_n^{\text{tot}} (E_i - V_n)}_{\text{synapse-local eligibility}} \cdot \underbrace{\frac{\partial L}{\partial V_n}}_{\text{compartment error}}, \quad (9)$$

167 where the eligibility term depends only on quantities available at synapse  $i$  (presynaptic activity  
168  $x_i$ , input resistance  $R_n^{\text{tot}}$ , driving force  $E_i - V_n$ ), and the compartment error  $\partial L / \partial V_n$  is the sole  
169 non-local quantity.

The factorization in Eq. (9) isolates the only non-local quantity as  $\partial L/\partial V_n = \alpha_n \delta_0$  under identity upward transfer. With post-voltage transfer, the activation-derivative factors in Theorem 1 enter the effective path gain, but the compartment error remains the only non-local term. Any practical local rule therefore depends entirely on how well a broadcast signal approximates that exact compartment error, which we evaluate in Sec. 4.2. In the models studied here, shunting improves this approximation by reducing and stabilizing  $R_n^{\text{tot}}$ , which narrows the spread of local sensitivities across the tree. The theorem is stated for the pre-activation compartment voltage  $V_n$  produced by the conductance stage itself. In the empirical networks, each dendritic branch layer may optionally apply a post-voltage reactivation  $f(V_n)$  before passing activity to the next stage; unless otherwise noted, the manuscript sweeps use a learnable bounded `param_tanh` reactivation. Our exact-error diagnostics are computed on the hooked pre-activation voltage, so the factorization check and compartment-error analyses target the conductance-stage quantity defined above rather than the post-activation activity.

### 3 Local Learning Rules

#### 3.1 Broadcast Error Approximation

We approximate the exact compartment error  $\partial L/\partial V_n$  (Corollary 1) with a broadcast signal  $e_n$  derived from the somatic/core error  $\delta_0$ . For a linear readout this reduces to  $\delta_0 = W_{\text{dec}}^\top (\partial L/\partial \hat{y})$ ; for a nonlinear decoder we use the corresponding decoder-input Jacobian product  $J_{\text{dec}}(h_{\text{core}})^\top (\partial L/\partial \hat{y})$  so that the broadcast remains defined in the same core-space coordinates. We distinguish five feedback objects:

- **Rank-1 shared broadcast:**  $e_n = \bar{\delta} \cdot \mathbf{1}$ , where  $\bar{\delta} = \text{mean}(\delta_0)$  is one scalar field delivered to all compartments.
- **Neuron-wise broadcast:**  $e_n = \delta_0$  when the layer width matches the core/output error dimension. In our headline feedforward settings this condition usually fails, so the default historical “per-soma” implementation reduces to rank-1 shared broadcast.
- **Rank- $K$  random broadcast:**  $c = P_K \delta_0 \in \mathbb{R}^K$  and  $e_n = \sum_{k=1}^K q_{n,k} c_k$ , where  $P_K$  is a fixed random projection and  $q_{n,k}$  mixes those  $K$  channels into compartment  $n$ .
- **Path-structured broadcast:**  $e_n = \lambda_{\text{mix}} \bar{\delta} \cdot \mathbf{1} + (1 - \lambda_{\text{mix}}) \Gamma_n(\delta_0)$ , where  $\Gamma_n$  projects the somatic error into router-inferred pathway roles, gates those roles by local pathway activity, and transports the resulting vector hierarchically through the block-structured dendritic tree.
- **Transported oracle:**  $e_n^{\text{transport}} = \alpha_n \delta_0 = \partial L/\partial V_n$ . This is an analysis upper bound, not a proposed local rule, because it supplies the exact path gain  $\alpha_n$  from Theorem 1.

We use the historical `per_soma` implementation for most experiments, but in the main feedforward architectures it is effectively rank-1 because hidden widths exceed the decoder error dimension (Appendix Table S18). Thus the headline rank-1 result should be read as the architecture-induced fallback of the default per-soma mode, not as a claim that neuron-wise broadcasts were unavailable in principle. Under shunting, that deliberately low-bandwidth field is often sufficient on standard classification tasks. Routed tasks and rank-bridge controls in the appendix test the opposite regime, where rank-1 broadcast collapses branch identity and richer feedback is needed. A local-mismatch heuristic is included only as an appendix control because it performs much worse than the somatic-error broadcasts (Appendix A).

#### 3.2 Three-Factor Rule (3F)

**Definition 1** (3F Update). *For synaptic and dendritic conductances:*

$$\Delta g_j^{\text{syn}} = \eta \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n \rangle_B, \quad \Delta g_j^{\text{den}} = \eta \langle R_n^{\text{tot}} (V_j - V_n) e_n \rangle_B, \quad (10)$$

where  $\langle \cdot \rangle_B$  denotes the batch average.

The three factors are presynaptic activity  $x_j$  (or a voltage difference), postsynaptic modulation through driving force and input resistance, and the broadcast error  $e_n$ . The same rule applies to excitatory and inhibitory synapses. The sign difference comes from the driving force  $(E_j - V_n)$ .

216 **Additive control.** Rule (10) is the local gradient of the *shunting* voltage  $V_n = \sum E_j g_j x_j / g_n^{\text{tot}}$ . For  
 217 the additive control ( $V_n = \sum g_j x_j$ , no divisive normalization), the correct local gradient is simpler:  
 218  $\Delta g_j^{\text{syn}} = \eta \langle x_j e_n \rangle_B$ , with no driving-force or  $R_n^{\text{tot}}$  terms ( $R_n^{\text{tot}} \equiv 1$  by definition since there is no  
 219 denominator). Throughout, each architecture uses the learning rule derived from its own forward-pass  
 220 dynamics. This keeps the comparison architecture-matched rather than applying a shunting-derived  
 221 rule to an additive model.

### 222 3.3 Practical Heuristic Wrappers: 4F and 5F

223 **4F (instructive attenuation proxy).** In the conductance-stage recursion (4), the path gain  $\alpha_n$   
 224 attenuates the somatic error differently at each compartment. To compensate without explicitly  
 225 computing that gain, 4F uses an online correlation proxy

$$r_n^{4F} = \text{Cov}(\bar{V}_n, \bar{V}_0) / (\sqrt{\text{Var}(\bar{V}_n) \text{Var}(\bar{V}_0)} + \varepsilon).$$

226 High  $r_n^{4F}$  up-weights branches whose voltages covary with the soma; low  $r_n^{4F}$  down-weights branches  
 227 where the broadcast is likely to be a poor proxy. We keep 4F as an instructive intermediate that tests  
 228 whether a scalar branch-soma relevance statistic helps rank-1 broadcast. Empirically, this factor alone  
 229 gives little improvement over 3F in the completed competence sweeps; the practical jump comes  
 230 from the 5F confidence factor below.

231 **5F (confidence-weighted 4F).** 5F adds a second heuristic factor

$$\phi_n = \text{Var}(V_n) / (\sigma_{\text{res}}^2 + \varepsilon),$$

232 where  $\sigma_{\text{res}}^2$  is the residual variance of  $V_n$  after linear regression on the parent voltage. Operationally,  
 233  $\phi_n$  is a branch-wise confidence gate: it becomes large when local voltages are strong and stable  
 234 relative to unexplained residual fluctuations, and small when local fluctuations dominate. We clamp  
 235 it to  $[0.25, 4.0]$  for stability. This factor is empirical, not theorem-derived. The paper’s core claim  
 236 rests on the exact factorization, path-gain compressibility, and the transported-oracle experiment; the  
 237 gain from  $\phi_n$  should be read as practical denoising rather than theoretical necessity.

**Definition 2** (5F Update).

$$\Delta g_j^{\text{syn}} = \eta r_n^{4F} \phi_n \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n \rangle_B. \quad (11)$$

238 Eq. (11) is therefore the headline practical LocalCA update, while the 3F rule is the theorem-facing  
 239 local gradient and 4F is the intermediate ablation. A feedback-alignment-style random-broadcast  
 240 argument is included in Appendix C; it is supportive but not central to the main mechanism claim.

### 241 3.4 Structured pathway feedback and branch-role plasticity

242 For tasks with latent pathway structure, scalar broadcast can erase branch identity. Motivated  
 243 by vectorized dendritic teaching signals [43], PV-LocalCA projects the somatic error into a low-  
 244 dimensional pathway basis and gates it by local branch roles,  $e_n^{\text{PV}} = \sum_{k=1}^K q_{n,k} c_k$ . This is a  
 245 structured rank- $K$  broadcast, not a separate learning principle. Current cue-routing results show that  
 246 moving beyond rank-1 helps, but the best structured pathway-vector implementation is not yet better  
 247 than the strongest unstructured low-rank control; we therefore treat it as a failure-mode diagnostic  
 248 and design direction rather than a solved headline mechanism.

## 249 4 Experiments

### 250 4.1 Setup

251 We evaluate capacity-calibrated supervised tasks and controlled sweeps that isolate inhibition strength,  
 252 rule family, broadcast mode, morphology, and architecture. The main text uses MNIST, Fashion-  
 253 MNIST [37], context gating, noise resilience, and a compressed morphology-regime summary;  
 254 CIFAR-10 and routed-feedback diagnostics are reported in the appendix. Main comparisons are  
 255 architecture-matched additive vs. shunting dendritic cores and local vs. backprop training; DFA is  
 256 a supplementary baseline. Unless otherwise stated, local runs use 5F with the historical `per_soma`  
 257 implementation, which reduces to rank-1 shared broadcast when hidden width exceeds the decoder




Figure 3: **Mechanistic chain from path gains to learning.** (A) Example dendritic path-gain field at matched inhibition, with measured MNIST CV annotated for additive and shunting networks. (B) Exact-error compressibility on MNIST: lower rank-1 residual and lower participation rank mean that a scalar broadcast is a better approximation. (C) Post-training inhibitory-conductance interventions (L: learned; 0/sh/m/u: zero, shuffled, mean-clamped, uniform matched); removing learned structure collapses performance. (D) Direct compartment-error fidelity on noise resilience. Solid lines compare rank-1 broadcast to the exact compartment error; dashed lines show the transported-error oracle. (E) Learning with the transported-error oracle. Improving the broadcast field recovers near-ceiling local learning where rank-1 broadcast fails. Error bars denote  $\pm 1$  s.d. across seeds.

error dimension. All main-claim dendritic runs use nonnegative conductance transforms and non-negative first-layer synaptic drive; image inputs are in  $[0, 1]$ , corrupted inputs are clipped back to  $[0, 1]$ , and context gating uses a nonnegative contextual figure-ground MNIST construction. The transported-error condition is an oracle upper bound derived from Theorem 1, not the proposed biological rule. Main mechanistic, oracle, and competence results use 5 seeds per condition unless noted otherwise.

## 4.2 From Exact Factorization to Credit-Signal Fidelity

To test whether these performance differences correspond to better credit signals, we compare local and backprop gradients on the same batch at the same parameter values. For each parameter tensor  $p$ , we compute directional alignment (cosine similarity) and scale mismatch, defined as  $|\log_{10}(\|g_p^{\text{local}}\|/\|g_p^{\text{bp}}\|)|$ , aggregated by parameter count.

Across the checkpointed MNIST gradient-dynamics diagnostic, shunting is more aligned with backprop and closer in norm than the additive control at the final checkpoint: weighted cosine rises from  $0.005 \pm 0.026$  to  $0.100 \pm 0.033$ , and scale mismatch falls from  $0.50 \pm 0.05$  to  $0.26 \pm 0.05$ . Figure 2A confirms that the theory is operational in code under the nonnegative-input setup: reconstructing gradients from Corollary 1 using exact compartment errors yields weighted relative reconstruction error below  $2 \times 10^{-7}$  and weighted scale mismatch below  $4 \times 10^{-8}$  across runs. In a few low-norm parameter blocks, cosine becomes numerically unstable even though the reconstruction itself is essentially exact, so we report the norm-based diagnostics in the figure. Figure 2D rules out a trivial zero-gradient explanation: in the 5F trajectory diagnostic, additive local and backprop gradient norms remain finite throughout learning. The only approximation made by LocalCA is therefore the broadcast substitute for  $\partial L/\partial V_n$ .

## 4.3 Mechanistic Chain: Path Gains, Broadcast Fidelity, and Oracle Transport

Figure 3 provides the central empirical chain. At matched inhibition ( $N_I=5$ ), shunting narrows the MNIST conductance-stage path-gain field (CV 0.23 vs. 1.16 for additive) and makes the exact compartment-error matrix more compressible: the best rank-1 residual drops from  $0.83 \pm 0.07$  to  $0.65 \pm 0.18$ , while participation rank drops from  $8.5 \pm 3.3$  to  $3.2 \pm 1.2$ . This is the diagnostic matched to the paper’s main claim: rank-1 broadcast is useful when the exact error field is closer to low rank.

The post-training inhibition intervention shows that the learned inhibitory field is necessary for the trained computation and is not interchangeable with a matched conductance load: zeroing,




Figure 4: **Capacity-calibrated competence and regime dependence.** (A) Matched backprop ceilings vs. rank-1 5F LocalCA on MNIST, Fashion-MNIST, and context gating, including the matched additive context-gating control. (B) Inhibitory dose-response on MNIST and noise resilience. (C) Morphology-dependent shunting advantage on noise resilience from a supportive 3-seed morphology sweep. (D) Rule-family ablation: 3F/4F are weak practical optimizers, while 5F is the empirical wrapper used for headline LocalCA. Error bars denote  $\pm 1$  s.d. across seeds where available.

shuffling, or clamping learned branch inhibition drops MNIST accuracy from about 90% to 20–26% and noise-resilience accuracy from about 81% to 10–19%. Thus inhibition carries sample-specific pathway structure rather than merely setting global scale. The fidelity and oracle panels test the same bottleneck directly: rank-1 broadcast tracks the exact compartment error better under shunting once inhibition is present (about 0.20–0.24 vs. 0.14–0.19 for additive), transported error reaches about 0.95–1.0 fidelity, and transported learning lifts both cores to 92–95% on noise resilience; MNIST transported-error values in the appendix remain in the 95–97% range. CIFAR-10 and direct-I versus explicit-I path-gain probes remain appendix diagnostics (Figs. S10, S14; Table S8).

#### 4.4 Capacity-Calibrated Competence

In the capacity-calibrated regime, shunting LocalCA stays within about 5–6 percentage points of matched backpropagation on MNIST, Fashion-MNIST, and context gating (Table S4; Fig. 4). Figure 4D keeps the rule-family ablation in the main text: 5F is the strongest practical wrapper but remains heuristic, while 3F/4F are instructive but weak optimizers. Broader rule and feedback-design controls are summarized in Appendix Fig. S1; matched DFA baselines remain much weaker than shunting LocalCA (Appendix Fig. S9).

#### 4.5 Regime Dependence of the Shunting Advantage

With matched rank-1 broadcast, shunting is modestly better on MNIST, Fashion-MNIST, and context gating, including  $0.803 \pm 0.006$  vs.  $0.787 \pm 0.007$  for additive on context gating. The gap grows on inhibition-sensitive noise resilience once inhibitory conductance is present (Fig. 4B). Morphology shifts the useful inhibitory operating regime: different tree depths and branch factors peak at different  $N_I$ , and the shunting advantage is not monotonic in inhibition (Fig. 4C).

### 5 Discussion and Limitations

The main result is a condition for faithful dendritic broadcast, not a claim that LocalCA universally replaces backpropagation. Exact conductance-tree credit separates a synapse-local eligibility term from one non-local compartment error, so local learning succeeds or fails according to how faithfully a low-dimensional broadcast approximates that exact path-specific error field.

The experiments support a coherent mechanism. Shunting inhibition changes the conductance denominator and thereby modulates input resistance along credit paths. In the regimes studied here, this makes the exact error field more compressible for rank-1 broadcast, improves local-backprop alignment, and reduces scale mismatch. Transported error closes much of the local-learning gap, while post-training inhibition interventions show that learned inhibitory conductance carries sample-

specific pathway structure rather than merely adding a matched conductance load. Morphology and inhibition interact, so more inhibition is not the mechanism; different trees shift the operating regime in which shunting is useful.

The boundary conditions are equally important. The practical 5F rule is an empirical stabilizer rather than a theorem-derived necessity, Eq. (8) is interpretive rather than a positive headline diagnostic, and activation-shape audits show that the mechanism depends on the learned bounded reactivation regime used in the main experiments. Rank-1 broadcast remains insufficient in routed and harder-data settings, where higher-rank or transported feedback is needed. The present models use input-driven shunting inhibition and a feedforward explicit inhibitory-cell probe, not recurrent lateral inhibition. More broadly, the results connect a biophysical operation, shunting control of input resistance, to a machine-learning object: the low-rank structure of the error field.

## References

- [1] Koch, C. (1999). *Biophysics of Computation*. Oxford University Press.
- [2] Dayan, P., & Abbott, L. F. (2001). *Theoretical Neuroscience*. MIT Press.
- [3] Poirazi, P., Brannon, T., & Mel, B. W. (2003). Pyramidal neuron as two-layer neural network. *Neuron*, 37(6), 989–999.
- [4] London, M., & Häusser, M. (2005). Dendritic computation. *Annual Review of Neuroscience*, 28, 503–532.
- [5] Urbanczik, R., & Senn, W. (2014). Learning by the dendritic prediction of somatic spiking. *Neuron*, 81(3), 521–528.
- [6] Carandini, M., & Heeger, D. J. (2012). Normalization as a canonical neural computation. *Nature Reviews Neuroscience*, 13(1), 51–62.
- [7] Holt, G. R., & Koch, C. (1997). Shunting inhibition does not have a divisive effect on firing rates. *Neural Computation*, 9(5), 1001–1013.
- [8] Vogels, T. P., Sprekeler, H., Zenke, F., Clopath, C., & Gerstner, W. (2011). Inhibitory plasticity balances excitation and inhibition in sensory pathways and memory networks. *Science*, 334(6062), 1569–1573.
- [9] Lillicrap, T. P., Cownden, D., Tweed, D. B., & Akerman, C. J. (2016). Random synaptic feedback weights support error backpropagation for deep learning. *Nature Communications*, 7, 13276.
- [10] Nøkland, A. (2016). Direct feedback alignment provides learning in deep neural networks. *NeurIPS*, 29.
- [11] Guerguiev, J., Lillicrap, T. P., & Richards, B. A. (2017). Towards deep learning with segregated dendrites. *eLife*, 6, e22901.
- [12] Sacramento, J., Costa, R. P., Bengio, Y., & Senn, W. (2018). Dendritic cortical microcircuits approximate the backpropagation algorithm. *NeurIPS*, 31.
- [13] Whittington, J. C., & Bogacz, R. (2019). Theories of error back-propagation in the brain. *Trends in Cognitive Sciences*, 23(3), 235–250.
- [14] Richards, B. A., & Lillicrap, T. P. (2019). Dendritic solutions to the credit assignment problem. *Current Opinion in Neurobiology*, 54, 28–36.
- [15] Scellier, B., & Bengio, Y. (2017). Equilibrium propagation. *Frontiers in Computational Neuroscience*, 11, 24.
- [16] Song, Y., Millidge, B., Salvatori, T., Lukasiewicz, T., Xu, Z., & Bogacz, R. (2024). Inferring neural activity before plasticity as a foundation for learning beyond backpropagation. *Nature Neuroscience*, 27, 348–358.

- [17] Gretton, A., Bousquet, O., Smola, A. J., & Schölkopf, B. (2005). Measuring statistical dependence with Hilbert-Schmidt norms. *Algorithmic Learning Theory*, 63–77.
- [18] Frémaux, N., & Gerstner, W. (2016). Neuromodulated spike-timing-dependent plasticity, and theory of three-factor learning rules. *Frontiers in Neural Circuits*, 9, 85.
- [19] Bellec, G., et al. (2020). A solution to the learning dilemma for recurrent networks of spiking neurons. *Nature Communications*, 11, 3625.
- [20] Welford, B. P. (1962). Note on a method for calculating corrected sums of squares and products. *Technometrics*, 4(3), 419–420.
- [21] Turrigiano, G. G. (2008). The self-tuning neuron: synaptic scaling of excitatory synapses. *Cell*, 135(3), 422–435.
- [22] Payeur, A., Guerguiev, J., Zenke, F., Richards, B. A., & Naud, R. (2021). Burst-dependent synaptic plasticity can coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24(7), 1010–1019.
- [23] Greedy, W., Zhu, H. W., Pemberton, J., Mellor, J., & Ponte Costa, R. (2022). Single-phase deep learning in cortico-cortical networks. *NeurIPS*, 35, 24213–24225.
- [24] Haider, P., Ellenberger, B., Kriener, L., Jordan, J., Senn, W., & Petrovici, M. A. (2021). Latent Equilibrium: A unified learning theory for arbitrarily fast computation with arbitrarily slow neurons. *NeurIPS*, 34.
- [25] Hinton, G. (2022). The forward-forward algorithm. *arXiv:2212.13345*.
- [26] Dellaferrera, G., & Kreiman, G. (2022). Error-driven input modulation: Solving the credit assignment problem without a backward pass. *Proceedings of the 39th International Conference on Machine Learning, Proceedings of Machine Learning Research*, 162, 4937–4955.
- [27] Lee, D.-H., et al. (2015). Difference target propagation. *ECML*, 498–515.
- [28] Meulemans, A., et al. (2021). Credit assignment in neural networks through deep feedback control. *NeurIPS*, 34.
- [29] Millidge, B., Seth, A. K., & Buckley, C. L. (2021). Predictive coding: A theoretical and experimental review. *arXiv:2107.12979*.
- [30] Koch, C., Poggio, T., & Torre, V. (1983). Nonlinear interactions in a dendritic tree. *PNAS*, 80(9), 2799–2802.
- [31] Silver, R. A. (2010). Neuronal arithmetic. *Nature Reviews Neuroscience*, 11(7), 474–489.
- [32] Max, K., Kriener, L., Pineda García, G. et al. (2024). Learning efficient backprojections across cortical hierarchies in real time. *Nature Machine Intelligence*, 6(6), 619–630.
- [33] Ma, C., Wu, J., Si, C., & Tan, K. C. (2024). Scaling supervised local learning with augmented auxiliary networks. *International Conference on Learning Representations (ICLR)*.
- [34] Lv, C., Xu, J., Lu, Y., Wang, X., Wang, Z., Xu, Z., Yu, D., Du, X., Zheng, X., & Huang, X. (2025). Dendritic Localized Learning: Toward Biologically Plausible Algorithm. *Proceedings of the 42nd International Conference on Machine Learning, Proceedings of Machine Learning Research*, 267, 41682–41700.
- [35] Erdogan, M., Pehlevan, C., & Erdogan, A. T. (2025). Error broadcast and decorrelation as a potential artificial and natural learning mechanism. *OpenReview*, NeurIPS 2025 spotlight.
- [36] Kao, C.-H., & Hariharan, B. (2024). Counter-Current Learning: A Biologically Plausible Dual Network Approach for Deep Learning. *Advances in Neural Information Processing Systems*, 37.
- [37] Xiao, H., Rasul, K., & Vollgraf, R. (2017). Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms. *arXiv:1708.07747*.

- 409 [38] Song, S., Sjöström, P. J., Reigl, M., Nelson, S., & Chklovskii, D. B. (2005). Highly nonrandom  
410 features of synaptic connectivity in local cortical circuits. *PLoS Biology*, 3(3), e68.
- 411 [39] Buzsáki, G., & Mizuseki, K. (2014). The log-dynamic brain: How skewed distributions affect  
412 network operations. *Nature Reviews Neuroscience*, 15(4), 264–278.
- 413 [40] The MICrONS Consortium (2025). Functional connectomics spanning multiple areas of mouse  
414 visual cortex. *Nature*, 640(8058), 435–447.
- 415 [41] Loewenstein, Y., Kuras, A., & Rumpel, S. (2011). Multiplicative Dynamics Underlie the  
416 Emergence of the Log-Normal Distribution of Spine Sizes in the Neocortex In Vivo. *Journal of*  
417 *Neuroscience*, 31(26), 9481–9488.
- 418 [42] Megias, M., et al. (2001). Total number and distribution of inhibitory and excitatory synapses  
419 on hippocampal CA1 pyramidal cells. *Neuroscience*, 102(3), 527–540.
- 420 [43] Francioni, V., Tang, V. D., Toloza, E. H. S. et al. (2026). Vectorized instructive signals in cortical  
421 dendrites. *Nature*. <https://doi.org/10.1038/s41586-026-10190-7>
- 422 [44] Iyer, A., Grewal, K., Velu, A., Souza, L. O., Forest, J., & Ahmad, S. (2022). Avoiding catas-  
423 trophe: Active dendrites enable multi-task learning in dynamic environments. *Frontiers in*  
424 *Neurorobotics*, 16, 846219.

## 425 A Supplementary Results

426 **Appendix overview.** The appendix records the empirical division of labor behind the main claims.  
427 It starts with rule-family, broadcast-mode, and component summaries; then reports gradient dynamics,  
428 CIFAR-10, inhibition/pathway diagnostics, rank-bridge controls, biological plausibility assumptions,  
429 implementation and asset details, and theory extensions. Table S2 is the quickest map of which  
430 components carry each result.

431 **Boundary diagnostics.** Several controls are intentionally negative or mixed. The log-gain co-  
432 variance margin from Eq. (8) is not a positive headline result, Double-CIFAR remains near chance  
433 even under matched backpropagation, positive cue-routing controls saturate too strongly to separate  
434 mechanisms, and the current pathway-vector implementation does not beat the best unstructured  
435 low-rank cue-routing control. We include these results to delimit the claim rather than to smooth over  
436 failure modes.

### 437 Mechanistic Support and Broadcast Controls

#### 438 Rule-Family Ranking

#### 439 Which Components Carry?




Figure S1: **Local-rule and feedback-design axes.** (A) The theorem-facing 3F update and the 4F attenuation proxy are weak practical optimizers in the completed local-competence sweeps; 5F is the strongest empirical wrapper. (B) The broadcast definition matters: on MNIST, rank-1 shared somatic error is strong, while a local-mismatch error surrogate fails. (C) On noise resilience, recursively propagating or increasing the rank of the error field closes much of the gap between rank-1 broadcast and exact path transport. (D) Cue routing probes a branch-specific regime: higher rank helps, while the current pathway-vector implementation is useful but not yet better than the best unstructured low-rank control. Error bars denote  $\pm 1$  s.d. where five-seed summaries are available.

| Dataset        | Rule | Top-10 valid | Top-10 test  |
|----------------|------|--------------|--------------|
| MNIST          | 3F   | 0.611        | 0.622        |
| MNIST          | 4F   | 0.620        | 0.628        |
| MNIST          | 5F   | <b>0.912</b> | <b>0.916</b> |
| Context gating | 3F   | 0.398        | 0.396        |
| Context gating | 4F   | 0.411        | 0.411        |
| Context gating | 5F   | <b>0.807</b> | <b>0.789</b> |

Table S1: **Rule-family ranking.** Top-10 mean across completed local-competence sweeps.

| Component                      | Representative effect                                                                                                                                                                                                                                                                                                                 | Interpretation                                                                                                                                                                                                                                                                                                                               |
|--------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Shunting architecture          | In the nonnegative-input noise-resilience sweep, additive performs better at $N_I=0$ (0.668 vs. 0.543), but shunting overtakes it once inhibition is present (0.837 vs. 0.818 at $N_I=10$ , 0.861 vs. 0.830 at $N_I=40$ ).                                                                                                            | The shunting gain is operating-regime dependent rather than uniform: it appears once inhibition shapes conductance scaling strongly enough to help the shared broadcast.                                                                                                                                                                     |
| Broadcast quality              | On MNIST (5F), shunting with the historical <code>per_soma</code> /rank-1 fallback reaches $0.912 \pm 0.005$ , while local-mismatch remains at $0.146 \pm 0.046$ .                                                                                                                                                                    | Correctly matching the broadcast field to the somatic error matters far more than arbitrary local mismatch heuristics.                                                                                                                                                                                                                       |
| Transported oracle             | Across inhibited noise-resilience regimes, additive local learning rises from roughly 81–84% under rank-1 fallback to 92–94% under transported broadcast; shunting rises from roughly 81–86% to 95–95.2%.                                                                                                                             | The oracle remains the strongest direct evidence that broadcast fidelity, not just architecture, is the dominant bottleneck.                                                                                                                                                                                                                 |
| Harder-data<br>CIFAR-10 ladder | In the nonnegative-input CIFAR-10 control ladder, shunting backprop reaches $49.5 \pm 0.4\%$ vs. $48.3 \pm 1.0\%$ for additive; shunting transported LocalCA reaches $47.2 \pm 0.6\%$ , while shunting rank-1 reaches $29.7 \pm 1.3\%$ . Removing learned I-to-E conductance drops shunting transported LocalCA to $38.6 \pm 1.0\%$ . | CIFAR-10 is not used as a benchmark claim, but it is a useful non-saturated stress test: the stronger result requires both shunting inhibition and a high-fidelity transported credit signal.                                                                                                                                                |
| Rule hierarchy                 | In the competence sweeps, MNIST improves from 0.622 (3F) and 0.628 (4F) to 0.916 (5F); context gating improves from 0.396/0.411 to 0.789.                                                                                                                                                                                             | 5F is the strongest practical rule in this paper, though its extra factor remains heuristic rather than theorem-derived.                                                                                                                                                                                                                     |
| 5F stability                   | Tightening the $\phi$ clamp to $[0.5, 2.0]$ changes MNIST by only about +0.4 pp; varying EMA from 0.05 to 0.20 changes it by less than 0.2 pp.                                                                                                                                                                                        | The practical gain from 5F is not especially brittle to its main tuning knobs.                                                                                                                                                                                                                                                               |
| Structured feedback            | In the cue-routing rank sweep, learned-routing additive local learning remains near ceiling while learned-routing shunting rises from 0.742 at random rank-1 to 0.846 at random rank-2; the current tuned pathway-vector variant reaches 0.815.                                                                                       | Routed tasks remain the main failure mode of rank-1 feedback, but the robust conclusion is that higher rank helps; the best structured fix is not yet settled enough to anchor a headline claim.                                                                                                                                             |
| Random low-rank channels       | On the shunting noise-resilience bridge, random low-rank broadcast improves from 0.453 at $K=1$ to 0.764 at $K=4$ and 0.795 at $K=8$ , while exact path transport reaches 0.847; on cue routing, the best unstructured low-rank setting is $K=2$ at 0.846.                                                                            | Higher-rank unstructured broadcast now closes a large fraction of the rank-1 gap, but it still does not recover the transported-oracle ceiling and does not by itself establish a structure-specific advantage.                                                                                                                              |
| Path propagation proxy         | In the shunting noise-resilience bridge, the current morphology-aware path-propagation proxy reaches $0.762 \pm 0.047$ , well above the rank-1 fallback ( $0.662 \pm 0.016$ ) and close to random low-rank $K=4$ ( $0.764 \pm 0.025$ ), though still below exact path transport ( $0.847 \pm 0.018$ ).                                | The present recursive attenuation proxy is a meaningful constructive baseline: it closes much of the oracle gap, though a stronger strictly local approximation to transported error remains open.                                                                                                                                           |
| Morphology regime              | In the exploratory noise-resilience regime map (Fig. S13), morphology shifts the useful inhibitory operating regime; the shunting advantage is non-monotonic and depends on tree depth and branch factor.                                                                                                                             | Conductance normalization and tree geometry interact rather than contributing independently, so the mechanism is not equivalent to simply adding more inhibition.                                                                                                                                                                            |
| Activation choice              | In the five-seed frozen activation audit, identity outperforms frozen <code>param_tanh</code> for LocalCA on noise resilience in both architectures (0.590/0.859 vs. 0.109/0.772 for additive at $N_I=0/10$ ; 0.411/0.769 vs. 0.331/0.701 for shunting), while the backprop effect remains strongly regime-dependent.                 | The post-voltage firing-rate transform is a real modeling variable, not just an implementation detail. Even a fixed saturating nonlinearity changes the operating point substantially, so the main paper’s mechanistic claims should be read in the learned bounded <code>param_tanh</code> regime used throughout the headline experiments. |

Table S2: **Which components carry the observed gains?** A compact summary of the empirical division of labor across architecture, broadcast design, and LocalCA rule variants.

| Core     | Broadcast      | Decoder  | Test (mean $\pm$ std) |
|----------|----------------|----------|-----------------------|
| Shunting | per-soma       | local    | 0.912 $\pm$ 0.005     |
| Shunting | per-soma       | backprop | 0.909 $\pm$ 0.008     |
| Shunting | local-mismatch | local    | 0.146 $\pm$ 0.046     |
| Shunting | local-mismatch | backprop | 0.146 $\pm$ 0.037     |
| Additive | per-soma       | local    | 0.894 $\pm$ 0.007     |
| Additive | per-soma       | backprop | 0.900 $\pm$ 0.001     |
| Additive | local-mismatch | local    | 0.342 $\pm$ 0.058     |
| Additive | local-mismatch | backprop | 0.348 $\pm$ 0.095     |

Table S3: **Local-mismatch recheck (MNIST, 5F)**. The historical per\_soma mode is consistently strong in this setting, where it functions as the rank-1 fallback described in Table S18; local-mismatch remains much weaker.




Figure S2: **Capacity calibration (supplementary)**. (A) Phase 1 backprop ceilings across all architectures and datasets. (B) Rule-family ranking: 5F consistently outperforms 4F and 3F. (C) Decoder mode comparison (local vs. backprop decoder, 5F MNIST). (D) Broadcast mode comparison: the historical per\_soma mode/rank-1 fallback is required for strong performance; local-mismatch fails.

#### 440 Broadcast Mode Comparison and Local-Mismatch Recheck

#### 441 Competence, Verification, and Stress Tests

#### 442 Phase 1 Capacity Ceilings

#### 443 Local Competence and Regime Dependence

#### 444 Ablation Results

#### 445 Extended Gradient Analysis and $N_I$ Sweep Detail

446 **Alignment dynamics are not a zero-gradient artifact.** To check whether low additive alignment  
447 could be caused by vanishing gradients, we analyzed a separate 3-seed checkpoint trajectory with  
448 parameter-level LocalCA and backprop gradients. The additive local gradients are nonzero throughout  
449 training (weighted norm range  $9.5 \times 10^{-4}$ – $9.5 \times 10^{-3}$ ), and the corresponding backprop gradients are

| Dataset        | BP ceiling | Shunting local    | Additive local    | BP gap |
|----------------|------------|-------------------|-------------------|--------|
| MNIST          | 0.971      | $0.912 \pm 0.005$ | $0.894 \pm 0.007$ | 5.9%   |
| Fashion-MNIST  | 0.889      | $0.838 \pm 0.003$ | $0.831 \pm 0.004$ | 5.1%   |
| Context gating | 0.861      | $0.803 \pm 0.006$ | $0.787 \pm 0.007$ | 5.8%   |

Table S4: **Local competence.** Backprop ceilings are matched-capacity shunting references from dedicated 5-seed standard-training refreshes for MNIST and context gating and from the dedicated 5-seed Fashion-MNIST competence sweep. Local values are 5F with rank-1 broadcast and local decoder updates. Context gating additionally uses an HSIC auxiliary objective (weight 0.01; Appendix E). Error bars on local values:  $\pm 1$  s.d. across 5 seeds.

| Condition                                               | Metric           | Value                   |
|---------------------------------------------------------|------------------|-------------------------|
| Decoder: local vs backprop vs frozen (MNIST, toy model) | test acc         | 0.379 vs 0.379 vs 0.176 |
| Shunting vs additive, matched (MNIST, toy model)        | $\Delta$ test    | +0.210                  |
| Per-soma, path on vs off (synthetic task, toy model)    | test / MI(E,I;C) | -0.003 / +0.017         |

Table S5: **Mechanistic ablations.** Results from controlled small-network experiments using minimal architectures to isolate individual factors. These runs use a small 2-layer dendritic network (branch factor [9]) trained for 30 epochs; they are not directly comparable to the headline capacity-calibrated results.

also nonzero ( $2.3 \times 10^{-5}$ – $5.7 \times 10^{-3}$ ). The low additive cosine therefore reflects directional and scale mismatch, not absent gradients: additive 5F cosine stays near zero from epoch 0 to epoch 50 (0.018 to 0.005), while its local/backprop norm ratio falls from about  $8.8 \times 10^2$  to 0.69.

#### Verification and Reproducibility

#### Additional Stress Tests

#### Broadcast Bandwidth (Quantization) Control

#### FA/DFA Baseline Comparison

#### CIFAR-10 Results

Table S6 provides the exact values for the appendix CIFAR-10 figure. In this stronger compact direct-I-stream family, shunting performs well under standard training (49.5% vs. 48.3% for additive), so the harder-data stress test does not unfairly penalize the shunting architecture. The broadcast-fidelity ladder also survives once decoder-aware soma mapping is used: additive LocalCA reaches 25.5% under rank-1 per-soma broadcast and 32.4% under transported error, while shunting rises from 29.7% under rank-1 broadcast to 35.0% with random low-rank  $K=4$  and 47.2% with transported error, nearly matching the 49.5% shunting backprop ceiling. Removing learned I-to-E conductance lowers the shunting ceiling by 5.7 pp and transported LocalCA by 8.6 pp, making this the clearest harder-data evidence that inhibitory conductance is part of the useful credit geometry rather than just an added architectural detail.

A routed Double-CIFAR follow-up with strictly nonnegative router outputs was also not included as positive evidence. The task remained near chance even under matched backprop: additive BP reached  $10.9 \pm 0.7\%$ , shunting BP with learned I-to-E reached  $11.6 \pm 0.4\%$ , and the strongest LocalCA condition, shunting transported error with learned I-to-E, reached only  $11.7 \pm 0.3\%$  (Fig. S15D). We therefore treat this as a failed hard routed-task screen and keep the stronger compact CIFAR-10 control ladder as the paper’s harder-data diagnostic.

**CIFAR-10 soma-on extension.** We also tested the soma-on extension on the same compact direct-I-stream depth-4 family with quantile init (Fig. S11; Table S7; 3 seeds). With `somatic_synapses=true`, path-transported LocalCA stays strong in both cores: additive reaches  $45.5 \pm 0.6\%$  and shunting reaches  $46.6 \pm 1.0\%$ . Relative to the soma-off control ladder, the additive path-transport condition improves substantially, while shunting remains within 1 pp of its soma-off path-transport value (47.2%) and within 3 pp of the matched shunting backprop ceiling (49.5%).




Figure S3: **Extended gradient and  $N_I$  analysis (supplementary).** (A) Scale mismatch bars from a separate matched-weight static diagnostic; these values should not be compared directly to the checkpointed final-state diagnostic in Fig. 2. Shunting achieves near-ideal scale (0.117); additive exhibits order-of-magnitude distortion ( $> 1.0$ ). (B) Noise-resilience  $N_I$  dose-response with error bands ( $\pm 1$  s.d.). (C) MNIST  $N_I$  dose-response detail with error bands. (D) Fashion-MNIST individual seeds for all conditions, showing consistency across runs.




Figure S4: **Alignment dynamics with explicit gradient-norm checks.** (A) Weighted LocalCA-vs.-backprop cosine over training for the 5F rule. (B,C) Weighted local and backprop gradient norms stay nonzero for both additive and shunting models. (D) The additive control begins with a much larger scale mismatch, then approaches the backprop scale without recovering strong directional alignment. Shading is s.d. across 3 seeds.

480 Per-soma rank-1 broadcast also becomes much more competitive for additive ( $41.6 \pm 0.6\%$ ) than for  
 481 shunting ( $35.9 \pm 1.2\%$ ), while low-rank  $K=4$  remains weak in both cores. Figure S11 makes the two  
 482 useful views explicit: the absolute soma-on ordering still favors path transport strongly, and shunting  
 483 transported error is already near its soma-off ceiling. On this harder dataset, the soma channel helps,  
 484 but it still does not eliminate the importance of broadcast quality.




Figure S5: **5F sensitivity analysis (MNIST, shunting, per-soma broadcast)**. (A) Sensitivity to the  $\phi_n$  clamp interval. The default clamp [0.25, 4.0] is already strong ( $92.1 \pm 0.8\%$ ); tightening it to [0.5, 2.0] slightly improves performance ( $92.5 \pm 0.3\%$ ), while widening it to [0.1, 8.0] is effectively neutral ( $92.1 \pm 0.9\%$ ). (B) Sensitivity to the 4F/5F EMA rate used for online statistics. Varying  $\eta_{\text{ema}}$  from 0.05 to 0.20 changes mean test accuracy by less than 0.2 pp. Together these results suggest that the 5F gain is not driven by a brittle single hyperparameter choice, even though the factor itself remains heuristic. Errors:  $\pm 1$  s.d. across 3 seeds.




Figure S6: **Verification and reproducibility (supplementary)**. (A) MNIST verification: main seeds (42–46) yield  $91.4 \pm 0.3\%$ ; held-out seeds (47–49) yield  $89.8 \pm 0.9\%$ , confirming generalization. (B) Context gating verification: main seeds (+HSIC)  $80.3 \pm 0.6\%$ ; held-out seeds (no HSIC)  $77.2 \pm 1.5\%$ . The roughly 3 pp gap reflects HSIC removal, not seed sensitivity. (C) HSIC weight ablation on context gating: moderate weights (0.01–0.1) perform best.




Figure S7: **Additional stress tests**. (A) Depth scaling: shunting local (green) degrades more gracefully from 63.5% to 57.4%; additive local (blue) drops from 54.9% to 29.7%. Backprop ceilings (dashed) remain stable. (B) Broadcast-noise robustness: shunting remains stable for  $\sigma \leq 0.1$ ; additive degrades rapidly and reaches chance at  $\sigma = 1$ . (C) Fashion-MNIST: the shunting gain is modest on this cleaner benchmark (83.8% local vs. 88.9% backprop), consistent with the regime-dependent story in the main text.




Figure S8: **Coarse shared broadcast still works.** Test accuracy as a function of effective broadcast bandwidth. Full-precision, 8-bit, 4-bit, and 2-bit quantization of the shared error signal all reach essentially the same accuracy (61–62%;  $p > 0.9$  for 4-bit vs. full), while a sign-only broadcast and a sparse top-30% broadcast degrade sharply. This shows that the shunting LocalCA signal carries most of its useful information in its direction rather than in fine-grained magnitude, consistent with the mechanism-summary picture in the main text. This panel is reported separately so the main figure can carry the full causal chain (path-gain compressibility → alignment → oracle → global summary).




Figure S9: **Feedback alignment baselines (MNIST).** (A) Grouped comparison: standard backprop (neutral gray), DFA (orange), and FA (purple). Hatched N/A markers indicate that FA is not implemented for the block-structured dendritic architectures because the random feedback matrices are not dimensionally compatible with the tree-structured branch updates. DFA achieves 66.1% on shunting vs. 62.2% additive vs. 51.8% point MLP. (B) Backprop–DFA gap: dendritic architectures (24.9–27.6 pp) show smaller gaps than point MLPs (39.5 pp), suggesting conductance-based architecture is partially compatible with random feedback.




Figure S10: **CIFAR-10 control ladder in the compact direct-I-stream family.** (A) Matched backprop ceilings. With nonnegative inputs and a nonlinear decoder, shunting slightly exceeds additive under standard training, and removing learned I-to-E conductance lowers the shunting ceiling. (B) Broadcast ladder. Rank-1 broadcast remains weak on this harder dataset, rank-4 improves partially, and exact transported error nearly reaches the shunting backprop ceiling. (C) Additive fairness controls. Input-dependent gain, running-stat gain, learned gain, and dendritic normalization do not rescue additive rank-1 broadcast. This figure is a harder-data diagnostic, not a claim of competitive CIFAR-10 benchmarking.

| Condition                                | Core     | Training / broadcast | I-to-E | Test acc.                        |
|------------------------------------------|----------|----------------------|--------|----------------------------------|
| standard                                 | additive | backprop             | yes    | 48.3 $\pm$ 1.0                   |
| standard                                 | shunting | backprop             | yes    | <b>49.5 <math>\pm</math> 0.4</b> |
| standard, no learned inhibition          | shunting | backprop             | no     | 43.9 $\pm$ 1.2                   |
| rank-1                                   | additive | 5F per-soma          | yes    | 25.5 $\pm$ 1.6                   |
| rank-1, input-dependent gain             | additive | 5F per-soma          | yes    | 25.9 $\pm$ 4.1                   |
| rank-1, running-stat gain                | additive | 5F per-soma          | yes    | 24.8 $\pm$ 0.9                   |
| rank-1, learned gain                     | additive | 5F per-soma          | yes    | 26.3 $\pm$ 1.5                   |
| rank-1, dendritic normalization          | additive | 5F per-soma          | yes    | 25.5 $\pm$ 1.2                   |
| rank-1                                   | shunting | 5F per-soma          | yes    | 29.7 $\pm$ 1.3                   |
| rank-1, no learned inhibition            | shunting | 5F per-soma          | no     | 22.3 $\pm$ 0.7                   |
| transported error                        | additive | 5F path transport    | yes    | 32.4 $\pm$ 1.5                   |
| rank-4                                   | shunting | 5F low-rank          | yes    | 35.0 $\pm$ 1.0                   |
| transported error                        | shunting | 5F path transport    | yes    | <b>47.2 <math>\pm</math> 0.6</b> |
| transported error, no learned inhibition | shunting | 5F path transport    | no     | 38.6 $\pm$ 1.0                   |

Table S6: **Exact CIFAR-10 values for the 70-run control ladder (5 seeds per condition, validation-best epoch).** All rows use flattened CIFAR-10 in  $[0, 1]$ , nonnegative transfer outputs, positive conductance transforms, a compact depth-4 dendritic core, and a nonlinear decoder. LocalCA maps output error back to decoder-input/soma coordinates through the decoder Jacobian. The useful conclusions are narrow but important: learned shunting slightly improves the matched backprop ceiling, learned I-to-E conductance is necessary for the strongest shunting results, transported error nearly closes the shunting LocalCA gap, and additive rank-1 fairness controls do not explain the transported-shunting result.

| Condition (soma-on)       | Additive       | Shunting                         |
|---------------------------|----------------|----------------------------------|
| LocalCA per-soma (rank-1) | 41.6 $\pm$ 0.6 | 35.9 $\pm$ 1.2                   |
| LocalCA low-rank $K=4$    | 27.5 $\pm$ 3.5 | 26.6 $\pm$ 1.7                   |
| LocalCA path transport    | 45.5 $\pm$ 0.6 | <b>46.6 <math>\pm</math> 1.0</b> |
| soma-off path transport   | 32.4 $\pm$ 1.5 | 47.2 $\pm$ 0.6                   |
| standard BP (ceiling)     | 48.3 $\pm$ 1.0 | 49.5 $\pm$ 0.4                   |

Table S7: **CIFAR-10 compact direct-I-stream depth-4 with `somatic_synapses=true`.** Adding direct somatic synapses and using quantile reactivation init substantially improves the additive transported-error condition and keeps the shunting transported-error condition within 1 pp of its soma-off counterpart. Per-soma (rank-1) becomes much more competitive for additive than for shunting, while low-rank  $K=4$  remains weak in both cores. The soma channel therefore helps, but does not remove the need for a strong credit signal.




Figure S11: **CIFAR-10 compact direct-I-stream depth-4 with somatic\_synapses=true.** **(A)** Absolute soma-on ladder on the compact direct-I-stream depth-4 family. Path transport stays strongest in both cores and remains near the matched backprop ceilings (dashed lines), while low-rank  $K=4$  remains weak. **(B)** Matched shift from soma-off to soma-on for the conditions with a clean paired baseline. Soma helps additive strongly under per-soma and path transport, improves shunting per-soma, leaves shunting transported error close to its soma-off value, and worsens shunting low-rank  $K=4$ . Additive low-rank  $K=4$  is omitted from panel B because the soma-off mechanism package did not include that exact matched condition.




Figure S12: **Additive + normalization control (MNIST).** (A) Standard backprop: normalization provides a small boost (89.7%  $\rightarrow$  90.9%). (B) Local learning: normalization improves additive from 46.5% to 56.0% (+9.5 pp), partially closing the gap to shunting (61.7%, dashed green). The remaining gap ( $-5.7$  pp) indicates that shunting provides benefits beyond simple voltage normalization.




Figure S13: **Morphology shifts the useful inhibitory regime.** (A) Shunting-minus-additive LocalCA accuracy gap across dendritic morphologies and inhibitory synapse counts. (B) Best inhibitory count for each morphology. (C) Representative path-gain diagnostic for selected morphology cells, used to connect the performance map to credit geometry. (D) Average shunting advantage by depth, with markers showing the peak inhibitory count for each depth.

#### 485 Additive Normalization Control

#### 486 Structured Feedback, Activation, and Supportive Extensions

#### 487 Morphology $\times$ Inhibition Regime Map

488 This supportive regime map tests whether tree geometry changes the inhibitory operating point at  
 489 which shunting helps. We treat it as evidence that morphology changes the useful inhibition regime;  
 490 a fuller mechanistic account would still require matched theory diagnostics for each cell in the grid.

#### 491 Input-Mode Probe: Direct Inhibitory Stream vs. Explicit Inhibitory Cells

492 The headline experiments use `input_mode=1`, where the transfer layer provides nonnegative exci-  
 493 tatory and inhibitory streams from the external input and the excitatory dendrites receive learned  
 494 I-to-E conductance banks directly. To test whether the same mechanism can be generated by an  
 495 explicit feedforward inhibitory population, we ran a one-feedforward-layer noise-resilience probe  
 496 that preserves the relevant dendritic depth ([3, 3] morphology) while removing an unnecessary second  
 497 feedforward population. In this clean setting, explicit inhibitory cells reach  $94.2 \pm 0.2\%$  under  
 498 path-transport LocalCA when their dendrites use the same update policy as excitatory dendrites  
 499 (`explicit_inhibitory_update_mode=local_ca`),  $93.8 \pm 0.2\%$  when those inhibitory-cell up-  
 500 dates are frozen (`freeze`), and  $95.1 \pm 0.3\%$  under standard backprop. Direct input-driven inhibition  
 501 gives the expected broadcast-fidelity ladder on the same one-layer architecture: per-soma rank-1  
 502 reaches  $85.2 \pm 0.7\%$ , while exact path transport reaches  $95.2 \pm 0.0\%$ . Thus explicit inhibitory cells




Figure S14: **Pathwise inhibition story and one-layer input-mode probe.** (A) In a branched conductance tree, the error path gain  $\alpha_n$  is the product of local input resistances and dendritic coupling conductances along the path to the soma. Inhibitory conductance on any compartment along that path lowers local input resistance and attenuates all upstream credit paths. (B) The same branch-level inhibitory conductance can be supplied directly from a nonnegative inhibitory input stream (`input_mode=1`) or through an explicit feedforward inhibitory cell (`input_mode=0`). (C) The clean mechanism probe keeps the dendritic tree depth ([3, 3]) but uses only one feedforward dendritic population; adding a second feedforward E/I layer tests a deeper graph-credit problem rather than the path-gain mechanism itself. (D) Noise-resilience results for the one-layer probe. Direct-I path transport reaches the explicit-I backprop ceiling, and explicit inhibitory cells with path-transport LocalCA also remain near ceiling. This supports the interpretation that explicit inhibitory cells can generate useful path-gain structure when the architecture isolates the dendritic-path question.

| Inhibitory input mode                                  | Training / broadcast                           | Test accuracy                       |
|--------------------------------------------------------|------------------------------------------------|-------------------------------------|
| Direct inhibitory stream ( <code>input_mode=1</code> ) | LocalCA, per-soma rank-1                       | $0.852 \pm 0.007$                   |
| Direct inhibitory stream ( <code>input_mode=1</code> ) | LocalCA, exact path transport                  | <b><math>0.952 \pm 0.000</math></b> |
| Explicit I cells ( <code>input_mode=0</code> )         | LocalCA, path transport, I-cell updates on     | $0.942 \pm 0.002$                   |
| Explicit I cells ( <code>input_mode=0</code> )         | LocalCA, path transport, I-cell updates frozen | $0.938 \pm 0.002$                   |
| Explicit I cells ( <code>input_mode=0</code> )         | standard backprop                              | $0.951 \pm 0.003$                   |

Table S8: **One-feedforward-layer input-mode path-gain probe on noise resilience (3 seeds).** All rows use one excitatory dendritic population with [3, 3] morphology and nonnegative inputs. The direct-I rows have no explicit inhibitory neurons; the input stream drives learned branch-level inhibitory conductance banks directly. The explicit-I rows use a matched inhibitory population with [3, 3] morphology. `local_ca` updates explicit inhibitory-cell dendrites with the same policy used for excitatory dendrites, while `freeze` holds those dendrites fixed but still lets excitatory I-to-E synapses learn. Exact path transport closes the direct-I rank-1 gap, and explicit inhibitory cells nearly match both direct-I path transport and the explicit-I backprop ceiling.

503 can generate useful path-gain structure once the architecture isolates the dendritic-path question. The  
504 older two-feedforward-layer explicit-I diagnostic was much weaker (about 58–63% under LocalCA  
505 despite a 94.7% backprop ceiling), so we treat it as evidence that deeper explicit  $E \rightarrow I \rightarrow E$  graphs  
506 need graph-aware local errors, not as evidence against inhibition-mediated path gains.

507 **Fitted inhibitory path-gain diagnostic.** We also asked whether the direct-I-stream models actually  
508 use inhibitory synapses in the way implied by Prop. 2. We reran the cue-routing pathway controls with  
509 strictly nonnegative router outputs and matched learned-I versus no-I conditions (50 complete runs).  
510 The task is now too easy to separate feedback methods: all conditions reach 99.0–99.1% test accuracy,  
511 and removing learned I-to-E conductance changes accuracy by less than the seed-to-seed variation.  
512 The useful information is therefore mechanistic rather than performance-based. In the learned-I  
513 LocalCA runs, the first-layer pathway-inhibition balance carries context information, and higher  
514 inhibitory conductance is consistently associated with lower local input resistance. The sign of the  
515 balance can flip across random seeds, so we report best-polarity context decoding as a sign-invariant  
516 measure. This does not prove that pathway-vector feedback has solved routed credit assignment,  
517 but it supports the narrower claim that learned inhibitory conductance can create input-dependent  
518 path-gain differences by modulating branch resistance.

| Cue feedback mode | Context decode    | MI                     | Corr( $I, R$ )     |
|-------------------|-------------------|------------------------|--------------------|
| rank-1            | $0.963 \pm 0.032$ | $0.867 \pm 0.100$ bits | $-0.608 \pm 0.085$ |
| rank-2            | $0.856 \pm 0.180$ | $0.810 \pm 0.213$ bits | $-0.602 \pm 0.041$ |
| pathway-vector    | $0.852 \pm 0.170$ | $0.694 \pm 0.300$ bits | $-0.627 \pm 0.028$ |
| path transport    | $0.816 \pm 0.153$ | $0.692 \pm 0.227$ bits | $-0.642 \pm 0.018$ |

Table S9: **Fitted direct-I cue-routing diagnostic (5 seeds per learned-I LocalCA mode).** In input\_mode=1, inhibitory synapse banks receive a nonnegative input stream directly. In trained cue-routing models, the balance of inhibitory activation across routed pathways carries context information and higher inhibitory conductance coincides with lower local input resistance. Context decoding is reported with best polarity because branch sign can flip across seeds; it should be read as evidence for context-dependent inhibitory path gating, not as a fixed anatomical assignment of one branch to one cue.




Figure S15: **Positive routed-task pathway and inhibition controls.** (A) Strictly nonnegative cue-routing controls saturate near 99% for learned-I and no-I variants, so this control is not a performance separation. (B) Fitted excitatory and inhibitory synapse banks retain branch-pathway structure above the 50% chance level. (C) In learned-I cue models, pathway-inhibition balance carries context information and inhibitory conductance is negatively coupled to local input resistance. (D) A nonnegative Double-CIFAR contextual screen remains near chance even under matched backprop, so it is not used as evidence for the paper’s local-learning claims. Error bars show s.d. across 5 seeds.

519 **Post-training inhibition intervention and exact-error rank.** We next tested whether learned  
520 inhibitory conductance carries task-specific structure rather than only changing global scale. We  
521 replayed trained shunting LocalCA checkpoints while replacing the branch-level inhibitory conduc-  
522 tance with four post-training interventions: zero inhibition, sample-shuffled inhibition, per-branch  
523 batch means, or one uniform matched mean. On a 512-example held-out subset, learned MNIST  
524 inhibition at  $N_I=5$  gives  $90.4 \pm 1.1\%$  accuracy, while the four interventions drop to 19.8–26.4%.  
525 On noise resilience at  $N_I=10$ , learned inhibition gives  $80.5 \pm 2.4\%$ , while the interventions drop  
526 to 10.3–19.3%. Thus the inhibitory field is not interchangeable with a matched global conductance  
527 load. We also computed the SVD of exact compartment-error matrices across samples and compart-  
528 ments. At matched MNIST  $N_I=5$ , shunting has lower rank-1 residual than additive ( $0.65 \pm 0.18$   
529 vs.  $0.83 \pm 0.07$ ) and lower participation rank ( $3.2 \pm 1.2$  vs.  $8.5 \pm 3.3$ ), supporting the path-gain  
530 compressibility interpretation.

531 We also computed the direct log-gain covariance margin in Eq. (8) on the selected morphology  
532 diagnostic checkpoints. The margin is zero when no inhibitory conductance is present and slightly  
533 negative in the inhibited shunting selections ( $-2.5 \times 10^{-3}$ ,  $-5.0 \times 10^{-4}$ , and  $-2.8 \times 10^{-4}$  at  
534  $N_I=5, 20, 40$ ). We therefore do not use this covariance condition as a positive headline result. Its  
535 role is to clarify when shunting should compress gains; the empirical claim rests on the measured  
536 comparative path-gain dispersion, exact-error rank, and inhibition-intervention diagnostics.

### 537 Cue-Routing Feedback-Rank Diagnostic

538 Cue routing tests the regime where one shared teaching signal is expected to fail: context determines  
539 which noisy cue is reliable, so different pathways should receive different credit. Figure S17 shows  
540 that the benchmark is learnable and that moving beyond rank-1 feedback helps, but the current  
541 path-structured variant does not yet beat the best unstructured low-rank setting. We therefore interpret  
542 this result as a failure-mode and rank-requirement diagnostic, not as a solved structured-feedback  
543 method.




Figure S16: **Post-training inhibition intervention and exact-error compressibility.** (A,B) Post-training inhibitory-conductance interventions on trained shunting LocalCA checkpoints. Accuracy is measured on a 512-example held-out subset; dots show seeds. Removing, shuffling, or clamping learned inhibition collapses performance, showing that the learned inhibitory field carries sample-specific task structure. (C,D) SVD diagnostics of the exact compartment-error matrix on matched MNIST  $N_I=5$  runs. Shunting has a smaller rank-1 residual and lower participation rank than additive, so the exact error field is more compressible. Error bars are s.d. across 5 seeds.




Figure S17: **Structured feedback when rank-1 broadcast fails.** (A) Cue-routing task schematic. Context determines which noisy cue is reliable on each trial, so different dendritic pathways should receive different teaching signals. (B) Five-seed comparison of rank-1, random low-rank, and pathway-structured feedback. Moving beyond rank-1 helps, while the best structured repair remains open. (C) Test accuracy vs. router specialization across learned-router conditions. High specialization alone does not guarantee success. (D) Recovered pathway assignments for a representative learned-router run. Cue-A and cue-B features segregate into separate latent pathways even when feedback quality, rather than router discreteness, is the limiting factor.

## Random Low-Rank Broadcast Channels

## Frozen Activation-Choice Audit

## B Implementation Details

## Biological Plausibility Assumptions

The theorems and rules above rest on a small set of explicit assumptions. We state them as a checklist so that readers can see exactly what is assumed and where each assumption is used:

**A1 (Steady-state voltage).** We use the steady-state solution of the passive cable equation (Eq. (1)). Temporal dynamics are not modeled; all gradients, path gains, and fidelity measurements refer to the conductance-stage voltage at steady state.

**A2 (Tree topology).** Each dendritic cell is a rooted tree with the soma as the root. There are no recurrent connections within the dendrite, which makes the path from any compartment to the soma unique and is what allows the closed-form path gain  $\alpha_n$  in Theorem 1.

**A3 (Nonnegative conductances).** Synaptic and dendritic conductance parameters are pushed through a positive transform (`softplus`), so  $g_j^{\text{syn}}, g_j^{\text{den}} \geq 0$  at every optimisation step. This is a biophysical constraint and is also what keeps the shunting denominator bounded away from zero.

**A4 (Nonnegative presynaptic drive).** The first synaptic drive  $x$  is nonnegative, implemented either by nonnegative inputs or an explicit ReLU transfer stage. Together with A3 this enforces the positive-input condition used for all headline sweeps.

| Setting                                                           | Broadcast                         | Test accuracy                       |
|-------------------------------------------------------------------|-----------------------------------|-------------------------------------|
| <i>Noise resilience, shunting, rank bridge</i>                    |                                   |                                     |
| per-soma                                                          | rank-1 per-soma                   | $0.662 \pm 0.016$                   |
| path propagation                                                  | recursive attenuation proxy       | $0.762 \pm 0.047$                   |
| low-rank                                                          | $K=1$                             | $0.453 \pm 0.126$                   |
| low-rank                                                          | $K=2$                             | $0.630 \pm 0.042$                   |
| low-rank                                                          | $K=4$                             | $0.764 \pm 0.025$                   |
| low-rank                                                          | $K=8$                             | $0.795 \pm 0.019$                   |
| path transport                                                    | exact conductance-stage transport | <b><math>0.847 \pm 0.018</math></b> |
| <i>Cue routing, learned shunting router, rank/structure sweep</i> |                                   |                                     |
| per-soma                                                          | tuned rank-1                      | $0.803 \pm 0.125$                   |
| low-rank                                                          | $K=1$                             | $0.742 \pm 0.022$                   |
| low-rank                                                          | $K=2$                             | <b><math>0.846 \pm 0.082</math></b> |
| low-rank                                                          | $K=4$                             | $0.748 \pm 0.085$                   |
| pathway-vector                                                    | tuned structured rank-2           | $0.815 \pm 0.126$                   |

Table S10: **Rank bridge and rank-structure comparison.** On shunting noise resilience, both the current path-propagation proxy and higher-rank random low-rank broadcast close a substantial fraction of the gap between rank-1 per-soma and exact path transport, with low-rank continuing to improve up to  $K=8$ . On cue routing, the learned-router sweep shows that moving beyond rank-1 helps, but the best current unstructured low-rank setting ( $K=2$ ) is still slightly better than the tuned structured pathway-vector variant. We therefore interpret routed credit assignment as an open higher-rank problem rather than as a solved pathway-vector rescue.

| Dataset / strategy          | Core     | $N_I$ | Identity (none)    | Frozen param_tanh |
|-----------------------------|----------|-------|--------------------|-------------------|
| MNIST / LocalCA             | additive | 0     | $0.780 \pm 0.014$  | $0.709 \pm 0.037$ |
| MNIST / LocalCA             | additive | 10    | $0.907 \pm 0.003$  | $0.881 \pm 0.005$ |
| MNIST / LocalCA             | shunting | 0     | $0.841 \pm 0.007$  | $0.778 \pm 0.020$ |
| MNIST / LocalCA             | shunting | 10    | $0.905 \pm 0.009$  | $0.840 \pm 0.016$ |
| Noise resilience / LocalCA  | additive | 0     | $0.590 \pm 0.162$  | $0.109 \pm 0.006$ |
| Noise resilience / LocalCA  | additive | 10    | $0.859 \pm 0.001$  | $0.772 \pm 0.009$ |
| Noise resilience / LocalCA  | shunting | 0     | $0.411 \pm 0.028$  | $0.331 \pm 0.020$ |
| Noise resilience / LocalCA  | shunting | 10    | $0.769 \pm 0.004$  | $0.701 \pm 0.006$ |
| MNIST / backprop            | additive | 0     | $0.925 \pm 0.001$  | $0.942 \pm 0.001$ |
| MNIST / backprop            | additive | 10    | $0.921 \pm 0.001$  | $0.950 \pm 0.001$ |
| MNIST / backprop            | shunting | 0     | $0.963 \pm 0.0004$ | $0.950 \pm 0.001$ |
| MNIST / backprop            | shunting | 10    | $0.967 \pm 0.001$  | $0.946 \pm 0.001$ |
| Noise resilience / backprop | additive | 0     | $0.896 \pm 0.003$  | $0.106 \pm 0.007$ |
| Noise resilience / backprop | additive | 10    | $0.893 \pm 0.002$  | $0.927 \pm 0.002$ |
| Noise resilience / backprop | shunting | 0     | $0.898 \pm 0.001$  | $0.869 \pm 0.002$ |
| Noise resilience / backprop | shunting | 10    | $0.926 \pm 0.001$  | $0.912 \pm 0.001$ |

Table S11: **Frozen activation-shape audit on matched dendritic architectures (5 seeds per condition).** Identity (none) and a frozen bounded post-voltage reactivation param\_tanh lead to markedly different operating regimes under nonnegative first-layer synaptic drive. In this audit, identity is the safer LocalCA choice on both MNIST and the noise-resilience task, while the backprop effect is strongly regime-dependent: frozen param\_tanh helps additive most clearly in the inhibited noise-resilience setting ( $N_I=10$ ) but collapses it at  $N_I=0$ . These runs match additive reactivation initialization to the shunting setting and hold the activation parameters fixed, isolating activation shape from both additive/shunting initialization differences and activation learning itself.

563 **A5 (Pre-reactivation upward transfer).** Theorem 1 is stated for the pre-reactivation compartment  
564 voltage. A post-voltage reactivation  $f$  (e.g. param\_tanh) is handled by multiplying the  
565 path gain by  $f'(V_n)$  along the recursion. The compartment error remains the sole non-local  
566 quantity; only the effective path gain acquires local activation-derivative factors.

| Dataset          | Core     | $N_I$ | Identity (none)        | Frozen param_tanh      |
|------------------|----------|-------|------------------------|------------------------|
| MNIST            | additive | 0     | CV 0.280, cosine 0.220 | CV 1.122, cosine 0.074 |
| MNIST            | additive | 10    | CV 0.310, cosine 0.212 | CV 1.079, cosine 0.094 |
| MNIST            | shunting | 0     | CV 0.393, cosine 0.170 | CV 0.770, cosine 0.114 |
| MNIST            | shunting | 10    | CV 0.385, cosine 0.281 | CV 1.177, cosine 0.106 |
| Noise resilience | additive | 0     | CV 0.124, cosine 0.079 | CV 0.114, cosine 0.052 |
| Noise resilience | additive | 10    | CV 0.414, cosine 0.206 | CV 0.840, cosine 0.094 |
| Noise resilience | shunting | 0     | CV 0.255, cosine 0.324 | CV 0.344, cosine 0.341 |
| Noise resilience | shunting | 10    | CV 0.762, cosine 0.140 | CV 2.497, cosine 0.079 |

Table S12: **Frozen activation-shape theory audit on matched LocalCA runs (5 seeds per condition).** Each entry reports conductance-stage path-gain dispersion (CV) and direct per-soma compartment-error fidelity (cosine between the per-soma broadcast and the exact pre-reactivation compartment error). The positive-input theory audit largely agrees with the empirical one: freezing param\_tanh usually broadens path-gain dispersion and weakens per-soma fidelity, especially in the inhibited settings where the main mechanistic story lives. The one mild exception is low-inhibition shunting on noise resilience, where frozen param\_tanh leaves fidelity roughly unchanged while still hurting final accuracy, suggesting that activation choice affects optimization and representation as well as direct credit fidelity.

567 **A6 (Single perisomatic broadcast).** Every synapse receives one non-local scalar or low-rank vector  
568 signal delivered from the soma; all other ingredients of the update rule (eligibility traces,  
569 driving force, input resistance, covariance modulators) are computable from quantities  
570 available at the synapse or compartment.

571 Each synapse therefore has access to: presynaptic activity  $x_j$  (local), compartment voltage  $V_n$   
572 (local membrane potential), reversal potential  $E_j$  (fixed by receptor type), and total input resistance  
573  $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$  (computable from local conductances). The 4F modulator  $r_n^{4F}$  requires online  
574 estimation of voltage covariance (biologically plausible via slow calcium signals); the 5F factor  $\phi_n$   
575 requires a linear regression proxy (implementable via eligibility traces). The only non-local quantity  
576 is the broadcast error  $e_n$ , which requires a top-down or neuromodulatory signal from the soma to  
577 dendritic compartments. We assume per-neuron resolution (one error per output neuron) for the  
578 per-soma broadcast mode.

## 579 Dendritic Structure in Code

580 We use the morphology notation  $[b_1, b_2, \dots, b_D]$  for rooted dendritic trees, where each factor gives  
581 the fan-in from one dendritic stage to the next on the path to the soma. For example,  $[3, 3]$  means that  
582 each soma aggregates 3 proximal branches and each proximal branch aggregates 3 distal branches,  
583 for 9 distal leaves per soma;  $[3, 3, 3]$  adds a third dendritic stage with 27 leaves per soma. In code,  
584 each excitatory unit is instantiated as a DendriNet, which is a stack of DendriticBranchLayers.  
585 Each branch layer contains an excitatory synapse bank, an optional inhibitory synapse bank, a  
586 branch-to-parent aggregation map, and an optional post-voltage reactivation. In the main feedfor-  
587 ward models studied here, somatic\_synapses=false, so synaptic inputs terminate on dendritic  
588 branches rather than directly on the soma. The headline runs use input\_mode=1: the transfer layer  
589 creates nonnegative excitatory and inhibitory streams from the same external input, and nonzero  
590 ie\_synapses\_per\_branch\_per\_layer attaches learned inhibitory conductance banks directly to  
591 excitatory dendritic branches. These models therefore have input-driven branch inhibition but no  
592 separate inhibitory neuron population. By contrast, input\_mode=0 instantiates explicit feedforward  
593 inhibitory cells when inhibitory\_layer\_sizes is nonempty; these cells can transform excitatory  
594 activity before projecting inhibitory conductance back onto excitatory dendrites. In that setting  
595 the inhibitory cells are also dendritic modules, so their E-to-I and I-to-I synaptic conductances are  
596 trained by the same local rule family rather than by a separate global inhibitory optimizer. We treat  
597 that explicit-I setting as an exploratory mechanism probe rather than a headline baseline, because  
598 the main claims were measured in the direct-I-stream regime. The additive and shunting cores  
599 are architecture-matched at the tree level; they differ in the branch-level voltage integration rule  
600 rather than in whether a dendritic tree exists. In all reported dendritic runs, synaptic and dendritic  
601 conductance parameters are pushed through positive transforms (softplus in the main paper runs),




Figure S18: **Soma-on extension for LocalCA.** Enabling direct somatic input (`somatic_synapses=true`) improves aggregate mean test accuracy across all four manuscript LocalCA families. Each family is summarized over the completed configurations in matched soma-off and soma-on sweeps. The largest gain appears in morphology scaling, where the family mean rises from 35.8% to 70.7% and the best run rises from 57.0% to 76.4%. Gains in the shunting-regime and error-shaping families are smaller but still substantial. Delta labels show soma-on minus soma-off. Because the main mechanistic analysis was developed in the dendrite-only feedforward regime, we report soma-on as a supportive extension rather than replacing the main operating regime.

| Family             | Soma-off | Soma-on | Δ mean  | Δ best  |
|--------------------|----------|---------|---------|---------|
| Phase-1 capacity   | 0.8784   | 0.8999  | +0.0215 | -0.0007 |
| Shunting regime    | 0.5641   | 0.6715  | +0.1074 | +0.0090 |
| Morphology scaling | 0.3578   | 0.7070  | +0.3492 | +0.1940 |
| Error shaping      | 0.2798   | 0.3906  | +0.1109 | +0.1210 |

Table S13: **Representative soma-on deltas across the four LocalCA families.** Values compare matched families with and without `somatic_synapses` and report both the change in family-mean test accuracy and the change in best-run test accuracy. Cue-routing soma and CIFAR-10 depth-4 soma are reported separately in Fig. S19, Fig. S11, and Table S7 because those extensions are summarized on different task families.

so the conductances themselves remain nonnegative at every optimization step; unconstrained decoder weights are separate readout parameters and are not interpreted as conductances.

Matched soma-on LocalCA sweeps suggest that direct somatic input is a meaningful extension rather than a minor implementation detail. Relative to matched non-somatic sweeps, enabling `somatic_synapses=true` raises the mean test accuracy in all four manuscript families: phase-1 capacity from 0.8784 to 0.8999, shunting-regime sweeps from 0.5641 to 0.6715, morphology-scaling sweeps from 0.3578 to 0.7070, and error-shaping sweeps from 0.2798 to 0.3906. These numbers are family-level aggregates over the completed configurations in each matched soma-off and soma-on sweep, rather than a single 5-seed cell. Because the main mechanistic analysis in this paper was developed in the dendrite-only feedforward regime, we keep the main text anchored to `somatic_synapses=false` and treat soma-on as a supportive extension rather than silently replacing the original operating regime. Figure S18 summarizes the mean and best-accuracy comparisons across the four families.

The soma-on benefit extends beyond the capacity/scaling families to the cue-routing benchmark that stress-tests rank-1 feedback. With dendrite-only pathways, the best learned-routing shunting LocalCA configuration was unstable at  $0.81 \pm 0.13$ ; with `somatic_synapses=true`, fixed-router shunting LocalCA reaches  $1.00 \pm 0.0$  test accuracy across 5 seeds, and learned-router additive LocalCA reaches  $0.99 \pm 0.003$  (Fig. S19). The soma channel provides a high-bandwidth direct input that bypasses the rank-1 bottleneck induced by per-soma broadcast.




Figure S19: **Soma-on rescues cue-routing LocalCA.** The left bars show the previously reported soma-off LocalCA results on the cue-integration benchmark. The right bars show the soma-on extension: fixed-router shunting LocalCA reaches 1.00 test accuracy, and learned-router additive LocalCA with soma reaches 0.99, matching the additive backprop ceiling with soma. The soma channel supplies a direct high-bandwidth input that bypasses the rank-1 bottleneck identified in the main text. Learned-router shunting with soma is omitted because this benchmark uses signed cues, which conflict with the shunting positive-input check in the learned-router pipeline; fixed-router shunting is therefore the clean shunting data point.

## 621 Post-Voltage Reactivation

622 The conductance stage produces a pre-reactivation compartment voltage  $V_n$ , after which each branch  
 623 layer may apply a monotone scalar nonlinearity. The implementation supports identity (none), fixed  
 624  $\tanh/\text{sigmoid}/\text{relu}$ , and several parametric families including `param_tanh`, `param_tanh_only_m`,  
 625 and piecewise linear-saturating variants. The main manuscript sweeps use `param_tanh`, a learnable  
 626 bounded firing-rate transform of the form  $(\tanh(m(V - b)) + 1)/2$ , together with an explicitly  
 627 configured nonnegative transfer stage so that the first synaptic drive remains nonnegative even  
 628 after input corruption or structured task noise. Most headline feedforward sweeps use a linear  
 629 decoder, while the CIFAR extension uses a nonlinear decoder and therefore maps output error back to  
 630 decoder-input/soma space through the decoder Jacobian before applying LocalCA. Unless explicitly  
 631 overridden, additive and shunting cores can start from different reactivation initializations; activation-  
 632 fairness ablations therefore match additive initialization to the shunting setting when comparing  
 633 alternative reactivation choices directly. The appendix activation audit additionally fixes `param_tanh`  
 634 after initialization to isolate activation shape from activation plasticity.

635 LocalCA uses occupancy-quantile initialization together with a calibration pass that floors tiny  
 636 quantile widths, clamps extreme fitted slopes, and reverts invalid per-layer fits rather than allowing  
 637 pathological intermediate values to propagate. In the non-somatic feedforward publication families,  
 638 this calibration leaves the results essentially unchanged relative to the same occupancy-quantile setup  
 639 without the additional validity checks: phase-1 capacity changes from 0.8786 to 0.8784 mean test  
 640 accuracy, shunting-regime sweeps from 0.5617 to 0.5641, morphology-scaling sweeps from 0.3600  
 641 to 0.3578, and error-shaping sweeps from 0.2860 to 0.2798. The calibration therefore improves  
 642 numerical safety without changing the qualitative LocalCA story.

643 We also compared two post-initialization policies for the reactivation parameters under quantile  
 644 initialization: *learned local* ( $b, m$ ), where LocalCA updates the gate parameters through the local  
 645 rule, and *quantile-maintained* ( $b, m$ ), where  $(b, m)$  are periodically refreshed from the current  
 646 voltage distribution instead of being updated by the local rule. This comparison matters because  
 647 quantile-maintained gates substantially simplify nonlinear reactivation learning: the rest of the  
 648 network still uses LocalCA, but the gate parameters themselves are set by a local distributional  
 649 rule rather than a learned plasticity rule. On MNIST classification sweeps, learned local ( $b, m$ )  
 650 remains clearly better (additive 0.9247 vs. 0.9061 mean test accuracy; shunting 0.9104 vs. 0.8766).  
 651 On CIFAR-10, the picture is mixed in a useful way: additive is essentially tied on mean (0.4223  
 652 learned-local vs. 0.4215 quantile-maintained) while shunting improves under quantile-maintained




Figure S20:  $(b, m)$  update policy under quantile initialization. Each cluster of bars compares the mean test accuracy of learned local  $(b, m)$  and quantile-maintained  $(b, m)$  under LocalCA for additive and shunting cores, averaged across a matched morphology/init grid  $([2, 2]$  and  $[3, 3, 3]$ ; analytical and quantile init variants). Dashed and dotted horizontal lines mark the matched standard backpropagation ceiling for additive and shunting respectively. **Left (MNIST):** Learned local  $(b, m)$  is consistently better; quantile-maintained lags by 1.9 pp (additive) and 3.4 pp (shunting) on mean accuracy. **Right (CIFAR-10):** The gap reverses for shunting: quantile-maintained reaches 43.6% mean vs. 40.3% for learned local, while additive is essentially tied. The pattern suggests that quantile-maintained  $(b, m)$  prevents over-fitting the reactivation gate on harder inputs, making it a useful alternative in harder visual regimes.

| Quantity         | Symbol                           | Convention              |
|------------------|----------------------------------|-------------------------|
| Voltage          | $V$                              | Normalized to $[-1, 1]$ |
| Conductances     | $g^{\text{syn}}, g^{\text{den}}$ | Nonneg. via softplus    |
| Leak conductance | $g^{\text{leak}}$                | Set to 1                |
| Input resistance | $R^{\text{tot}}$                 | $\leq 1$                |

Table S14: Units and normalization.

gates (0.4028 to 0.4356 mean test accuracy). These bars aggregate over a matched morphology/init grid rather than a dedicated 5-seed replicate, so we treat them as a supportive policy comparison rather than a headline result. Accordingly, the main manuscript continues to use occupancy-quantile initialization followed by LocalCA updates of  $(b, m)$  as the default local-learning setting, while quantile-maintained  $(b, m)$  is reported as a supplementary alternative. Figure S20 shows the full MNIST and CIFAR-10 comparison.

## Component Ablation

To isolate the contribution of each ingredient in our LocalCA recipe, we ran a per-component ablation on MNIST starting from the manuscript default and removing one component at a time. The default combines quantile init, learned local  $(b, m)$ , `somatic_synapses=false`, and `param_tanh` reactivation. We also include `somatic_synapses=true` as a positive extension and matched standard backpropagation as the upper reference. Figure S21 reports test accuracy (5 seeds per cell) for both shunting and additive dendritic cores on clean MNIST with morphology  $[3, 3]$ .

## Units and Parameterization

### Decoder Update Modes

$W_{\text{dec}}$  maps  $V_L \rightarrow \hat{y}$ . Modes: **backprop** ( $\nabla_W L$  via autograd), **local** ( $\Delta W = \eta \langle \delta_0 V_L^T \rangle_B$ ), **frozen** ( $\Delta W = 0$ ).




Figure S21: **Per-component ablation for LocalCA on MNIST.** Bars show mean test accuracy ( $\pm$  s.d. across 5 seeds); coloured bars highlight the manuscript default (blue/green) and the soma-on extension (green). **(A) Shunting:** The full default reaches 90.7%; removing quantile init leaves accuracy unchanged at 90.8% (quantile init is a mild calibration nudge on this clean task), removing learned ( $b, m$ ) drops by about 1 pp, and removing reactivation entirely drops accuracy to 87.3%; the reactivation transform is the largest single LocalCA ingredient. **(B) Additive:** Quantile init contributes +2.1 pp (90.1  $\rightarrow$  92.2%), learned ( $b, m$ ) contributes +1.7 pp, and reactivation contributes +1.1 pp; each ingredient helps on the additive core. **Soma-on** is the largest positive change in both cores (+3 to +3.5 pp over full config), while matched standard backpropagation reaches 97.0% in both cores and sets the upper reference. The fixed-relu reactivation condition did not satisfy the same initialization checks and is omitted.

| Category                    | Implemented modes                                                                                | Role in this paper                                                                                                                                                                                  |
|-----------------------------|--------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Core models                 | point MLP; dendritic MLP; dendritic additive; dendritic shunting; pathway-routed dendritic cores | Main comparisons use dendritic additive vs. dendritic shunting; routed cores are used for cue integration.                                                                                          |
| Training strategies         | standard backprop; LocalCA; DFA; FA; transported-broadcast oracle                                | Main text uses standard and LocalCA; DFA/FA remain appendix baselines; path transport is an oracle upper bound.                                                                                     |
| Broadcast modes             | scalar; per-soma; local mismatch; low-rank; path transport; pathway vector                       | Main text centers on per-soma, low-rank, and transported-error controls; scalar and local-mismatch are supporting controls, while pathway vector remains an appendix structured-feedback extension. |
| Recent dendritic extensions | path propagation; router-derived branch roles; pathway-vector feedback; inhibitory homeostasis   | These were added to test morphology-aware credit transport and structured branch-specific plasticity beyond the original 3F/4F/5F family.                                                           |

Table S15: Implemented model families and LocalCA operating modes. Only a subset is used for the main claims; the remaining modes serve as controls, ablations, or structured extensions.

## 670 Model and Mode Inventory

## 671 Representative Manuscript Architectures

## 672 Hyperparameters

673 Table S17 collects the core training hyperparameters for each headline experiment. Learning rates  
674 are applied through parameter groups (top- $k$ , dendritic block-linear, reactivation, decoder) with the  
675 values given in the representative config files; where a single LR is reported, every group uses that  
676 value.

## 677 Effective Broadcast Dimensionality in Headline Architectures

### 678 LocalCA Extension Compatibility

### 679 Seed Protocol and Evidence Map

680 The seed depth is not uniform across every sweep, and we therefore separate main evidence from  
 681 supporting diagnostics. Table S20 pins each figure or table in this paper to its seed count and evidence  
 682 tier.

### 683 Compute Resources

684 All experiments were run on an institutional GPU cluster. Most MNIST, Fashion-MNIST, and  
 685 synthetic-task runs used a single NVIDIA A100 or H100 GPU and finished in minutes to tens of  
 686 minutes per seed. The flattened CIFAR-10 runs required longer single-GPU jobs but still did not use  
 687 distributed or multi-GPU training. Reported headline metrics are aggregated over 5 seeds unless noted  
 688 otherwise. The remaining 3-seed results are confined to selected supportive sweeps, and the broader  
 689 exploratory screens used to tune and diagnose the method required additional compute beyond the  
 690 final reported runs.

### 691 Code, Data, and Asset Availability

692 MNIST, Fashion-MNIST, CIFAR-10, and MICrONS are public datasets or public benchmark assets  
 693 cited in the paper; we use them under their standard published access conditions and cite their original  
 694 sources in the bibliography. The present anonymous submission does not include a public code URL  
 695 in order to preserve double-blind review, but the paper specifies the model family, training modes, key  
 696 hyperparameters, and representative configurations, and we plan to release code and configuration  
 697 files after review. The paper does not release a new dataset or a stand-alone pretrained model asset.

### 698 Algorithm

---

#### Algorithm 1 Headline 5F Rank-1 LocalCA Update

---

- 1: **Input:** Dendritic model, batch  $(x, y)$ , learning rate  $\eta$
  - 2: Forward pass; loss  $L$ , output error  $\delta^y$
  - 3: Somatic/core error  $\delta_0 = \text{decoder-input Jacobian product } J_{\text{dec}}(h_{\text{core}})^\top \delta^y$  (or  $W_{\text{dec}}^\top \delta^y$  for a linear decoder)
  - 4: **for** each layer  $n$  (reverse) **do**
  - 5:   Rank-1 broadcast  $e_n = \text{mean}(\delta_0) \mathbf{1}_n$
  - 6:   Update EMA estimates for  $r_n^{4F}$  and  $\phi_n$  from batch voltages
  - 7:    $\Delta g_j^{\text{syn}} \leftarrow \eta r_n^{4F} \phi_n \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n \rangle_B$
  - 8:    $\Delta g_j^{\text{den}} \leftarrow \eta r_n^{4F} \phi_n \langle R_n^{\text{tot}} (V_j - V_n) e_n \rangle_B$
  - 9: **end for**
  - 10: Optional controls replace line 4 with neuron-wise, rank- $K$ , path-structured, or transported-oracle broadcasts.
  - 11: Clip gradients; optimizer step
- 

## 699 C Theoretical Details

### 700 Variant Taxonomy

#### 701 Random-Broadcast Alignment

702 **Proposition 3** (Supportive random-broadcast alignment). *Let the broadcast matrix  $B_n$  have i.i.d.*  
 703 *zero-mean entries with  $\mathbb{E}[B_n^\top B_n] = \sigma_B^2 I$ . Then the expected cosine between the resulting local*  
 704 *gradient and the exact gradient is positive,  $\mathbb{E}[\cos \angle(g^{\text{local}}, g^{\text{exact}})] \geq c_n > 0$ , by an argument*  
 705 *analogous to feedback alignment [9]. The constant  $c_n$  depends on how strongly local eligibility*  
 706 *factors correlate with the conductance-stage path gain in Eq. (4); shunting improves this correlation*  
 707 *by narrowing the spread of intermediate sensitivities.*

We treat this argument as supportive rather than central. The main theoretical object in the paper is still the exact pre-reactivation compartment error  $\partial L / \partial V_n$ ; under identity upward transfer this reduces to  $\alpha_n \delta_0$ , while with post-voltage activation the corresponding effective path gain additionally includes local  $f'_n(V_n)$  factors. The main empirical question is how faithfully different broadcast fields approximate that quantity.

## Biological Analogs

## D Morphology-Aware Extensions

**Path-integrated propagation.** Modulate broadcast error by  $\pi_n = \pi_{n-1} \cdot R_{n-1}^{\text{tot}} \cdot \bar{g}_{n-1}^{\text{den}}$ , approximating depth attenuation from Eq. (4).

**Depth modulation.** Per-branch scaling  $\kappa_j = \kappa_{\text{base}} / (d_j + d_0)$ , mirroring cable attenuation.

**Dendritic normalization.**  $\Delta g_j^{\text{den}} \leftarrow \Delta g_j^{\text{den}} / (\sum_k g_k^{\text{den}} + \varepsilon)$ , analogous to homeostatic scaling [21].

**Pathway-vector feedback.** For tasks with latent pathway structure, the broadcast becomes a role vector inferred from the router. Local pathway activity gates this vector. Upward block transport then passes it to earlier branches, aligning their feedback with downstream descendants.

**Router-derived branch roles.** Each branch receives a role profile inferred from router assignments and incoming block weights rather than from a hand-coded apical/basal label. Plasticity is then modulated by branch selectivity and a synapse-specific role-alignment factor, emphasizing pathway-consistent updates without imposing a heuristic branch taxonomy.

## E HSIC Auxiliary Objectives

Following [17], we use kernel-based HSIC objectives:

$$\mathcal{L}^{\text{self}} = B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Z \mathbf{H}), \quad \mathcal{L}^{\text{target}} = -B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Y \mathbf{H}).$$

Moderate weights (0.01–0.1) help on context gating; negligible on MNIST. Online statistics ( $r_n^{\text{4F}}$ ,  $\phi_n$ ) use Welford’s algorithm [20].

## F Online Variant with Eligibility Traces

The continuous-time eligibility trace is

$$\tau_e \dot{e}_j^{\text{syn}} = -e_j^{\text{syn}} + x_j(E_j - V_n)R_n^{\text{tot}},$$

with update

$$\Delta g_j^{\text{syn}} \propto \int e_j^{\text{syn}}(t) m_n(t) dt$$

[18, 19].

## G Weight Statistics and Biological References

As a secondary analysis, we examine whether the conductance-based update naturally induces characteristic weight statistics. Because synaptic updates scale with input resistance  $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$ , the voltage equation implies multiplicative weight dynamics, creating a regime analogous to exponentiated-gradient updates. Such multiplicative dynamics suggest log-normal weight distributions in trained networks [38, 39]. We compare the resulting distributions against two biological references: electrophysiological data from Song et al. [38] (rat V1 L5 pyramidal,  $\sigma = 0.94$ ) and the MICrONS mm<sup>3</sup> connectome [40] (mouse V1 synapse sizes).

**Model weight distributions are log-normal.** We fit log-normal distributions to the active (top- $K$  masked) excitatory synaptic weights after training on MNIST and Fashion-MNIST. Table S24 shows the fitted  $\sigma$  (standard deviation of  $\log w$ ) for each condition. Shunting consistently produces narrower distributions than additive integration, and backpropagation produces narrower distributions than local learning. These two factors combine approximately additively: BP+Shunting  $\rightarrow$  narrowest, Local+Additive  $\rightarrow$  broadest.

**Shunting is closest to the biological references considered here.** BP+Shunting produces  $\sigma = 0.94$  on MNIST, very close to the Song et al. reference ( $\Delta = 0.008$ ). Local+Shunting produces  $\sigma = 1.13$ – $1.22$ , close to MICrONS E $\rightarrow$ E synapse sizes ( $\sigma = 1.14$ ;  $\Delta \leq 0.08$ ). By contrast, Local+Additive produces  $\sigma = 1.43$ – $1.69$ , broader than any biological reference considered. This pattern is consistent across datasets (MNIST and Fashion-MNIST), suggesting it reflects structural properties of the learning rule rather than task-specific features.

**The match depends on inhibitory conductance.** We test sensitivity to the E/I synapse ratio using a sweep over  $N_I \in \{0, 2, 5, 10, 20, 30, 40\}$  inhibitory synapses per branch (Appendix Fig. S22). Without inhibition ( $N_I=0$ ), all conditions produce  $\sigma > 2.0$ , far from biology. With biologically constrained ratios ( $N_I=10$ – $20$ , corresponding to E/I ratios of 4:1 to 2:1; [42]), shunting networks enter the biological range ( $\sigma \approx 0.9$ – $1.2$ ), while the effect of  $N_I$  on additive networks is less systematic. The weight distribution shape is also robust to dendritic depth (branch factors [9] through [3,3,3,3]; see Appendix Fig. S22A): BP+Shunting converges toward Song et al. with increasing depth, while Local+Shunting remains near MICrONS E $\rightarrow$ E across depths.

## H Weight Distribution Sensitivity Analysis

**Side-by-side BP vs. LocalCA across depths.** To complement the per-architecture analysis above, we ran a matched BP vs. LocalCA sweep over three branch-factor settings ([2, 2], [3, 3], [3, 3, 3]) under both shunting and additive cores on MNIST (5 seeds per cell; quantile init; learned ( $b, m$ ) for LocalCA). Figure S23 shows that at shallow and medium depths the local rule and backprop reach broadly similar excitatory weight means ( $\approx 0.8$ – $0.9$ ), while the coefficient of variation (CV) under LocalCA is roughly twice as large as under BP at these depths (e.g. shunting [2, 2]:  $CV_{BP} = 0.15$  vs.  $CV_{Loc.} = 0.31$ ; shunting [3, 3]:  $0.11$  vs.  $0.26$ ). At the deepest setting [3, 3, 3] the weight means drop and the CV jumps substantially for both optimizers ( $CV \approx 0.66$ – $0.88$ ), and test accuracy drops by  $\sim 5$  pp for BP and  $\sim 15$ – $20$  pp for LocalCA. Across the explored depth range, shunting LocalCA stays closer to shunting BP than additive LocalCA does to additive BP on both weight and accuracy axes, consistent with the manuscript claim that the conductance denominator is the key mechanistic ingredient.

**Weight distributions under soma-on.** To check whether the direct somatic synapse extension (Fig. S18) alters the weight-distribution picture, we ran a matched 5-seed sweep at depth [3, 3] on MNIST with `somatic_synapses=true` across all four (core, strategy) cells; results are summarized in Fig. S24. Two effects emerge. First, shunting + BP becomes *tighter* than in the soma-off case: excitatory-weight CV drops from  $\approx 0.11$  (soma-off, [3, 3]) to  $\approx 0.10$  (soma-on), consistent with the soma sink absorbing part of the drive and reducing per-branch variance. Second, shunting + LocalCA goes in the opposite direction: CV grows from  $\approx 0.26$  (soma-off) to  $\approx 0.40$  (soma-on), while MNIST test accuracy drops  $\sim 3$  pp (from 96.5% to 93.6%). Additive cores sit between these extremes ( $CV \approx 0.23$ – $0.30$ , accuracy 95.7–97.0%) and are only weakly affected. The ordering is consistent with the mechanism in the main text: the somatic synapse is an additional BP-visible path that backprop exploits cleanly, but under the local rule it adds broadcast variance that the scalar/per-soma signal cannot sharpen, widening the weight distribution without a matching accuracy gain.

**MICrONS connectome analysis.** We extracted synapse size distributions from the MICrONS mm<sup>3</sup> condensed connectome (v1181; 9.0M connections, 71.9K neurons; 89% excitatory). Splitting by pre/post cell type: E $\rightarrow$ E connections ( $n=3.85$ M) show  $\sigma_{size}=1.14$  and  $CV=1.01$ ; I $\rightarrow$ E ( $n=3.53$ M) show  $\sigma=0.92$  and  $CV=1.03$ ; E $\rightarrow$ I ( $n=1.11$ M) show  $\sigma=0.83$ ; I $\rightarrow$ I ( $n=0.51$ M) show  $\sigma=0.85$ . Inhibitory connections are more frequently multi-synaptic (I $\rightarrow$ E: 37.5% multi-synapse; E $\rightarrow$ E: 11.5%), and inhibitory synapses are on average smaller (median 3532 vs. 4676 for E $\rightarrow$ E).




Figure S22: **Synaptic weight distribution sensitivity to architecture.** (A) Log-normal  $\sigma$  vs. dendritic depth (branch factors [9] to [3,3,3,3]). Both backprop and local shunting remain closer to the biological-width references than their additive counterparts across depth. (B)  $\sigma$  vs. inhibitory synapses per branch ( $N_I$ ). Without inhibition ( $N_I=0$ ), all conditions are overly broad; with biologically constrained  $N_I=10$ –20, shunting enters the biological range. Shaded bands show the sweep over  $N_E=20$  to 80. (C) Summary of MNIST and Fashion-MNIST results with biological references. The shunting ordering is consistent across datasets. (D) Excitatory vs. inhibitory weight  $\sigma$  on MNIST with MICrONS E→E and I→E reference bands.

794 **Sensitivity to ee\_synapses.** Using the E/I grid sweep ( $N_E \in \{20, 40, 80\}$  and  $N_I \in$   
795  $\{0, 2, 5, 10, 20, 30, 40\}$ ; 168 models), excitatory  $\sigma$  increases modestly with  $N_E$  in both architec-  
796 tures, but the shunting-versus-additive ordering is preserved at every  $N_E$  value. The dominant driver  
797 of  $\sigma$  is still  $N_I$ : moving from  $N_I=0$  to  $N_I=10$  reduces shunting  $\sigma$  by roughly 1.0–1.5 units, while  
798 the effect on additive is smaller and less systematic.

799 **Sensitivity to depth.** From the depth-scaling sweep ([9], [3,3], [3,3,3], [3,3,3,3]; 80 models, 5 seeds  
800 each), BP+Shunting excitatory  $\sigma$  decreases with depth (1.06  $\rightarrow$  0.90), converging toward Song  
801 et al. (0.94). Local+Shunting  $\sigma$  increases slightly (1.18  $\rightarrow$  1.33) but remains in the biological  
802 range at shallow depths. BP+Additive shows the strongest depth sensitivity (1.38  $\rightarrow$  0.79), while  
803 Local+Additive diverges with depth (2.01  $\rightarrow$  2.44). The morphology-aware modulators (path  
804 propagation, depth/centrality weighting) have negligible effect on  $\sigma$ .

805 **Mechanistic interpretation.** The shunting conductance  $g_n^{\text{tot}} = g_n^{\text{leak}} + \sum g_n^{\text{exc}} + \sum g_n^{\text{inh}}$  appears  
806 in the update factor  $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$ . As inhibitory conductance increases,  $g_n^{\text{tot}}$  grows and  $R_n^{\text{tot}}$   
807 shrinks, compressing the multiplicative scale of weight updates. This is analogous to a learning rate  
808 that self-adjusts via the total conductance: strong synapses contribute more to  $g_n^{\text{tot}}$ , reducing their  
809 own future update magnitude. This acts as a natural form of weight-dependent regularization. The  
810 resulting equilibrium distribution is log-normal, consistent with Loewenstein et al. [41] and related  
811 multiplicative spine-dynamics analyses. Additive integration lacks this self-normalizing mechanism,  
812 producing broader distributions that fall outside the biological range.




Figure S23: **Synaptic weight statistics vs. dendritic depth (MNIST, 5 seeds/cell).** (A) Excitatory weight mean (error bars = std). At [2, 2] and [3, 3], all four conditions cluster around 0.8–0.9; at [3, 3, 3] the means drop to  $\approx 0.5$ –0.75 and the cross-seed spread grows, especially for additive LocalCA. (B) Coefficient of variation (CV = std/mean). At shallow/medium depth, LocalCA produces  $\sim 2\times$  higher CV than matched BP, but both remain well-behaved ( $\leq 0.33$ ). At the deepest network all four conditions jump to CV  $\approx 0.66$ –0.88, with additive LocalCA the broadest. (C) MNIST test accuracy. BP tracks  $\sim 96$ –97% at [2, 2]/[3, 3] and  $\sim 92$ –95% at [3, 3, 3]. LocalCA stays within  $\sim 5$ –6 pp of BP at shallow/medium depth but trails by  $\sim 15$ –20 pp at [3, 3, 3], identifying depth as the remaining open challenge for this local rule. The qualitative ordering (shunting LocalCA tracks shunting BP more closely than additive LocalCA tracks additive BP) reinforces the manuscript claim that the conductance denominator is the key mechanistic ingredient.




Figure S24: **Excitatory-weight CV and MNIST accuracy under soma-off vs. soma-on (depth [3, 3], 5 seeds/cell).** (A) Coefficient of variation of excitatory weights for the four (core, strategy) cells; paired bars show soma-off (from Fig. S23) and soma-on (this sweep). Shunting + BP tightens slightly; shunting + LocalCA broadens; additive cells change little. (B) Matched MNIST test accuracy (sanity panel). BP accuracy is preserved or marginally improved; LocalCA accuracy is preserved for additive but drops  $\sim 3$  pp for shunting, consistent with the CV widening in panel A.

## 813 I Synthetic Benchmark Descriptions

814 Each synthetic task is generated procedurally from a fixed random seed and uses fixed train/val/test  
815 splits for reproducibility. The MNIST-derived synthetic tasks use flattened  $28 \times 28$  inputs; cue  
816 integration instead uses a small vector input with two output classes. Output dimensionality therefore  
817 varies by task and is stated in the task description when it differs from 10-way digit classification.  
818 Figure S25 gives a visual summary of each benchmark.

819 **Context gating.** Inputs are contextual figure-ground MNIST examples. The right half of the image  
820 carries the original MNIST digit, while the left half is replaced by low-amplitude uniform noise; all  
821 pixels therefore remain nonnegative. This task tests whether credit assignment can exploit a simple  
822 context-dependent spatial structure without leaving the nonnegative-drive regime of the conductance  
823 model. The HSIC auxiliary objective (weight 0.01) is used to decorrelate layer representations,  
824 adding roughly 3 pp under local learning.




Figure S25: **Synthetic benchmark schematics.** (A) Context gating: the right half carries the MNIST digit; the left half is filled with low-amplitude uniform noise. (B) Noise resilience: MNIST image corrupted by structured channel noise (fixed 50-D latent projection,  $\sigma_{\text{task}} = 1.5$ , clipped to  $[0, 1]$ ). (C) Cue integration: a context cue selects which of two noisy cue streams is reliable; dendritic integration routes the reliable cue to the binary readout. (D) Fashion-MNIST: 10-class apparel classification sharing the MNIST input shape.

**Noise resilience.** Inputs are flattened MNIST images corrupted by structured channel noise: a fixed 50-dimensional latent Gaussian perturbation is projected into input space and added to each example, with  $\sigma_{\text{task}} = 1.5$ , then clipped to  $[0, 1]$ . The projection is fixed across training, so the network must filter correlated interference to recover the underlying class signal. Credit assignment must remain coherent under the corrupted signal, making this task particularly sensitive to path-gain dispersion.

**Info shunting.** Inputs contain two additive streams: a low-frequency signal and a high-frequency distractor drawn with equal energy, together with a separate context channel indicating which stream is relevant on each trial. This benchmark uses signed synthetic inputs, so it violates the nonnegative presynaptic-drive assumption enforced in the main experiments. We therefore retain it only as an exploratory stress test and exclude it from the main-text claims.

**Cue integration.** Two noisy cue streams ( $A$  and  $B$ ) are presented simultaneously, each carrying partial information about one of 2 classes. A binary one-hot context signal indicates which cue is reliable on each trial (cue  $A$  or cue  $B$ ), and each cue vector is clipped to  $[0, 1]$  after noise is added. The fixed-pathway variant duplicates the context into both cue branches; the learned-routing variant requires the router to discover the cue separation from data. This task is used in Sec. A to test structured pathway-vector broadcast while staying inside the nonnegative-input regime.

## J Depth Scaling and Noise Robustness

**Depth scaling.** Varying dendritic depth from 1–4 layers (branch factors  $[9]$  to  $[3, 3, 3, 3]$ ): shunting local degrades from 63.5% to 57.4%; additive local degrades more steeply from 54.9% to 29.7%. The shunting advantage grows with depth ( $+8.5 \rightarrow +27.7$  pp). Backprop ceilings remain stable at about 90–92%. See Fig. S7A.

**Noise robustness.** Gaussian noise  $\mathcal{N}(0, \sigma^2)$  on broadcast error: shunting is robust to  $\sigma \leq 0.1$  (about 62%) while additive drops from 46.5% to chance at  $\sigma=1.0$ , confirming that shunting credit signals carry useful learning information beyond the broadcast magnitude. See Fig. S7B.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The abstract and introduction state the paper’s main mechanistic claim, the scope of the conductance-based setting, and the regime-dependent nature of the empirical results; see Sec. 1 and Sec. 5.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

## 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The discussion explicitly narrows the scope of the claims and discusses limitations involving benchmark strength, broadcast assumptions, depth scaling, and the exploratory stress tests; see Sec. 5.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

## 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The voltage model, assumptions, theorem statements, and proof details are given in Sec. 2 and Appendix C; proof sketches and the exact factorization sanity check are cross-referenced in the main text.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The main paper and appendices specify the datasets, architectures, local-rule family, broadcast modes, representative hyperparameters, and compute regime needed to reproduce the headline results; see Sec. 4 and Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

#### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

962 Answer: [No]

963 Justification: To preserve double-blind review, the anonymous submission does not include

964 a public repository or anonymized code bundle. The paper instead provides architectural,

965 algorithmic, and configuration details and states that code and configs will be released after

966 review; see Appendix B.

967 Guidelines:

- 968 • The answer [N/A] means that paper does not include experiments requiring code.
- 969 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 970
- 971 • While we encourage the release of code and data, we understand that this might not
- 972 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
- 973 including code, unless this is central to the contribution (e.g., for a new open-source
- 974 benchmark).
- 975 • The instructions should contain the exact command and environment needed to run to
- 976 reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- 977
- 978 • The authors should provide instructions on data access and preparation, including how
- 979 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 980 • The authors should provide scripts to reproduce all experimental results for the new
- 981 proposed method and baselines. If only a subset of experiments are reproducible, they
- 982 should state which ones are omitted from the script and why.
- 983 • At submission time, to preserve anonymity, the authors should release anonymized
- 984 versions (if applicable).
- 985 • Providing as much information as possible in supplemental material (appended to the
- 986 paper) is recommended, but including URLs to data and code is permitted.

987 **6. Experimental setting/details**

988 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-

989 rameters, how they were chosen, type of optimizer) necessary to understand the results?

990 Answer: [Yes]

991 Justification: The main text states the benchmark families, seed protocol, representative

992 local-rule settings, and major training choices, while the appendices provide implementation

993 details and operating-mode inventories; see Sec. 4 and Appendix B.

994 Guidelines:

- 995 • The answer [N/A] means that the paper does not include experiments.
- 996 • The experimental setting should be presented in the core of the paper to a level of detail
- 997 that is necessary to appreciate the results and make sense of them.
- 998 • The full details can be provided either with the code, in appendix, or as supplemental
- 999 material.

1000 **7. Experiment statistical significance**

1001 Question: Does the paper report error bars suitably and correctly defined or other appropriate

1002 information about the statistical significance of the experiments?

1003 Answer: [Yes]

1004 Justification: The paper reports means and standard deviations across 5 seeds for the headline

1005 results and across 3 seeds for selected appendix-only support sweeps, states this explicitly

1006 in the text, and shows error bars or seed variability in the relevant tables and figures; see

1007 Sec. 4 and Appendix A.

1008 Guidelines:

- 1009 • The answer [N/A] means that the paper does not include experiments.
- 1010 • The authors should answer [Yes] if the results are accompanied by error bars, confidence
- 1011 intervals, or statistical significance tests, at least for the experiments that support the
- 1012 main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: Appendix B describes the institutional GPU cluster setting, the use of single A100/H100-class GPUs, rough per-run runtime scales, and the fact that broader exploratory sweeps required additional compute beyond the final reported runs.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses standard public benchmarks and mechanistic simulations, does not involve human-subjects experimentation or restricted data collection, and we are not aware of any aspect that violates the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The discussion includes a broader-relevance and potential-harms paragraph that distinguishes mechanistic neuroscience relevance from possible misuse through overclaiming biological plausibility or overstating robustness; see Sec. 5.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The submission does not release a high-risk generative model, foundation model, or scraped dataset, so the safeguard question is not applicable here.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The paper credits all benchmark and connectomic assets, does not redistribute third-party raw assets, and includes an asset/access table with license or access terms where available; see Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, `paperswithcode.com/datasets` has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: The paper does not release a new dataset, benchmark, or pretrained model asset in this anonymous submission.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The work does not involve crowdsourcing or direct research with human subjects.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

1169 Answer: [N/A]  
 1170 Justification: The work does not involve crowdsourcing or human-subjects research requiring  
 1171 IRB review.  
 1172 Guidelines:  
 1173 • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
 1174 with human subjects.  
 1175 • Depending on the country in which research is conducted, IRB approval (or equivalent)  
 1176 may be required for any human subjects research. If you obtained IRB approval, you  
 1177 should clearly state this in the paper.  
 1178 • We recognize that the procedures for this may vary significantly between institutions  
 1179 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
 1180 guidelines for their institution.  
 1181 • For initial submissions, do not include any information that would break anonymity (if  
 1182 applicable), such as the institution conducting the review.  
 1183 **16. Declaration of LLM usage**  
 1184 Question: Does the paper describe the usage of LLMs if it is an important, original, or  
 1185 non-standard component of the core methods in this research? Note that if the LLM is used  
 1186 only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
 1187 scientific rigor, or originality of the research, declaration is not required.  
 1188 Answer: [N/A]  
 1189 Justification: LLMs are not part of the proposed method, theorem, experiments, or analysis  
 1190 pipeline, so no declaration is required for the core scientific contribution.  
 1191 Guidelines:  
 1192 • The answer [N/A] means that the core method development in this research does not  
 1193 involve LLMs as any important, original, or non-standard components.  
 1194 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not  
 1195 be described.

| Setting                                                    | Encoder  | E layers           | Tree                                         | Synapses / branch                                          | Seeds                | Notes                                                                                                                                         |
|------------------------------------------------------------|----------|--------------------|----------------------------------------------|------------------------------------------------------------|----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| Gradient fidelity / path-gain / noise-resilience mechanism | id.+ReLU | [128, 128]         | [3, 3]                                       | $N_E=40, N_I \in \{0, 5, 10, 20, 40\}$                     | 5 fidelity, 5 oracle | Main mechanism sweeps in Figs. 2 and 3; reactivation <code>param_tanh</code> .                                                                |
| MNIST / context-gating local competence                    | id.+ReLU | [128]              | [3, 3]                                       | $N_E=40, N_I=20$                                           | 5                    | Best local 5F per-soma runs in Table S4; reactivation <code>param_tanh</code> .                                                               |
| Fashion-MNIST competence                                   | id.+ReLU | [128]              | [3, 3]                                       | $N_E=40, N_I=20$                                           | 5                    | Dedicated five-seed standard-vs-local competence sweep; reactivation <code>param_tanh</code> .                                                |
| Cue-routing PV-LocalCA                                     | router   | [64]               | [2]                                          | $N_E=10, N_I=4$                                            | 5                    | Learned router with two pathway groups and rank-2 structured feedback; reactivation <code>param_tanh</code> .                                 |
| Positive cue-routing inhibition controls                   | router   | [64]               | [2]                                          | $N_E=10, N_I=4$ direct I-to-E or no I-to-E                 | 5 per condition      | Strictly nonnegative router outputs; learned-I vs. no-I controls for rank-1, rank-2, pathway-vector, path transport, and BP.                  |
| Double-CIFAR routed screen                                 | router   | [256]              | [2]                                          | $N_E=20, N_I=6$ direct I-to-E or no I-to-E                 | 5 per condition      | Strictly nonnegative routed CIFAR streams; retained as a boundary screen because matched BP remains near chance.                              |
| CIFAR-10 compact direct-I-stream control ladder            | id.+ReLU | [20] E, no I cells | [3, 3, 3, 3]                                 | $N_E=25, N_I=25$ direct I-to-E, with matched no-I controls | 5 per condition      | Harder-data extension with input-driven inhibitory conductance, decoder [32, 16], decoder-aware soma mapping, and additive fairness controls. |
| Morphology $\times$ inhibition appendix map                | id.+ReLU | [128, 128]         | [2, 2], [3, 3], [4, 4], [2, 2, 2], [3, 3, 3] | $N_E=40, N_I \in \{0, 5, 10, 20, 40\}$                     | 3                    | Supportive 3-seed sweep; reactivation <code>param_tanh</code> .                                                                               |

Table S16: **Representative architectures used in the manuscript.** All dendritic rows use explicit branch-level excitatory and inhibitory synapse banks with `somatic_synapses=false`, nonnegative conductances, and a nonnegative first-layer drive enforced by an explicit ReLU transfer stage in the main runs. The tree column gives dendritic morphology per excitatory neuron, and the synapse counts report excitatory and inhibitory synapses per branch.

| Setting                                                      | Opt. | LR                                            | Batch | Epochs    | W-dec.   | Notes                                                            |
|--------------------------------------------------------------|------|-----------------------------------------------|-------|-----------|----------|------------------------------------------------------------------|
| MNIST / Fashion-MNIST / context gating (5F per-soma LocalCA) | Adam | $10^{-3} / 5 \cdot 10^{-4}$ (block / react.)  | 256   | 100       | 0        | fixed-epoch, no early stopping; param_tanh reactivation          |
| MNIST / Fashion-MNIST / context gating (BP ceiling)          | Adam | $10^{-3} / 5 \cdot 10^{-4}$                   | 256   | 100       | 0        | matched to LocalCA setup                                         |
| Gradient fidelity / path-gain / noise resilience             | Adam | $10^{-3}$                                     | 256   | 50        | 0        | 5 seeds, hooks enabled for diagnostic capture                    |
| Morphology $\times$ inhibition regime map (exploratory)      | Adam | $10^{-3}$                                     | 256   | 50        | 0        | 3 seeds                                                          |
| Cue routing (soma-off & soma-on)                             | Adam | $10^{-3}$                                     | 256   | 90        | 0        | early stopping with patience 30                                  |
| Positive cue-routing inhibition controls                     | Adam | $1.5 \cdot 10^{-3} / 7 \cdot 10^{-4}$ (block) | 256   | 140       | 0        | early stopping with patience 35; 50 complete runs                |
| Double-CIFAR routed screen                                   | Adam | $1.2 \cdot 10^{-3} / 6 \cdot 10^{-4}$ (block) | 256   | 220       | 0        | no early stopping; 70 complete runs; retained as boundary screen |
| CIFAR-10 compact direct-I-stream depth-4 LocalCA             | Adam | $10^{-3} / 10^{-4}$ (block / react.)          | 256   | 400       | 0 / 0.01 | grad-clip 5.0, early stop patience 50, nonlinear decoder         |
| CIFAR-10 compact direct-I-stream depth-4 BP                  | Adam | $10^{-3}$                                     | 256   | 200       | 0.01     | early stop patience 40, matched ceiling                          |
| Weight-distribution vs. depth                                | Adam | $10^{-3}$                                     | 256   | 100       | 0        | weight_analysis hook captures final-epoch statistics             |
| Component ablation (MNIST)                                   | Adam | $10^{-3}$                                     | 256   | 100       | 0        | 5 seeds, single-layer dendrinet [3, 3]                           |
| ( $b, m$ ) policy comparison (MNIST / CIFAR-10)              | Adam | $10^{-3} / 10^{-4}$                           | 256   | 100 / 400 | 0 / 0.01 | quantile refresh every 5 epochs; EMA rate in {0.25, 0.5, 1.0}    |

Table S17: **Training hyperparameters by experiment.** All runs use Adam with default  $\beta_1=0.9$ ,  $\beta_2=0.999$ ,  $\varepsilon=10^{-8}$ . Reactivation and block-linear parameter groups use a slower LR to keep conductance-level dynamics stable. Full configs are provided under drafts/dendritic-local-learning/configs/.

| Setting                                          | Hidden width(s) | Output dim           | Effective default per-soma field                                                                              |
|--------------------------------------------------|-----------------|----------------------|---------------------------------------------------------------------------------------------------------------|
| Gradient fidelity / path-gain / noise resilience | 128, 128        | 10                   | Scalar / rank-1 at both dendritic layers                                                                      |
| MNIST / Fashion-MNIST competence                 | 128             | 10                   | Scalar / rank-1 at the dendritic hidden layer                                                                 |
| Context gating competence                        | 128             | 2                    | Scalar / rank-1 at the dendritic hidden layer                                                                 |
| Cue routing                                      | 64              | 2                    | Scalar / rank-1 by default; higher-rank tests use explicit low-rank or pathway-vector feedback                |
| Compact CIFAR-10 harder-data extension           | 20              | decoder-input dim 20 | Vector-valued per-soma after decoder-aware mapping; higher-rank tests use explicit low-rank or path transport |

Table S18: **Effective dimensionality of the default “per-soma” broadcast in the headline architectures.** In code, per-soma uses the output error vector only when layer width matches decoder output dimension; otherwise it falls back to a shared scalar. The main feedforward settings in this paper therefore test a genuinely rank-1 broadcast regime by default, which is why the higher-rank low-rank, path-transport, and pathway-vector controls are important.

| Extension              | Meaning                                                              | Compatible with | Typical use here                                                                               |
|------------------------|----------------------------------------------------------------------|-----------------|------------------------------------------------------------------------------------------------|
| Per-soma broadcast     | Somatic error vector when widths match, otherwise scalar fallback    | 3F, 4F, 5F      | Default main-text LocalCA setting; in the headline feedforward runs this is effectively rank-1 |
| Low-rank broadcast     | Small number of broadcast channels instead of a scalar or full field | 3F, 4F, 5F      | Intermediate control when rank-1 is too restrictive                                            |
| Path transport         | Recursive tree-transported error using conductance-stage path gains  | 3F, 4F, 5F      | Oracle upper-bound broadcast                                                                   |
| Path propagation       | Approximate path-gain attenuation applied to broadcast               | 3F, 4F, 5F      | Morphology-aware modifier for deeper trees                                                     |
| Pathway vector         | Structured branch/pathway-specific broadcast field                   | 3F, 4F, 5F      | Used mainly in PV-LocalCA on routed tasks                                                      |
| Branch-role rules      | Scales updates by router- or pathway-derived branch specialization   | 3F, 4F, 5F      | Used in the structured pathway extension                                                       |
| Inhibitory homeostasis | Auxiliary local inhibitory balancing term                            | 3F, 4F, 5F      | Optional stabilizer in routed or shunting settings                                             |

Table S19: Compatibility of the newer LocalCA extensions with the base 3F/4F/5F family. These extensions mostly modify the error field or branch scaling rather than defining a new “6F” or “7F” rule.

| Figure / Table             | Sweep                                                                  | Replication                                                                                                      | Tier                    |
|----------------------------|------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|-------------------------|
| Fig. 1<br>Fig. 2           | Model, rule, and broadcast definitions<br>Gradient-fidelity diagnostic | N/A<br>3 checkpointed seeds<br>for dynamics;<br>exact-factorization<br>summary over all<br>diagnostic conditions | Definitions<br>Headline |
| Fig. S4                    | Alignment dynamics and<br>nonzero-gradient check                       | 3 seeds                                                                                                          | Supportive              |
| Fig. 3                     | Mechanistic chain + oracle + summary                                   | 5 seeds                                                                                                          | Headline                |
| Fig. 4                     | Competence / regime dose-response                                      | 5 seeds                                                                                                          | Headline                |
| Fig. S10 / Table S6        | CIFAR-10 compact direct-I-stream<br>control ladder                     | 5 seeds per condition;<br>70 complete runs                                                                       | Supportive              |
| Fig. S17                   | Cue-routing (soma off)                                                 | 5 seeds                                                                                                          | Supportive              |
| Table S4                   | Competence gaps (BP ceiling vs. 5F)                                    | 5 seeds                                                                                                          | Headline                |
| Fig. S18                   | Soma-on extension across manuscript<br>families                        | Family aggregates<br>(144–420 cfgs)                                                                              | Supportive              |
| Table S13                  | Soma-on delta summary across<br>manuscript families                    | Family aggregates                                                                                                | Supportive              |
| Fig. S20                   | $(b, m)$ update-policy comparison                                      | Matched 4-cell grid /<br>bar                                                                                     | Supportive              |
| Fig. S21                   | Component ablation (MNIST)                                             | 5 seeds                                                                                                          | Supportive              |
| Fig. S19                   | Cue-routing soma-on (stable subset)                                    | 5 seeds                                                                                                          | Supportive              |
| Table S9                   | Fitted direct-I inhibitory path-gain<br>diagnostic                     | 5 runs                                                                                                           | Supportive              |
| Fig. S15                   | Positive cue controls and routed<br>Double-CIFAR screen                | 50 cue runs; 70<br>Double-CIFAR runs                                                                             | Supportive / boundary   |
| Fig. S16                   | Inhibition interventions and exact-error<br>rank                       | 5 seeds                                                                                                          | Supportive              |
| Fig. S22                   | Weight-distribution sensitivity                                        | 5 seeds                                                                                                          | Supportive              |
| Fig. S23                   | Weight-distribution vs. depth                                          | 5 seeds                                                                                                          | Supportive              |
| Fig. S24                   | Weight-distribution soma-off vs. soma-on                               | 5 seeds                                                                                                          | Supportive              |
| Fig. S11 / Table S7        | CIFAR-10 depth-4 soma extension                                        | 3 seeds                                                                                                          | Supportive              |
| Fig. S9                    | FA/DFA baseline comparison                                             | 5 seeds                                                                                                          | Supportive              |
| Fig. S8                    | Broadcast-bandwidth / quantization                                     | 5 seeds                                                                                                          | Supportive              |
| Fig. S12                   | Additive + normalization control                                       | 5 seeds                                                                                                          | Supportive              |
| Fig. S7 / S6               | Stress tests; verification                                             | 5 seeds                                                                                                          | Supportive              |
| Fig. S5                    | 5F sensitivity                                                         | 3 seeds                                                                                                          | Supportive              |
| Table S10                  | Rank-bridge                                                            | 5 seeds                                                                                                          | Supportive              |
| Fig. S13                   | Morphology $\times$ inhibition map                                     | 3 seeds                                                                                                          | Supportive              |
| Fig. S14 / Table S8        | One-layer direct-I vs. explicit-I<br>input-mode probe                  | 3 seeds                                                                                                          | Supportive              |
| Signed-input info-shunting | Signed-input stress test                                               | 1–3 seeds                                                                                                        | Exploratory             |

**Table S20: Evidence map.** Each main-paper and appendix figure or table is pinned to its sweep and replication type. Most main claims come from dedicated 5-seed sweeps; the supporting soma-extension and  $(b, m)$  policy appendices instead summarize matched sweep families or matched config grids, and are labeled accordingly. The BP ceilings in Table S4 come from dedicated matched-capacity 5-seed standard-training refreshes.

| Asset           | Use in this paper                                                   | Access / license note                                                                                                                        |
|-----------------|---------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| MNIST           | Digit classification and derived nonnegative synthetic tasks        | Public academic benchmark credited to the original dataset authors; used through standard public access, with no raw-data redistribution.    |
| Fashion-MNIST   | Apparel classification stress test                                  | Public benchmark released with the Fashion-MNIST paper and repository; used through standard public access, with no raw-data redistribution. |
| CIFAR-10        | Flattened harder-data diagnostic                                    | Public research benchmark from the CIFAR dataset collection; used through standard public access, with no raw-data redistribution.           |
| MICrONS         | Biological reference statistics in the weight-distribution appendix | Public connectomics resource cited to the MICrONS Consortium; only aggregate comparison statistics are reported here.                        |
| Synthetic tasks | Context gating, noise resilience, cue integration                   | Generated procedurally from public benchmarks or random seeds described in Appendix I; no new third-party asset is introduced.               |

Table S21: **Existing assets used in the manuscript.** We credit the original sources in the bibliography, use public benchmark assets under their published access conditions, and do not redistribute third-party raw data in the anonymous submission.

| Rule       | Factors                                | Cost               | Best regime                   |
|------------|----------------------------------------|--------------------|-------------------------------|
| 3F         | $x, (E - V), e$                        | $\mathcal{O}(1)$   | Baseline                      |
| 4F         | $3F + r^{4F}$                          | $\mathcal{O}(1)$   | Instructive attenuation proxy |
| 5F         | $4F + \phi$                            | $\mathcal{O}(d_n)$ | Strongest overall             |
| 5F+VH      | 5F + local voltage homeostasis         | $\mathcal{O}(d_n)$ | Stabilized LocalCA / warm-up  |
| PV-LocalCA | 3F + structured $e_n$ + role alignment | $\mathcal{O}(r_n)$ | Pathway discovery             |

Table S22: Variant taxonomy. ‘5F+VH’ denotes the same base 5F rule with an additional strictly local voltage-centering auxiliary that can also be scheduled through multi-stage LocalCA warm-up; the codebase also supports this as a shorthand ‘5f\_vh’ alias in configuration files.

| Component          | Analog                         | Interpretation                         |
|--------------------|--------------------------------|----------------------------------------|
| $R_n^{\text{tot}}$ | Input resistance               | Sensitivity modulation                 |
| $(E_j - V_n)$      | Synaptic driving force         | Local gradient factor                  |
| Shunting           | Divisive normalization         | $\partial V / \partial g_I \propto -V$ |
| $r_n^{4F}$         | Layer relevance                | Output correlation                     |
| $\phi_n$           | Signal propagation             | Conditional predictability             |
| Pathway vector     | Branch-specific top-down drive | Structured teaching signal             |
| Role alignment     | Pathway-selective plasticity   | Branch specialization                  |

Table S23: Biological analogs.

| Condition                 | MNIST                                                      |                 | Fashion-MNIST   |                 |
|---------------------------|------------------------------------------------------------|-----------------|-----------------|-----------------|
|                           | Exc $\sigma$                                               | Inh $\sigma$    | Exc $\sigma$    | Inh $\sigma$    |
| BP + Shunting             | 0.94 $\pm$ 0.01                                            | 0.87 $\pm$ 0.01 | 0.92 $\pm$ 0.01 | 0.84 $\pm$ 0.01 |
| BP + Additive             | 1.30 $\pm$ 0.01                                            | 1.22 $\pm$ 0.01 | 1.16 $\pm$ 0.01 | 1.05 $\pm$ 0.01 |
| Local + Shunting          | 1.22 $\pm$ 0.01                                            | 1.71 $\pm$ 0.02 | 1.13 $\pm$ 0.06 | 1.36 $\pm$ 0.06 |
| Local + Additive          | 1.69 $\pm$ 0.16                                            | 1.57 $\pm$ 0.07 | 1.43 $\pm$ 0.16 | 1.11 $\pm$ 0.12 |
| Song et al. 2005          | $\sigma = 0.94$ (rat V1 L5, EPSP amplitude)                |                 |                 |                 |
| MICrONS E $\rightarrow$ E | $\sigma = 1.14$ (mouse V1, synapse size, $n=3.9\text{M}$ ) |                 |                 |                 |
| MICrONS I $\rightarrow$ E | $\sigma = 0.92$ (mouse V1, synapse size, $n=3.5\text{M}$ ) |                 |                 |                 |

Table S24: **Weight distribution width.** Log-normal  $\sigma$  of excitatory and inhibitory synaptic weights (mean $\pm$ std across 5 seeds). Biological references are included only as descriptive anchors. With biologically constrained E/I ratios (4:1), shunting self-normalization produces weight heterogeneity in the range spanned by those references, whereas additive integration is broader and less stable.