
Shunting Inhibition Improves Local Credit Assignment in Dendritic Networks

Anonymous Author(s)

Affiliation

Address

email

Abstract

Conductance-based shunting inhibition, long studied as a gain-control mechanism, also creates favorable conditions for local credit assignment in dendritic networks. Starting from steady-state conductance equations, we derive exact gradients for compartmental dendritic trees and show that each synaptic gradient factorizes into a synapse-local eligibility term and a single non-local compartment error. This factorization motivates a family of local learning rules (3F/4F/5F) that combine synapse-local variables with a somatic error broadcast. On matched MNIST architectures, shunting raises the local-vs-backprop cosine from 0.006 to 0.202 and lowers scale mismatch from 1.053 to 0.117, corresponding to about $30\times$ relative improvement in direction and about $10\times$ lower scale distortion. The advantage is regime-dependent, growing with inhibitory conductance and becoming largest on tasks that require robust credit signals under noise. On a hard cue-integration benchmark where the model must discover separate dendritic pathways, per-soma broadcast remains unstable under shunting local learning, but a hierarchical pathway-vector teaching signal with router-derived branch-role plasticity stabilizes performance at $99.1 \pm 0.3\%$ test accuracy over 5 seeds. Together, these results identify divisive normalization as a useful mechanism for improving local credit assignment in conductance-based dendritic networks.

1 Introduction

Credit assignment in deep networks usually depends on backpropagation: global error transport through symmetric pathways with no obvious biological substrate. Dendritic compartments suggest an alternative. Each synapse has access to rich local state such as driving force, membrane conductance, and branch voltage, while supervisory information could arrive as a low-bandwidth broadcast from the soma [15]. The question in this paper is whether such local information is sufficient for useful learning.

We show that it can be, but only in a specific biophysical regime. Starting from conductance-based dendritic voltage equations [1], we derive exact gradients for dendritic trees (Theorem 1) and show that each synaptic gradient factorizes into purely local terms and a single non-local compartment error. Approximating that compartment error with a broadcast signal yields a family of local learning rules (3F, 4F, and 5F) that use only branch-local variables together with the broadcast. For tasks where different branches must be credited differently, we further introduce a structured pathway-vector signal that preserves latent pathway identity during credit assignment.

The central empirical finding is that these local rules work much better with *shunting inhibition* than with an additive control. In the conductance model, synaptic sensitivity depends on both driving force




Figure 1: **Model and credit assignment.** (A) A compartmental dendritic neuron with excitatory (blue triangles) and inhibitory (red circles; shunting) synapses. Voltages propagate to the soma through learned dendritic conductances (green arrows). A somatic broadcast error e_n (dashed red) reaches all compartments; each synapse combines e_n with local quantities (presynaptic drive x_j , input resistance R_n^{tot} , and driving force $E_j - V_n$). (B) Compact LocalCA hierarchy: 3F uses local eligibility and a broadcast field; 4F adds the morphology-aware factor ρ ; 5F further adds the confidence-like factor ϕ . (C) Broadcast-mode schematic. Scalar broadcast sends one shared field to all branches; per-soma broadcast sends a vector-valued somatic error to each branch; structured pathway feedback uses separate channels to preserve pathway identity when different branches should receive different teaching signals. (D) MNIST rule-family summary for shunting LocalCA. Bars show the rule progression from 3F to 4F to 5F using the same summary values reported in Appendix Table S1; the dashed green line marks the MNIST backprop ceiling (96.5%). This panel makes clear that the method is a family, not only its best-performing 5F member.

and compartment input resistance R_n^{tot} . When inhibitory reversal is near rest, shunting increases total conductance and reduces both the magnitude and variability of R_n^{tot} across compartments. In our experiments, this improves both the directional alignment and the scale stability of local credit signals relative to additive controls. The gains are largest in regimes that stress credit assignment, particularly under strong inhibition and noisy inputs. These results suggest a mechanistic link between divisive normalization [6] and the quality of local credit signals.

We quantify these effects using a gradient-fidelity diagnostic based on directional alignment (cosine similarity) and scale mismatch between local and backprop gradients (Table 1, Fig. 2). The absolute cosine values remain modest, but they are consistently more positive under shunting and predict which local-learning regimes ultimately succeed. The shunting advantage is regime-dependent: it grows with inhibitory conductance and is largest on tasks where credit signals must remain robust under noise (Fig. 4). On standard classification benchmarks, simple per-soma broadcast is often sufficient. On the hard cue-routing benchmark, however, scalar broadcast collapses distinctions between competing branches, motivating the structured pathway-vector rule studied in Sec. 4.7.

Contributions

- Exact gradients for conductance-based dendritic trees.** We derive exact loss gradients for dendritic trees and show that each synaptic gradient factorizes into a synapse-local eligibility term and a compartment-specific error.
- A family of local learning rules.** We use this factorization to define 3F/4F/5F local updates based on synapse-local variables and a somatic broadcast, and extend this family with PV-LocalCA when pathway identity must be preserved.
- A regime-dependent advantage for shunting.** In matched architectures, shunting improves gradient fidelity and local learning relative to an additive control, with the largest gains in noisy or inhibition-dependent settings.
- A gradient-fidelity diagnostic.** We use a component-wise local-vs-backprop diagnostic to link dendritic architecture and broadcast design to the quality of local credit signals.

61 2 Compartmental Voltage Model and Gradient Derivation

62 We use a steady-state conductance model derived from discretized passive cable dynamics [1, 2]. In
 63 normalized units, leak and inhibitory reversal are set to 0, the excitatory reversal is set to 1, and leak
 64 conductance is fixed at 1. Each compartment voltage is therefore a conductance-weighted average.
 65 Two properties of this form matter for what follows: local sensitivities depend on both driving force
 66 $(E - V)$ and input resistance R^{tot} , and shunting inhibition corresponds to adding conductance when
 67 $E_{\text{inh}} \approx 0$.

68 2.1 Voltage Equation and Local Sensitivities

69 Consider compartment n with synaptic inputs j (activity x_j , reversal E_j , conductance $g_j^{\text{syn}} \geq 0$) and
 70 dendritic inputs from children (voltage V_j , conductance $g_j^{\text{den}} \geq 0$). The steady-state voltage is:

$$V_n = \frac{\sum_j E_j x_j g_j^{\text{syn}} + \sum_j V_j g_j^{\text{den}}}{\underbrace{\sum_j x_j g_j^{\text{syn}} + \sum_j g_j^{\text{den}} + 1}_{g_n^{\text{tot}}}}, \quad R_n^{\text{tot}} = 1/g_n^{\text{tot}}. \quad (1)$$

71 V_n is a convex combination of reversal potentials, child voltages, and leak. Writing $\mathcal{S}_n = \{0\} \cup$
 72 $\{E_j\}_j \cup \{V_j\}_j$, we have $\min \mathcal{S}_n \leq V_n \leq \max \mathcal{S}_n$ and $0 < R_n^{\text{tot}} \leq 1$. The local sensitivities follow
 73 directly:

Proposition 1 (Local Sensitivities).

$$\frac{\partial V_n}{\partial g_i^{\text{syn}}} = x_i R_n^{\text{tot}} (E_i - V_n), \quad \frac{\partial V_n}{\partial V_i} = g_i^{\text{den}} R_n^{\text{tot}}, \quad \frac{\partial V_n}{\partial g_i^{\text{den}}} = R_n^{\text{tot}} (V_i - V_n). \quad (2)$$

74 2.2 Shunting Inhibition as Divisive Gain Control

75 An inhibitory synapse with $E_{\text{inh}} \approx 0$ contributes current $(0 - V_n)x_j g_j^{\text{syn}}$ and increases g_n^{tot} . Its
 76 sensitivity is $\partial V_n / \partial g_j^{\text{syn}} = -x_j R_n^{\text{tot}} V_n$, which corresponds to multiplicative attenuation (divisive
 77 normalization). Shunting is divisive at the voltage level, but its effect on firing rates can be subtractive
 78 in some regimes [7]. Inhibitory plasticity can balance excitation dynamically [8]; our learned
 79 inhibitory conductances serve an analogous role.

80 2.3 Exact Gradients for Dendritic Trees

81 Let V_0 be the somatic output, $\hat{y} = W_{\text{dec}} V_0$ the decoder, and $\delta_0 = W_{\text{dec}}^\top (\partial L / \partial \hat{y})$ the somatic error.

82 **Theorem 1** (Backpropagation on a Dendritic Tree). *For a rooted dendritic tree with soma at node 0*
 83 *and unique parent $p(n)$ for each non-somatic compartment, the loss gradient satisfies*

$$\frac{\partial L}{\partial V_n} = \frac{\partial L}{\partial V_{p(n)}} R_{p(n)}^{\text{tot}} g_{n \rightarrow p(n)}^{\text{den}}, \quad (3)$$

84 *with boundary condition $\partial L / \partial V_0 = \delta_0$. Because the graph is a tree, each compartment has a unique*
 85 *path to the soma. Defining the path gain*

$$\alpha_n = \prod_{(i \rightarrow k) \in \text{path}(n \rightsquigarrow 0)} R_k^{\text{tot}} g_{i \rightarrow k}^{\text{den}}, \quad \frac{\partial L}{\partial V_n} = \alpha_n \delta_0, \quad (4)$$

86 *with $\alpha_0 = 1$.*

87 *Proof.* Apply the chain rule on the tree-structured computation graph using Prop. 1. $\square$

88 **Corollary 1** (Local–Global Factorization). *The exact synaptic gradient at compartment n factorizes*
 89 *as:*

$$\frac{\partial L}{\partial g_i^{\text{syn}}} = \underbrace{x_i R_n^{\text{tot}} (E_i - V_n)}_{\text{synapse-local eligibility}} \cdot \underbrace{\frac{\partial L}{\partial V_n}}_{\text{compartment error}}, \quad (5)$$

where the eligibility term depends only on quantities available at synapse i (presynaptic activity x_i , input resistance R_n^{tot} , driving force $E_i - V_n$), and the compartment error $\partial L / \partial V_n$ is the sole non-local quantity.

This factorization isolates the only non-local quantity: $\partial L / \partial V_n$. Any practical local rule therefore depends on how well a broadcast signal approximates that compartment error, which we evaluate in Sec. 4.2. In the models studied here, shunting improves this approximation by reducing and stabilizing R_n^{tot} , which narrows the spread of local sensitivities across the tree.

3 Local Learning Rules

3.1 Broadcast Error Approximation

We approximate the exact compartment error $\partial L / \partial V_n$ (Corollary 1) with a broadcast signal e_n derived from the somatic error $\delta_0 = W_{\text{dec}}^\top (\partial L / \partial \hat{y})$. We consider four broadcast modes of increasing locality:

- **Scalar:** $e_n = \bar{\delta} \cdot \mathbf{1}$, where $\bar{\delta} = \text{mean}(\delta_0)$ is a single scalar broadcast to all compartments.
- **Per-soma:** $e_n = \delta_0$ when the layer dimension matches the output; layers with mismatched dimensions fall back to scalar.
- **Local mismatch:** $e_n = (1 - \alpha) \bar{\delta} \cdot \tilde{m}_n + \alpha \bar{\delta}$, where $\tilde{m}_n$ is the RMS-normalized, batch-centered parent-child voltage difference $P_n - V_n$ and $\alpha = 0.2$.
- **Pathway-vector:** $e_n = \alpha \bar{\delta} \cdot \mathbf{1} + (1 - \alpha) \Gamma_n(\delta_0)$, where Γ_n projects the somatic error into router-inferred pathway roles, gates those roles by local pathway activity, and transports the resulting vector hierarchically through the block-structured dendritic tree.

We use per-soma broadcast by default. In our experiments, the local-mismatch variant performs much worse (Appendix A), indicating that broadcast quality is critical. Under shunting, this simple broadcast is often sufficient on standard classification tasks, but the cue-routing benchmark in Sec. 4.7 requires the structured pathway-vector form. Appendix Table S5 summarizes the broader set of implemented model families, training strategies, and broadcast modes used throughout the study.

3.2 Three-Factor Rule (3F)

Definition 1 (3F Update). *For synaptic and dendritic conductances:*

$$\Delta g_j^{\text{syn}} = \eta \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n \rangle_B, \quad \Delta g_j^{\text{den}} = \eta \langle R_n^{\text{tot}} (V_j - V_n) e_n \rangle_B, \quad (6)$$

where $\langle \cdot \rangle_B$ denotes the batch average.

The three factors are presynaptic activity x_j (or a voltage difference), postsynaptic modulation through driving force and input resistance, and the broadcast error e_n . The same rule applies to excitatory and inhibitory synapses. The sign difference comes from the driving force $(E_j - V_n)$.

Additive control. Rule (6) is the local gradient of the *shunting* voltage $V_n = \sum E_j g_j x_j / g_n^{\text{tot}}$. For the additive control ($V_n = \sum g_j x_j$, no divisive normalization), the correct local gradient is simpler: $\Delta g_j^{\text{syn}} = \eta \langle x_j e_n \rangle_B$, with no driving-force or R_n^{tot} terms ($R_n^{\text{tot}} \equiv 1$ by definition since there is no denominator). Throughout, each architecture uses the learning rule derived from its own forward-pass dynamics. This keeps the comparison architecture-matched rather than applying a shunting-derived rule to an additive model.

3.3 Higher-Order Rules: 4F and 5F

4F (heuristic attenuation proxy). In the exact gradient (4), the path gain α_n attenuates the somatic error differently at each compartment. To compensate without explicitly computing that gain, we use an online correlation proxy

$$\rho_n = \text{Cov}(\bar{V}_n, \bar{V}_0) / (\sqrt{\text{Var}(\bar{V}_n) \text{Var}(\bar{V}_0)} + \varepsilon).$$

High ρ_n indicates that compartment and somatic activity covary strongly, making the broadcast δ_0 a more plausible proxy for the local compartment error; low ρ_n down-weights updates where that proxy is likely to be poor. ρ_n is estimated online via exponential moving average ($\alpha = 0.1$).

5F (confidence-weighted 4F). The current 5F implementation adds a second heuristic factor

$$\phi_n = \text{Var}(V_n) / (\sigma_{\text{res}}^2 + \varepsilon),$$

where σ_{res}^2 is the residual variance of V_n after linear regression on the parent voltage. This is an inverse residual-fraction or signal-to-noise proxy: it becomes large when compartment activity is strong and stable relative to residual fluctuations unexplained by the parent, and small when local fluctuations dominate. Operationally, ϕ_n behaves as a branch-wise confidence gate rather than as a pure unique-information estimator. The large empirical gap between 4F and 5F in our sweeps indicates that this confidence reweighting removes a major source of update noise: 3F/4F treat all compartments with similar broadcast quality similarly, whereas 5F suppresses branches whose local voltage is dominated by residual variability. When σ_{res}^2 is very small, ϕ_n would otherwise become degenerate, so we clamp it to a finite interval and interpret saturated values simply as “high-confidence” branches rather than as precise estimates of information content. Clamped to $[0.25, 4.0]$ for stability.

Proposition 2 (5F Update).

$$\Delta g_j^{\text{syn}} = \eta \rho_n \phi_n \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n \rangle_B. \quad (7)$$

Gradient alignment under random broadcast. When the broadcast matrix B_n has i.i.d. zero-mean entries with $\mathbb{E}[B_n^\top B_n] = \alpha I$, the expected cosine between local and exact gradients is positive: $\mathbb{E}[\cos \angle(g^{\text{local}}, g^{\text{exact}})] \geq c_n > 0$, by an argument analogous to feedback alignment [9]. The constant c_n depends on the correlation between local factors and the exact path gain (4); shunting architecture increases this correlation by normalizing the scale of intermediate signals.

3.4 Structured pathway feedback and branch-role plasticity

For tasks with latent pathway structure, scalar broadcast can erase the distinction between different dendritic branches that should be credited differently. We therefore augment the local rule with a structured feedback variant used only on the cue-routing benchmark. We refer to this structured cue-routing rule as *PV-LocalCA* (pathway-vector LocalCA).

Hierarchical pathway-vector broadcast. The encoder/router defines a low-dimensional role basis over latent pathways (for example, one role per cue stream together with shared context). Instead of collapsing δ_0 to a scalar or per-soma vector, we project it into this role basis, gate each role by the local pathway activity on the target branch, and propagate the resulting vector through the proximal-to-distal block tree. Earlier dendritic layers therefore receive a pathway-specific teaching signal that remains aligned with their downstream descendants. Writing the low-rank role coefficients as c_k and the local branch role weights as $q_{n,k}$, the structured broadcast takes the form

$$e_n^{\text{PV}} = \sum_{k=1}^K q_{n,k} c_k,$$

with K latent pathways. Scalar broadcast is the special case $K=1$, while per-soma broadcast corresponds to a dense output-aligned basis.

Router-derived branch roles. Rather than assigning hand-coded apical/basal labels, each branch j receives a normalized role profile r_j inferred from router assignments and the branch’s incoming block weights. Plasticity then depends on two additional local quantities: a branch selectivity factor that up-weights branches with concentrated role mass, and a synapse-specific alignment factor proportional to $\langle r_i, r_j \rangle$ for pre/post role profiles. This biases learning toward pathway-consistent afferents without imposing a heuristic branch taxonomy.

4 Experiments

4.1 Setup

We evaluate the model in two experimental settings: a *capacity-calibrated* regime, where the same architectures achieve strong performance under backpropagation, and a set of *controlled sweeps*

Dataset	Core	Weighted cosine $\uparrow$	Scale mismatch $\downarrow$
MNIST	Shunting	0.202	0.117
MNIST	Additive	0.006	1.053
Context gating	Shunting	0.108	0.036
Context gating	Additive	-0.007	2.154

Table 1: **Gradient fidelity (5F + per-soma)**. Local vs. backprop gradients on matched weights. Shunting shows higher directional alignment and lower scale mismatch (about 30 $\times$ and 10 $\times$ relative to additive in this setting).

that isolate the effects of inhibition strength, broadcast mode, and architecture on credit-signal quality. Our main datasets are MNIST, Fashion-MNIST [39], and four synthetic tasks: *context gating* (context-dependent category boundaries), *noise resilience* (learning under structured input noise), *info shunting* (a task designed to require inhibition-mediated processing), and *cue integration* (two competing sensory cues plus context indicating which cue is reliable). The cue-integration benchmark is specifically designed to test whether dendritic pathway structure can support context-dependent reliability weighting rather than simple averaging: a fixed-pathway control duplicates context into each cue pathway, while a learned router must discover that separation from data. Architectures include point MLP baselines and dendritic cores with either additive integration or shunting (conductance-based) inhibition. Our primary credit-assignment comparisons are architecture-matched additive vs. shunting cores and local vs. backprop training; DFA is included as a supplementary feedback baseline where the layer geometry permits it. Most headline MNIST, Fashion-MNIST, and context-gating results use the 5F rule with per-soma broadcast. The hard cue-routing benchmark instead uses PV-LocalCA: a structured 3F variant with pathway-vector broadcast, path propagation, inhibitory homeostasis, and router-derived branch-role plasticity. For the inhibition sweeps, we additionally evaluate a transported-error oracle that recursively applies the exact path gains from Theorem 1; this oracle is used only as an upper-bound control, not as the main biological proposal. Because context gating benefits from a small HSIC auxiliary objective (about 3 percentage points in Appendix Fig. S4C), we treat it as a supporting structured benchmark rather than as the cleanest test of the base LocalCA rule. We report means $\pm$ s.d. across 3–5 seeds for all headline results. We will release code and configuration files after the review process. The experiments are organized in three stages: we first evaluate how well different broadcasts track exact gradients, then test whether those fidelity differences translate into learning, and finally study a structured-feedback extension for settings where scalar broadcast fails.

4.2 Shunting improves gradient fidelity

To test whether these performance differences correspond to better credit signals, we compare local and backprop gradients on the same batch at the same parameter values. For each parameter tensor p , we compute directional alignment (cosine similarity) and scale mismatch, defined as $|\log_{10}(\|g_p^{\text{local}}\|/\|g_p^{\text{bp}}\|)|$, aggregated by parameter count.

Alignment dynamics over training. Figure 2B tracks per-layer cosine similarity over epochs. In shunting networks, proximal layers approach high positive alignment over training, whereas distal layers show smaller but still positive alignment. Additive networks show near-zero or negative alignment at all layers and epochs. Component-wise decomposition (Fig. 2C) reveals that dendritic conductances and excitatory synapses carry the strongest alignment signal in shunting networks, consistent with the biophysical role of conductance-based driving forces. Although the absolute cosines are still far from 1, they are strongly predictive of local-learning competence: the same inhibition regimes that raise these cosines are the regimes in which local training succeeds.

Exact factorization sanity check. Using the exact compartment errors $\partial L/\partial V_n$ extracted by autograd, we reconstructed the corresponding synaptic and dendritic conductance gradients directly from the factorization in Corollary 1, using the exact piecewise denominator derivative of the implemented shunting layer. Across the full 100-run inhibition sweep used for the mechanistic analysis, the reconstructed gradients match autograd to numerical precision (weighted cosine 1.000001 ± 0.000001 , weighted scale mismatch $< 4 \times 10^{-8}$). This confirms that the theory is operational in code: the




Figure 2: **Gradient-fidelity mechanism.** (A) Weighted cosine similarity between local and backprop gradients: shunting (green) shows about $30\times$ better directional alignment than additive (blue) on both MNIST and context gating. (B) Per-layer cosine similarity over training epochs. Shunting proximal layers approach 1.0; additive layers remain near zero. (C) Component-wise alignment: E-synapses and dendritic conductances carry the strongest alignment signal in shunting networks.

only approximation made by LocalCA is in the broadcast substitute for $\partial L/\partial V_n$, not in the local eligibility factors.

4.3 Mechanistic evidence for the shunting advantage

We provide evidence for why shunting improves gradient fidelity and test broadcast robustness (Fig. 3).

Path-gain concentration. For a tree, the exact compartment error is the somatic error multiplied by a unique path gain $\alpha_n = \prod_{k \in \text{path}(n \rightarrow 0)} R_k^{\text{tot}} g_k^{\text{den}}$. Panel A measures the coefficient of variation of these path gains across compartments. On MNIST, shunting keeps the path-gain distribution tightly concentrated across inhibition levels (CV about 0.21–0.27), whereas the additive control shows much broader dispersion (CV about 0.53–1.08). This is the deterministic quantity that scalar or per-soma broadcast is implicitly trying to ignore: when the omitted path gains vary less, one shared broadcast is a better proxy for the true compartment-specific error.

Direct compartment-error fidelity. The theorem identifies $\partial L/\partial V_n$ as the sole non-local term, so we measure it directly. Panel B compares the local broadcast e_n with the exact compartment error on the noise-resilience sweep, using cosine similarity at the compartment level. Per-soma broadcast tracks the exact compartment error much better under shunting than under additive integration, rising from 0.20 at $N_I=0$ to about 0.32 near the best-performing inhibition range, while additive stays near zero or negative. The transported-error oracle obtained by recursively applying the exact path gains improves fidelity further in both architectures, showing that the theorem is not merely descriptive: it defines an operational upper bound on low-bandwidth broadcast. The large gap between the additive per-soma broadcast and its transported upper bound shows that omitted path gains are a major part of the additive failure mode.

Transported broadcast is an operational upper bound. Panel C trains the same local rule using the transported-error oracle as its teaching signal. On noise resilience, additive learning rises from 33–47% under per-soma broadcast to 75.6% at $N_I=0$ and 98.5% at $N_I=40$ under transported broadcast. Shunting improves from 71.7–90.3% under per-soma to 98.5–99.3% under transported broadcast across the same sweep. On MNIST, transported broadcast reaches 94.6% for additive and 96.8% for shunting. These oracle results show that broadcast fidelity is a major bottleneck for LocalCA, while the remaining additive–shunting gap under matched transported broadcast points to an additional role for conductance-shaped local sensitivities.




Figure 3: **Mechanistic evidence.** (A) Path-gain dispersion on MNIST, measured as the coefficient of variation of the exact tree attenuation α_n across compartments. Shunting keeps path gains much more concentrated than the additive control as inhibition increases. (B) Direct compartment-error fidelity on noise resilience. Solid lines compare the per-soma broadcast e_n to the exact compartment error $\partial L / \partial V_n$; dashed lines show the transported-error oracle obtained by recursively applying exact path gains. Shunting substantially improves the fidelity of the simple per-soma broadcast, while the transported oracle defines the remaining approximation gap. (C) Learning with the transported-error oracle on noise resilience. Improving the broadcast field is enough to recover near-ceiling local learning, especially in the additive architecture where per-soma broadcast fails most severely. (D) Broadcast bandwidth. Four-bit quantization preserves full performance; only sign-only and sparse modes degrade sharply.

247 **Low-bandwidth broadcast suffices.** Testing degraded broadcast modes (Fig. 3D), 4-bit quantiza-
 248 tion preserves full performance (61.6% vs. 61.7%; $p > 0.9$), while sign-only (37.5%) and sparse
 249 top-30% (38.2%) degrade sharply. Here “effective bits per neuron” means the precision of the
 250 broadcast channel after compression: 32 for full-precision, b for b -bit quantization, 1 for sign-only,
 251 and a separate sparse code for the top-30% condition. These results suggest that broadcast direction
 252 matters more than exact magnitude once local sensitivities are well scaled.

253 **Normalization-only controls remain insufficient.** Explicitly normalizing additive voltages im-
 254 proves local learning but does not close the gap to shunting (Appendix Fig. S7). DFA baselines
 255 (Appendix Fig. S5) likewise improve over point-MLP random-feedback baselines without matching
 256 the shunting conductance model. These controls support the same conclusion as panel B: useful local
 257 learning requires more than a broadcast signal with the right overall scale.

258 4.4 Local learning under capacity-calibrated settings

259 Having established that shunting improves gradient fidelity, we next test whether the same conditions
 260 improve end-task learning. In the capacity-calibrated regime, shunting dendritic cores reach 0.965 on
 261 MNIST and 0.864 on context gating under backpropagation. Within the same architectures, the best
 262 local configuration (5F, per-soma broadcast, local decoder) reaches the accuracies reported in Table 2.

263 Within the local-rule family, 5F consistently outperforms 4F and 3F (Appendix Table S1), and
 264 per-soma broadcast strongly outperforms scalar and local-mismatch modes (Appendix Table S2).
 265 The large 4F→5F jump is important: in our current implementation, most of that gain comes
 266 from confidence-weighting updates by ϕ_n , which suppresses noisy branch fluctuations that would
 267 otherwise receive the same broadcast-driven update as reliable branches. We therefore interpret 5F as
 268 the strongest practical rule in this paper, but not as a uniquely derived optimum; its clamp bounds and
 269 residual estimator are still heuristic design choices. In a focused MNIST sensitivity sweep, however,
 270 those heuristics are not especially fragile: tightening the clamp from $[0.25, 4.0]$ to $[0.5, 2.0]$ changes

Dataset	BP ceiling	Best local (5F)	Gap
MNIST	0.965	0.914 ± 0.003	5.1%
Fashion-MNIST	0.879	0.811 ± 0.012	6.8%
Context gating	0.864	0.803 ± 0.006	6.1%

Table 2: **Local competence.** Backprop ceilings from capacity sweeps; local values are 5F with per-soma broadcast on shunting dendritic cores. Context gating additionally uses HSIC auxiliary objective (weight 0.01; Appendix E). Errors: ± 1 s.d. across 5 seeds.




Figure 4: **Local competence and regime dependence.** (A) Backprop ceiling (gray) vs. best local rule (5F per-soma) for shunting (green) and additive (blue) cores on MNIST, Fashion-MNIST, and context gating. (B) N_I dose-response: test accuracy vs. inhibitory synapses per branch (N_I) for shunting and additive cores on MNIST (solid) and noise resilience (dashed). Shunting needs $N_I \geq 5$ to reach high accuracy; additive learning fails on noise resilience regardless of N_I . (C) Shunting advantage (Δ , in percentage points) vs. N_I . The shunting advantage grows with inhibitory synapses per branch and is larger on the noise-resilience task than on MNIST.

mean test accuracy by only +0.4 pp, widening it to $[0.1, 8.0]$ is effectively neutral, and varying the 4F/5F EMA rate from 0.05 to 0.20 shifts accuracy by less than 0.2 pp (Appendix Fig. S3).

4.5 Regime dependence of the shunting advantage

The benefit of shunting varies across tasks. On MNIST with matched per-soma broadcast, shunting outperforms additive by about 2 percentage points; a similar pattern holds on Fashion-MNIST (81.1% vs. 79.4%). On tasks that require credit signals robust to noise, the gap is much larger: +50.3 pp on noise resilience (with $N_I=10$ inhibitory synapses per branch) and +24.8 pp on info shunting ($N_I=0$). Figure 4B–C shows that the advantage grows with inhibitory conductance strength. Additive cores can still perform reasonably on easier settings, reaching about 89% on MNIST after tuning. However, they fail in regimes where inhibition-mediated normalization matters for gradient propagation.

4.6 Scalability and generalization

We test whether the shunting advantage persists under three stress tests: increased dendritic depth, noisy broadcast signals, and a second real-world benchmark.

Depth scaling. Varying dendritic depth from 1 to 4 layers (branch factors $[9]$ to $[3, 3, 3, 3]$), shunting local learning degrades from 63.5% to 57.4% while additive local degrades more steeply from 54.9% to 29.7% (Fig. 5A). Backprop ceilings remain stable at about 90–92% for both architectures. The shunting advantage grows with depth (+8.5 pp at depth 1 to +27.7 pp at depth 4), consistent with the path-gain attenuation in Theorem 1: each additional layer introduces an attenuation factor of the




Figure 5: **Scalability and generalization.** (A) Depth scaling: shunting local (green) degrades gracefully from 63.5% to 57.4%; additive local (blue) drops from 54.9% to 29.7%. Backprop ceilings (dashed) remain stable. (B) Noise robustness: shunting is robust to $\sigma \leq 0.1$; additive degrades rapidly and reaches chance at $\sigma = 1.0$. (C) Fashion-MNIST: shunting local 81.1% vs. 87.9% backprop; additive local 79.4% vs. 87.4% backprop.

289 form $R_k^{\text{tot}} g_k^{\text{den}} < 1$. In practice, shunting changes the distribution of these factors and yields better
 290 depth scaling under local learning.

291 **Noise robustness.** Adding Gaussian noise $\mathcal{N}(0, \sigma^2)$ to the broadcast error (Fig. 5B), shunting local
 292 learning is robust to $\sigma \leq 0.1$ (about 62%) and degrades gracefully beyond, while additive drops
 293 from 46.5% to chance at $\sigma = 1.0$. This pattern is consistent with the proposed mechanism: shunting
 294 produces better-scaled local sensitivities, so gradient direction is preserved more often even when the
 295 broadcast magnitude is corrupted.

296 **Fashion-MNIST.** On Fashion-MNIST (Fig. 5C), shunting local reaches 81.1% vs. a 87.9% back-
 297 prop ceiling (gap 6.8%); additive local reaches 79.4% vs. 87.4% backprop (gap 8.0%). The shunting
 298 advantage is modest (about 1.7 pp) on this clean classification task, consistent with the regime-
 299 dependent pattern above: the benefit is largest on tasks that require robust credit signals under
 300 noise.

301 4.7 Extension: structured feedback when scalar broadcast fails

302 The main paper focuses on low-bandwidth scalar or per-soma broadcast. We now consider a harder
 303 cue-integration benchmark where that approximation is insufficient because different dendritic
 304 pathways must be credited differently. The task asks the model to combine two noisy sensory cues
 305 using a context signal that indicates which cue is reliable on each trial. In the fixed-pathway control,
 306 context is copied into both cue pathways. In the learned-pathway setting, the router must instead
 307 discover that separation and align one latent pathway with each cue stream.

308 Figure 6 summarizes five-seed runs for the hard benchmark together with a focused tuning sweep.
 309 Fixed-pathway additive local learning remains near ceiling ($99.1 \pm 0.3\%$), fixed-pathway shunting
 310 local learning is strong but more variable ($92.8 \pm 6.1\%$), learned-routing additive local learning also
 311 succeeds ($98.1 \pm 2.4\%$), and standard learned-router training stays near $99.2 \pm 0.4\%$ for both additive
 312 and shunting cores. Taken together, these controls show that the benchmark and router are learnable,
 313 while also indicating that shunting LocalCA is already somewhat stressed on this task even before
 314 learned routing is introduced.

315 The specific failure appears when shunting local learning uses the default per-soma broadcast. At
 316 the original 90-epoch setting, learned-routing accuracy is only $82.6 \pm 14.2\%$. Lowering router
 317 temperature to 0.3 or freezing the encoder for the first 20 epochs changes little ($87.6 \pm 13.3\%$ and




Figure 6: Structured feedback for pathway discovery on cue integration. (A) Task schematic. Context determines which noisy cue is reliable on each trial. Scalar or per-soma broadcast sends the same teaching signal to both pathways, whereas PV-LocalCA preserves pathway identity through a low-rank pathway-vector feedback field. (B) Seed-averaged diagnosis and focused tuning (5 seeds per condition). Learned-routing additive local learning remains high, but learned-routing shunting with per-soma broadcast stays unstable; low-temperature, freeze-20, and per-soma tuning help only modestly, whereas tuned pathway-vector feedback reaches $99.1 \pm 0.3\%$. (C) Test accuracy vs. router specialization across learned-router conditions. High specialization alone does not guarantee success: per-soma, low-temperature, and freeze-20 runs remain inaccurate despite near-discrete assignments. (D) Recovered pathway assignments for a representative tuned pathway-vector run. Cue-A and cue-B features segregate cleanly into separate latent pathways.

83.8 $\pm$ 13.9%), even though the router already becomes highly specialized (mean max assignment 0.98–1.00). A focused optimization sweep shows that training budget alone is not the issue: extending training to 140 epochs leaves per-soma performance unchanged (82.5 $\pm$ 14.1%), and even lowering the learning rate to 10^{-3} reaches only 88.2 $\pm$ 13.9%.

Replacing per-soma broadcast with PV-LocalCA changes the regime entirely. The untuned 90-epoch version already improves to 94.0 $\pm$ 11.0%, and the tuned setting (140 epochs, learning rate 10^{-3}) reaches 99.1 $\pm$ 0.3% test and 99.4 $\pm$ 0.2% validation accuracy while maintaining strong specialization (mean max assignment 0.995 $\pm$ 0.008). The failure is therefore not caused by pathway discovery itself; it arises because scalar broadcast does not preserve pathway identity once different cue branches must be credited differently. When the feedback preserves that structure, learned-routing shunting local learning becomes reliable.

5 Related Work

Dendritic models of credit assignment. Dendritic trees support nonlinear computation [3, 4, 31] and have inspired biologically plausible learning schemes: segregated dendrites [11], dendritic prediction errors [5, 12] (reported about 98% on MNIST in an abstract-compartment setting), burst-dependent plasticity [23, 24], latent equilibrium [25] (98.9% MNIST with abstract compartments), and dendritic localized learning [36]. We differ in deriving rules from conductance-based equations rather than abstract compartment surrogates and in using gradient-fidelity diagnostics to ask when a low-bandwidth somatic broadcast is accurate enough to support learning.

Feedback alignment and local learning. Random feedback [9], DFA [10], forward-forward [26], PEPITA [27], and PAL [34] established strong local-learning baselines on standard MLP benchmarks. Recent supervised local-learning systems have also raised the empirical bar in deeper settings: AugLocal scales auxiliary-network local learning to deep CNNs including ImageNet-sized models [35], DLL reports strong MLP/CNN/RNN performance under explicit plausibility criteria [36], CCL introduces

a biologically plausible dual-network training framework [38], and EBD provides a complementary theory of direct error broadcast through decorrelation [37]. Our aim is therefore not to claim the strongest benchmark accuracy among all local-learning methods, but to identify a biophysical mechanism that explains when low-bandwidth credit signals become useful in conductance-based dendritic networks. Our broadcast modes generalize FA/DFA to dendrites while additionally exploiting conductance-based signals unavailable to standard architectures. Standard FA fails entirely on dendritic layers, while DFA is compatible and shunting yields a smaller backprop–DFA gap (24.9 pp) than point MLPs (39.5 pp; Appendix Fig. S5).

Target propagation and energy-based methods. DTP [28], DFC [29], equilibrium propagation [16], and predictive coding [14, 30] solve weight transport through diverse mechanisms. Our contribution is orthogonal: conductance-based biophysics provides an additional route to local credit. We do not benchmark these methods directly here because their iterative inference or controller dynamics are not architecture-matched to the steady-state conductance cores studied in this paper. We therefore treat them as complementary solutions to locality rather than as drop-in baselines for the same model family.

Divisive normalization. Shunting implements divisive normalization [6], though its effect on rates can be subtractive [7]. Silver [32] showed that inhibitory conductance modulates gain and SNR. Beniaguev et al. [33] showed single neurons are computationally equivalent to 5–8 layer DNNs. To our knowledge, prior work has not directly examined gradient alignment between local updates and backpropagation in the context of shunting inhibition.

6 Discussion

We started from exact gradients for a conductance-based dendritic tree and asked when a somatic broadcast can serve as a useful proxy for compartment-specific errors. In the models studied here, the answer depends strongly on the inhibitory regime. Shunting changes the conductance denominator and stabilizes effective local sensitivities, improves agreement between local and backprop gradients, and leads to better local learning than an additive control, especially in noisy or inhibition-dependent settings. Explicit normalization recovers part of this benefit but not all of it, suggesting that shunting changes more than just overall scale.

Limitations. Our best accuracy (91.4% MNIST, 81.1% Fashion-MNIST) is below several recent local-learning methods that operate in abstract compartments or standard activation spaces [35–38]. This reflects the constraints of operating in conductance-based voltage space: bounded voltages, positive conductances, and a denominator-heavy computation graph. In the conservative flattened-CIFAR-10 sweep shown in Appendix Fig. S6, shunting local learning reaches 21.2% vs. 40.1% backprop, indicating that the gradient-fidelity mechanism alone is insufficient for harder tasks without architectural advances (e.g., convolutional encoders). That trade-off reflects our emphasis on mechanism rather than absolute benchmark performance: the gradient-fidelity analysis helps explain why some architectures support local learning better than others. Second, naive local broadcasts remain much weaker than per-soma on MNIST, and the hard cue-routing benchmark requires an even more structured pathway-vector signal (Sec. 4.7). Our claims are therefore specific to the broadcast family studied here rather than arbitrary local mismatch heuristics. Per-soma broadcast should be understood as low-bandwidth but still neuron-global: it assumes access to the full somatic error vector and therefore sets the current locality boundary of the method. Third, scaling to deeper architectures degrades local learning more than backprop (Fig. 5A; Appendix I), indicating that depth-dependent credit attenuation remains challenging. At depth 4, shunting local learning is still roughly 33 percentage points below backprop, so the present rules improve deep credit assignment but do not solve it. Fourth, standard FA is incompatible with dendritic layer geometry (Appendix Fig. S5), limiting the space of applicable feedback methods. We also do not directly compare against iterative target-propagation or equilibrium-style methods, because a clean comparison would require re-implementing their inference dynamics inside the same conductance-based dendritic architecture rather than benchmarking against separate abstract models. Finally, the weight-statistics analysis is secondary to the main credit-assignment claim and is reported in Appendix G; it should be read as descriptive consistency rather than as core evidence for the learning mechanism.

Broader relevance. More broadly, the results are relevant to neuroscience, machine learning, and neuromorphic hardware. For neuroscience, they suggest a connection between divisive normalization and local credit fidelity, extending its known roles in gain control [32] and sensory coding [6]. For machine learning, they show that conductance-based inductive biases can shape gradient geometry in ways that benefit local learning. For neuromorphic implementations, the update rules remain strictly local and parallelizable.

Potential harmful consequences. The work is primarily methodological and does not introduce a deployable decision system or a new human dataset. Still, more efficient local learning rules for adaptive hardware could eventually be incorporated into surveillance, autonomous, or other safety-critical systems. We therefore view the most plausible risks as downstream misuse of efficient on-device learning rather than harm from the present experimental setup itself.

Testable predictions. Our framework predicts that (1) stronger perisomatic inhibition yields more precise synaptic plasticity; (2) GABA_A blockade selectively impairs multi-layer credit tasks; (3) the dendritic–somatic voltage correlation (ρ_n) increases during learning; (4) learned excitatory weight distributions follow log-normal statistics whose width depends systematically on inhibitory conductance (Appendix G). These predictions can help distinguish our account from models where dendrites mainly serve as computational substrates.

References

- [1] Koch, C. (1999). *Biophysics of Computation*. Oxford University Press.
- [2] Dayan, P., & Abbott, L. F. (2001). *Theoretical Neuroscience*. MIT Press.
- [3] Poirazi, P., Brannon, T., & Mel, B. W. (2003). Pyramidal neuron as two-layer neural network. *Neuron*, 37(6), 989–999.
- [4] London, M., & Häusser, M. (2005). Dendritic computation. *Annual Review of Neuroscience*, 28, 503–532.
- [5] Urbanczik, R., & Senn, W. (2014). Learning by the dendritic prediction of somatic spiking. *Neuron*, 81(3), 521–528.
- [6] Carandini, M., & Heeger, D. J. (2012). Normalization as a canonical neural computation. *Nature Reviews Neuroscience*, 13(1), 51–62.
- [7] Holt, G. R., & Koch, C. (1997). Shunting inhibition does not have a divisive effect on firing rates. *Neural Computation*, 9(5), 1001–1013.
- [8] Vogels, T. P., et al. (2011). Inhibitory plasticity balances excitation and inhibition. *Science*, 334(6062), 1569–1573.
- [9] Lillicrap, T. P., et al. (2016). Random synaptic feedback weights support error backpropagation. *Nature Communications*, 7, 13276.
- [10] Nøkland, A. (2016). Direct feedback alignment provides learning in deep neural networks. *NeurIPS*, 29.
- [11] Guerguiev, J., Lillicrap, T. P., & Richards, B. A. (2017). Towards deep learning with segregated dendrites. *eLife*, 6, e22901.
- [12] Sacramento, J., et al. (2018). Dendritic cortical microcircuits approximate the backpropagation algorithm. *NeurIPS*, 31.
- [13] Bartunov, S., et al. (2018). Assessing the scalability of biologically-motivated deep learning algorithms and architectures. *NeurIPS*, 31.
- [14] Whittington, J. C., & Bogacz, R. (2019). Theories of error back-propagation in the brain. *Trends in Cognitive Sciences*, 23(3), 235–250.

- [15] Richards, B. A., & Lillicrap, T. P. (2019). Dendritic solutions to the credit assignment problem. *Current Opinion in Neurobiology*, 54, 28–36.
- [16] Scellier, B., & Bengio, Y. (2017). Equilibrium propagation. *Frontiers in Computational Neuroscience*, 11, 24.
- [17] Gretton, A., et al. (2005). Measuring statistical dependence with Hilbert-Schmidt norms. *ALT*, 63–77.
- [18] Frémaux, N., & Gerstner, W. (2016). Neuromodulated spike-timing-dependent plasticity, and theory of three-factor learning rules. *Frontiers in Neural Circuits*, 9, 85.
- [19] Bellec, G., et al. (2020). A solution to the learning dilemma for recurrent networks of spiking neurons. *Nature Communications*, 11, 3625.
- [20] Larkum, M. (2013). A cellular mechanism for cortical associations. *Trends in Neurosciences*, 36(3), 141–151.
- [21] Welford, B. P. (1962). Note on a method for calculating corrected sums of squares and products. *Technometrics*, 4(3), 419–420.
- [22] Turrigiano, G. G. (2008). The self-tuning neuron: synaptic scaling of excitatory synapses. *Cell*, 135(3), 422–435.
- [23] Payeur, A., et al. (2021). Burst-dependent synaptic plasticity can coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24(7), 1010–1019.
- [24] Greedy, W., et al. (2022). Single-phase deep learning in cortico-cortical networks. *NeurIPS*, 35.
- [25] Haider, P., et al. (2021). Latent equilibrium. *NeurIPS*, 34.
- [26] Hinton, G. (2022). The forward-forward algorithm. *arXiv:2212.13345*.
- [27] Dellaferrera, G., & Bhatt, D. (2022). Error-driven input modulation: Solving the credit assignment problem without a backward pass. *ICML*, 4937–4955.
- [28] Lee, D.-H., et al. (2015). Difference target propagation. *ECML*, 498–515.
- [29] Meulemans, A., et al. (2021). Credit assignment in neural networks through deep feedback control. *NeurIPS*, 34.
- [30] Millidge, B., Seth, A., & Buckley, C. L. (2022). Predictive coding: A theoretical and experimental review. *arXiv:2107.12979*.
- [31] Koch, C., Poggio, T., & Torre, V. (1983). Nonlinear interactions in a dendritic tree. *PNAS*, 80(9), 2799–2802.
- [32] Silver, R. A. (2010). Neuronal arithmetic. *Nature Reviews Neuroscience*, 11(7), 474–489.
- [33] Beniaguev, D., Segev, I., & London, M. (2021). Single cortical neurons as deep artificial neural networks. *Neuron*, 109(17), 2727–2739.
- [34] Bhatt, D., et al. (2024). Parallel local learning with alignment. *Nature Machine Intelligence*, 6, 1–12.
- [35] Ma, C., Wu, J., Si, C., & Tan, K. C. (2024). Scaling supervised local learning with augmented auxiliary networks. *ICLR*.
- [36] Lv, C., Xu, J., Lu, Y., Wang, X., Wang, Z., Xu, Z., Yu, D., & Du, X. (2025). Dendritic Localized Learning: Toward Biologically Plausible Algorithm. *Proceedings of Machine Learning Research*, 267.
- [37] Erdogan, M., Pehlevan, C., & Erdogan, A. (2025). Error broadcast and decorrelation as a potential artificial and natural learning mechanism. *NeurIPS*.

Dataset	Rule	Top-10 valid	Top-10 test
MNIST	3F	0.611	0.622
MNIST	4F	0.620	0.628
MNIST	5F	0.912	0.916
Context gating	3F	0.398	0.396
Context gating	4F	0.411	0.411
Context gating	5F	0.807	0.789

Table S1: **Rule-family ranking.** Top-10 mean across completed local-competence sweeps.

Core	Broadcast	Decoder	Test (mean $\pm$ std)
Shunting	per-soma	local	0.912 $\pm$ 0.005
Shunting	per-soma	backprop	0.909 $\pm$ 0.008
Shunting	local-mismatch	local	0.146 $\pm$ 0.046
Shunting	local-mismatch	backprop	0.146 $\pm$ 0.037
Additive	per-soma	local	0.894 $\pm$ 0.007
Additive	per-soma	backprop	0.900 $\pm$ 0.001
Additive	local-mismatch	local	0.342 $\pm$ 0.058
Additive	local-mismatch	backprop	0.348 $\pm$ 0.095

Table S2: **Local-mismatch recheck (MNIST, 5F).** Per-soma is consistently strong; local-mismatch remains much weaker.

- 480 [38] Kao, C.-H., & Hariharan, B. (2024). Counter-Current Learning: A biologically plausible dual
481 network approach for deep learning. *NeurIPS*.
- 482 [39] Xiao, H., Rasul, K., & Vollgraf, R. (2017). Fashion-MNIST: A novel image dataset for bench-
483 marking machine learning algorithms. *arXiv:1708.07747*.
- 484 [40] Song, S., Sjöström, P. J., Reigl, M., Nelson, S., & Bhatt, D. B. (2005). Highly nonrandom
485 features of synaptic connectivity in local cortical circuits. *PLoS Biology*, 3(3), e68.
- 486 [41] Buzsáki, G., & Mizuseki, K. (2014). The log-dynamic brain: How skewed distributions affect
487 network operations. *Nature Reviews Neuroscience*, 15(4), 264–278.
- 488 [42] Saxena, S., et al. (2024). Lognormal recurrent networks for biological plausibility. *ICLR*.
- 489 [43] MICrONS Consortium (2021). Functional connectomics spanning multiple areas of mouse
490 visual cortex. *bioRxiv*, 2021.07.28.454025.
- 491 [44] Loewenstein, Y., Kuras, A., & Bhatt, D. (2011). Multiplicative dynamics underlie the emergence
492 of the log-normal distribution of spine sizes in the neocortex in vivo. *Journal of Neuroscience*,
493 31(26), 9481–9488.
- 494 [45] Megías, M., et al. (2001). Total number and distribution of inhibitory and excitatory synapses
495 on hippocampal CA1 pyramidal cells. *Neuroscience*, 102(3), 527–540.




Figure S1: **Capacity calibration (supplementary).** (A) Phase 1 backprop ceilings across all architectures and datasets. (B) Rule-family ranking: 5F consistently outperforms 4F and 3F. (C) Decoder mode comparison (local vs. backprop decoder, 5F MNIST). (D) Broadcast mode comparison: per-soma is required for strong performance; local-mismatch fails.

Condition	Metric	Value
Decoder: local vs backprop vs frozen (MNIST)	test acc	0.379 vs 0.379 vs 0.176
Shunting vs additive, matched (MNIST)	Δ test	+0.210
Per-soma, path on vs off	test / MI(E,I;C)	-0.003 / +0.017

Table S3: **Mechanistic ablations** in controlled small-network architectures.




Figure S2: **Extended gradient and N_I analysis (supplementary).** (A) Scale mismatch bars: shunting achieves near-ideal scale (0.117); additive exhibits order-of-magnitude distortion (> 1.0). (B) Noise resilience N_I dose-response with error bands (± 1 s.d.). (C) MNIST N_I dose-response detail with error bands. (D) Fashion-MNIST individual seeds for all conditions, showing consistency across runs.

A Supplementary Results

Rule-Family Ranking

Broadcast Mode Comparison and Local-Mismatch Recheck

Phase 1 Capacity Ceilings

Ablation Results

Extended Gradient Analysis and N_I Sweep Detail

Verification and Reproducibility

FA/DFA Baseline Comparison

CIFAR-10 Results

Additive Normalization Control

B Implementation Details

Biological Plausibility Assumptions

Each synapse has access to: presynaptic activity x_j (local), compartment voltage V_n (local membrane potential), reversal potential E_j (fixed by receptor type), and total input resistance $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$




Figure S3: **5F sensitivity analysis (MNIST, shunting, per-soma broadcast).** (A) Sensitivity to the ϕ_n clamp interval. The default clamp [0.25, 4.0] is already strong ($92.1 \pm 0.8\%$); tightening it to [0.5, 2.0] slightly improves performance ($92.5 \pm 0.3\%$), while widening it to [0.1, 8.0] is effectively neutral ($92.1 \pm 0.9\%$). (B) Sensitivity to the 4F/5F EMA rate used for online statistics. Varying α from 0.05 to 0.20 changes mean test accuracy by less than 0.2 pp. Together these results suggest that the 5F gain is not driven by a brittle single hyperparameter choice, even though the factor itself remains heuristic. Errors: ± 1 s.d. across 3 seeds.




Figure S4: **Verification and reproducibility (supplementary).** (A) MNIST verification: main seeds (42–46) yield $91.4 \pm 0.3\%$; held-out seeds (47–49) yield $89.8 \pm 0.9\%$, confirming generalization. (B) Context gating verification: main seeds (+HSIC) $80.3 \pm 0.6\%$; held-out seeds (no HSIC) $77.2 \pm 1.5\%$. The roughly 3 pp gap reflects HSIC removal, not seed sensitivity. (C) HSIC weight ablation on context gating: moderate weights (0.01–0.1) perform best.

(computable from local conductances). The 4F modulator ρ_n requires online estimation of voltage covariance (biologically plausible via slow calcium signals); the 5F factor ϕ_n requires a linear regression proxy (implementable via eligibility traces). The only non-local quantity is the broadcast error e_n , which requires a top-down or neuromodulatory signal from the soma to dendritic compartments. We assume per-neuron resolution (one error per output neuron) for the per-soma broadcast mode.

Units and Parameterization

Decoder Update Modes

W_{dec} maps $V_L \rightarrow \hat{y}$. Modes: **backprop** ($\nabla_W L$ via autograd), **local** ($\Delta W = \eta \langle \delta_0 V_L^T \rangle_B$), **frozen** ($\Delta W = 0$).




Figure S5: **Feedback alignment baselines (MNIST).** (A) Grouped comparison: standard backprop (gray), DFA (orange), and FA (purple). Red X marks indicate FA failure on dendritic architectures (dimensional incompatibility between random feedback matrices and block-structured dendritic layers). DFA achieves 66.1% on shunting vs. 62.2% additive vs. 51.8% point MLP. (B) Backprop–DFA gap: dendritic architectures (24.9–27.6 pp) show smaller gaps than point MLPs (39.5 pp), suggesting conductance-based architecture is partially compatible with random feedback.




Figure S6: **Conservative flattened CIFAR-10 reference sweep.** In this appendix stress test, shunting backprop reaches 40.1%, additive backprop 37.8%, shunting local 21.2%, and additive local 10.0% (chance). The backprop–local gap (about 19 pp for shunting) is larger than on MNIST, reflecting the increased difficulty of propagating credit through conductance-based layers for complex visual features. This panel should be read as a conservative flattened-image benchmark rather than as the tuned CIFAR ceiling for the architecture family.

519 Model and Mode Inventory

520 LocalCA Extension Compatibility

521 Compute Resources

522 All experiments were run on an internal GPU cluster. Most MNIST, Fashion-MNIST, and synthetic-
 523 task runs used a single NVIDIA A100 or H100 GPU and finished in minutes to tens of minutes per
 524 seed; the flattened CIFAR-10 runs required longer single-GPU jobs but still did not use distributed or
 525 multi-GPU training. Reported headline metrics are aggregated over 3–5 seeds as stated in the text;




Figure S7: **Additive + normalization control (MNIST)**. (A) Standard backprop: normalization provides a small boost (89.7% $\rightarrow$ 90.9%). (B) Local learning: normalization improves additive from 46.5% to 56.0% (+9.5 pp), partially closing the gap to shunting (61.7%, dashed green). The remaining gap (-5.7 pp) indicates shunting provides benefits beyond simple voltage normalization, including bounded activations, conductance-dependent input weighting, and biophysically constrained sensitivity structure.

Quantity	Symbol	Convention
Voltage	V	Normalized to $[-1, 1]$
Conductances	$g^{\text{syn}}, g^{\text{den}}$	Nonneg. via softplus
Leak conductance	g^{leak}	Set to 1
Input resistance	R^{tot}	≤ 1

Table S4: Units and normalization.

526 the broader exploratory sweeps used to tune and diagnose the method required additional compute
527 beyond the final reported runs.

528 Code, Data, and Asset Availability

529 MNIST, Fashion-MNIST, CIFAR-10, and MICrONS are public datasets or public benchmark assets
530 cited in the paper; we use them under their standard published access conditions and cite their original
531 sources in the bibliography. The present anonymous submission does not include a public code URL
532 in order to preserve double-blind review, but the paper specifies the model family, training modes, key
533 hyperparameters, and representative configurations, and we plan to release code and configuration
534 files after review. The paper does not release a new dataset or a stand-alone pretrained model asset.

535 Algorithm

Algorithm 1 Local Credit Assignment

- 1: **Input:** Model, batch (x, y) , config $\mathcal{C}$
- 2: Forward pass; loss L , output error δ^y
- 3: Somatic error $\delta_0 = W_{\text{dec}}^\top \delta^y$
- 4: **for** each layer n (reverse) **do**
- 5: $e_n = \text{broadcast}(\delta_0, \mathcal{C})$; optionally gate by pathway roles
- 6: Compute ρ_n, ϕ_n (EMA estimators, if enabled)
- 7: Apply the configured local update (Eq. 6 or 7)
- 8: **end for**
- 9: Clip gradients; optimizer step

Category	Implemented modes	Role in this paper
Core models	point MLP; dendritic MLP; dendritic additive; dendritic shunting; pathway-routed dendritic cores	Main comparisons use dendritic additive vs. dendritic shunting; routed cores are used for cue integration.
Training strategies	standard backprop; LocalCA; DFA; FA; transported-broadcast oracle	Main text uses standard and LocalCA; DFA/FA remain appendix baselines; path transport is an oracle upper bound.
Broadcast modes	scalar; per-soma; local mismatch; low-rank; path transport; pathway vector	Main text centers on per-soma and pathway-vector; scalar and local-mismatch are supporting controls; path transport is diagnostic only.
Recent dendritic extensions	path propagation; router-derived branch roles; pathway-vector feedback; inhibitory homeostasis	These were added to test morphology-aware credit transport and structured branch-specific plasticity beyond the original 3F/4F/5F family.

Table S5: Implemented model families and LocalCA operating modes. Only a subset is used for the main claims; the remaining modes serve as controls, ablations, or structured extensions.

Extension	Meaning	Compatible with	Typical use here
Per-soma broadcast	Somatic error vector shared across compartments	3F, 4F, 5F	Default main-text LocalCA setting
Low-rank broadcast	Small number of broadcast channels instead of a scalar or full field	3F, 4F, 5F	Intermediate control when rank-1 is too restrictive
Path transport	Recursive tree-transported error using path gains	3F, 4F, 5F	Oracle upper-bound broadcast
Path propagation	Approximate path-gain attenuation applied to broadcast	3F, 4F, 5F	Morphology-aware modifier for deeper trees
Pathway vector	Structured branch/pathway-specific broadcast field	3F, 4F, 5F	Used mainly in PV-LocalCA on routed tasks
Branch-role rules	Scales updates by router- or pathway-derived branch specialization	3F, 4F, 5F	Structured pathway extension
Inhibitory homeostasis	Auxiliary local inhibitory balancing term	3F, 4F, 5F	Stabilizer for routed/shunting settings

Table S6: Compatibility of the newer LocalCA extensions with the base 3F/4F/5F family. These extensions mostly modify the error field or branch scaling rather than defining a new “6F” or “7F” rule.

536 C Theoretical Details

537 Variant Taxonomy

538 Biological Analogs

539 D Morphology-Aware Extensions

540 **Path-integrated propagation.** Modulate broadcast error by $\pi_n = \pi_{n-1} \cdot R_{n-1}^{\text{tot}} \cdot \bar{g}_{n-1}^{\text{den}}$, approxi-
541 mating depth attenuation from Eq. 4.

542 **Depth modulation.** Per-branch scaling $\rho_j = \rho_{\text{base}} / (d_j + \alpha)$, mirroring cable attenuation.

543 **Dendritic normalization.** $\Delta g_j^{\text{den}} \leftarrow \Delta g_j^{\text{den}} / (\sum_k g_k^{\text{den}} + \varepsilon)$, analogous to homeostatic scaling
544 [22].

545 **Pathway-vector broadcast.** For tasks with latent pathway structure, the broadcast becomes a
546 low-dimensional vector over router-inferred pathway roles. This vector is gated by local pathway

Rule	Factors	Cost	Best regime
3F	$x, (E - V), e$	$\mathcal{O}(1)$	Baseline
4F	$3F + \rho$	$\mathcal{O}(1)$	Improved conditioning
5F	$4F + \phi$	$\mathcal{O}(d_n)$	Strongest overall
PV-LocalCA	$3F + \text{structured } e_n + \text{role alignment}$	$\mathcal{O}(r_n)$	Pathway discovery

Table S7: Variant taxonomy.

Component	Analog	Interpretation
R_n^{tot}	Input resistance	Sensitivity modulation
$(E_j - V_n)$	Synaptic driving force	Local gradient factor
Shunting	Divisive normalization	$\partial V / \partial g_I \propto -V$
ρ_n	Layer relevance	Output correlation
ϕ_n	Signal propagation	Conditional predictability
Pathway vector	Branch-specific top-down drive	Structured teaching signal
Role alignment	Pathway-selective plasticity	Branch specialization

Table S8: Biological analogs.

activity and propagated through the block-structured tree so that earlier branches receive feedback aligned with their downstream descendants.

Router-derived branch roles. Each branch receives a role profile inferred from router assignments and incoming block weights rather than from a hand-coded apical/basal label. Plasticity is then modulated by branch selectivity and a synapse-specific role-alignment factor, emphasizing pathway-consistent updates without imposing a heuristic branch taxonomy.

E HSIC Auxiliary Objectives

Following [17], we use kernel-based HSIC objectives. Self-decorrelation: $\mathcal{L}^{\text{self}} = B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Z \mathbf{H})$. Target-correlation: $\mathcal{L}^{\text{target}} = -B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Y \mathbf{H})$. Moderate weights (0.01–0.1) help on context gating; negligible on MNIST. Online statistics (ρ_n, ϕ_n) use Welford’s algorithm [21].

F Online Variant with Eligibility Traces

Continuous-time eligibility: $\tau_e \dot{e}_j^{\text{syn}} = -e_j^{\text{syn}} + x_j(E_j - V_n)R_n^{\text{tot}}$. Update: $\Delta g_j^{\text{syn}} \propto \int e_j^{\text{syn}}(t)m_n(t)dt$ [18, 19].

G Weight Statistics and Biological References

As a secondary analysis, we examine whether the conductance-based update naturally induces characteristic weight statistics. Because synaptic updates scale with input resistance $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$, the voltage equation implies multiplicative weight dynamics, creating a regime analogous to exponentiated gradient descent [42]. Such multiplicative dynamics suggest log-normal weight distributions in trained networks [40, 41]. We compare the resulting distributions against two biological references: electrophysiological data from Song et al. [40] (rat V1 L5 pyramidal, $\sigma = 0.94$) and the MICrONS mm³ connectome [43] (mouse V1 synapse sizes).

Model weight distributions are log-normal. We fit log-normal distributions to the active (top- K masked) excitatory synaptic weights after training on MNIST and Fashion-MNIST. Table S9 shows the fitted σ (standard deviation of $\log w$) for each condition. Shunting consistently produces narrower distributions than additive integration, and backpropagation produces narrower distributions than local learning. These two factors combine approximately additively: BP+Shunting $\rightarrow$ narrowest, Local+Additive $\rightarrow$ broadest.

Condition	MNIST		Fashion-MNIST	
	Exc σ	Inh σ	Exc σ	Inh σ
BP + Shunting	0.94±0.01	0.87±0.01	0.92±0.01	0.84±0.01
BP + Additive	1.30±0.01	1.22±0.01	1.16±0.01	1.05±0.01
Local + Shunting	1.22±0.01	1.71±0.02	1.13±0.06	1.36±0.06
Local + Additive	1.69±0.16	1.57±0.07	1.43±0.16	1.11±0.12
Song et al. 2005	$\sigma = 0.94$ (rat V1 L5, EPSP amplitude)			
MICrONS E→E	$\sigma = 1.14$ (mouse V1, synapse size, $n=3.9M$)			
MICrONS I→E	$\sigma = 0.92$ (mouse V1, synapse size, $n=3.5M$)			

Table S9: **Weight distribution width.** Log-normal σ of excitatory and inhibitory synaptic weights (mean±std across 5 seeds). Biological references are included only as descriptive anchors. With biologically constrained E/I ratios (4:1), shunting self-normalization produces weight heterogeneity in the range spanned by those references, whereas additive integration is broader and less stable.

Shunting is closest to the biological references considered here. BP+Shunting produces $\sigma = 0.94$ on MNIST, very close to the Song et al. reference ($\Delta = 0.008$). Local+Shunting produces $\sigma = 1.13$ – 1.22 , close to MICrONS E→E synapse sizes ($\sigma = 1.14$; $\Delta \leq 0.08$). By contrast, Local+Additive produces $\sigma = 1.43$ – 1.69 , broader than any biological reference considered. This pattern is consistent across datasets (MNIST and Fashion-MNIST), suggesting it reflects structural properties of the learning rule rather than task-specific features.

The match depends on inhibitory conductance. We test sensitivity to the E/I synapse ratio using a sweep over $N_I \in \{0, 2, 5, 10, 20, 30, 40\}$ inhibitory synapses per branch (Appendix Fig. S8). Without inhibition ($N_I=0$), all conditions produce $\sigma > 2.0$, far from biology. With biologically constrained ratios ($N_I=10$ – 20 , corresponding to E/I ratios of 4:1 to 2:1; [45]), shunting networks enter the biological range ($\sigma \approx 0.9$ – 1.2), while the effect of N_I on additive networks is less systematic. The weight distribution shape is also robust to dendritic depth (branch factors [9] through [3,3,3,3]; see Appendix Fig. S8A): BP+Shunting converges toward Song et al. with increasing depth, while Local+Shunting remains near MICrONS E→E across depths.

H Weight Distribution Sensitivity Analysis

MICrONS connectome analysis. We extracted synapse size distributions from the MICrONS mm³ condensed connectome (v1181; 9.0M connections, 71.9K neurons; 89% excitatory). Splitting by pre/post cell type: E→E connections ($n=3.85M$) show $\sigma_{\text{size}}=1.14$ and CV=1.01; I→E ($n=3.53M$) show $\sigma=0.92$ and CV=1.03; E→I ($n=1.11M$) show $\sigma=0.83$; I→I ($n=0.51M$) show $\sigma=0.85$. Inhibitory connections are more frequently multi-synaptic (I→E: 37.5% multi-synapse; E→E: 11.5%), and inhibitory synapses are on average smaller (median 3532 vs. 4676 for E→E).

Sensitivity to ee.synapses. Using the E/I grid sweep ($N_E \in \{20, 40, 80\} \times N_I \in \{0, 2, 5, 10, 20, 30, 40\}$; 168 models), excitatory σ increases modestly with N_E in both architectures (+0.4–0.6 from $N_E=20$ to 80), but the shunting–additive ordering is preserved at all N_E values. The dominant driver of σ is N_I : going from $N_I=0$ to $N_I=10$ reduces shunting σ by approximately 1.0–1.5 units, while the effect on additive is smaller and less systematic.

Sensitivity to depth. From the depth-scaling sweep ([9], [3,3], [3,3,3], [3,3,3,3]; 80 models, 5 seeds each), BP+Shunting excitatory σ decreases with depth ($1.06 \rightarrow 0.90$), converging toward Song et al. (0.94). Local+Shunting σ increases slightly ($1.18 \rightarrow 1.33$) but remains in the biological range at shallow depths. BP+Additive shows the strongest depth sensitivity ($1.38 \rightarrow 0.79$), while Local+Additive diverges with depth ($2.01 \rightarrow 2.44$). The morphology-aware modulators (path propagation, depth/centrality weighting) have negligible effect on σ .

Mechanistic interpretation. The shunting conductance $g_n^{\text{tot}} = g_n^{\text{leak}} + \sum g_n^{\text{exc}} + \sum g_n^{\text{inh}}$ appears in the update factor $R_n^{\text{tot}} = 1/g_n^{\text{tot}}$. As inhibitory conductance increases, g_n^{tot} grows and R_n^{tot} shrinks, compressing the multiplicative scale of weight updates. This is analogous to a learning rate




Figure S8: **Synaptic weight distribution sensitivity to architecture.** (A) Log-normal σ vs. dendritic depth (branch factors [9] to [3,3,3,3]). Both backprop and local shunting remain closer to the biological-width references than their additive counterparts across depth. (B) σ vs. inhibitory synapses per branch (N_I). Without inhibition ($N_I=0$), all conditions are overly broad; with biologically constrained $N_I=10$ –20, shunting enters the biological range. Shaded bands show the sweep over $N_E=20$ to 80. (C) Summary of MNIST and Fashion-MNIST results with biological references. The shunting ordering is consistent across datasets. (D) Excitatory vs. inhibitory weight σ on MNIST with MICrONS E→E and I→E reference bands.

that self-adjusts via the total conductance: strong synapses contribute more to g_n^{tot} , reducing their own future update magnitude. This acts as a natural form of weight-dependent regularization. The resulting equilibrium distribution is log-normal, consistent with Loewenstein et al. [44] (multiplicative spine dynamics) and Saxena et al. [42] (exponentiated gradient). Additive integration lacks this self-normalizing mechanism, producing broader distributions that fall outside the biological range.

I Depth Scaling and Noise Robustness

Depth scaling. Varying dendritic depth from 1–4 layers (branch factors [9] to [3, 3, 3, 3]): shunting local degrades from 63.5% to 57.4%; additive local degrades more steeply from 54.9% to 29.7%. The shunting advantage grows with depth (+8.5 → +27.7 pp). Backprop ceilings remain stable at about 90–92%. See Fig. 5A.

Noise robustness. Gaussian noise $\mathcal{N}(0, \sigma^2)$ on broadcast error: shunting is robust to $\sigma \leq 0.1$ (about 62%) while additive drops from 46.5% to chance at $\sigma=1.0$, confirming that shunting credit signals carry useful learning information beyond the broadcast magnitude. See Fig. 5B.

NeurIPS Paper Checklist

1. Claims [Yes] The abstract and introduction state the paper’s main claims, and the discussion explicitly narrows their scope to the conductance-based setting studied here; see Secs. 1 and 6.

2. Limitations [Yes] The paper includes an explicit limitations paragraph discussing benchmark performance, broadcast assumptions, depth scaling, and the absence of iterative target-propagation-style baselines; see Sec. 6.

- 629 **3. Theory, Assumptions and Proofs** [Yes] The theoretical assumptions are stated with the voltage
630 model and gradient theorem in Sec. 2, and complete proofs or proof sketches are provided in the
631 appendix where needed; see also Appendix C.
- 632 **4. Experimental Result Reproducibility** [Yes] The model family, update rules, benchmark
633 settings, and representative hyperparameters are described in the main text and appendices, with
634 exact configurations summarized in Appendix B.
- 635 **5. Open Access to Data and Code** [No] The anonymous submission does not include a public
636 code release or anonymized repository URL. Public datasets are cited, and we plan to release code
637 and configuration files after review; see Appendix B.
- 638 **6. Experimental Setting / Details** [Yes] The paper specifies the datasets, architectural families,
639 learning modes, broadcast modes, and seed averaging protocol in Sec. 4 and Appendix B.
- 640 **7. Experiment Statistical Significance** [Yes] Headline experimental results report means and
641 standard deviations across 3–5 seeds, and figures/tables indicate when error bars or seed variability
642 are shown; see Sec. 4 and Appendix A.
- 643 **8. Experiments Compute Resource** [Yes] Appendix B reports the compute setting at the level of
644 single-GPU cluster runs, worker type, and rough runtime scale for the experiments.
- 645 **9. Code of Ethics** [Yes] We have read the NeurIPS Code of Ethics and are not aware of any
646 aspect of the work that violates it. The paper uses standard public benchmarks and does not involve
647 prohibited applications.
- 648 **10. Broader Impacts** [Yes] Sec. 6 discusses broader relevance and a short potential-harms
649 paragraph tailored to this work’s most plausible downstream misuse modes.
- 650 **11. Safeguards** [NA] The paper does not release a high-risk generative model, foundation model,
651 or other asset that requires access control or misuse safeguards.
- 652 **12. Licenses** [Yes] The paper cites the public benchmark and connectomic assets it uses, and
653 Appendix B states the asset-usage and release posture for this submission.
- 654 **13. Assets** [NA] The submission does not release a new dataset, benchmark, or pretrained model
655 asset.
- 656 **14. Crowdsourcing and Research with Human Subjects** [NA] The work does not use crowd-
657 sourcing or direct research with human subjects.
- 658 **15. IRB Approvals** [NA] The work does not involve human-subjects research requiring IRB
659 approval.