

---

# Shunting Inhibition and Dendritic Branching Shape Local Credit Assignment

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Biological neurons assign credit across branching dendrites, where synaptic drive,  
2 reversal potential, local voltage, and input resistance jointly determine plasticity.  
3 We study conductance-based dendritic networks with E/I synapse banks, shunting  
4 inhibition, and tree-structured branch-to-soma coupling, and ask when restricted  
5 somatic feedback can approximate compartment-specific backpropagated errors.  
6 Exact gradients factor into *local eligibility*  $\times$  *compartment error*: the eligibility uses  
7 presynaptic activity, driving force, and input resistance, while the fast non-local  
8 term is a path-specific error obtained by transporting a soma error through dendritic  
9 gains. This factorization turns local learning into a credit-signal compression  
10 problem. We test the hypothesis that shunting inhibition helps when it reshapes the  
11 compartment-error field to better match global-scalar, per-soma, low-rank, or path-  
12 structured feedback. Exact-gradient reconstruction verifies the factorization; path-  
13 gain, rank, broadcast-fidelity, inhibition-intervention, and transported-error-oracle  
14 diagnostics support the proposed mechanism. Under nonnegative conductances  
15 and per-soma 5F feedback, shunting LocalCA remains 5–6 percentage points  
16 below matched backpropagation on MNIST, Fashion-MNIST, and figure-ground  
17 MNIST, indicating that feedback-field fidelity remains a major bottleneck. These  
18 results show how E/I conductance, shunting inhibition, and dendritic branching  
19 can reshape credit-signal geometry in restricted local learning.

## 20 1 Introduction

21 Unlike point units in standard neural networks, biological neurons assign credit across branching  
22 dendritic trees. A synapse receives more than presynaptic activity and a scalar postsynaptic activation:  
23 it also has access to local membrane voltage, synaptic reversal potential, total conductance, and input  
24 resistance. Inhibitory synapses can therefore do more than subtract activity: by increasing total local  
25 conductance, inhibition can shunt voltage and change the gain of every dendritic path passing through  
26 that compartment.

27 We examine whether these biophysical ingredients can facilitate local credit assignment (LocalCA).  
28 Backpropagation [16] would assign a distinct error to every dendritic compartment, whereas a  
29 biologically plausible supervisory signal is more likely to arrive as a low-bandwidth somatic or  
30 modulatory broadcast [21]. The central question is therefore not whether such a broadcast can  
31 reproduce unconstrained backpropagation in every setting, but when the exact dendritic error field  
32 is simple enough for it to approximate. Our contribution is not simply to add dendritic structure to  
33 a neural network, but to show how conductance-based synapses, shunting inhibition, and dendritic  
34 topology change the geometry of the credit signal itself.




Figure 1: **Model and credit assignment.** (A) Dendritic E/I unit with branch synapse banks. (B) Layer of such units projecting to a task readout;  $\delta_0 = \partial L / \partial V_0$  is the soma/core error used for feedback. (C) Exact updates multiply local eligibility by path-specific errors  $\delta_{u,n}$ ; LocalCA replaces them with broadcast estimates  $e_{u,n}$  of increasing fidelity.

Starting from conductance-based dendritic voltage equations [1], we derive exact gradients for dendritic trees (Theorem 1). Each synaptic gradient factorizes into a synapse-local eligibility term and a single non-local compartment error:

$$\text{gradient} = \underbrace{\text{local eligibility}}_{\text{presynaptic drive, } E - V, R^{\text{tot}}} \times \underbrace{\text{compartment error}}_{\text{path-specific non-local term}}.$$

The compartment error is path-specific: under identity upward transfer it is the corresponding soma error multiplied by a conductance-stage path gain through that dendritic tree. LocalCA preserves the exact local eligibility and replaces this path-specific error with a broadcast field.

**What is local in LocalCA?** The fast eligibility variables are synapse- or compartment-local: presynaptic drive, local voltage, driving force, and input resistance. The 4F and 5F variants additionally use slowly estimated branch-level statistics. The fast task-dependent non-local quantity is the compartment error, approximated by a broadcast field  $e_n$ . Thus LocalCA variants differ by the broadcast estimator and optional branch-level modulators, while preserving the same local eligibility mechanism.

This factorization turns local credit assignment into a feedback-compatibility problem. Intrinsic low rank of the exact compartment-error matrix is useful, but it is not sufficient: the error field must also align with the particular feedback object available to the rule. We therefore distinguish global scalar broadcast, neuron-wise/per-soma broadcast, random low-rank feedback, path-structured feedback, and exact transported-error oracles. Dendritic branching creates the paths; E/I synapse banks set the local conductance state; and shunting inhibition changes the input resistance along each path. Shunting is not assumed to help monotonically. It helps restricted feedback when inhibitory conductance attenuates high-gain paths or otherwise makes the exact error field better aligned with the available broadcast.

**Related work.** This work builds on models of dendritic signal integration, branch-level nonlinear computation, divisive normalization, three-factor plasticity, and biologically plausible teaching signals [1, 3, 4, 8, 9, 12, 18–20, 23, 27–29, 35, 36, 46]. It also connects to machine-learning approaches that replace exact backpropagation with local, random, predictive, equilibrium, perturbation, or broadcast feedback signals [14–17, 22, 30–34, 37–41]. Our contribution is complementary: we derive the exact conductance-tree credit object and test when shunting inhibition and dendritic morphology make that object broadcast-compatible.

We test this chain directly. Exact-gradient reconstruction verifies the factorization numerically. Gradient-fidelity diagnostics are used as diagnostics, not as a claim of backpropagation equivalence: shunting improves directional alignment relative to an additive control, but the absolute cosine remains modest, while scale mismatch is substantially reduced. Exact-error rank and direct broadcast-fidelity diagnostics show that shunting can make the compartment-error field easier for the feedback field used by the learning rule to approximate. Transported-error oracle experiments show that supplying the exact path-transported compartment error substantially reduces the local-learning gap. Together, these results identify feedback fidelity as an important bottleneck in these experiments.

**Contributions.** This paper makes three contributions. First, it derives the conductance-stage path gain, its local-eligibility and path-error factorization, and a closed-form inhibitory path-gain ratio. Second, it identifies path-gain and broadcast-compatibility metrics that mediate whether restricted somatic feedback can approximate exact compartment errors. Third, it shows empirically that shunting inhibition can improve this approximation while also identifying regimes where scalar or low-rank feedback fails.

**Scope of evidence.** Analytical statements concern the conductance-tree factorization, path gains, and inhibitory path-gain derivatives under the stated steady-state assumptions. Simulation controls test exact reconstruction, broadcast fidelity, transported-oracle upper bounds, and additive/shunting comparisons. Biological interpretations are restricted to branch-local inhibition and morphology-dependent credit routing; transported-error experiments are oracle diagnostics, not proposed biological learning rules.

## 2 Compartmental Voltage Model and Gradient Derivation

For each dendritic compartment, we use a steady-state conductance model derived from discretized passive cable dynamics [1, 2]. In normalized units, leak and inhibitory reversal potentials are set to 0,

the excitatory reversal potential is set to 1, and leak conductance is fixed at 1 (i.e., all conductances are expressed relative to leak). Each compartment computes a conductance-weighted voltage: excitatory conductance pulls the voltage toward the excitatory reversal potential, leak and shunting inhibition pull it toward zero, and dendritic coupling transfers child voltages upward toward the soma. This form makes local sensitivity depend on both driving force ( $E - V$ ) and total input resistance  $R^{\text{tot}}$ . Throughout the paper we write dendritic morphologies as rooted trees  $[b_1, b_2, \dots, b_D]$ , where each factor gives the fan-in from one dendritic stage to the next on the path to the soma. Thus [3, 3] means that each soma aggregates 3 proximal branches and each proximal branch aggregates 3 distal branches, for 9 distal leaves per soma. In the main feedforward experiments, synaptic inputs terminate on dendritic branches rather than directly on the soma. Each branch layer carries explicit excitatory conductance-based synapse banks and, when enabled, learned inhibitory branch-level conductance-based synapse banks driven by nonnegative inputs. This is input-driven shunting inhibition, not a recurrent lateral inhibitory circuit; explicit inhibitory-cell controls are reported in the appendix. The additive and shunting cores are architecture-matched at the tree level and differ only in the branch-level synaptic integration rule. These ingredients play distinct roles in the credit-assignment argument. Branching creates the path structure. E/I synapse banks determine the local conductance state. Shunting turns inhibitory conductance into a multiplicative change in path gain through input resistance. The local learning rule then asks whether a low-bandwidth teaching signal can stand in for the exact path-specific error induced by that tree. **Notation.** We use  $\delta_{0,u}^{(\ell),V} = \partial L / \partial V_{0,u}^{(\ell)}$  for the exact voltage-space error of soma  $u$  in layer  $\ell$ ,  $\delta_{0,u}^{(\ell),a} = \partial L / \partial a_{0,u}^{(\ell)}$  for the corresponding activation-space error when a soma activity  $a_0 = f_0(V_0)$  is used,  $\delta_0$  for the layer-wide soma/core teaching vector used operationally,  $\bar{\delta}$  for the global scalar obtained by averaging soma coordinates within an example, and  $\delta_{u,n}^{(\ell)} = \partial L / \partial V_{u,n}^{(\ell)}$  for compartment  $n$  of neuron  $u$  in layer  $\ell$ . We use  $\alpha_n^{\text{cond}}$  for the conductance-stage path gain from compartment  $n$  to the soma and  $\tilde{\alpha}_n$  for the effective path gain after including any post-voltage activation derivatives; when upward transfer is identity,  $\tilde{\alpha}_n = \alpha_n^{\text{cond}}$ . We use  $N_E$  and  $N_I$  for the number of excitatory and inhibitory synapses per branch,  $r_n^{\text{4F}}$  for the empirical 4F branch-soma covariance proxy,  $\phi_n$  for the 5F bounded branch-level preconditioner,  $\lambda_{\text{mix}}$  for the pathway-broadcast mixing weight, and  $\kappa_j$  for depth-dependent modulation in morphology extensions. All conductances and presynaptic activities in the main experimental setting are nonnegative; additive voltages may be signed comparators, but the conductance variables that define the shunting model are nonnegative.

*Voltage equation and local sensitivities.* Consider compartment  $n$  with synaptic inputs  $j$  (activity  $x_j$ , reversal  $E_j$ , conductance  $g_j^{\text{syn}} \geq 0$ ) and dendritic inputs from children (voltage  $V_j$ , conductance  $g_j^{\text{den}} \geq 0$ ). The steady-state voltage is:

$$V_n = \frac{\sum_j E_j x_j g_j^{\text{syn}} + \sum_j V_j g_j^{\text{den}}}{\underbrace{\sum_j x_j g_j^{\text{syn}} + \sum_j g_j^{\text{den}} + 1}_{g_n^{\text{tot}}}}, \quad R_n^{\text{tot}} = 1/g_n^{\text{tot}}. \quad (1)$$

$V_n$  is a convex combination of reversal potentials, child voltages, and leak. Writing  $\mathcal{S}_n = \{0\} \cup \{E_j\}_j \cup \{V_j\}_j$ , we have  $\min \mathcal{S}_n \leq V_n \leq \max \mathcal{S}_n$  and  $0 < R_n^{\text{tot}} \leq 1$ . The local sensitivities follow directly:

**Proposition 1** (Local Sensitivities).

$$\frac{\partial V_n}{\partial g_i^{\text{syn}}} = x_i R_n^{\text{tot}} (E_i - V_n), \quad \frac{\partial V_n}{\partial V_i} = g_i^{\text{den}} R_n^{\text{tot}}, \quad \frac{\partial V_n}{\partial g_i^{\text{den}}} = R_n^{\text{tot}} (V_i - V_n). \quad (2)$$

Eq. (2) gives the eligibility factors used by the local rules below: presynaptic activity or voltage difference, driving force, and input resistance. If a child transmits post-voltage activity  $a_i = f_i(V_i)$  rather than raw voltage, the dendritic-conductance derivative is  $R_n^{\text{tot}}(a_i - V_n)$  and the child-voltage error includes the local factor  $f'_i(V_i)$ .

*Shunting inhibition as divisive gain control.* An inhibitory synapse with  $E_{\text{inh}} \approx 0$  contributes current  $(0 - V_n)x_j g_j^{\text{syn}}$  and increases  $g_n^{\text{tot}}$ . Its sensitivity is  $\partial V_n / \partial g_j^{\text{syn}} = -x_j R_n^{\text{tot}} V_n$ , which corresponds to multiplicative attenuation (divisive normalization). Shunting is divisive at the voltage level, but its effect on firing rates can be subtractive in some regimes [10]. Inhibitory plasticity can balance

excitation dynamically [11]; our learned inhibitory conductances provide a trainable balancing mechanism. The same denominator also gives the direct connection to path gain. Writing  $G_n^E(x)$  and  $G_n^I(x)$  for the total excitatory and inhibitory synaptic conductance impinging on compartment  $n$ ,

$$R_n^{\text{tot}}(x) = (1 + G_n^E(x) + G_n^I(x) + G_n^{\text{den}})^{-1}.$$

Increasing  $G_n^I(x)$  lowers  $R_n^{\text{tot}}(x)$  and therefore multiplicatively suppresses every upstream path whose error must pass through compartment  $n$ . In this sense inhibition can gate credit flow by changing the path gain, even when the inhibitory drive is feedforward rather than lateral. Prop. 2 formalizes this statement after the tree path gain is defined.

**Exact gradients for dendritic trees.** For one neuron/tree in layer  $\ell$ , let  $V_{0,u}^{(\ell)}$  be the somatic voltage and define the exact voltage-space soma error as  $\delta_{0,u}^{(\ell),V} := \partial L / \partial V_{0,u}^{(\ell)}$ . If a decoder receives soma activity  $a_0 = f_0(V_0)$  rather than raw voltage, the decoder supplies an activation-space error  $\delta_0^a$  and the voltage-space boundary is  $\delta_0^V = f'_0(V_0)\delta_0^a$ . The theoretical derivation is stated for the conductance-stage voltage before any post-voltage reactivation. When a branch applies a monotone reactivation before transmitting activity upward, the same recursion holds with additional local derivative factors along the path. All exact-error diagnostics below are computed as pre-rectivation voltage errors, including activation-derivative factors when reactivation is enabled.

**Theorem 1** (Conductance-Stage Backpropagation on a Dendritic Tree). *For a rooted dendritic tree with soma at node 0 and unique parent  $p(n)$  for each non-somatic compartment, suppose the parent receives either the pre-rectivation voltage  $V_n$  or a local post-voltage activity  $a_n = f_n(V_n)$ . For identity upward transfer, the loss gradient satisfies*

$$\frac{\partial L}{\partial V_n} = \frac{\partial L}{\partial V_{p(n)}} R_{p(n)}^{\text{tot}} g_{n \rightarrow p(n)}^{\text{den}}, \quad (3)$$

with boundary condition  $\partial L / \partial V_0 = \delta_0^V$ . With post-voltage transfer  $a_n = f_n(V_n)$ , the local factor is the product of  $f'_n(V_n)$ ,  $R_{p(n)}^{\text{tot}}$ , and  $g_{n \rightarrow p(n)}^{\text{den}}$ . Because the graph is a tree, each compartment has a unique path to the soma. Defining the conductance-stage path gain

$$\alpha_n^{\text{cond}} = \prod_{(i \rightarrow k) \in \text{path}(n \rightsquigarrow 0)} R_k^{\text{tot}} g_{i \rightarrow k}^{\text{den}}, \quad \frac{\partial L}{\partial V_n} = \alpha_n^{\text{cond}} \delta_0^V, \quad (4)$$

with  $\alpha_0^{\text{cond}} = 1$  under identity transfer. In the post-voltage case, the exact effective path gain is

$$\tilde{\alpha}_n = \prod_{(i \rightarrow k) \in \text{path}(n \rightsquigarrow 0)} f'_i(V_i) R_k^{\text{tot}} g_{i \rightarrow k}^{\text{den}}, \quad \frac{\partial L}{\partial V_n} = \tilde{\alpha}_n \delta_0^V, \quad (5)$$

where  $f'_i(V_i)$  is omitted for identity transfer.

*Proof.* Apply the chain rule on the tree-structured computation graph using Prop. 1.  $\square$

The one-step recursion in Eq. (3) and the path product in Eq. (4) separate local conductance transfer from the non-local somatic error for one soma/tree and one layer. Supplying an approximate layer-soma teaching vector, reusing a decoder-derived vector at another layer, or compressing  $\delta_{0,u}^{(\ell),V}$  into a per-soma or global scalar broadcast is an additional feedback-approximation step introduced by LocalCA, not a consequence of the theorem itself.

**Proposition 2** (Inhibitory control of conductance-stage path gain). *Let  $\mathcal{A}(n)$  be the set of compartments on the path from compartment  $n$  to the soma, excluding  $n$  and including the parent compartments through which the error must pass. Holding dendritic coupling conductances fixed, an added inhibitory conductance  $\Delta G_k^I(x) \geq 0$  at any  $k \in \mathcal{A}(n)$  changes the conductance-stage path gain by*

$$\frac{\alpha_n^{\text{cond}}(x; \Delta G^I)}{\alpha_n^{\text{cond}}(x; 0)} = \prod_{k \in \mathcal{A}(n)} \frac{g_k^{\text{tot}}(x)}{g_k^{\text{tot}}(x) + \Delta G_k^I(x)}, \quad \frac{\partial \log \alpha_n^{\text{cond}}}{\partial G_k^I} = \begin{cases} -R_k^{\text{tot}}, & k \in \mathcal{A}(n), \\ 0, & k \notin \mathcal{A}(n). \end{cases} \quad (6)$$

*Proof.* Substitute  $R_k^{\text{tot}} = 1/g_k^{\text{tot}}$  into the path-gain product in Eq. (4). Adding inhibitory conductance changes only the denominator at compartments on the path from  $n$  to the soma, which gives the product ratio and the log derivative in Eq. (6).  $\square$

Prop. 2 makes two directed predictions. First, inhibition is a path gate: it attenuates all credit paths below the inhibited compartment. Second, inhibition is useful for restricted-feedback local learning only when this attenuation makes the distribution of  $\alpha_n^{\text{cond}}$  more uniform or better aligned with the task. If inhibition is too strong, too heterogeneous, or generated by a poorly calibrated inhibitory population, it can suppress useful signals and hurt learning. In the additive control the same inhibitory input changes voltages and local eligibilities, but it does not enter the conductance-stage path recursion as a multiplicative resistance gate.

**When does shunting compress path gains?** The attenuation identity above is not by itself a concentration result. Let  $z_n = \log \alpha_n^{\text{cond}}$  be the log conductance-stage path gain for compartment  $n$ . If added inhibitory conductance contributes a path-dependent attenuation

$$\beta_n(x) = \sum_{k \in \mathcal{A}(n)} \log \left( 1 + \frac{\Delta G_k^I(x)}{g_k^{\text{tot}}(x)} \right), \quad z'_n = z_n - \beta_n, \quad (7)$$

then, over compartments or examples,

$$\text{Var}(z') = \text{Var}(z) + \text{Var}(\beta) - 2 \text{Cov}(z, \beta). \quad (8)$$

Thus shunting narrows the log path-gain field exactly when

$$\text{Cov}(z, \beta) > \frac{1}{2} \text{Var}(\beta). \quad (9)$$

Eq. (7) follows by taking the logarithm of the product ratio in Prop. 2; Eq. (8) then applies the variance identity for  $z - \beta$ . Eq. (9) is a diagnostic condition under one averaging measure, not the empirical mechanism claimed below. In finite trained networks, the covariance margin can disagree with CV and exact-error-rank summaries because these statistics average over different compartment and sample axes. We therefore use direct path-gain dispersion, exact-error compressibility, and inhibition-intervention measurements as the empirical tests, and report the covariance margin as a boundary check in the appendix.

**Corollary 1** (Local–Global Factorization). *The exact synaptic gradient at compartment  $n$  factorizes as:*

$$\frac{\partial L}{\partial g_i^{\text{syn}}} = \underbrace{x_i R_n^{\text{tot}} (E_i - V_n)}_{\text{synapse-local eligibility}} \cdot \underbrace{\frac{\partial L}{\partial V_n}}_{\text{compartment error}}, \quad (10)$$

where the eligibility term depends only on quantities available at synapse  $i$  (presynaptic activity  $x_i$ , input resistance  $R_n^{\text{tot}}$ , driving force  $E_i - V_n$ ), and the compartment error  $\partial L / \partial V_n$  is the sole non-local quantity.

**Proposition 3** (Path-gain dispersion controls 3F alignment). *For a single example before batch averaging, consider one soma/tree with the exact voltage-space soma error  $\delta_0^V$  supplied, and write the local eligibility for conductance parameter  $i$  as  $z_i$ . Let  $n(i)$  be the compartment containing parameter  $i$ . Under identity transfer replace  $\tilde{\alpha}_{n(i)}$  below by  $\alpha_{n(i)}^{\text{cond}}$ ; with post-voltage transfer use the effective gain  $\tilde{\alpha}_{n(i)}$ . Assume both gradient vectors are nonzero, equivalently  $\sum_i w_i > 0$  and  $\sum_i w_i \tilde{\alpha}_{n(i)}^2 > 0$ . The exact and 3F per-soma gradients are*

$$g_i^{\text{BP}} = z_i \tilde{\alpha}_{n(i)} \delta_0^V, \quad \hat{g}_i^{\text{3F}} = z_i \delta_0^V. \quad (11)$$

If transfer derivatives and conductance couplings are nonnegative,  $\tilde{\alpha}_{n(i)} \geq 0$ , so corresponding nonzero components have the same sign. Their cosine is

$$\cos(\hat{g}^{\text{3F}}, g^{\text{BP}}) = \frac{\sum_i w_i \tilde{\alpha}_{n(i)}}{\sqrt{\sum_i w_i} \sqrt{\sum_i w_i \tilde{\alpha}_{n(i)}^2}} = \frac{1}{\sqrt{1 + \text{CV}_w(\tilde{\alpha}_{n(i)}^2)}}, \quad w_i = (z_i \delta_0^V)^2. \quad (12)$$

*Proof.* Substitute Eq. (11) into the cosine formula and cancel the common signed eligibility factors through  $w_i$ . The final equality is the weighted identity  $E_w[\tilde{\alpha}^2] = E_w[\tilde{\alpha}]^2(1 + \text{CV}_w(\tilde{\alpha})^2)$ .  $\square$

Prop. 3 is conditional on the exact layer-soma error. It explains why low effective path-gain dispersion is a useful within-tree credit geometry, while also making clear that practical multi-layer LocalCA must approximate the layer-soma error before this within-tree argument applies. The factorization in

Eq. (10) isolates the fast task-error-dependent non-local quantity as  $\partial L/\partial V_n = \alpha_n^{\text{cond}} \delta_0^V$  under identity upward transfer and  $\partial L/\partial V_n = \tilde{\alpha}_n \delta_0^V$  when post-voltage transfer is enabled. The implemented updates in Eqs. (13)–(14) are batch averages of products, so Prop. 3 should be read as the pre-batch geometry of per-example gradient contributions. The 4F and 5F factors introduced below use slower voltage statistics, so they are empirical reliability preconditioners rather than additional error signals. Any practical local rule therefore depends on how well its broadcast signal approximates that exact compartment error, which we evaluate in the gradient-fidelity diagnostic below. In the models studied here, shunting improves this approximation by reducing and stabilizing  $R_n^{\text{tot}}$ , which narrows the spread of local sensitivities across the tree. In the empirical networks, each dendritic branch layer may optionally apply a post-voltage reactivation  $f(V_n)$  before passing activity to the next stage; unless otherwise noted, the manuscript sweeps use a learnable bounded tanh reactivation.

### 3 Local Learning Rules

**Broadcast error approximation.** We approximate the exact compartment error  $\partial L/\partial V_n$  (Corollary 1) with a broadcast signal  $e_n^V$  derived from a soma/core teaching vector. For a linear readout this vector is  $W_{\text{dec}}^\top (\partial L/\partial \hat{y})$ ; for a nonlinear decoder we use the corresponding decoder-input Jacobian product  $J_{\text{dec}}(h_{\text{core}})^\top (\partial L/\partial \hat{y})$  so that feedback is defined in soma/core coordinates. If this vector is used at an earlier matched-width dendritic layer, practical LocalCA assumes neuron-index correspondence and reuses the coordinate as an approximate layer-soma error. This inter-layer teaching approximation is distinct from the within-tree question addressed by Theorem 1: given a layer-soma error, how faithfully can a restricted branch-level field approximate  $\partial L/\partial V_n$ ? The main competence and mechanism runs use per-soma feedback: when soma/core dimension matches a layer’s neuron dimension, each neuron receives one soma-specific teaching coordinate and the same value is shared across that neuron’s dendritic compartments. Conditional on an exact layer-soma error, per-soma sharing isolates the within-tree approximation. Practical multi-layer LocalCA additionally approximates earlier-layer soma errors by reusing final-core coordinates at matched-width layers, as separated explicitly in Fig. 2D. A stricter scalar mode compresses the soma/core vector to one number per example,  $\bar{\delta}_b = d_0^{-1} \sum_{c=1}^{d_0} \delta_{0,bc}^a$ , and broadcasts that number to all neurons and compartments. Thus scalar feedback also discards heterogeneity across soma-error coordinates. For higher-bandwidth controls,  $P_K$  is a fixed random projection from soma-error coordinates to  $K$  feedback channels,  $q_{n,k}$  are fixed compartment-specific mixing coefficients, and  $\Gamma_n$  denotes the pathway-structured transport map for compartment  $n$ . We distinguish five feedback objects: **global scalar broadcast**,  $e_n = \bar{\delta} \mathbf{1}$ ; **per-soma shared broadcast**,  $e_{u,n}$  is one coordinate shared across the compartments of neuron  $u$ ; **rank- $K$  random broadcast**,  $c = P_K \delta_0^a$  and  $e_n = \sum_k q_{n,k} c_k$ ; **path-structured broadcast**,  $e_n = \lambda_{\text{mix}} \bar{\delta} \mathbf{1} + (1 - \lambda_{\text{mix}}) \Gamma_n(\delta_0^a)$  with pathway roles motivated by vectorized dendritic teaching signals [45]; and the **transported oracle**,  $e_n^{\text{transport}} = \tilde{\alpha}_n \delta_0^V = \partial L/\partial V_n$ , an analysis upper bound because it supplies the exact path-transported compartment error from Theorem 1. Operationally, path-transport runs with reactivation enabled transport activation-space errors through parent reactivation slopes and multiply by the local  $f'_n(V_n)$  before conductance updates; with identity transfer this reduces to  $\alpha_n^{\text{cond}} \delta_0^V$ . Thus the main feedforward result tests neuron-wise/per-soma feedback rather than a single layer-wide scalar. Under shunting, this restricted field supports nontrivial learning on the standard classification tasks, while scalar, routed, and rank-bridge controls in the appendix expose regimes where lower-bandwidth feedback collapses branch identity and richer feedback is needed. A local-mismatch heuristic is included only as an appendix control because it performs much worse than the somatic-error broadcasts (Appendix A).

*Three-factor learning rule (3F).*

**Definition 1** (3F Update). *For synaptic and dendritic conductances, with  $a_j = V_j$  under identity upward transfer and  $a_j = f_j(V_j)$  when post-voltage transfer is enabled:*

$$\hat{\nabla}_{g_j^{\text{syn}}} L = \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n^V \rangle_B, \quad \hat{\nabla}_{g_j^{\text{den}}} L = \langle R_n^{\text{tot}} (a_j - V_n) e_n^V \rangle_B, \quad (13)$$

where  $\langle \cdot \rangle_B$  denotes the batch average.

The three factors are presynaptic activity  $x_j$  (or a transmitted child activity difference), postsynaptic modulation through driving force and input resistance, and the broadcast voltage-space error  $e_n^V$ . The same rule applies to excitatory and inhibitory synapses. The sign difference comes from the




Figure 2: **Exact reconstruction and gradient diagnostics.** (A) Exact-factorization reconstruction matches autograd. (B,C) Final LocalCA-backpropagation cosine and scale mismatch; dots show seeds. (D) Layer-soma factorial diagnostic separating exact vs. reused soma teaching signals and exact path transport vs. per-soma sharing. Metrics/protocol: Appendix A. Error bars in (B,C) are  $\pm 1$  s.d.

257 driving force ( $E_j - V_n$ ). We write  $\widehat{\nabla}_g L$  as the local gradient estimate supplied to the optimizer; the  
 258 optimizer applies the usual learning rate and descent step. Equivalently, one could absorb the sign  
 259 convention into  $e_n^V$  and treat it as a descent signal.

260 **Additive control.** Rule (13) is the local gradient of the *shunting* voltage  $V_n = \sum E_j g_j x_j / g_n^{\text{tot}}$ . For  
 261 the additive E/I control, inhibition enters as a fixed signed voltage contribution rather than through  
 262 the conductance denominator:

$$V_n^{\text{add}} = \sum_{j \in E} g_j^E x_j^E - \sum_{j \in I} g_j^I x_j^I + \sum_k g_k^{\text{den}} a_k.$$

263 Equivalently,  $\partial V_n^{\text{add}} / \partial g_j^{\text{syn}} = s_j x_j$  with  $s_j = +1$  for excitatory and  $s_j = -1$  for inhibitory synapses,  
 264 and  $\partial V_n^{\text{add}} / \partial g_k^{\text{den}} = a_k$ . The additive local gradient therefore has no driving-force or  $R_n^{\text{tot}}$  terms  
 265 ( $R_n^{\text{tot}} \equiv 1$  by definition since there is no denominator). For the identity-transfer additive tree, the exact  
 266 upward gain is  $\alpha_n^{\text{add}} = \prod_{(i \rightarrow k) \in \text{path}(n \rightsquigarrow 0)} g_{i \rightarrow k}^{\text{den}}$ ; with reactivation,  $\tilde{\alpha}_n^{\text{add}} = \prod_{(i \rightarrow k)} f'_i(V_i) g_{i \rightarrow k}^{\text{den}}$ .  
 267 Thus additive controls have an exact tree gain, but inhibition does not enter that gain through a  
 268 conductance denominator. Throughout, each architecture uses the learning rule derived from its  
 269 own forward-pass dynamics. This keeps the comparison architecture-matched rather than applying a  
 270 shunting-derived rule to an additive model. In implementations with unconstrained raw parameters  $\theta$   
 271 and  $g = \text{softplus}(\theta)$ , Eqs. (13)–(14) give conductance-space gradients; before assigning gradients to  
 272  $\theta$ , we multiply by  $\partial g / \partial \theta$ . The exact-gradient reconstruction diagnostics include this parameterization  
 273 factor; Appendix B gives the compact shunting/additive comparison.






Figure 3: **Mechanistic chain from path gains to learning.** (A) Measured path-gain field. (B) MNIST exact-error geometry: unrestricted rank-1 SVD residual  $\rho_1$  and actual per-soma broadcast cosine. (C) Post-training inhibitory-conductance interventions. (D,E) Broadcast fidelity and transported-error learning on noise resilience. CV,  $\rho_1$ , and interventions: Appendix A. Error bars are  $\pm 1$  s.d.

274 *Practical heuristic wrappers: 4F and 5F.* The 4F and 5F variants add slow branch-level preconditioners to the same fast eligibility and broadcast error; they are empirical stabilizers, not additional  
 275 task-error channels. 4F uses a branch-soma covariance proxy and 5F adds a bounded predictability  
 276 factor,  
 277

$$r_n^{4F} = \frac{\text{Cov}_B(\bar{V}_{n,b}, \bar{V}_{0,b})}{\sqrt{\text{Var}_B(\bar{V}_{n,b}) \text{Var}_B(\bar{V}_{0,b})} + \varepsilon}, \quad \phi_n = \frac{\text{Var}(V_n)}{\sigma_{\text{res}}^2 + \varepsilon},$$

278 with implementation clamps  $r_n^{4F} \in [0.1, 2.0]$  and  $\phi_n \in [0.25, 4.0]$ . The appendix specifies the  
 279 averaging axes, EMA initialization, detaching, and clamp rationale. Empirically, 4F alone gives little  
 280 improvement over 3F; the practical rule is the 5F update below.

**Definition 2** (5F Update).

$$\hat{\nabla}_{g_j^{\text{syn}}} L = r_n^{4F} \phi_n \langle x_j R_n^{\text{tot}} (E_j - V_n) e_n^V \rangle_B, \quad \hat{\nabla}_{g_j^{\text{den}}} L = r_n^{4F} \phi_n \langle R_n^{\text{tot}} (a_j - V_n) e_n^V \rangle_B. \quad (14)$$

281 Eq. (14) is therefore the main practical LocalCA update. The 3F rule uses the theorem-derived local  
 282 eligibility with an approximate broadcast error, while 4F and 5F add empirical reliability factors. A  
 283 feedback-alignment-style random-broadcast argument is included in Appendix C; it is supportive but  
 284 not central to the main mechanistic claim.

## 285 4 Experiments




Figure 4: **Matched-capacity performance and regime dependence.** (A) Matched backprop vs. per-soma 5F LocalCA. (B) Inhibitory dose-response. (C) Morphology-dependent shunting advantage. (D) MNIST controls: BP, exact path transport (PT), activation choice, and additive gain/normalization variants. Values/protocols: Table S3; Appendix A. Error bars denote  $\pm 1$  s.d.

286 *Setup.* The experiments test factorization, gradient fidelity, exact-error geometry, inhibition necessity,  
 287 and transported-error oracles using architecture-matched additive/shunting cores and local/backprop  
 288 training on MNIST [43], Fashion-MNIST [42], figure-ground MNIST, and noise resilience; CIFAR-  
 289 10 [44], cue routing, DFA, and detailed protocols are in the appendix. Unless stated otherwise,  
 290 local runs use 5F with per-soma feedback, conductance runs use nonnegative first-layer drive, and  
 291 main mechanism/oracle/competence results use five seeds. Transported error is an oracle upper  
 292 bound from Theorem 1, not the proposed biological rule. Figure shorthand: BP backpropagation, PS  
 293 per-soma feedback, PT exact path transport, PV pathway-vector feedback, Or. transported oracle,  $K$   
 294 random low-rank channels,  $N_I$  inhibitory synapses/branch, pp percentage points, and MN/FM/FG =  
 295 MNIST/Fashion-MNIST/figure-ground MNIST. Metrics, implementation, and tasks are defined in  
 296 Appendices A, B, and D.

297 *From exact factorization to credit-signal fidelity.* We compare local and backprop gradients  
 298 on matched batches and parameters using parameter-count-weighted cosine and scale mismatch  
 299  $|\log_{10}(\|g_p^{\text{local}}\|/\|g_p^{\text{bp}}\|)|$ . Shunting improves the final diagnostic, though absolute directional agree-  
 300 ment remains weak: cosine rises from  $0.005 \pm 0.026$  to  $0.100 \pm 0.033$ , and scale mismatch falls from  
 301  $0.50 \pm 0.05$  to  $0.26 \pm 0.05$ . Exact-gradient reconstruction matches autograd to numerical precision  
 302 (Fig. 2A). The layer-soma factorial diagnostic separates the two practical approximations: exact layer-  
 303 soma errors plus path transport recover autograd, while per-soma compression and reused final-core  
 304 coordinates mainly damage the first layer (Fig. 2D). Thus the central non-local approximation is the  
 305 broadcast substitute for  $\partial L/\partial V_n$ , not the tree recursion.

306 *Mechanistic chain: path gains, broadcast fidelity, and oracle transport.*




Figure 5: **Rule and feedback controls.** (A) Rule ablation. (B) Error-source ablation. (C) Noise-resilience feedback ladder: PS, path proxy, low-rank  $K$ , and transported oracle. (D) Cue-routing rank diagnostic with pathway-vector feedback (PV). Details: Appendix Fig. S1; Table S7.

Figure 3 gives the central empirical chain. At matched inhibition ( $N_I=5$ ), shunting narrows the MNIST conductance-stage path-gain component (CV 0.23 vs. 1.16 for additive), lowers exact-error rank (rank-1 residual  $0.70 \pm 0.21$  vs.  $0.81 \pm 0.11$ ), and improves compatibility with the implemented feedback field (actual per-soma broadcast cosine  $0.74 \pm 0.13$  vs.  $0.37 \pm 0.18$ ). Panel A is an architectural path-gain diagnostic rather than a direct validation of Prop. 3; the SVD and broadcast-cosine panels test intrinsic geometry and implemented-feedback compatibility, respectively.

Learned branch inhibition is functionally important: zeroing, shuffling, or clamping it drops MNIST from about 90% to 20–26% and noise resilience from about 81% to 10–19%. Direct fidelity and oracle panels probe the same bottleneck: under shunting, per-soma feedback better tracks exact compartment errors, transported-error fidelity approaches 0.95–1.0, and transported learning lifts both architectures to 92–95% on noise resilience.

*Matched-capacity performance.* In the matched-capacity regime, shunting LocalCA remains within 5–6 percentage points of matched backpropagation on MNIST, Fashion-MNIST, and figure-ground MNIST (Table S3; Fig. 4). Controls in Fig. 4D identify a feedback-field bottleneck: exact path transport reaches the matched shunting backprop reference, activation choice has little effect, and additive gain/normalization variants do not rescue the additive baseline. Fig. 5 shows that 5F is the strongest practical LocalCA variant, soma-derived feedback is essential, and higher-rank/path feedback improves the noise-resilience bridge.

*Regime dependence of the shunting advantage.* With matched per-soma feedback, shunting is modestly better on MNIST, Fashion-MNIST, and figure-ground MNIST, and the gap grows on inhibition-sensitive noise resilience. Morphology shifts the useful inhibitory regime, so the shunting advantage is not monotonic in  $N_I$  or tree depth (Fig. 4B,C; Appendix Fig. S6).

## 5 Discussion

We studied local credit assignment in conductance-based dendritic networks with explicit E/I conductance synapses, shunting inhibition, and tree-structured branching. Exact gradients separate synapse-local eligibility–presynaptic activity, driving force, and input resistance—from one path-specific compartment error, making LocalCA an error-approximation problem. The experiments support a coherent mechanism: branching creates credit paths, E/I synapse banks set the conductance state, and shunting modulates input resistance along those paths. In the regimes studied here this lowers exact-error rank, improves feedback-fidelity diagnostics, reduces gradient-scale mismatch, and lets transported-error oracles substantially reduce the local-learning gap. Post-training interventions further show that learned inhibitory conductance is sample-dependent and functionally important to the trained forward computation.

The operating regime matters. Shunting is not universally beneficial, more inhibition is not always better, and the practical 5F rule is an empirical bounded preconditioner. Exact decoder-to-soma feedback requires the decoder Jacobian or linear readout to supply the correct soma/core coordinate; practical multi-layer LocalCA also approximates layer-soma errors by reusing coordinates at earlier matched-width layers. The core theorem covers conductance-stage branch parameters under steady-state, non-spiking dynamics; reactivation, decoder, router, top- $k$ , 4F/5F batch-time statistics, and HSIC terms are implementation-specific updates. Eq. (9) is interpretive, and activation shape, morphology, task structure, and feedback rank can change broadcast compatibility. The inhibition studied here is input-driven rather than recurrent lateral inhibition, and post-training interventions establish forward necessity rather than isolating credit-assignment causality. We therefore claim a credit-geometry mechanism supported by exact reconstruction, fidelity diagnostics, controls, and oracle transport, not competitive CIFAR-10 performance or sufficiency of a single global scalar broadcast for routed/pathway-specific tasks.

## References

- [1] Koch, C. (1999). *Biophysics of Computation: Information Processing in Single Neurons*. Oxford University Press.
- [2] Abbott, L. F., & Dayan, P. (2001). *Theoretical Neuroscience: Computational and Mathematical Modeling of Neural Systems*. MIT Press.
- [3] Poirazi, P., Brannon, T., & Mel, B. W. (2003). Pyramidal neuron as two-layer neural network. *Neuron*, 37(6), 989–999.
- [4] London, M., & Häusser, M. (2005). Dendritic computation. *Annual Review of Neuroscience*, 28, 503–532.
- [5] Spruston, N. (2008). Pyramidal neurons: dendritic structure and synaptic integration. *Nature Reviews Neuroscience*, 9, 206–221. <https://doi.org/10.1038/nrn2286>
- [6] Branco, T., & Häusser, M. (2010). The single dendritic branch as a fundamental functional unit in the nervous system. *Current Opinion in Neurobiology*, 20(4), 494–502. <https://doi.org/10.1016/j.conb.2010.07.009>
- [7] Larkum, M. (2013). A cellular mechanism for cortical associations: an organizing principle for the cerebral cortex. *Trends in Neurosciences*, 36(3), 141–151. <https://doi.org/10.1016/j.tins.2012.11.006>
- [8] Urbanczik, R., & Senn, W. (2014). Learning by the dendritic prediction of somatic spiking. *Neuron*, 81(3), 521–528.
- [9] Carandini, M., & Heeger, D. J. (2012). Normalization as a canonical neural computation. *Nature Reviews Neuroscience*, 13(1), 51–62.

- [10] Holt, G. R., & Koch, C. (1997). Shunting inhibition does not have a divisive effect on firing rates. *Neural Computation*, 9(5), 1001–1013.
- [11] Vogels, T. P., Sprekeler, H., Zenke, F., Clopath, C., & Gerstner, W. (2011). Inhibitory plasticity balances excitation and inhibition in sensory pathways and memory networks. *Science*, 334(6062), 1569–1573.
- [12] Frémaux, N., & Gerstner, W. (2016). Neuromodulated spike-timing-dependent plasticity, and theory of three-factor learning rules. *Frontiers in Neural Circuits*, 9, 85. <https://doi.org/10.3389/fncir.2015.00085>
- [13] Hennequin, G., Agnes, E. J., & Vogels, T. P. (2017). Inhibitory plasticity: balance, control, and codependence. *Annual Review of Neuroscience*, 40, 557–579. <https://doi.org/10.1146/annurev-neuro-072116-031005>
- [14] Lillicrap, T. P., Cownden, D., Tweed, D. B., & Akerman, C. J. (2016). Random synaptic feedback weights support error backpropagation for deep learning. *Nature Communications*, 7, 13276.
- [15] Lillicrap, T. P., Santoro, A., Marris, L., Akerman, C. J., & Hinton, G. (2020). Backpropagation and the brain. *Nature Reviews Neuroscience*, 21(6), 335–346. <https://doi.org/10.1038/s41583-020-0277-3>
- [16] Rumelhart, D. E., Hinton, G. E., & Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323, 533–536.
- [17] Nøkland, A. (2016). Direct feedback alignment provides learning in deep neural networks. *NeurIPS*, 29, 1037–1045.
- [18] Guerguiev, J., Lillicrap, T. P., & Richards, B. A. (2017). Towards deep learning with segregated dendrites. *eLife*, 6, e22901.
- [19] Sacramento, J., Costa, R. P., Bengio, Y., & Senn, W. (2018). Dendritic cortical microcircuits approximate the backpropagation algorithm. *NeurIPS*, 31.
- [20] Whittington, J. C., & Bogacz, R. (2019). Theories of error back-propagation in the brain. *Trends in Cognitive Sciences*, 23(3), 235–250.
- [21] Richards, B. A., & Lillicrap, T. P. (2019). Dendritic solutions to the credit assignment problem. *Current Opinion in Neurobiology*, 54, 28–36.
- [22] Scellier, B., & Bengio, Y. (2017). Equilibrium Propagation: Bridging the Gap between Energy-Based Models and Backpropagation. *Frontiers in Computational Neuroscience*, 11, 24. <https://doi.org/10.3389/fncom.2017.00024>
- [23] Song, Y., Millidge, B., Salvatori, T., Lukasiewicz, T., Xu, Z., & Bogacz, R. (2024). Inferring neural activity before plasticity as a foundation for learning beyond backpropagation. *Nature Neuroscience*, 27, 348–358. <https://doi.org/10.1038/s41593-023-01514-1>
- [24] Gretton, A., Bousquet, O., Smola, A. J., & Schölkopf, B. (2005). Measuring statistical dependence with Hilbert-Schmidt norms. *Lecture Notes in Computer Science*, 3734, 63–77. [https://doi.org/10.1007/11564089\\_7](https://doi.org/10.1007/11564089_7)
- [25] Welford, B. P. (1962). Note on a method for calculating corrected sums of squares and products. *Technometrics*, 4(3), 419–420. <https://doi.org/10.1080/00401706.1962.10490022>
- [26] Turrigiano, G. G. (2008). The self-tuning neuron: synaptic scaling of excitatory synapses. *Cell*, 135(3), 422–435.
- [27] Payeur, A., Guerguiev, J., Zenke, F., Richards, B. A., & Naud, R. (2021). Burst-dependent synaptic plasticity can coordinate learning in hierarchical circuits. *Nature Neuroscience*, 24(7), 1010–1019.
- [28] Greedy, W., Zhu, H. W., Pemberton, J., Mellor, J., & Ponte Costa, R. (2022). Single-phase deep learning in cortico-cortical networks. *NeurIPS*, 35, 24213–24225.

- [29] Haider, P., Ellenberger, B., Kriener, L., Jordan, J., Senn, W., & Petrovici, M. A. (2021). Latent Equilibrium: A unified learning theory for arbitrarily fast computation with arbitrarily slow neurons. *NeurIPS*, 34, 17839–17851.
- [30] Hinton, G. (2022). The Forward-Forward Algorithm: Some Preliminary Investigations. *arXiv:2212.13345*.
- [31] Dellaferrera, G., & Kreiman, G. (2022). Error-driven input modulation: Solving the credit assignment problem without a backward pass. *Proceedings of the 39th International Conference on Machine Learning, Proceedings of Machine Learning Research*, 162, 4937–4955.
- [32] Lee, D.-H., Zhang, S., Fischer, A., & Bengio, Y. (2015). Difference target propagation. *Machine Learning and Knowledge Discovery in Databases: ECML PKDD 2015*, 498–515. [https://doi.org/10.1007/978-3-319-23528-8\\_31](https://doi.org/10.1007/978-3-319-23528-8_31)
- [33] Meulemans, A., Tristany Farinha, M., García Ordóñez, J., Vilimelis Aceituno, P., Sacramento, J., & Grewe, B. F. (2021). Credit assignment in neural networks through deep feedback control. *NeurIPS*, 34, 4674–4687.
- [34] Millidge, B., Seth, A. K., & Buckley, C. L. (2021). Predictive coding: A theoretical and experimental review. *arXiv:2107.12979*.
- [35] Koch, C., Poggio, T., & Torre, V. (1983). Nonlinear interactions in a dendritic tree: localization, timing, and role in information processing. *PNAS*, 80(9), 2799–2802. <https://doi.org/10.1073/pnas.80.9.2799>
- [36] Silver, R. A. (2010). Neuronal arithmetic. *Nature Reviews Neuroscience*, 11(7), 474–489.
- [37] Max, K., Kriener, L., Pineda García, G., Nowotny, T., Jaras, I., Senn, W., & Petrovici, M. A. (2024). Learning efficient backprojections across cortical hierarchies in real time. *Nature Machine Intelligence*, 6, 619–630. <https://doi.org/10.1038/s42256-024-00845-3>
- [38] Ma, C., Wu, J., Si, C., & Tan, K. C. (2024). Scaling supervised local learning with augmented auxiliary networks. *International Conference on Learning Representations (ICLR)*.
- [39] Lv, C., Xu, J., Lu, Y., Wang, X., Wang, Z., Xu, Z., Yu, D., Du, X., Zheng, X., & Huang, X. (2025). Dendritic Localized Learning: Toward Biologically Plausible Algorithm. *Proceedings of the 42nd International Conference on Machine Learning, Proceedings of Machine Learning Research*, 267, 41682–41700.
- [40] Erdogan, M., Pehlevan, C., & Erdogan, A. T. (2025). Error Broadcast and Decorrelation as a Potential Artificial and Natural Learning Mechanism. *Advances in Neural Information Processing Systems*, 38.
- [41] Kao, C.-H., & Hariharan, B. (2024). Counter-Current Learning: A Biologically Plausible Dual Network Approach for Deep Learning. *Advances in Neural Information Processing Systems*, 37.
- [42] Xiao, H., Rasul, K., & Vollgraf, R. (2017). Fashion-MNIST: A novel image dataset for benchmarking machine learning algorithms. *arXiv:1708.07747*.
- [43] LeCun, Y., Bottou, L., Bengio, Y., & Haffner, P. (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278–2324.
- [44] Krizhevsky, A. (2009). Learning multiple layers of features from tiny images. Technical report, University of Toronto.
- [45] Francioni, V., Tang, V. D., Toloza, E. H. S., Ding, Z., Brown, N. J., & Harnett, M. T. (2026). Vectorized instructive signals in cortical dendrites. *Nature*, 652, 1254–1263. <https://doi.org/10.1038/s41586-026-10190-7>
- [46] Iyer, A., Grewal, K., Velu, A., Souza, L. O., Forest, J., & Ahmad, S. (2022). Avoiding catastrophe: Active dendrites enable multi-task learning in dynamic environments. *Frontiers in Neurobotics*, 16, 846219.

## A Supplementary Results and Diagnostics

This appendix records the supporting diagnostics, implementation details, and task protocols needed to interpret and reproduce the paper’s main claims.

**Appendix structure.** The appendix is organized to make the paper reproducible from mechanism to implementation. Appendix A first reports the supplementary diagnostics that support or delimit the main empirical claims. Appendix B then specifies the implemented dendritic networks, including morphology, synapse banks, shunting and additive forward equations, positive parameterization, reactivation, decoder errors, and optimizer-level gradient assignment. Appendix C collects learning-rule variants and auxiliary theory details that are useful for reproducing the local updates but are not needed for the main proof. Appendix D gives task construction and robustness protocols.

**Boundary diagnostics.** Several controls are intentionally negative or mixed. The log-gain covariance margin from Eq. (9) is not used as a positive main result, and the pathway-vector feedback variant does not beat the best unstructured low-rank cue-routing control. These diagnostics delimit the claim rather than smoothing over failure modes.

### Reproducibility Summary

Detailed run sources and commands are maintained in `REPRODUCIBILITY.md`. The table below summarizes the experiments, architectures, seed protocols, selection rules, and reproduction commands used for the manuscript figures.

| Experiment family                                           | Architecture / condition                                                                            | Seeds and selection                                                              | Reproduction command                                               |
|-------------------------------------------------------------|-----------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|--------------------------------------------------------------------|
| Exact-factorization and gradient diagnostics                | Nonnegative-input E/I dendritic cores, [128, 128] soma layers, [3, 3] tree, 3F–5F local rules       | Main summaries use 5 seeds; gradient-dynamics panels use 3 seeds where indicated | <code>python scripts/generate_neurips_figures.py</code>            |
| Layer-soma factorial diagnostic                             | Two-layer shunting checkpoints with exact/reused soma errors and transported/per-soma branch fields | Six fixed checkpoints; no validation selection                                   | <code>python scripts/measure_layer_soma_factorial.py</code>        |
| Path-gain, exact-error rank, and inhibition intervention    | Matched additive and shunting cores at fixed inhibition; post-training intervention states          | Five seeds per main condition; interventions use trained checkpoints             | <code>python scripts/generate_theory_diagnostics_figures.py</code> |
| Transported-error oracle and feedback construction controls | Per-soma, random low-rank, path-structured, and exact transported-error feedback                    | Five seeds for main oracle and feedback-ladder summaries                         | <code>python scripts/summarize_path_transport_sweep.py</code>      |
| Matched-capacity performance                                | MNIST, Fashion-MNIST, figure-ground MNIST; per-soma 5F LocalCA and matched backprop references      | Five seeds; validation-selected configurations report held-out test accuracy     | <code>python scripts/summarize_neurips_claim_sweeps.py</code>      |

Table S1: **Reproducibility summary.** Commands are run from the draft directory after the corresponding training or diagnostic outputs have been generated.

### Diagnostic Metrics

The gradient-fidelity diagnostics compare LocalCA and backpropagation gradients at matched parameters and batches. For a parameter block  $p$ , let  $g_p^{\text{local}}$  and  $g_p^{\text{bp}}$  be the two gradients and let  $w_p = \text{numel}(p) / \sum_q \text{numel}(q)$ . The reported weighted cosine is

$$C_w = \sum_p w_p \frac{\langle g_p^{\text{local}}, g_p^{\text{bp}} \rangle}{\|g_p^{\text{local}}\|_2 \|g_p^{\text{bp}}\|_2 + \varepsilon},$$

and weighted scale mismatch is

$$S_w = \sum_p w_p \left| \log_{10} \frac{\|g_p^{\text{local}}\|_2 + \varepsilon}{\|g_p^{\text{bp}}\|_2 + \varepsilon} \right|.$$

For exact-error compressibility, we form a matrix  $E$  whose rows are examples and whose columns are compartment–neuron coordinates of the exact pre-activation compartment error. If  $E = U\Sigma V^\top$




Figure S1: **Detailed rule and feedback controls.** (A) Eligibility-rule ablation: 3F, 4F, and 5F. (B) Error-source ablation: per-soma/shared feedback vs. local mismatch. (C) Feedback-construction ladder on noise resilience: PS per-soma/shared, recursive Path, random low-rank K2/K4/K8, and transported Oracle. (D) Cue routing: PS per-soma/shared, random low-rank K1/K2/K4, and pathway-vector feedback.

and  $\sigma_i$  are the singular values, the unrestricted rank-1 residual and participation rank are

$$\rho_1 = \frac{\sqrt{\sum_{i>1} \sigma_i^2}}{\sqrt{\sum_i \sigma_i^2}}, \quad r_{\text{part}} = \frac{(\sum_i \sigma_i^2)^2}{\sum_i \sigma_i^4}.$$

This is an intrinsic geometry metric: both the sample coefficients and spatial template are chosen in hindsight by SVD. The constrained shared-field residual instead fixes the spatial template to all ones and allows only the best per-example scalar,

$$\rho_1 = \frac{\|E - \text{rowmean}(E)\mathbf{1}^\top\|_F}{\|E\|_F}.$$

The per-soma/blockwise residual matches the practical feedback object more closely. Let  $\Pi_{\text{PS}}$  replace, for each sample and soma, all compartments of that soma by their optimal shared coefficient. Then

$$\rho_{\text{PS}} = \frac{\|E - \Pi_{\text{PS}}E\|_F}{\|E\|_F}.$$

The broadcast-conditioned residual uses the actual feedback field  $B$  generated by the implemented mode,

$$\rho_{\text{bc}} = \frac{\|E - B\|_F}{\|E\|_F},$$

with a corresponding flattened cosine and, where useful, a globally rescaled version of  $B$ . These constrained metrics are emitted by `scripts/measure_error_rank_diagnostics.py` alongside the unrestricted SVD metrics, so intrinsic error geometry and compatibility with the implemented broadcast field can be reported separately. Path-gain dispersion is the coefficient of variation of  $\alpha_n^{\text{cond}}$  or  $\tilde{\alpha}_n$  over the measured sample and compartment field, depending on whether reactivation is enabled. This unweighted architectural statistic differs from the eligibility-weighted  $\text{CV}_w(\tilde{\alpha}_{n(i)})$  in Prop. 3. Compartment-error fidelity is the cosine between the tested broadcast field and the exact compartment-error field after flattening over the same sample and compartment axes.

## Rule and Feedback Controls

These controls separate three questions that are deliberately collapsed in the main competence summary: which local-rule wrapper is practical, whether the somatic error source matters, and how performance changes as the feedback field gains bandwidth or path information.



514 **Rule-Family Ranking**

515 This table keeps the 3F/4F/5F optimizer distinction explicit: the theorem-derived eligibility is central  
 516 to the proposed mechanism, but the reported LocalCA performance results use the empirical 5F  
 517 stabilizer.

| Dataset             | Rule | Mean valid, 10 validation-selected configs | Mean test, same 10 configs |
|---------------------|------|--------------------------------------------|----------------------------|
| MNIST               | 3F   | 0.611                                      | 0.622                      |
| MNIST               | 4F   | 0.620                                      | 0.628                      |
| MNIST               | 5F   | <b>0.912</b>                               | <b>0.916</b>               |
| Figure-ground MNIST | 3F   | 0.398                                      | 0.396                      |
| Figure-ground MNIST | 4F   | 0.411                                      | 0.411                      |
| Figure-ground MNIST | 5F   | <b>0.807</b>                               | <b>0.789</b>               |

Table S2: **Rule-family ranking.** Mean over the ten best validation-selected configurations from local-competence sweeps; this is not top-10 classification accuracy, and test scores are reported only after validation selection.

518 **5F Stabilizer Sensitivity**

519 Because the main LocalCA runs use the empirical 5F stabilizer, we also tested whether performance  
 520 depends sharply on the preconditioner clamp or the EMA rate used to estimate branch statistics.  
 521 Across three-seed MNIST sensitivity checks, tightening or widening the clamp and varying the EMA  
 522 rate leaves test accuracy near the default setting. This does not make 5F theorem-derived; it supports  
 523 interpreting it as a bounded branch-level reliability preconditioner rather than a brittle tuned error  
 524 source.




Figure S2: **5F stabilizer sensitivity.** (A) Changing the 5F preconditioner clamp from the default [0.25, 4.0] to tighter [0.5, 2.0] or wider [0.1, 8.0] bounds leaves MNIST test accuracy near 92%. (B) Changing the 4F/5F EMA smoothing rate over  $\alpha \in \{0.05, 0.10, 0.20\}$  gives similarly stable performance. Error bars and bands are  $\pm 1$  s.d. across 3 seeds.

| Dataset             | Shunt. BP | Add. BP           | Shunt. local      | Add. local        | Shunt. gap (pp) |
|---------------------|-----------|-------------------|-------------------|-------------------|-----------------|
| MNIST               | 0.971     | $0.970 \pm 0.001$ | $0.911 \pm 0.005$ | $0.894 \pm 0.007$ | 6.0             |
| Fashion-MNIST       | 0.889     | $0.876 \pm 0.001$ | $0.838 \pm 0.003$ | $0.831 \pm 0.004$ | 5.1             |
| Figure-ground MNIST | 0.861     | –                 | $0.803 \pm 0.006$ | $0.787 \pm 0.007$ | 5.8             |

Table S3: **Local performance.** Backpropagation references are matched-capacity standard-training runs. Additive backpropagation is reported where a matched five-seed reference is available; the figure-ground additive backpropagation reference was not part of the matched competence source table. Local values are 5F with per-soma feedback and local decoder updates. Figure-ground MNIST additionally uses an HSIC auxiliary objective (weight 0.01; see the HSIC auxiliary-objective appendix paragraph). Error bars are  $\pm 1$  s.d. across 5 seeds where available; gaps are percentage points.

## 525 Local Performance and Regime Dependence

526 Table S3 reports the matched-capacity performance numbers used in the main text. The point  
527 of this table is not to claim state-of-the-art benchmark performance, but to show the remaining  
528 LocalCA-to-backprop gap under matched dendritic capacity and per-soma 5F feedback.

## 529 Extended Gradient Analysis and $N_I$ Sweep Detail

530 Figure 2 in the main text summarizes final-state gradient diagnostics and the layer-soma factorial  
531 check. The supplementary view below separates static scale checks, inhibitory dose-response curves,  
532 and individual Fashion-MNIST seed behavior so that the aggregate claims are not driven by a single  
533 summary panel.




Figure S3: **Extended gradient and  $N_I$  analysis (supplementary).** (A) Log-scale mismatch,  $|\log_{10}(\|g^{\text{local}}\| / \|g^{\text{bp}}\|)|$ , from a separate matched-weight static diagnostic; these values should not be compared directly to the checkpointed final-state diagnostic in Fig. 2. The ideal value is 0; shunting is near ideal (0.117), whereas additive shows order-of-magnitude distortion ( $> 1.0$ ). (B) Noise-resilience  $N_I$  dose-response with error bands ( $\pm 1$  s.d.). (C) MNIST  $N_I$  dose-response detail with error bands. (D) Fashion-MNIST individual seeds for all conditions, showing consistency across runs.




Figure S4: **Alignment dynamics with explicit gradient-norm checks.** (A) Weighted cosine similarity between LocalCA and backpropagation gradients over training for the 5F rule. (B,C) Weighted local and backprop gradient norms stay nonzero for both additive and shunting models. (D) The additive control begins with a much larger scale mismatch, then approaches the backprop scale without recovering strong directional alignment. Shading is s.d. across 3 seeds.

534 **Alignment dynamics are not a zero-gradient artifact.** To check whether low additive alignment  
535 could be caused by vanishing gradients, we analyzed a separate 3-seed checkpoint trajectory with  
536 parameter-level LocalCA and backprop gradients. The additive local gradients are nonzero throughout  
537 training (weighted norm range  $9.5 \times 10^{-4}$ – $9.5 \times 10^{-3}$ ), and the corresponding backprop gradients are  
538 also nonzero ( $2.3 \times 10^{-5}$ – $5.7 \times 10^{-3}$ ). The low additive cosine therefore reflects directional and scale  
539 mismatch, not absent gradients: additive 5F cosine stays near zero from epoch 0 to epoch 50 (0.018  
540 to 0.005), while its local/backprop norm ratio falls from about  $8.8 \times 10^2$  to 0.69.




Figure S5: **Verification and reproducibility (supplementary).** (A) MNIST verification: the current-code rerun of main seeds (42–46) yields  $91.1 \pm 0.5\%$ ; held-out seeds (47–49) yield  $89.8 \pm 0.9\%$ , showing robustness across random seeds. (B) Matched figure-ground MNIST HSIC control: both main seeds (42–46) and held-out seeds (47–49) compare HSIC weight 0 against 0.01. (C) HSIC weight ablation on figure-ground MNIST: moderate weights (0.01–0.1) perform best.

#### 541 Verification and Reproducibility

542 These checks compare the main seed set with held-out seeds and report the HSIC-weight sensitivity  
 543 for figure-ground MNIST. The additional matched experiments use the same objective-weight  
 544 comparison on both seed sets, so seed robustness and HSIC contribution can be summarized as a  
 545 matched two-factor control.




Figure S6: **Additional stress tests.** (A) Depth scaling: shunting local (green) degrades more gracefully from 63.5% to 57.4%; additive local (blue) drops from 54.9% to 29.7%. Matched backpropagation references (dashed) remain stable. (B) Broadcast-noise robustness: shunting remains stable for  $\sigma \leq 0.1$ ; additive degrades rapidly and reaches chance at  $\sigma = 1$ . (C) Fashion-MNIST: the shunting gain is modest on this cleaner benchmark (83.8% local vs. 88.9% backprop), consistent with the regime-dependent interpretation in the main text.

#### 546 Additional Stress Tests

547 These stress tests probe two ways in which low-bandwidth broadcast can become unreliable: deeper  
 548 dendritic paths and noisy teaching signals. They are not separate main benchmarks; they delimit  
 549 where shunting remains useful and where additive local learning degrades.




Figure S7: **Feedback alignment baselines (MNIST)**. (A) Grouped comparison: standard backprop (neutral gray), DFA (orange), and FA (purple). Hatched N/A markers indicate that FA is not implemented for the block-structured dendritic architectures because the random feedback matrices are not dimensionally compatible with the tree-structured branch updates. DFA achieves 66.1% on shunting vs. 62.2% additive vs. 51.8% point MLP. (B) Backprop–DFA gap: dendritic architectures (24.9–27.6 pp) show smaller gaps than point MLPs (39.5 pp), suggesting conductance-based architecture is partially compatible with random feedback.

#### 550 FA/DFA Baseline Comparison

551 The feedback-alignment baseline is included as an adjacent local-learning reference point. DFA can  
 552 be evaluated on the dendritic cores, whereas the FA condition is marked unavailable for the block-  
 553 structured dendritic updates because the required random feedback matrices are not dimensionally  
 554 compatible with the branch-level parameterization.






Figure S8: **CIFAR-10 control ladder in the compact direct-I-stream family.** (A) Matched back-propagation references. With nonnegative inputs and a nonlinear decoder, shunting slightly exceeds additive under standard training, and removing learned I-to-E conductance lowers the shunting reference. (B) Broadcast ladder. Per-soma feedback remains weak on this harder dataset, random low-rank  $K=4$  improves partially, and exact effective transported error nearly reaches the matched shunting backpropagation reference. (C) Additive fairness controls. Input-dependent gain, running-stat gain, learned gain, and dendritic normalization do not rescue additive per-soma feedback. This figure is a harder-dataset stress test, not a claim of competitive CIFAR-10 benchmarking.

## 555 CIFAR-10 Results

| Condition                                | Core     | Training / broadcast | I-to-E | Test acc.                        |
|------------------------------------------|----------|----------------------|--------|----------------------------------|
| standard                                 | additive | backprop             | yes    | $48.3 \pm 1.0$                   |
| standard                                 | shunting | backprop             | yes    | <b><math>49.5 \pm 0.4</math></b> |
| standard, no learned inhibition          | shunting | backprop             | no     | $43.9 \pm 1.2$                   |
| per-soma                                 | additive | 5F per-soma          | yes    | $25.5 \pm 1.6$                   |
| per-soma, input-dependent gain           | additive | 5F per-soma          | yes    | $25.9 \pm 4.1$                   |
| per-soma, running-stat gain              | additive | 5F per-soma          | yes    | $24.8 \pm 0.9$                   |
| per-soma, learned gain                   | additive | 5F per-soma          | yes    | $26.3 \pm 1.5$                   |
| per-soma, dendritic normalization        | additive | 5F per-soma          | yes    | $25.5 \pm 1.2$                   |
| per-soma                                 | shunting | 5F per-soma          | yes    | $29.7 \pm 1.3$                   |
| per-soma, no learned inhibition          | shunting | 5F per-soma          | no     | $22.3 \pm 0.7$                   |
| transported error                        | additive | 5F path transport    | yes    | $32.4 \pm 1.5$                   |
| rank-4                                   | shunting | 5F low-rank          | yes    | $35.0 \pm 1.0$                   |
| transported error                        | shunting | 5F path transport    | yes    | <b><math>47.2 \pm 0.6</math></b> |
| transported error, no learned inhibition | shunting | 5F path transport    | no     | $38.6 \pm 1.0$                   |

Table S4: **Exact CIFAR-10 values for the 70-run control ladder (5 seeds per condition, validation-best epoch).** All rows use flattened CIFAR-10 in  $[0, 1]$ , nonnegative transfer outputs, positive conductance transforms, a compact depth-4 dendritic core, and a nonlinear decoder. LocalCA maps output error back to decoder-input/soma coordinates through the decoder Jacobian. The useful conclusions are narrow but important: learned shunting slightly improves the matched backpropagation reference, learned I-to-E conductance is necessary for the strongest shunting results, transported error nearly closes the shunting LocalCA gap, and additive per-soma fairness controls do not explain the transported-shunting result.

Table S4 provides the exact values for the appendix CIFAR-10 figure. CIFAR-10 is used only as a harder-dataset mechanistic stress test: under flattened inputs, nonnegative conductance constraints, and a compact passive dendritic core, the matched backpropagation reference is the relevant comparison rather than modern vision accuracy. In this stronger compact direct-I-stream family, shunting performs well under standard training (49.5% vs. 48.3% for additive), so the harder-dataset stress test does not unfairly penalize the shunting architecture. The broadcast-fidelity ladder also survives once decoder-aware soma mapping is used: additive LocalCA reaches 25.5% under per-soma feedback and 32.4% under transported error, while shunting rises from 29.7% under per-soma feedback to 35.0% with random low-rank  $K=4$  and 47.2% with transported error, nearly matching the 49.5% shunting backpropagation reference. Removing learned I-to-E conductance lowers the shunting reference by 5.7 pp and transported LocalCA by 8.6 pp, making this a harder-dataset stress test for inhibitory conductance as part of the useful credit geometry rather than just an added architectural detail.

#### Additive Normalization Control

This control separates generic additive voltage normalization from shunting conductance dynamics in a lower-capacity MNIST stress regime.




Figure S9: **Additive + normalization control (MNIST stress-regime diagnostic).** (A) Standard backprop: normalization provides a small boost (89.7% → 90.9%). (B) Local learning: normalization improves additive from 46.5% to 56.0% (+9.5 pp), partially closing the gap to shunting (61.7%, dashed green). These lower local accuracies come from a separate stress-regime diagnostic, not from the matched-capacity performance setting in Table S3.




Figure S10: **Morphology shifts the useful inhibitory regime.** (A) Shunting-minus-additive LocalCA accuracy gap across dendritic morphologies and inhibitory synapse counts. (B) Best inhibitory count for each morphology. (C) Representative path-gain diagnostic for selected morphology cells, used to connect the performance map to credit geometry. (D) Average shunting advantage by depth, with markers showing the peak inhibitory count for each depth.

#### 571 Morphology $\times$ Inhibition Regime Map

572 This regime map tests whether tree geometry changes the inhibitory operating point at which shunting  
 573 helps. It supports the claim that morphology changes the useful inhibition regime; a fuller mechanistic  
 574 account would require matched theory diagnostics for every cell in the grid.

## Input-Mode Probe: Direct Inhibitory Stream vs. Explicit Inhibitory Cells

The main experiments use input-driven inhibitory streams: the transfer layer provides nonnegative excitatory and inhibitory channels from the external input, and excitatory dendrites receive learned I-to-E conductance banks directly. To test whether the same mechanism can be generated by an explicit feedforward inhibitory population, we ran a one-feedforward-layer noise-resilience probe that preserves the relevant dendritic depth ( $[3, 3]$  morphology). In this setting, explicit inhibitory cells reach  $94.2 \pm 0.2\%$  under path-transport LocalCA when their dendrites use the same update policy as excitatory dendrites,  $93.8 \pm 0.2\%$  when those inhibitory-cell updates are frozen, and  $95.1 \pm 0.3\%$  under standard backprop. Direct input-driven inhibition gives the expected broadcast-fidelity ladder on the same one-layer architecture: per-soma feedback reaches  $85.2 \pm 0.7\%$ , while exact effective path transport reaches  $95.2 \pm 0.0\%$ . Thus explicit inhibitory cells can generate useful path-gain structure once the architecture isolates the dendritic-path question.

| Inhibitory input mode    | Training / broadcast                             | Test accuracy                       |
|--------------------------|--------------------------------------------------|-------------------------------------|
| Direct inhibitory stream | LocalCA, per-soma feedback                       | $0.852 \pm 0.007$                   |
| Direct inhibitory stream | LocalCA, exact effective path transport          | <b><math>0.952 \pm 0.000</math></b> |
| Explicit I cells         | LocalCA, path transport, train I-cell dendrites  | $0.942 \pm 0.002$                   |
| Explicit I cells         | LocalCA, path transport, freeze I-cell dendrites | $0.938 \pm 0.002$                   |
| Explicit I cells         | standard backprop                                | $0.951 \pm 0.003$                   |

Table S5: **One-feedforward-layer input-mode path-gain probe on noise resilience (3 seeds)**. All rows use one excitatory dendritic population with  $[3, 3]$  morphology and nonnegative inputs. The direct-stream rows have no explicit inhibitory neurons; the input stream drives learned branch-level inhibitory conductance banks directly. The explicit-I rows use a matched inhibitory population with  $[3, 3]$  morphology. Exact effective path transport closes the direct-stream per-soma gap, and explicit inhibitory cells nearly match both direct-stream path transport and the explicit-I matched backpropagation reference.

**Post-training inhibition intervention and exact-error rank.** We next tested whether learned inhibitory conductance carries sample-dependent structure rather than only changing global scale. We replayed trained shunting LocalCA checkpoints while replacing the branch-level inhibitory conductance with four post-training interventions: zero inhibition, sample-shuffled inhibition, per-branch batch means, or one uniform matched mean. On a 512-example held-out subset, learned MNIST inhibition at  $N_I=5$  gives  $90.4 \pm 1.1\%$  accuracy, while the four interventions drop to 19.8–26.4%. On noise resilience at  $N_I=10$ , learned inhibition gives  $80.5 \pm 2.4\%$ , while the interventions drop to 10.3–19.3%. Thus the inhibitory field is not interchangeable with a matched global conductance load. We also computed the SVD of exact compartment-error matrices across samples and compartments. At matched MNIST  $N_I=5$ , shunting has lower unrestricted rank-1 residual than additive ( $0.70 \pm 0.21$  vs.  $0.81 \pm 0.11$ ), lower participation rank ( $5.7 \pm 4.0$  vs.  $8.0 \pm 3.7$ ), and higher actual per-soma broadcast cosine ( $0.74 \pm 0.13$  vs.  $0.37 \pm 0.18$ ), supporting but not by itself proving the path-gain compressibility interpretation.

We also computed the direct log-gain covariance margin in Eq. (9) on the selected morphology diagnostic checkpoints. The margin is zero when no inhibitory conductance is present and slightly negative in the inhibited shunting selections ( $-2.5 \times 10^{-3}$ ,  $-5.0 \times 10^{-4}$ , and  $-2.8 \times 10^{-4}$  at  $N_I=5, 20, 40$ ). We therefore do not use this covariance condition as a positive main result. Its role is to clarify when shunting should compress gains; the empirical claim rests on the measured comparative path-gain dispersion, exact-error rank, and inhibition-intervention diagnostics.




Figure S11: **Post-training inhibition intervention and exact-error compressibility.** (A,B) Post-training inhibitory-conductance interventions on trained shunting LocalCA checkpoints. Accuracy is measured on a 512-example held-out subset; dots show seeds. Removing, shuffling, or clamping learned inhibition collapses performance, showing that the learned inhibitory field is sample-dependent and functionally important. (C,D) Exact-error diagnostics on matched MNIST  $N_I=5$  runs. Shunting has a smaller unrestricted rank-1 residual, lower participation rank, and higher actual per-soma broadcast cosine than additive, so the exact error field is both more intrinsically compressible and better aligned with the implemented feedback field. Error bars are s.d. across 5 seeds.

| Model implication                                                                                                                                                                                                                                                                                                                                                                                                                                                        | Biological prediction                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Path-local inhibitory conductance changes $\alpha_n^{\text{cond}}$ only for descendants whose credit path crosses the inhibited compartment. The log derivative $\partial \log \alpha_n^{\text{cond}} / \partial G_k^I = -R_k^{\text{tot}}$ predicts stronger credit gating in high-resistance compartments. Single-channel broadcast succeeds when exact compartment errors are effectively shared across the relevant branches and fails when branch identity matters. | Branch-local inhibitory perturbation should distort plasticity or teaching-signal efficacy for synapses below the perturbed compartment more than for sister branches.<br><br>The same inhibitory conductance change should have larger effects on low-conductance/high-resistance branches than on already high-conductance branches.<br><br>Tasks requiring route- or context-specific branch credit should need richer dendritic teaching signals than a single scalar modulatory broadcast. |

Table S6: **Experimental predictions from the path-gain mechanism.** These are qualitative biological predictions, not claims tested directly in this paper.

## 606 Experimental Predictions

607 The model is not a detailed cell-type circuit model, but its path-gain identity gives falsifiable dendritic  
608 predictions. These predictions are naturally aligned with branch-level dendritic computation and  
609 compartment-targeted inhibition in pyramidal neurons [5–7, 13].

## 610 Cue-Routing Feedback-Rank Diagnostic

611 Cue routing tests a regime where one shared teaching signal is expected to fail: context determines  
612 which noisy cue is reliable, so different pathways should receive different credit. Figure S12 shows  
613 that the benchmark is learnable and that moving beyond the restricted per-soma / single-channel  
614 baselines helps, but the path-structured variant does not yet beat the best unstructured low-rank  
615 setting. We therefore interpret this result as a failure-mode and rank-requirement diagnostic, not as a  
616 solved structured-feedback method.




Figure S12: **Structured feedback when restricted feedback fails.** (A) Cue-routing task schematic. Context determines which noisy cue is reliable on each trial, so different dendritic pathways should receive different teaching signals. (B) Five-seed comparison of per-soma, random low-rank, and pathway-structured feedback. Moving beyond the per-soma / single-channel baselines helps, while the best structured repair remains open. (C) Test accuracy vs. router specialization across learned-router conditions. High specialization alone does not guarantee success. (D) Recovered pathway assignments for a representative learned-router run. Cue-A and cue-B features segregate into separate latent pathways even when feedback quality, rather than router discreteness, is the limiting factor.



## 617 Random Low-Rank Broadcast Channels

618 These controls test whether failures of the restricted per-soma baseline reflect insufficient feedback  
619 bandwidth rather than a failure of local eligibility itself.

| Setting                                                           | Broadcast                      | Test accuracy                       |
|-------------------------------------------------------------------|--------------------------------|-------------------------------------|
| <i>Noise resilience, shunting, rank bridge</i>                    |                                |                                     |
| per-soma                                                          | per-soma feedback              | $0.662 \pm 0.016$                   |
| path propagation                                                  | recursive attenuation proxy    | $0.762 \pm 0.047$                   |
| low-rank                                                          | $K=1$                          | $0.453 \pm 0.126$                   |
| low-rank                                                          | $K=2$                          | $0.630 \pm 0.042$                   |
| low-rank                                                          | $K=4$                          | $0.764 \pm 0.025$                   |
| low-rank                                                          | $K=8$                          | $0.795 \pm 0.019$                   |
| path transport                                                    | exact effective path transport | <b><math>0.847 \pm 0.018</math></b> |
| <i>Cue routing, learned shunting router, rank/structure sweep</i> |                                |                                     |
| per-soma                                                          | tuned per-soma feedback        | $0.803 \pm 0.125$                   |
| low-rank                                                          | $K=1$                          | $0.742 \pm 0.022$                   |
| low-rank                                                          | $K=2$                          | <b><math>0.846 \pm 0.082</math></b> |
| low-rank                                                          | $K=4$                          | $0.748 \pm 0.085$                   |
| pathway-vector                                                    | tuned structured rank-2        | $0.815 \pm 0.126$                   |

Table S7: **Rank bridge and rank-structure comparison.** On shunting noise resilience, both recursive path propagation and higher-rank random low-rank broadcast close a substantial fraction of the gap between per-soma feedback and exact effective path transport, with low-rank continuing to improve up to  $K=8$ . On cue routing, the learned-router sweep shows that moving beyond the per-soma baseline helps, but the best unstructured low-rank setting ( $K=2$ ) is still slightly better than the tuned structured pathway-vector variant. We therefore interpret routed credit assignment as an open higher-rank problem rather than as a solved pathway-vector rescue.

## 620 B Dendritic Network Implementation

621 This section gives the implementation-level specification needed to reproduce the dendritic archi-  
622 tectures and LocalCA updates. It separates the biological modeling assumptions from software  
623 choices such as positive raw-parameter transforms, decoder-error coordinates, and optimizer gradient  
624 assignment.

### 625 Implementation Overview

626 Each feedforward dendritic layer is a population of dendritic neurons. A neuron contains a rooted  
627 tree of compartments; each compartment has an excitatory synapse bank, optionally an inhibitory  
628 synapse bank, and dendritic coupling conductances from its children. The main experiments use  
629 branch-level input-driven inhibition: a nonnegative transfer stage produces excitatory and inhibitory  
630 input streams from the external input, and the inhibitory stream drives learned I-to-E conductance  
631 banks on excitatory dendritic branches. The explicit-I probe in Table S5 instead instantiates a separate  
632 feedforward inhibitory population, whose output projects inhibitory conductance onto excitatory  
633 dendrites.

634 The additive and shunting cores are architecture-matched: they use the same tree, synapse counts,  
635 nonnegative transfer stage, positive conductance parameterization, optimizer protocol, and seed sets.  
636 They differ only in how a branch combines excitatory, inhibitory, and dendritic inputs. Decoder  
637 weights are unconstrained readout parameters and are not interpreted as conductances.

| Paper object              | Code anchor                                                                                                                                                       | Role                                                                                                             |
|---------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| Shunting branch voltage   | <code>DendriticBranchLayer.forward</code> with <code>use_shunting=True</code>                                                                                     | Implements numerator/denominator conductance integration and stores total conductance for diagnostics.           |
| Additive control voltage  | <code>DendriticBranchLayer.forward</code> with <code>use_shunting=False</code> ; <code>normalize_additive_voltage</code> for the normalization control            | Uses the same branch tree and synapse banks, but inhibitory input is subtractive rather than a denominator load. |
| Broadcast-mode switch     | <code>LocalCreditAssignment._apply_local_rule_gradients</code> ; <code>LocalRuleConfig.error_broadcast_mode</code>                                                | Chooses scalar/per-soma, local-mismatch, low-rank, path-transport, and pathway-vector feedback fields.           |
| Low-rank feedback         | <code>LocalCreditAssignment._compute_low_rank_broadcast</code> ; <code>broadcast_rank</code>                                                                      | Fixed random low-rank broadcast controls used in the bandwidth ladder.                                           |
| Path-transport oracle     | <code>LocalCreditAssignment._precompute_path_transport_errors</code>                                                                                              | Transported-error upper bound using the effective tree recursion.                                                |
| Path-propagation proxy    | <code>LocalCreditAssignment._precompute_path_propagation_factors</code>                                                                                           | Approximate morphology-aware attenuation control, distinct from exact path transport.                            |
| Shunting eligibility      | <code>LocalCreditAssignment._apply_local_rule_gradients</code> conductance branch                                                                                 | Assigns $xR^{\text{tot}}(E - V)e$ gradients, then multiplies by the positive-transform derivative.               |
| Additive eligibility      | <code>LocalCreditAssignment._apply_local_rule_gradients</code> additive branch                                                                                    | Assigns the matched signed additive derivative, with optional additive-gain controls.                            |
| 4F / 5F reliability gates | <code>_compute_layer_rho</code> , <code>_compute_layer_phi*</code> , <code>clamp_phi</code> , <code>FiveFactorConfig</code>                                       | Slow reliability preconditioners used by 4F/5F, not additional fast task-error channels.                         |
| Experiment knobs          | <code>LocalRuleConfig</code> ; YAML fields <code>rule_variant</code> , <code>error_broadcast_mode</code> , <code>broadcast_rank</code> , <code>five_factor</code> | Reproduces the rule family, feedback bandwidth, and 5F stabilizer settings.                                      |

Table S8: **Equation-to-code map**. Code anchors are symbolic names rather than line numbers so that minor refactors do not invalidate the reproducibility map.

| Parameter family                                | Forward role                                | Update source                                                                 | Feedback / auxiliary signal                                                                                            |
|-------------------------------------------------|---------------------------------------------|-------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------|
| Excitatory and inhibitory synaptic conductances | Branch E/I drive                            | LocalCA conductance eligibility                                               | Scalar, per-soma, low-rank, path-structured, or transported feedback plus local presynaptic drive and voltage          |
| Dendritic coupling conductances                 | Child-to-parent tree transfer               | LocalCA conductance eligibility                                               | Same branch feedback field, child activity, parent voltage, and input resistance                                       |
| Reactivation slope/bias                         | Post-voltage transfer $f(V)$                | Configured local/autograd schedule                                            | Local activation derivative and available batch objective gradients; not part of the conductance-factorization theorem |
| Top- $k$ synapse banks                          | Sparse input selection and selected weights | Masked LocalCA on selected banks; inactive updates only when enabled          | Same branch feedback field, multiplied by top- $k$ masks and local modulators                                          |
| Linear decoder                                  | Readout from final core activity            | Local output-space decoder gradient or autograd                               | $\delta^y$ for local decoder mode; autograd for backprop mode                                                          |
| Nonlinear decoder                               | Compact CIFAR readout                       | Autograd for decoder parameters; Jacobian product for dendritic error mapping | $J_{\text{dec}}(h)^\top \delta^y$ maps output error into soma/core coordinates                                         |
| HSIC auxiliary                                  | Batch dependence regularizer                | Auxiliary gradient on selected local representations when configured          | Batch-global kernel gradient; reported as an auxiliary objective, not a local theorem term                             |
| Router / pathway roles                          | Structured cue-routing pathway selection    | Configuration-dependent router/pathway updates                                | Pathway-vector controls use router-derived role signals; this remains an exploratory higher-rank mechanism             |

Table S9: **Update scope by parameter family**. The exact factorization covers conductance-stage branch and dendritic-coupling parameters. Decoder, reactivation, HSIC, and router mechanisms are implementation-level training components and are therefore documented separately from the theorem.

### 638 Morphology and Synapse Banks

639 The morphology notation  $[b_1, b_2, \dots, b_D]$  specifies a rooted dendritic tree. The factor  $b_d$  is the fan-in  
 640 from dendritic stage  $d$  to the next more proximal stage. Thus  $[3, 3]$  gives 3 proximal branches per  
 641 soma and 3 distal branches per proximal branch, for 9 distal leaves per soma. In the main feedforward  
 642 models, synaptic inputs terminate on dendritic branches rather than directly on the soma. Unless a  
 643 control explicitly changes this, each branch has  $N_E$  excitatory synapses and  $N_I$  inhibitory synapses.  
 644 The branch input is sparse through a learned top- $k$  or banked synapse selection module; inactive  
 645 synapses are masked out in the local update unless the corresponding control enables inactive-weight  
 646 updates.

647 For a branch  $n$ , write  $x_{nj}^E, x_{nj}^I \geq 0$  for the selected excitatory and inhibitory presynaptic drives and  
 648  $a_c$  for the post-activation activity of child compartment  $c$ . The branch-level excitatory, inhibitory,  
 649 and dendritic conductance currents are

$$U_n^E = \sum_{j=1}^{N_E} g_{nj}^E x_{nj}^E, \quad U_n^I = \sum_{j=1}^{N_I} g_{nj}^I x_{nj}^I, \quad D_n = \sum_{c \in \text{child}(n)} g_{c \rightarrow n}^{\text{den}} a_c.$$

650 The corresponding dendritic coupling load is

$$G_n^{\text{den}} = \sum_{c \in \text{child}(n)} g_{c \rightarrow n}^{\text{den}}.$$

651 For leaf branches,  $D_n = G_n^{\text{den}} = 0$ . For direct input-driven inhibition,  $x^I$  is produced by the  
 652 inhibitory transfer stream. For explicit-I controls,  $x^I$  is the activity of the inhibitory-cell population.

### 653 Shunting and Additive Forward Equations

654 In the shunting core, excitatory reversal is 1, leak and inhibitory reversal are 0, and leak conductance  
 655 is fixed to 1. The implemented branch voltage is therefore

$$V_n^{\text{shunt}} = \frac{U_n^E + D_n}{1 + U_n^E + U_n^I + G_n^{\text{den}} + \varepsilon}.$$

656 The small numerical  $\varepsilon$  is used only for floating-point safety; analytically the leak conductance keeps  
 657 the denominator positive. In implementation equations, the same stabilized denominator is used for  
 658 the recorded input resistance:

$$g_n^{\text{tot}} = 1 + U_n^E + U_n^I + G_n^{\text{den}}, \quad g_{n,\varepsilon}^{\text{tot}} = g_n^{\text{tot}} + \varepsilon, \quad R_{n,\varepsilon}^{\text{tot}} = (g_{n,\varepsilon}^{\text{tot}})^{-1}.$$

659 The main analytical equations write  $R_n^{\text{tot}} = (g_n^{\text{tot}})^{-1}$ ; exact-gradient reconstruction tests use the  
 660 implementation convention above. In the additive control, inhibition is subtractive rather than  
 661 shunting. With the same nonnegative conductance variables and nonnegative inputs, the implemented  
 662 comparator is

$$V_n^{\text{add}} = U_n^E + D_n - U_n^I,$$

663 with optional additive-normalization controls reported separately in Fig. S9. Thus additive inhibition  
 664 changes the signed voltage contribution and local additive eligibility, but it does not enter the  
 665 denominator and does not gate conductance-stage path gain through  $R_n^{\text{tot}}$ .

| Property              | Shunting / conductance core                               | Additive control                                      |
|-----------------------|-----------------------------------------------------------|-------------------------------------------------------|
| Voltage               | Normalized by total conductance                           | Signed E/I voltage sum                                |
| Inhibition            | Divisive conductance load                                 | Subtractive signed term                               |
| Synaptic eligibility  | $x R^{\text{tot}} (E - V) e$                              | $s x e, s \in \{+1, -1\}$                             |
| Driving force         | Present through $E - V$                                   | Absent from additive derivative                       |
| Input resistance      | State-dependent, $R^{\text{tot}} = (g^{\text{tot}})^{-1}$ | Effectively fixed                                     |
| Credit interpretation | Inhibition changes path gain and voltage sensitivity      | Inhibition shifts voltage but does not gate path gain |

Table S10: **Matched shunting/additive comparison.** Both architectures use their own local derivative, so the comparison changes the forward-pass integration rule rather than applying a shunting-derived update to an additive model.

## 666 Positive Parameterization and Raw-Parameter Gradients

667 Trainable synaptic and dendritic conductances are stored as unconstrained raw parameters  $\theta$  and  
 668 transformed by a positive map  $g = \psi(\theta)$ , with `softplus` used in the main reported runs:

$$\psi(\theta) = \log(1 + \exp \theta), \quad \psi'(\theta) = \sigma(\theta).$$

669 Equations in the main text give conductance-space gradients  $\partial L / \partial g$ . Before assigning gradients to  
 670 the optimizer variables, the implementation applies the chain rule

$$\frac{\partial L}{\partial \theta} = \frac{\partial L}{\partial g} \psi'(\theta).$$

671 The exact-gradient reconstruction diagnostic includes this parameterization factor, as well as top- $k$   
 672 activity masks and dendritic block-linear parameterization. This is why the reconstruction comparison  
 673 is against raw autograd gradients rather than only against abstract conductance variables.

## 674 Post-Voltage Reactivation

675 The conductance stage produces a pre-reactivation voltage  $V_n$ , after which each branch may transmit a  
 676 monotone post-voltage activity  $a_n = f_n(V_n)$ . The main sweeps use the learnable bounded transform

$$f_n(V) = \frac{\tanh(m_n(V - b_n)) + 1}{2}, \quad m_n = \exp(\ell_{m,n}),$$

677 with derivative

$$f'_n(V) = \frac{m_n}{2} [1 - \tanh^2(m_n(V - b_n))].$$

678 When reactivation is enabled, exact path transport uses the effective gain  $\tilde{\alpha}_n$ : parent-to-child transport  
 679 includes the parent activation derivative, and the local update converts activation-space error into  
 680 pre-reactivation voltage error by multiplying by the child derivative. With identity transfer, this  
 681 reduces to conductance-stage transport through  $\alpha_n^{\text{cond}}$ .

## 682 Decoder and Error Coordinates

683 Let  $h$  denote the final soma/core activity passed to the decoder and let  $\delta^y = \partial L / \partial \hat{y}$ . For a linear  
 684 decoder  $\hat{y} = W_{\text{dec}} h$ , the soma/core activation-space teaching error is

$$\delta_0^a = W_{\text{dec}}^\top \delta^y.$$

685 For a nonlinear decoder, used in the CIFAR-10 compact control ladder, the implementation uses the  
 686 decoder-input Jacobian product

$$\delta_0^a = J_{\text{dec}}(h)^\top \delta^y.$$

687 All LocalCA broadcast modes are defined in this soma/core activation-error coordinate system before  
 688 local conversion to voltage-space errors. Decoder-update modes are: **backprop**, which uses autograd  
 689 for decoder parameters; **local**, which is implemented for the linear readout and assigns

$$\hat{\nabla}_{W_{\text{dec}}} L = \langle \delta^y h^\top \rangle_B, \quad \hat{\nabla}_{b_{\text{dec}}} L = \langle \delta^y \rangle_B;$$

690 and **frozen**, which holds the decoder fixed. The local decoder gradient is output-space;  $\delta_0^a$  is the  
 691 mapped soma/core activation-space error used by dendritic broadcasts.

## 692 Biological Plausibility Assumptions

693 The model makes the following explicit assumptions. First, it uses the steady-state solution of  
 694 the passive cable equation rather than temporal membrane dynamics; all path gains and gradient  
 695 diagnostics refer to conductance-stage voltages at steady state. Second, each dendritic cell is a rooted  
 696 tree, so each compartment has a unique path to the soma and the closed-form path gains  $\alpha_n^{\text{cond}}$   
 697 and  $\tilde{\alpha}_n$  are well-defined. Third, synaptic and dendritic conductance parameters are nonnegative by  
 698 construction. Fourth, the main sweeps enforce nonnegative first-layer presynaptic drive through  
 699 nonnegative inputs or a ReLU transfer stage. Fifth, post-voltage reactivation is handled through  
 700 local  $f'(V)$  factors along the path. Sixth, each dendritic synapse receives a restricted non-local  
 701 feedback field-scalar, per-soma, low-rank, path-structured, or transported-oracle depending on the  
 702 condition-while presynaptic activity, voltage, reversal potential, input resistance, and slowly estimated  
 703 branch-level modulators are local or branch-local quantities.

| Quantity         | Symbol                           | Convention                                                                                                             |
|------------------|----------------------------------|------------------------------------------------------------------------------------------------------------------------|
| Voltage          | $V$                              | Conductance-stage shunting voltage in $[0, 1]$ under nonnegative drive; additive controls may be signed                |
| Conductances     | $g^{\text{syn}}, g^{\text{den}}$ | Nonnegative via $g = \text{softplus}(\theta)$ ; local gradients are mapped to raw $\theta$ by the transform derivative |
| Leak conductance | $g^{\text{leak}}$                | Set to 1                                                                                                               |
| Input resistance | $R^{\text{tot}}$                 | $\leq 1$                                                                                                               |

Table S11: Units and normalization.

#### 704 Units and Parameterization

705 Table S11 summarizes the numerical conventions used throughout the implementation and diagnostics,  
706 including which quantities are constrained conductances and which are unconstrained readout  
707 parameters.

## 708 Representative Manuscript Architectures

| Setting                                                    | Encoder  | E layers           | Tree                                         | Synapses / branch                                          | Seeds                | Notes                                                                                                                                              |
|------------------------------------------------------------|----------|--------------------|----------------------------------------------|------------------------------------------------------------|----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|
| Gradient fidelity / path-gain / noise-resilience mechanism | id.+ReLU | [128, 128]         | [3, 3]                                       | $N_E=40, N_I \in \{0, 5, 10, 20, 40\}$                     | 5 fidelity, 5 oracle | Main mechanism sweeps in Figs. 2 and 3; reactivation <code>param_tanh</code> .                                                                     |
| MNIST / figure-ground local competence                     | id.+ReLU | [128]              | [3, 3]                                       | $N_E=40, N_I=20$                                           | 5                    | Best local 5F per-soma runs in Table S3; reactivation <code>param_tanh</code> .                                                                    |
| Fashion-MNIST competence                                   | id.+ReLU | [128]              | [3, 3]                                       | $N_E=40, N_I=20$                                           | 5                    | Dedicated five-seed standard-vs-local competence sweep; reactivation <code>param_tanh</code> .                                                     |
| Cue-routing PV-LocalCA                                     | router   | [64]               | [2]                                          | $N_E=10, N_I=4$                                            | 5                    | Learned router with two pathway groups and rank-2 structured feedback; reactivation <code>param_tanh</code> .                                      |
| CIFAR-10 compact direct-I-stream control ladder            | id.+ReLU | [20] E, no I cells | [3, 3, 3, 3]                                 | $N_E=25, N_I=25$ direct I-to-E, with matched no-I controls | 5 per condition      | Harder-dataset stress test with input-driven inhibitory conductance, decoder [32, 16], decoder-aware soma mapping, and additive fairness controls. |
| Morphology $\times$ inhibition appendix map                | id.+ReLU | [128, 128]         | [2, 2], [3, 3], [4, 4], [2, 2, 2], [3, 3, 3] | $N_E=40, N_I \in \{0, 5, 10, 20, 40\}$                     | 3                    | Supportive 3-seed sweep; reactivation <code>param_tanh</code> .                                                                                    |

**Table S12: Representative architectures used in the manuscript.** All dendritic rows use explicit branch-level excitatory and inhibitory synapse banks without direct somatic synapses, nonnegative conductances, and a nonnegative first-layer drive enforced by an explicit ReLU transfer stage in the main runs. The tree column gives dendritic morphology per excitatory neuron, and the synapse counts report excitatory and inhibitory synapses per branch.

## 709 Hyperparameters

710 Table S13 collects the core training hyperparameters for each main experiment. Learning rates are  
711 applied through parameter groups (top- $k$ , dendritic block-linear, reactivation, decoder) with the values  
712 given in the representative config files; where a single LR is reported, every group uses that value.

| Setting                                                                  | Opt. | LR                                           | Batch | Epochs | Weight decay | Notes                                                                |
|--------------------------------------------------------------------------|------|----------------------------------------------|-------|--------|--------------|----------------------------------------------------------------------|
| MNIST / Fashion-MNIST / figure-ground MNIST (5F per-soma LocalCA)        | Adam | $10^{-3} / 5 \cdot 10^{-4}$ (block / react.) | 256   | 100    | 0            | fixed-epoch, no early stopping; <code>param_tanh</code> reactivation |
| MNIST / Fashion-MNIST / figure-ground MNIST (matched backprop reference) | Adam | $10^{-3} / 5 \cdot 10^{-4}$                  | 256   | 100    | 0            | matched to LocalCA setup                                             |
| Gradient fidelity / path-gain / noise resilience                         | Adam | $10^{-3}$                                    | 256   | 50     | 0            | 5 seeds, hooks enabled for diagnostic capture                        |
| Morphology $\times$ inhibition regime map (supportive)                   | Adam | $10^{-3}$                                    | 256   | 50     | 0            | 3 seeds                                                              |
| Cue routing                                                              | Adam | $10^{-3}$                                    | 256   | 90     | 0            | early stopping with patience 30                                      |
| CIFAR-10 compact direct-I-stream depth-4 LocalCA                         | Adam | $10^{-3} / 10^{-4}$ (block / react.)         | 256   | 400    | 0 / 0.01     | grad-clip 5.0, early stop patience 50, nonlinear decoder             |
| CIFAR-10 compact direct-I-stream depth-4 BP                              | Adam | $10^{-3}$                                    | 256   | 200    | 0.01         | early stop patience 40, matched backpropagation reference            |

**Table S13: Training hyperparameters by experiment.** All runs use Adam with default  $\beta_1=0.9$ ,  $\beta_2=0.999$ ,  $\epsilon=10^{-8}$ . Reactivation and block-linear parameter groups use a slower LR to keep conductance-level dynamics stable. Representative configs are provided in the accompanying repository.

## 713 Effective Broadcast Dimensionality in Main Architectures

| Setting                                          | Soma/core dim. | Output classes | Main feedback field                                                                                           |
|--------------------------------------------------|----------------|----------------|---------------------------------------------------------------------------------------------------------------|
| Gradient fidelity / path-gain / noise resilience | 128, 128       | 10             | Per-soma soma/core error at matched-width dendritic layers                                                    |
| MNIST / Fashion-MNIST competence                 | 128            | 10             | Per-soma soma/core error                                                                                      |
| Figure-ground MNIST competence                   | 128            | 10             | Per-soma soma/core error                                                                                      |
| Cue routing                                      | 64             | 2              | Tuned per-soma baseline; higher-rank tests use low-rank or pathway-vector feedback                            |
| Compact CIFAR-10 harder-dataset stress test      | 20             | 10             | Vector-valued per-soma after decoder-aware mapping; higher-rank tests use explicit low-rank or path transport |

Table S14: **Effective dimensionality of the main feedback in the reported architectures.** Output loss gradients are first mapped into decoder-input/soma coordinates. The main competence and mechanism runs use the implemented `per_soma` mode when the soma/core dimension matches layer width; scalar, low-rank, pathway-vector, and transported-error conditions are labeled separately as controls or oracles.

## 714 Compute Resources

715 All experiments were run on an institutional GPU cluster. Reproduction runs use one NVIDIA A100  
716 or H100-class GPU (40–80GB memory) per seed and do not require distributed training. Table S15  
717 gives approximate single-GPU runtime ranges for the reported experiment families; unreported  
718 exploratory runs used additional cluster time and are not counted in the reproduction estimate.

| Experiment family                                       | Runs | Time / run  | Approx. GPU-h |
|---------------------------------------------------------|------|-------------|---------------|
| Performance and verification                            | ~25  | 5–20 min    | 4–9           |
| Gradient, path-gain, inhibition, and oracle diagnostics | ~80  | 10–30 min   | 15–40         |
| Morphology, stress, rule, and feedback controls         | ~120 | 5–35 min    | 20–70         |
| CIFAR-10 control ladder                                 | 70   | 0.8–2.5 h   | 60–175        |
| Figure generation and CPU-side summaries                | –    | < 5 h total | < 5           |

Table S15: **Approximate compute for reported results.** Ranges reflect single-seed wall-clock variation across A100/H100-class GPUs, data-loading overhead, and early stopping.

## 719 Code, Data, and Asset Availability

720 MNIST, Fashion-MNIST, and CIFAR-10 are public datasets or public benchmark assets cited in the  
721 paper; we use them under their standard published access conditions and cite their original sources in  
722 the bibliography. The accompanying repository contains the core model and training code, selected  
723 diagnostic utilities, representative configuration files, tests, and reproduction instructions for the main  
724 experimental families. A public release should identify the immutable commit or release tag used for  
725 the posted manuscript. The paper does not release a new dataset or a stand-alone pretrained model  
726 asset. Dataset access pages are the original MNIST site / UCI entry (<https://archive.ics.uci.edu/dataset/683/mnist+database+of+handwritten+digits>), the Fashion-MNIST repository  
727 (<https://github.com/zalandoresearch/fashion-mnist>), and the CIFAR page (<https://www.cs.toronto.edu/~kriz/cifar.html>).  
728  
729

| Asset           | Use in this paper                                            | Access / license or terms note                                                                                                                                                     |
|-----------------|--------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MNIST           | Digit classification and derived nonnegative synthetic tasks | Public benchmark from the original MNIST site / UCI entry; UCI lists DOI 10.24432/C53K8Q and asks users to follow the original acknowledgement policy; no raw-data redistribution. |
| Fashion-MNIST   | Apparel classification stress test                           | Public Zalando Research benchmark; repository is MIT licensed; no raw-data redistribution beyond standard dataset loaders.                                                         |
| CIFAR-10        | Flattened harder-dataset stress test                         | Public CIFAR dataset site asks users to cite Krizhevsky’s technical report; no raw-data redistribution.                                                                            |
| Synthetic tasks | Figure-ground MNIST, noise resilience, cue integration       | Generated procedurally from public benchmarks or random seeds described in Appendix D; no new third-party asset is introduced.                                                     |

Table S16: **Existing assets used in the manuscript.** We credit the original sources in the bibliography, use public benchmark assets under their published access conditions, and do not redistribute third-party raw data in the manuscript package.

### 730 LocalCA Broadcast and Gradient Assignment

731 LocalCA first maps the output loss derivative to the soma/core activation-space error  $\delta_0^a$  and then  
732 constructs a branch voltage-error field. In the final dendritic layer this is the decoder-input error; in  
733 earlier matched-width layers the practical implementation reuses the same coordinate by neuron index  
734 as an approximate layer-soma teaching signal. This is an inter-layer feedback approximation, separate  
735 from the within-tree path-transport approximation studied by the main theory. In the main reported  
736 per\_soma mode, if the layer width matches the soma/core error dimension, each soma-specific error  
737 component is broadcast to the corresponding dendritic compartments:

$$e_{n,b}^a = \delta_{0,b}^a \in \mathbb{R}^{d_0}.$$

738 The stricter scalar control compresses the soma/core error to one value per example,

$$\bar{\delta}_b = \frac{1}{d_0} \sum_{c=1}^{d_0} \delta_{0,bc}^a, \quad e_{n,b}^a = \bar{\delta}_b \mathbf{1}_n,$$

739 where  $d_0$  is the dimension of the soma/core error. If dimensions do not match, the implementation  
740 falls back to scalar-style expansion rather than silently inventing an unmatched per-soma map. The  
741 low-rank controls use fixed random projection and mixing matrices  $P_K$  and  $Q_n$ ,

$$c_b = P_K \delta_{0,b}^a \in \mathbb{R}^K, \quad e_{n,b}^a = Q_n c_b.$$

742 The pathway-vector controls use router-inferred pathway roles to gate and transport a low-rank  
743 role vector. The transported oracle uses the effective tree recursion. If  $p$  is the parent of child  $c$ ,  
744 activation-space error is propagated by

$$e_c^a = e_p^a f'_p(V_p) R_{p,\varepsilon}^{\text{tot}} g_{c \rightarrow p}^{\text{den}},$$

745 and the local pre-activation voltage error used in conductance eligibility is

$$e_n^V = \begin{cases} e_n^a f'_n(V_n), & \text{post-voltage reactivation enabled,} \\ e_n^a, & \text{identity transfer.} \end{cases}$$

746 This operational definition is the implementation counterpart of the effective path gain  $\tilde{\alpha}_n$  in Theo-  
747 rem 1.

748 For shunting branches, the conductance-space LocalCA gradients assigned before raw-parameter  
749 transformation are

$$\hat{\nabla}_{g_{nj}^E} L = \langle e_n^V x_{nj}^E R_{n,\varepsilon}^{\text{tot}} (1 - V_n) \rangle_B, \quad \hat{\nabla}_{g_{nj}^I} L = \langle e_n^V x_{nj}^I R_{n,\varepsilon}^{\text{tot}} (0 - V_n) \rangle_B,$$



and for a dendritic coupling from child  $c$  to parent  $n$ ,

$$\widehat{\nabla}_{g_{c \rightarrow n}^{\text{den}}} L = \langle e_n^V R_{n,\varepsilon}^{\text{tot}} (a_c - V_n) \rangle_B.$$

For additive controls, the matching local derivatives are

$$\widehat{\nabla}_{g_{nj}^E} L = \langle e_n^V x_{nj}^E \rangle_B, \quad \widehat{\nabla}_{g_{nj}^I} L = -\langle e_n^V x_{nj}^I \rangle_B, \quad \widehat{\nabla}_{g_{c \rightarrow n}^{\text{den}}} L = \langle e_n^V a_c \rangle_B.$$

The optimizer applies the usual descent step; equivalently, the sign can be absorbed into the definition of the broadcast error. After these conductance-space gradients are formed, they are multiplied by the positive-transform derivative  $\psi'(\theta)$  and by any active synapse masks before assignment to raw parameters.

#### 4F and 5F Local Modulators

The 3F rule uses only the local eligibility and broadcast error above. The 4F and 5F variants multiply the conductance-space gradient by branch-level reliability factors before raw-parameter transformation. For 4F, the implementation records  $V_n$  from each branch layer and the corresponding soma/core activity  $V_0$ . For each batch item  $b$ , it first averages over the non-batch coordinates of the recorded layer tensor to obtain scalar summaries  $\bar{V}_{n,b}$  and  $\bar{V}_{0,b}$ . It then estimates covariance over the batch axis and smooths the resulting scalar with an EMA:

$$r_n^{4F} = \frac{\text{Cov}_B(\bar{V}_{n,b}, \bar{V}_{0,b})}{\sqrt{\text{Var}_B(\bar{V}_{n,b}) \text{Var}_B(\bar{V}_{0,b})} + \varepsilon}.$$

The EMA is initialized from the first observed batch statistic; the recorded tensors are detached from autograd, and the same mini-batch supplies both the statistics and the gradient update. This proxy is clamped to the positive stability range  $[0.1, 2.0]$  before it multiplies the local gradient; sign-opposed covariance therefore reduces the multiplier rather than reversing update direction. Because the smoothed numerator and denominator can be updated separately, and because single-sample fallbacks use online covariance moments, the clamped value can exceed one and should be read as a preconditioner rather than a literal Pearson correlation. The 5F factor adds a bounded branch-level preconditioner based on how much branch voltage variance remains predictable from parent or soma-level activity:

$$\phi_n = \text{clip}_{[0.25, 4.0]} \left( \frac{\text{Var}(V_n)}{\sigma_{\text{res},n}^2 + \varepsilon} \right),$$

where  $\sigma_{\text{res},n}^2$  is estimated online from a scalar ridge regression of the branch-layer voltage on the configured parent proxy. Batch-mode estimates center over examples and flatten remaining layer coordinates before updating the smoothed variance and covariance moments; single-sample fallback uses scalar online moments. Since these moments are smoothed and clamped rather than recomputed as a single ordinary least-squares fit on a fixed batch,  $\phi_n$  can transiently fall below one or rise above one. The lower clamp keeps noisy residual estimates from suppressing a branch-layer update entirely, and the upper clamp prevents highly predictable branches from dominating optimization. In the main 5F runs, the multiplier is  $r_n^{4F} \phi_n$ . The 5F sensitivity diagnostic in Fig. S2 varies the clamp and EMA rate to verify that the reported MNIST performance is not a single brittle clamp setting.

#### Algorithm

Algorithm 1 gives the update order for the main 5F per-soma LocalCA condition and shows where scalar, higher-bandwidth, and transported-error controls enter.

---

**Algorithm 1** Main 5F Per-Soma LocalCA Update

---

```
1: Input: Dendritic model, batch  $(x, y)$ , optimizer
2: Forward pass; loss  $L$ , output error  $\delta^y$ 
3: Somatic/core activation error  $\delta_0^a = \text{decoder-input Jacobian product } J_{\text{dec}}(h_{\text{core}})^\top \delta^y$  (or  $W_{\text{dec}}^\top \delta^y$  for a linear decoder)
4: for each layer  $n$  (reverse) do
5:   Per-soma feedback  $e_n^a = \delta_0^a$  when dimensions match; scalar/low-rank/path controls replace this line
6:   Update EMA estimates for  $r_n^{4F}$  and  $\phi_n$  from batch voltages
7:   Convert to pre-reactivation voltage error  $e_n^V$  using the local activation derivative when needed
8:    $\widehat{\nabla}_{g_j^{\text{syn}}} L \leftarrow r_n^{4F} \phi_n \langle x_j R_{n,\varepsilon}^{\text{tot}} (E_j - V_n) e_n^V \rangle_B$ 
9:    $\widehat{\nabla}_{g_{c \rightarrow n}^{\text{den}}} L \leftarrow r_n^{4F} \phi_n \langle R_{n,\varepsilon}^{\text{tot}} (a_c - V_n) e_n^V \rangle_B$ 
10:  Map conductance gradients to raw parameters by multiplying by  $\partial g / \partial \theta$ 
11: end for
12: Optional controls replace line 5 with scalar, rank- $K$ , path-structured, or transported-oracle feedback.
13: Clip gradients; optimizer step
```

---

## C Theoretical Details

This section collects theory-adjacent material that supports implementation and interpretation but is not part of the main proof. The central derivation remains Theorem 1, Corollary 1, and Prop. 2.

### Random-Broadcast Alignment Intuition

Feedback-alignment arguments suggest that random or low-rank feedback can be useful when the induced local update remains positively correlated with the exact gradient [14]. In this model, that correlation depends on how local eligibility factors co-vary with the conductance-stage path gain in Eq. (4). We treat this as intuition only. The central theoretical object in the paper is still the exact pre-reactivation compartment error  $\partial L / \partial V_n$ ; under identity upward transfer this reduces to  $\alpha_n^{\text{cond}} \delta_0^V$ , while post-voltage activation uses  $\tilde{\alpha}_n \delta_0^V$ . The empirical question is how faithfully different broadcast fields approximate that quantity.

### Morphology-Aware Extensions

The following variants are implemented controls or architectural extensions used to test whether coarse morphology-aware feedback can approximate exact path transport. They are not required for the main 5F per-soma results, but they define the path-propagation, depth, normalization, and pathway-vector controls reported in the appendix figures and tables.

**Path-integrated propagation.** Modulate broadcast error by  $\pi_n = \pi_{n-1} \cdot R_{n-1}^{\text{tot}} \cdot \bar{g}_{n-1}^{\text{den}}$ , approximating depth attenuation from Eq. (4) without computing the exact sample-specific path gain.

**Depth modulation.** Per-branch scaling  $\kappa_j = \kappa_{\text{base}} / (d_j + d_0)$ , mirroring cable attenuation and testing whether a simple depth prior can stabilize distal updates.

**Dendritic normalization.**  $\Delta g_j^{\text{den}} \leftarrow \Delta g_j^{\text{den}} / (\sum_k g_k^{\text{den}} + \varepsilon)$ , analogous to homeostatic scaling [26] and used as an additive-control normalization comparison.

**Pathway-vector feedback.** For tasks with latent pathway structure, the broadcast becomes a role vector inferred from the router. Local pathway activity gates this vector. Upward block transport then passes it to earlier branches, aligning their feedback with downstream descendants.

**Router-derived branch roles.** Each branch receives a role profile inferred from router assignments and incoming block weights rather than from a hand-coded apical/basal label. Plasticity is then modulated by branch selectivity and a synapse-specific role-alignment factor, emphasizing pathway-consistent updates without imposing a heuristic branch taxonomy.

## 813 HSIC Auxiliary Objectives

814 Following [24], the figure-ground MNIST LocalCA runs optionally apply an HSIC-style auxiliary  
815 gradient to selected layer representations  $Z$ :

$$\mathcal{L}^{\text{self}} = B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Z \mathbf{H}), \quad \mathcal{L}^{\text{target}} = -B^{-2} \text{tr}(\mathbf{K}_Z \mathbf{H} \mathbf{K}_Y \mathbf{H}).$$

816 Here  $B$  is batch size,  $\mathbf{H}$  is the centering matrix,  $\mathbf{K}_Z$  uses the configured representation kernel, and  
817  $\mathbf{K}_Y$  uses one-hot class labels when `target_source=labels`. The reported figure-ground runs  
818 use an RBF kernel with fixed bandwidth  $\sigma = 1.0$ , HSIC weight 0.01, self and target weights 0.3,  
819 five-epoch warmup, and gradient clipping at 0.1; the code applies the resulting auxiliary gradient  
820 through the configured local-learning representation pathway rather than through the conductance-  
821 factorization theorem. Moderate weights (0.01–0.1) help on figure-ground MNIST; negligible on  
822 MNIST. Online statistics ( $r_n^{4F}$ ,  $\phi_n$ ) use Welford’s algorithm [25].

## 823 D Task Construction and Robustness Protocols

824 This section specifies the benchmark, synthetic-task, depth, and feedback-noise protocols used in the  
825 experiments, with emphasis on which inputs preserve the nonnegative-drive regime assumed by the  
826 conductance model.

### 827 General Data Handling

828 MNIST, Fashion-MNIST, and CIFAR-10 are loaded from standard public dataset loaders, flattened  
829 before entering the dendritic core, scaled to  $[0, 1]$ , and passed through the configured nonnegative  
830 transfer stage in the main experiments. MNIST-derived synthetic tasks preserve the flattened  $28 \times 28$   
831 geometry. Fixed random seeds define task generators such as corruption projections and context  
832 masks; generator parameters are shared across train, validation, and test splits, while sample-level  
833 corruptions are drawn per example. Validation accuracy selects checkpoints, and test accuracy is  
834 reported. CIFAR-10 is used as a flattened harder-dataset stress test under the same conductance  
835 constraints rather than as a competitive vision benchmark; no data augmentation is used as a  
836 performance device.

### 837 Supervised Benchmark Tasks

838 **MNIST and Fashion-MNIST.** These ten-class tasks use nonnegative pixel inputs. MNIST tests  
839 whether per-soma 5F LocalCA approaches matched backpropagation on clean supervised data;  
840 Fashion-MNIST adds a same-dimensional distribution shift.

841 **CIFAR-10 harder-dataset stress test.** CIFAR-10 inputs are flattened RGB images in  $[0, 1]$ . The  
842 compact CIFAR family uses a smaller dendritic core, nonlinear decoder, and decoder-aware LocalCA  
843 mapping to test whether the broadcast-fidelity ladder persists on harder data: per-soma feedback is  
844 weak, higher-rank feedback helps, and exact effective path transport approaches the matched shunting  
845 backpropagation reference.

846 **Figure-ground MNIST.** For each reshaped  $28 \times 28$  MNIST image, the right half carries the digit  
847 and the left half is replaced by independent distractors  $u_{r,c} \sim \text{Uniform}(0, 0.25)$ , clipped to  $[0, 1]$   
848 when needed. The task tests whether local credit assignment exploits spatial signal/distractor structure  
849 without leaving the nonnegative-drive regime. The main LocalCA competence runs use the HSIC  
850 auxiliary objective with weight 0.01, which adds roughly 3 pp in the ablation.

851 **Noise resilience.** Flattened MNIST images are corrupted as  $\tilde{x} = \text{clip}(x + \sigma_{\text{task}} A z, 0, 1)$  with  
852  $\sigma_{\text{task}} = 1.5$ , a fixed task-seeded projection  $A \in \mathbb{R}^{784 \times 50}$ , and fresh  $z \sim \mathcal{N}(0, I_{50})$  for each  
853 corrupted example. The fixed projection creates correlated interference without reducing the task to  
854 memorizing one corruption pattern.

855 **Cue integration.** Two noisy cue streams ( $A$  and  $B$ ) are presented simultaneously for two-class  
856 classification, with a one-hot context indicating which cue is reliable on each trial. Cue vectors are  
857 clipped to  $[0, 1]$  after noise. The fixed-pathway variant duplicates context into both cue branches; the

858 learned-routing variant requires the router to discover cue separation from data. This feedback-rank  
859 diagnostic tests structured pathway-vector broadcast inside the nonnegative-input regime.

## 860 Depth Scaling and Noise Robustness

861 These stress diagnostics keep the same broad model family while changing dendritic depth or  
862 corrupting the broadcast signal. For depth scaling, dendritic depth varies from 1–4 layers (branch  
863 factors [9] to [3, 3, 3, 3]): shunting local degrades from 63.5% to 57.4%, additive local falls from  
864 54.9% to 29.7%, the shunting advantage grows from +8.5 to +27.7 pp, and matched backpropagation  
865 references remain near 90–92% (Fig. S6A). For noise robustness, Gaussian noise  $\mathcal{N}(0, \sigma^2)$  is added  
866 to the broadcast error; shunting remains robust to  $\sigma \leq 0.1$  (about 62%), while additive drops from  
867 46.5% to chance at  $\sigma=1.0$ , indicating that shunting credit signals carry useful learning information  
868 beyond broadcast magnitude alone (Fig. S6B).

## 869 NeurIPS Paper Checklist

### 870 1. Claims

871 Question: Do the main claims made in the abstract and introduction accurately reflect the  
872 paper’s contributions and scope?

873 Answer: [Yes]

874 Justification: The abstract and introduction state the paper’s main mechanistic claim, the  
875 scope of the conductance-based setting, and the regime-dependent nature of the empirical  
876 results; see Sec. 1 and Sec. 5.

877 Guidelines:

- 878 • The answer [N/A] means that the abstract and introduction do not include the claims  
879 made in the paper.
- 880 • The abstract and/or introduction should clearly state the claims made, including the  
881 contributions made in the paper and important assumptions and limitations. A [No] or  
882 [N/A] answer to this question will not be perceived well by the reviewers.
- 883 • The claims made should match theoretical and experimental results, and reflect how  
884 much the results can be expected to generalize to other settings.
- 885 • It is fine to include aspirational goals as motivation as long as it is clear that these goals  
886 are not attained by the paper.

### 887 2. Limitations

888 Question: Does the paper discuss the limitations of the work performed by the authors?

889 Answer: [Yes]

890 Justification: The discussion explicitly narrows the scope of the claims and discusses  
891 limitations involving benchmark strength, broadcast assumptions, depth scaling, morphology,  
892 activation shape, and feedback rank; see Sec. 5.

893 Guidelines:

- 894 • The answer [N/A] means that the paper has no limitation while the answer [No] means  
895 that the paper has limitations, but those are not discussed in the paper.
- 896 • The authors are encouraged to create a separate “Limitations” section in their paper.
- 897 • The paper should point out any strong assumptions and how robust the results are to  
898 violations of these assumptions (e.g., independence assumptions, noiseless settings,  
899 model well-specification, asymptotic approximations only holding locally). The authors  
900 should reflect on how these assumptions might be violated in practice and what the  
901 implications would be.
- 902 • The authors should reflect on the scope of the claims made, e.g., if the approach was  
903 only tested on a few datasets or with a few runs. In general, empirical results often  
904 depend on implicit assumptions, which should be articulated.
- 905 • The authors should reflect on the factors that influence the performance of the approach.  
906 For example, a facial recognition algorithm may perform poorly when image resolution  
907 is low or images are taken in low lighting. Or a speech-to-text system might not be

used reliably to provide closed captions for online lectures because it fails to handle technical jargon.

- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: The voltage model, assumptions, theorem statements, and short chain-rule proofs are given in Sec. 2; Appendix C labels supportive feedback-alignment material as intuition rather than a formal theorem. The exact factorization sanity check is cross-referenced in the main text.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The main paper and appendices specify the datasets, architectures, local-rule family, broadcast modes, representative hyperparameters, and compute regime needed to reproduce the main results, and the supplementary code archive provides code, configuration files, and reproduction instructions; see Sec. 4 and Appendix B.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed

instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The submission includes a supplementary code archive containing the core model and training code, selected diagnostic utilities, representative configuration files, tests, and reproduction instructions for the main experimental families. The public benchmarks are cited and accessed through their standard public sources; see Appendix B.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The main text states the benchmark families, seed protocol, representative local-rule settings, and major training choices, while the appendices and supplementary code

1015 archive provide implementation details, configuration files, and operating-mode inventories;  
1016 see Sec. 4 and Appendix B.

1017 Guidelines:

- 1018 • The answer [N/A] means that the paper does not include experiments.
- 1019 • The experimental setting should be presented in the core of the paper to a level of detail  
1020 that is necessary to appreciate the results and make sense of them.
- 1021 • The full details can be provided either with the code, in appendix, or as supplemental  
1022 material.

## 1023 7. Experiment statistical significance

1024 Question: Does the paper report error bars suitably and correctly defined or other appropriate  
1025 information about the statistical significance of the experiments?

1026 Answer: [Yes]

1027 Justification: The paper reports means and standard deviations across 5 seeds for the main  
1028 results and across 3 seeds for selected appendix-only support sweeps, states this explicitly  
1029 in the text, and shows error bars or seed variability in the relevant tables and figures; see  
1030 Sec. 4 and Appendix A.

1031 Guidelines:

- 1032 • The answer [N/A] means that the paper does not include experiments.
- 1033 • The authors should answer [Yes] if the results are accompanied by error bars, confidence  
1034 intervals, or statistical significance tests, at least for the experiments that support the  
1035 main claims of the paper.
- 1036 • The factors of variability that the error bars are capturing should be clearly stated (for  
1037 example, train/test split, initialization, random drawing of some parameter, or overall  
1038 run with given experimental conditions).
- 1039 • The method for calculating the error bars should be explained (closed form formula,  
1040 call to a library function, bootstrap, etc.)
- 1041 • The assumptions made should be given (e.g., Normally distributed errors).
- 1042 • It should be clear whether the error bar is the standard deviation or the standard error  
1043 of the mean.
- 1044 • It is OK to report 1-sigma error bars, but one should state it. The authors should  
1045 preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis  
1046 of Normality of errors is not verified.
- 1047 • For asymmetric distributions, the authors should be careful not to show in tables or  
1048 figures symmetric error bars that would yield results that are out of range (e.g., negative  
1049 error rates).
- 1050 • If error bars are reported in tables or plots, the authors should explain in the text how  
1051 they were calculated and reference the corresponding figures or tables in the text.

## 1052 8. Experiments compute resources

1053 Question: For each experiment, does the paper provide sufficient information on the com-  
1054 puter resources (type of compute workers, memory, time of execution) needed to reproduce  
1055 the experiments?

1056 Answer: [Yes]

1057 Justification: Appendix B describes the institutional GPU cluster setting, single A100/H100-  
1058 class GPUs, approximate memory class, per-family runtime ranges, and approximate GPU-  
1059 hours for the reported experiments.

1060 Guidelines:

- 1061 • The answer [N/A] means that the paper does not include experiments.
- 1062 • The paper should indicate the type of compute workers CPU or GPU, internal cluster,  
1063 or cloud provider, including relevant memory and storage.
- 1064 • The paper should provide the amount of compute required for each of the individual  
1065 experimental runs as well as estimate the total compute.

- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work uses standard public benchmarks and mechanistic simulations, does not involve human-subjects experimentation or restricted data collection, and we are not aware of any aspect that violates the NeurIPS Code of Ethics.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: The paper is a mechanistic study of dendritic credit assignment using public benchmarks and synthetic diagnostics, with no deployed system, human-subject data, scraped dataset, or direct societal application. The limitations section nevertheless warns against overgeneralizing the mechanism beyond the tested regimes; see Sec. 5.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]



1119 Justification: The submission does not release a high-risk generative model, foundation  
1120 model, or scraped dataset, so the safeguard question is not applicable here.

1121 Guidelines:

- 1122 • The answer [N/A] means that the paper poses no such risks.
- 1123 • Released models that have a high risk for misuse or dual-use should be released with  
1124 necessary safeguards to allow for controlled use of the model, for example by requiring  
1125 that users adhere to usage guidelines or restrictions to access the model or implementing  
1126 safety filters.
- 1127 • Datasets that have been scraped from the Internet could pose safety risks. The authors  
1128 should describe how they avoided releasing unsafe images.
- 1129 • We recognize that providing effective safeguards is challenging, and many papers do  
1130 not require this, but we encourage authors to take this into account and make a best  
1131 faith effort.

## 1132 12. Licenses for existing assets

1133 Question: Are the creators or original owners of assets (e.g., code, data, models), used in  
1134 the paper, properly credited and are the license and terms of use explicitly mentioned and  
1135 properly respected?

1136 Answer: [Yes]

1137 Justification: The paper credits all benchmark assets, does not redistribute third-party raw  
1138 assets, and includes an asset/access table with URLs, citation requirements, and license or  
1139 access terms where available; see Appendix B.

1140 Guidelines:

- 1141 • The answer [N/A] means that the paper does not use existing assets.
- 1142 • The authors should cite the original paper that produced the code package or dataset.
- 1143 • The authors should state which version of the asset is used and, if possible, include a  
1144 URL.
- 1145 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- 1146 • For scraped data from a particular source (e.g., website), the copyright and terms of  
1147 service of that source should be provided.
- 1148 • If assets are released, the license, copyright information, and terms of use in the  
1149 package should be provided. For popular datasets, `paperswithcode.com/datasets`  
1150 has curated licenses for some datasets. Their licensing guide can help determine the  
1151 license of a dataset.
- 1152 • For existing datasets that are re-packaged, both the original license and the license of  
1153 the derived asset (if it has changed) should be provided.
- 1154 • If this information is not available online, the authors are encouraged to reach out to  
1155 the asset's creators.

## 1156 13. New assets

1157 Question: Are new assets introduced in the paper well documented and is the documentation  
1158 provided alongside the assets?

1159 Answer: [Yes]

1160 Justification: The only new released asset is a supplementary code archive. It includes an  
1161 MIT license, README, manifest, reproduction instructions, representative configuration  
1162 files, tests, and the core model/training code needed for the main experimental families; the  
1163 paper does not release a new dataset, benchmark, or pretrained model.

1164 Guidelines:

- 1165 • The answer [N/A] means that the paper does not release new assets.
- 1166 • Researchers should communicate the details of the dataset/code/model as part of their  
1167 submissions via structured templates. This includes details about training, license,  
1168 limitations, etc.
- 1169 • The paper should discuss whether and how consent was obtained from people whose  
1170 asset is used.

- 1171                   • At submission time, remember to anonymize your assets (if applicable). You can either  
1172                   create an anonymized URL or include an anonymized zip file.

1173 **14. Crowdsourcing and research with human subjects**

1174           Question: For crowdsourcing experiments and research with human subjects, does the paper  
1175           include the full text of instructions given to participants and screenshots, if applicable, as  
1176           well as details about compensation (if any)?

1177           Answer: [N/A]

1178           Justification: The work does not involve crowdsourcing or direct research with human  
1179           subjects.

1180           Guidelines:

- 1181           • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
1182           with human subjects.
- 1183           • Including this information in the supplemental material is fine, but if the main contribu-  
1184           tion of the paper involves human subjects, then as much detail as possible should be  
1185           included in the main paper.
- 1186           • According to the NeurIPS Code of Ethics, workers involved in data collection, curation,  
1187           or other labor should be paid at least the minimum wage in the country of the data  
1188           collector.

1189 **15. Institutional review board (IRB) approvals or equivalent for research with human**  
1190 **subjects**

1191           Question: Does the paper describe potential risks incurred by study participants, whether  
1192           such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)  
1193           approvals (or an equivalent approval/review based on the requirements of your country or  
1194           institution) were obtained?

1195           Answer: [N/A]

1196           Justification: The work does not involve crowdsourcing or human-subjects research requiring  
1197           IRB review.

1198           Guidelines:

- 1199           • The answer [N/A] means that the paper does not involve crowdsourcing nor research  
1200           with human subjects.
- 1201           • Depending on the country in which research is conducted, IRB approval (or equivalent)  
1202           may be required for any human subjects research. If you obtained IRB approval, you  
1203           should clearly state this in the paper.
- 1204           • We recognize that the procedures for this may vary significantly between institutions  
1205           and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
1206           guidelines for their institution.
- 1207           • For initial submissions, do not include any information that would break anonymity (if  
1208           applicable), such as the institution conducting the review.

1209 **16. Declaration of LLM usage**

1210           Question: Does the paper describe the usage of LLMs if it is an important, original, or  
1211           non-standard component of the core methods in this research? Note that if the LLM is used  
1212           only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
1213           scientific rigor, or originality of the research, declaration is not required.

1214           Answer: [N/A]

1215           Justification: LLMs are not part of the proposed method, theorem, experiments, or analysis  
1216           pipeline, so no declaration is required for the core scientific contribution.

1217           Guidelines:

- 1218           • The answer [N/A] means that the core method development in this research does not  
1219           involve LLMs as any important, original, or non-standard components.
- 1220           • Please refer to our LLM policy in the NeurIPS handbook for what should or should not  
1221           be described.